

Basics of Computer Vision

Dmitry Ryabokon, github.com/dryabokon

Content

1. CNNs

- 1.1 CNNs out of box
- 1.2 Basic operations
- 1.3 Feature extraction
- 1.4 Transfer learning
- 1.5 Auto-encoders
- 1.6. GANs
- 1.7 Fine Tuning
- 1.8 Benchmarking the CNNs

2. Projective Geometry

- 2.1 KeyPoints detection
- 2.2 Matching the KeyPoints
- 2.3 Homography
- 2.4 Blending
- 2.5 Camera calibration
- 2.6 Stereopair rectification
- 2.7 Pose estimation

3. OpenGL

- 3.1 Aruco
- 3.2 Augmented reality

4. Object detection

- 4.1 YOLO
- 4.2 SSD
- 4.3 Mask RCNN
- 4.4 CAM
- 4.5 Captioning

5. Computer Vision Applications

- 5.1 Stereo Vision
- 5.2 Face Swap
- 5.3 Face Filter
- 5.4 Style Transfer

Part 1: CNNs

CNNs out of box

Basic operations

Feature extraction

Transfer learning

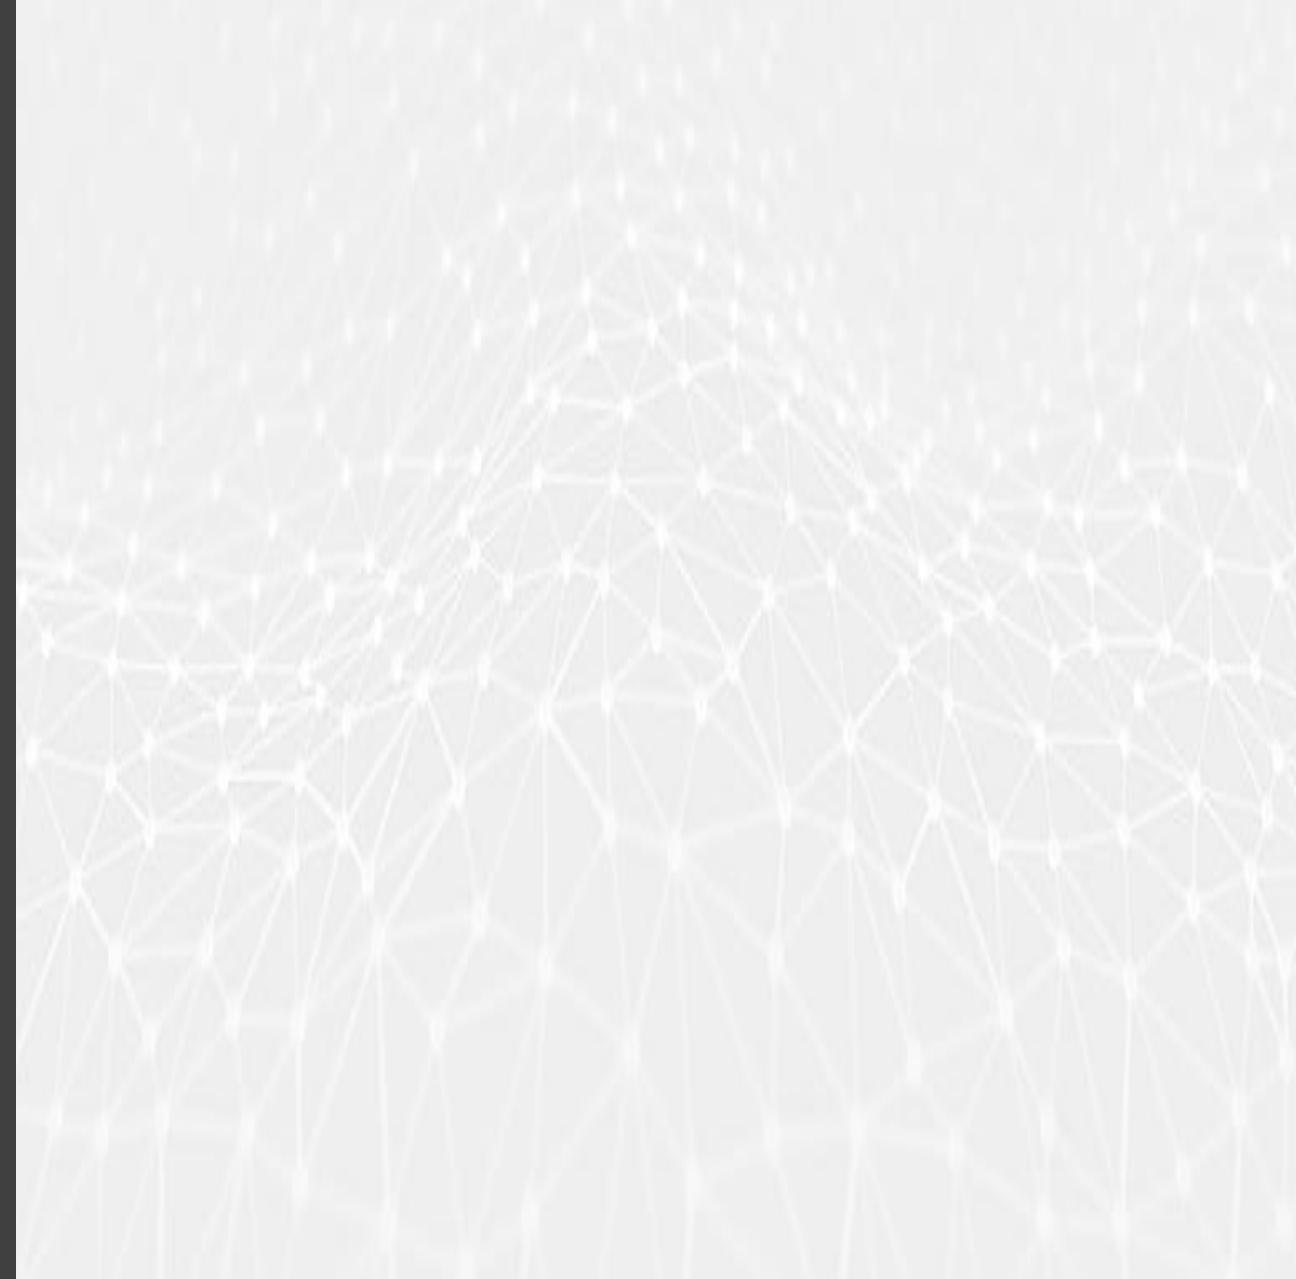
Auto-encoders

GANs

Fine Tuning

Benchmarking the CNNs

CNNs out of
box



CNNs out of box

The popular networks

Classification

- LeNet [Model](#)
- AlexNet [Model](#)
- VGG [Model](#)
- ResNet [Paper](#)
- YOLO9000 [Paper](#)
- DenseNet [Paper](#)

Segmentation

- FCN8 [Paper](#)
- SegNet [Paper](#)
- U-Net [Paper](#)
- E-Net [Paper](#)
- ResNetFCN [Paper](#)
- PSPNet [Paper](#)
- Mask RCNN [Paper](#)

Detection

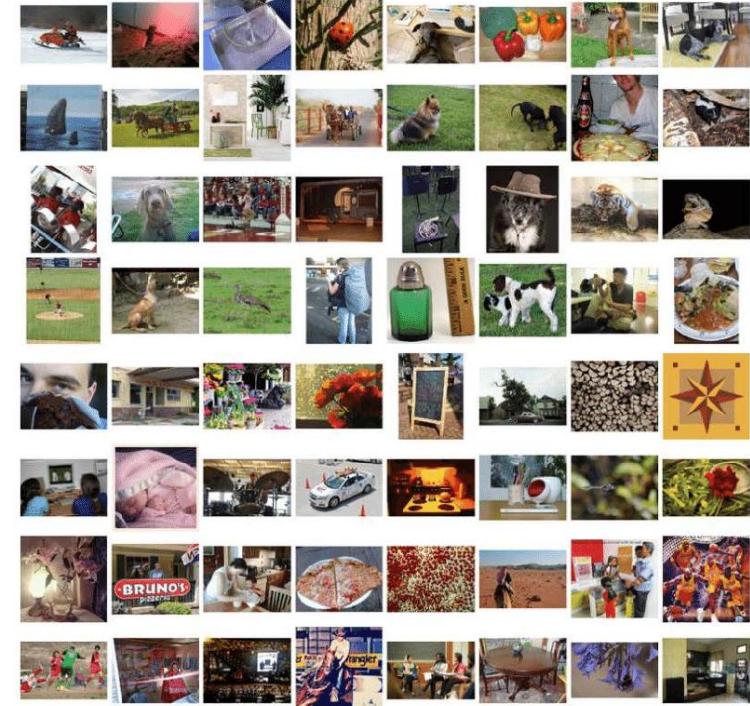
- Faster RCNN [Paper](#)
- SSD [Paper](#)
- YOLOv2 [Paper](#)
- R-FCN [Paper](#)

CNNs out of box

Some datasets available for research



MNIST: 10 classes, ~7000 ex. per class



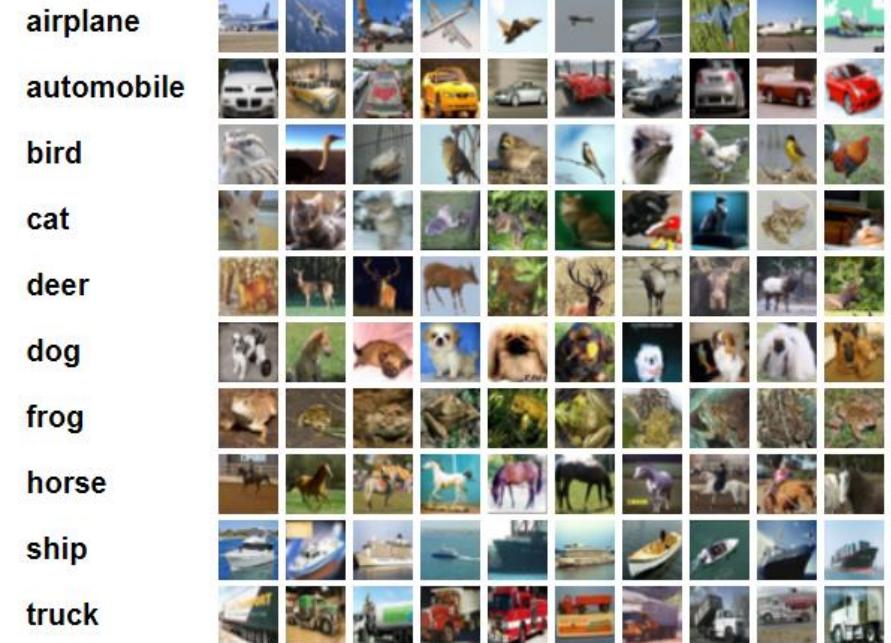
ImageNet: 1000 classes, ~100 ex per class

CNNs out of box

Some datasets available for research



The Street View House Numbers
10 classes, ~2000 ex. per class



CIFAR: 10 classes, 6000 ex. per class
100 classes, 600 ex per class

CNNs out of box

Some datasets available for research



Olivetti database: 40 classes

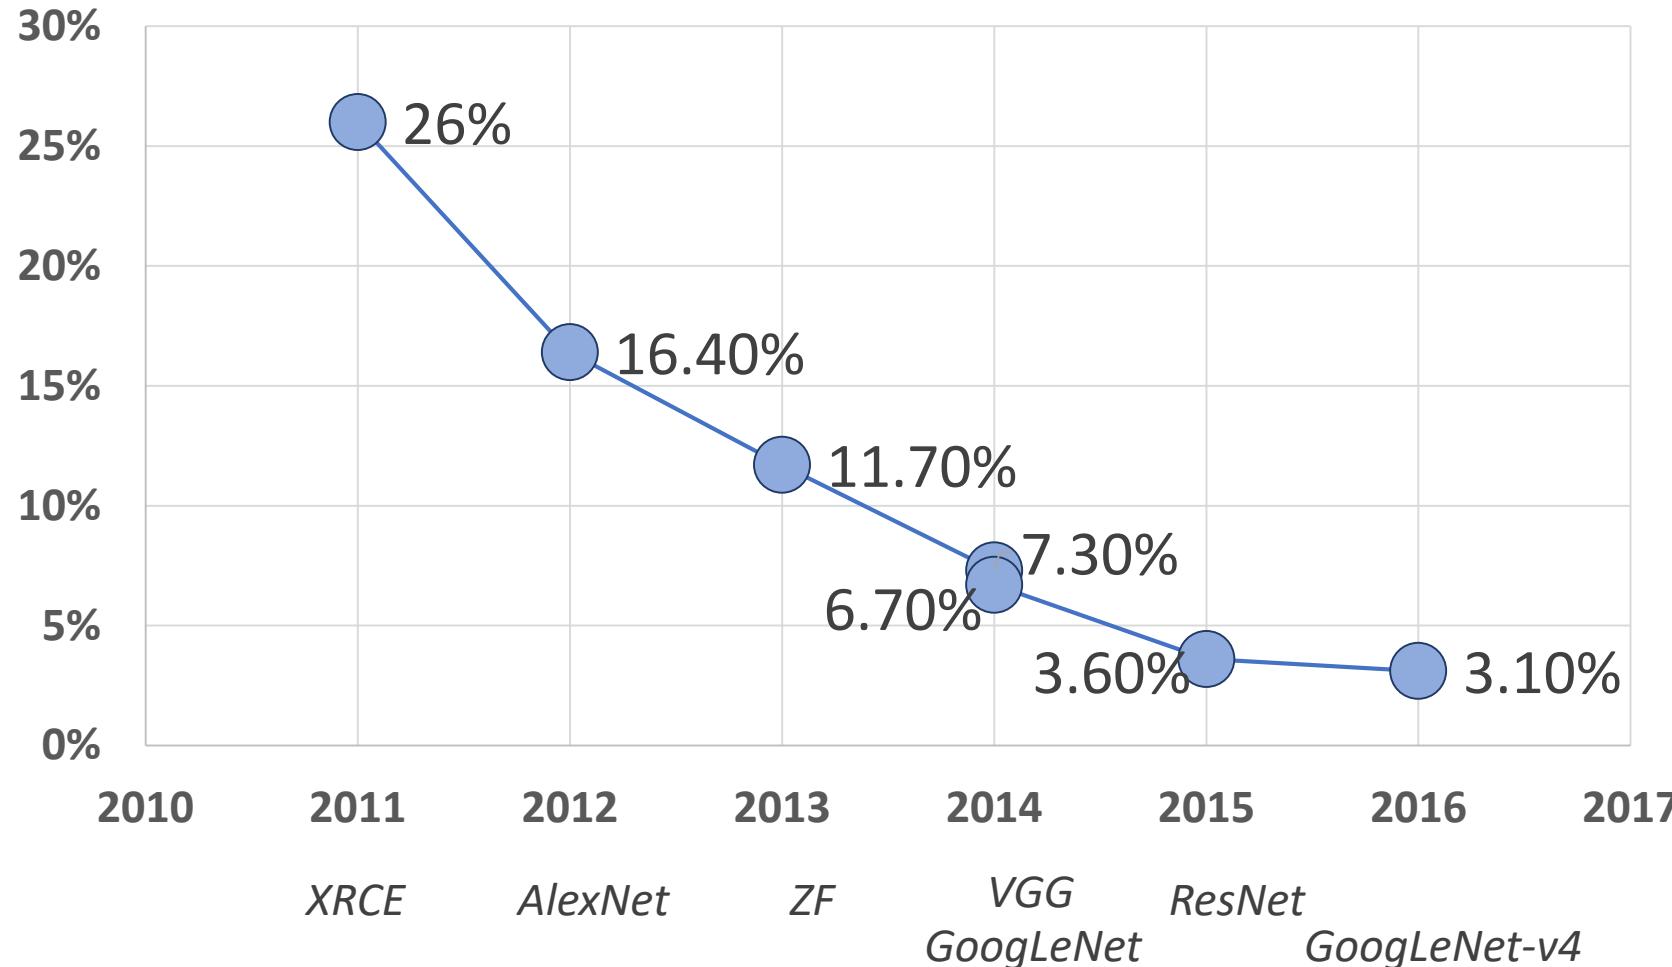
The screenshot shows the Kaggle datasets interface. At the top, there's a navigation bar with back, forward, and refresh buttons, and a URL field showing <https://www.kaggle.com/datasets>. Below the navigation is a blue header bar with the word "Datasets". To the right of the header is a "Documentation" link. Underneath the header, there are three tabs: "Public" (selected), "Your Datasets", and "Favorites". Below these tabs are filters for "Sizes", "File types", "Licenses", and "Tags". The main content area displays four dataset cards:

- Heart Disease UCI**: 101 datasets. Description: "https://archive.ics.uci.edu/ml/datasets/Heart+Disease". Last updated: 7 months ago (Version 1). Tags: biology, health, classification, binary clas...
- Graduate Admissions**: 424 datasets. Description: "Predicting admission from important parameters". Last updated: a month ago (Version 2). Tags: regression, model com..., random for..., + 2 more...
- FIFA 19 complete player dataset**: 334 datasets. Description: "18k+ FIFA 19 players, ~90 attributes extracted from the late...". Last updated: a month ago. Tags: sports, data visuali..., regression, + 2 more...
- Suicide Rates Overview 1985 to 2016**: 54 datasets. Description: "Compares socio-economic info with suicide rates by year a...". Tags: world, demograph..., economics

.. and much more @ kaggle

CNNs out of box

ImageNet Classification error



CNNs out of box

Usage of the neural networks

identify the name of a street (in France) from an image

compressing and decompressing images

semantic image segmentation

real-world image text extraction.

image classification

image-to-text

unsupervised learning

3D object reconstruction

image matching and retrieval

automatic speech recognition

localizing and identifying multiple objects in a single image

computer vision

discovery of latent 3D keypoints

predicting future video frames

<https://github.com/keras-team/keras/tree/master/examples>
<https://github.com/tensorflow/models/tree/master/research>

CNNs out of box

The screenshot shows a browser window with the URL https://tfhub.dev/google/imagenet/resnet_v2_50/feature_vector/1. The page has a header "TensorFlow Hub" with a search icon. Below the header, there's a breadcrumb navigation: "← imagenet/resnet_v2_50/feature_vector".

Usage

This module implements the common signature for computing [image feature vectors](#). It can be used like

```
module = hub.Module("https://tfhub.dev/google/imagenet/resnet_v2_50/feature_vector/1")
height, width = hub.get_expected_image_size(module)
images = ... # A batch of images with shape [batch_size, height, width, 3].
features = module(images) # Features with shape [batch_size, num_features].
```

...or using the signature name `image_feature_vector`. The output for each image in the batch is a feature vector of size `num_features = 2048`.

For this module, the size of the input image is fixed to `height x width = 224 x 224` pixels. The input `images` are expected to have color values in the range `[0,1]`, following the [common image input](#) conventions.

CNNs out of box

```
import tensorflow_hub as hub
import urllib.request
import cv2
import numpy
import json
import tensorflow as tf
# -----
URL = 'https://s3.amazonaws.com/deep-learning-models/image-models/imagenet_class_index.json'
data = json.loads(urllib.request.urlopen(URL).read().decode())
class_names = [data['%d'%i][1] for i in range(0,999)]
# -----
def example_predict():

    module = hub.Module("https://tfhub.dev/google/imagenet/resnet_v2_50/classification/1")
    img = cv2.imread('data/ex-natural/dog/dog_0000.jpg')
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = cv2.resize(img, (224, 224)).astype(numpy.float32)
    img = numpy.array([img]).astype(numpy.float32) / 255.0
    sess = tf.Session()
    sess.run(tf.global_variables_initializer())
    sess.run(tf.tables_initializer())

    outputs = module(dict(images=img), signature="image_classification", as_dict=True)
    prob = outputs["default"].eval(session=sess)[0]
    idx = numpy.argsort(-prob)[0]

    print(class_names[idx], prob[idx])
    sess.close()

    return
# -----
if __name__ == '__main__':
    example_predict()
```

CNNs out of box

← → ⌂ https://keras.io/getting-started/faq/#how-can-i-use-pre-trained-models-in-keras

Keras Documentation

Search docs

- Home
- Why use Keras
- GETTING STARTED
 - Guide to the Sequential model
 - Guide to the Functional API

FAQ

- How should I cite Keras?
- How can I run Keras on GPU?
- How can I run a Keras model on multiple GPUs?
- What does "sample", "batch", "epoch" mean?
- How can I save a Keras model?
- Why is the training loss much higher than the testing loss?
- How can I obtain the output of an intermediate layer?
- How can I use Keras with datasets that don't fit in memory?
- How can I interrupt training when the validation loss isn't decreasing

How can I use pre-trained models in Keras?

Code and pre-trained weights are available for the following image classification models:

- Xception
- VGG16
- VGG19
- ResNet50
- Inception v3
- Inception-ResNet v2
- MobileNet v1
- DenseNet
- NASNet
- MobileNet v2

They can be imported from the module `keras.applications`:

```
from keras.applications.xception import Xception
from keras.applications.vgg16 import VGG16
from keras.applications.vgg19 import VGG19
from keras.applications.resnet50 import ResNet50
from keras.applications.inception_v3 import InceptionV3
from keras.applications.inception_resnet_v2 import InceptionResNetV2
from keras.applications.mobilenet import MobileNet
from keras.applications.densenet import DenseNet121
from keras.applications.densenet import DenseNet169
from keras.applications.densenet import DenseNet201
from keras.applications.nasnet import NASNetLarge
from keras.applications.nasnet import NASNetMobile
from keras.applications.mobilenet_v2 import MobileNetV2

model = VGG16(weights='imagenet', include_top=True)
```

CNNs out of box

```
from keras.applications import MobileNet
from keras.applications.xception import Xception
from keras.applications.mobilenet import preprocess_input
import urllib.request
import cv2
import numpy
import json
from keras import backend as K
K.set_image_dim_ordering('tf')
# -----
URL = 'https://s3.amazonaws.com/deep-learning-models/image-models/imagenet_class_index.json'
data = json.loads(urllib.request.urlopen(URL).read().decode())
class_names = [data['%d'%i][1] for i in range(0,999)]
# -----
def example_predict():

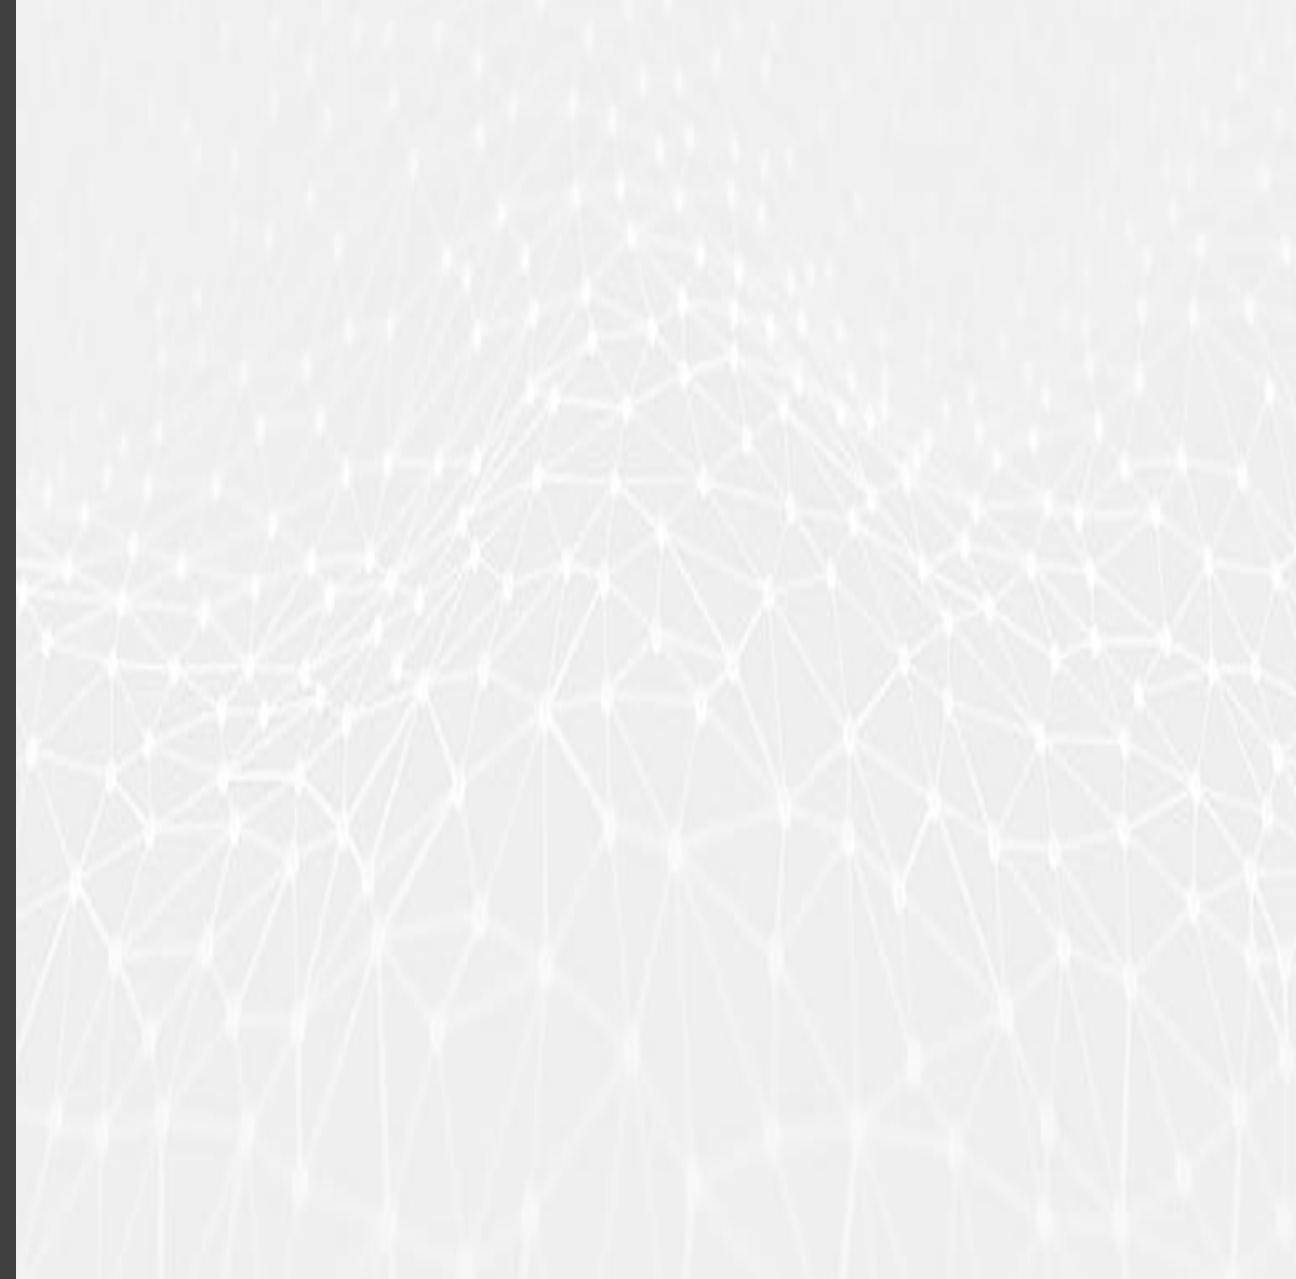
    CNN = MobileNet()
    #CNN = Xception()

    img = cv2.imread('data/ex-natural/dog/dog_0000.jpg')
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = cv2.resize(img, (224, 224)).astype(numpy.float32)

    prob = CNN.predict(preprocess_input(numpy.array([img])))
    idx = numpy.argsort(-prob[0])[0]
    print(class_names[idx], prob[0, idx])

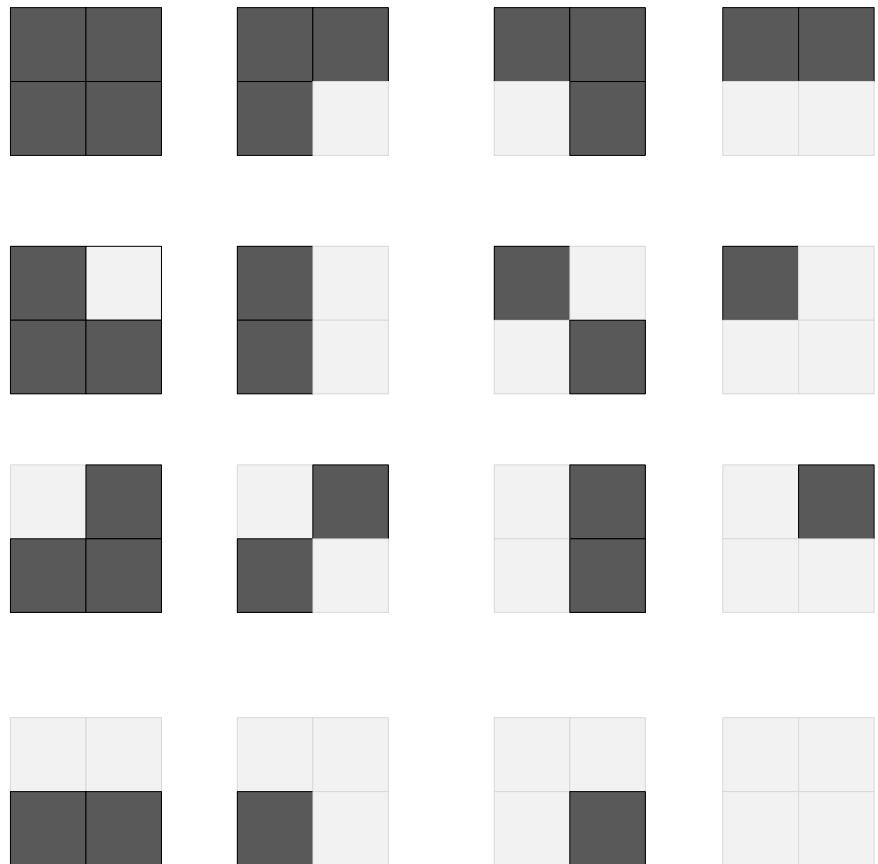
    return
# -----
if __name__ == '__main__':
    example_predict()
```

Basic operations



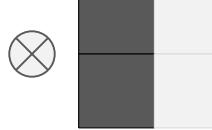
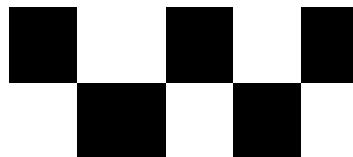
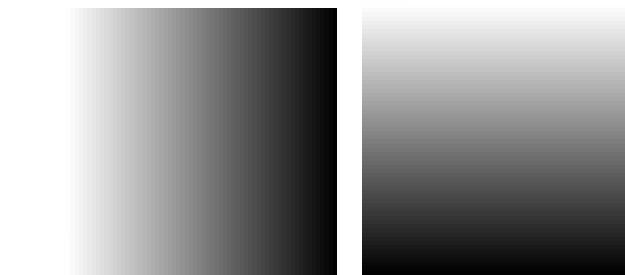
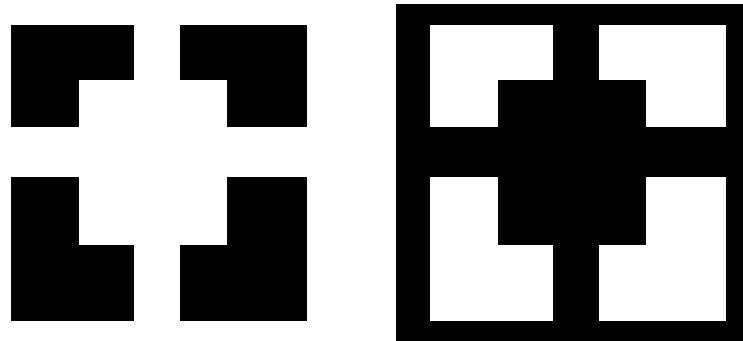
Basic operations

Convolution



Basic operations

Convolution

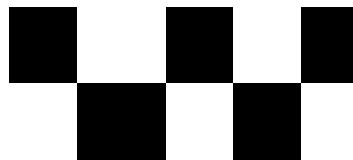
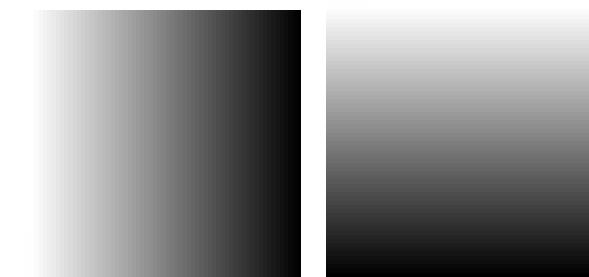
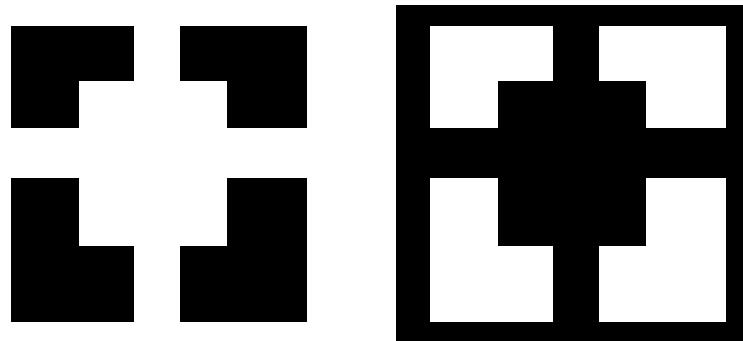


=



Basic operations

Convolution

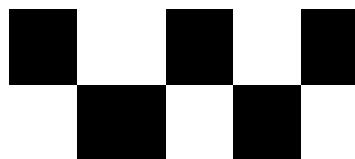
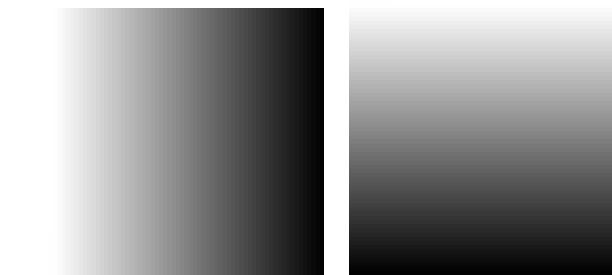
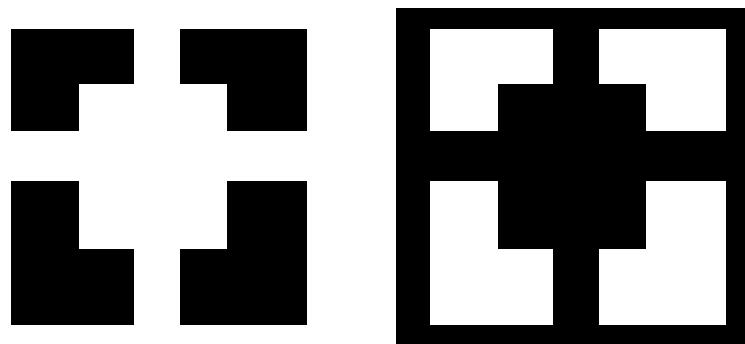


=



Basic operations

Convolution

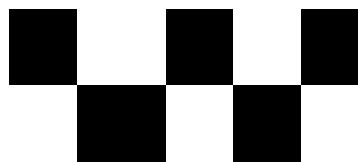
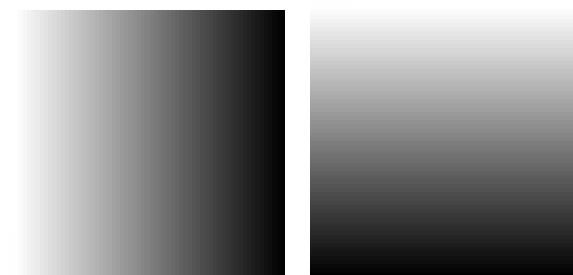
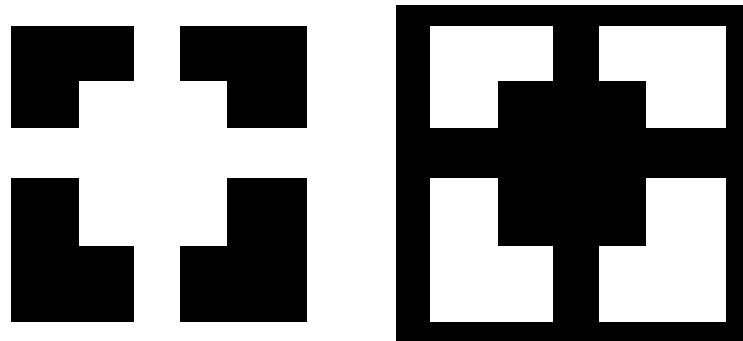


=

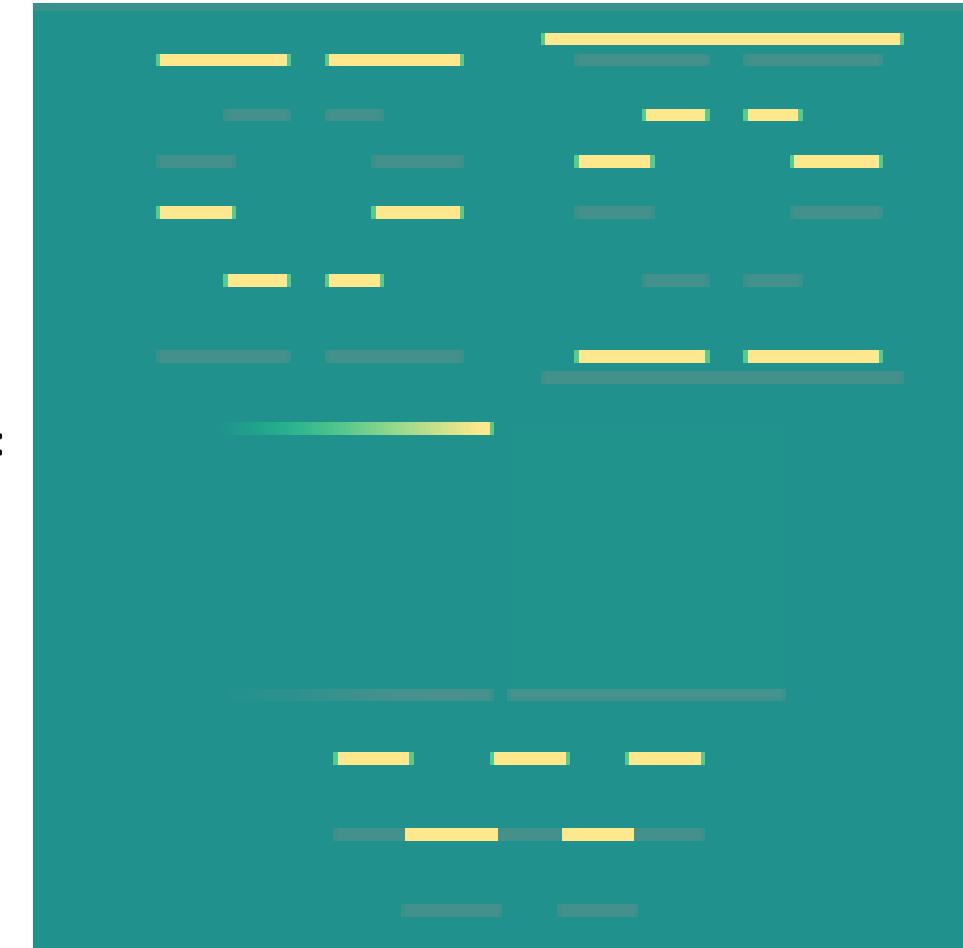


Basic operations

Convolution

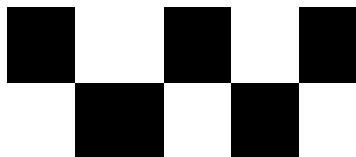
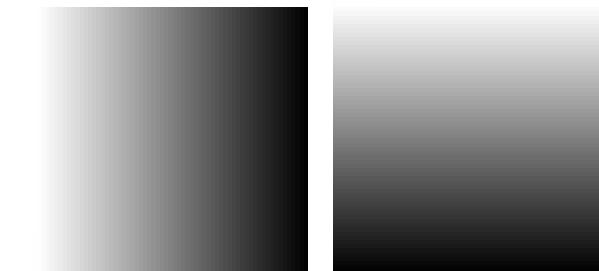


$$\otimes \begin{matrix} & & \\ & & \\ \text{---} & \text{---} & \text{---} \\ & & \end{matrix} =$$

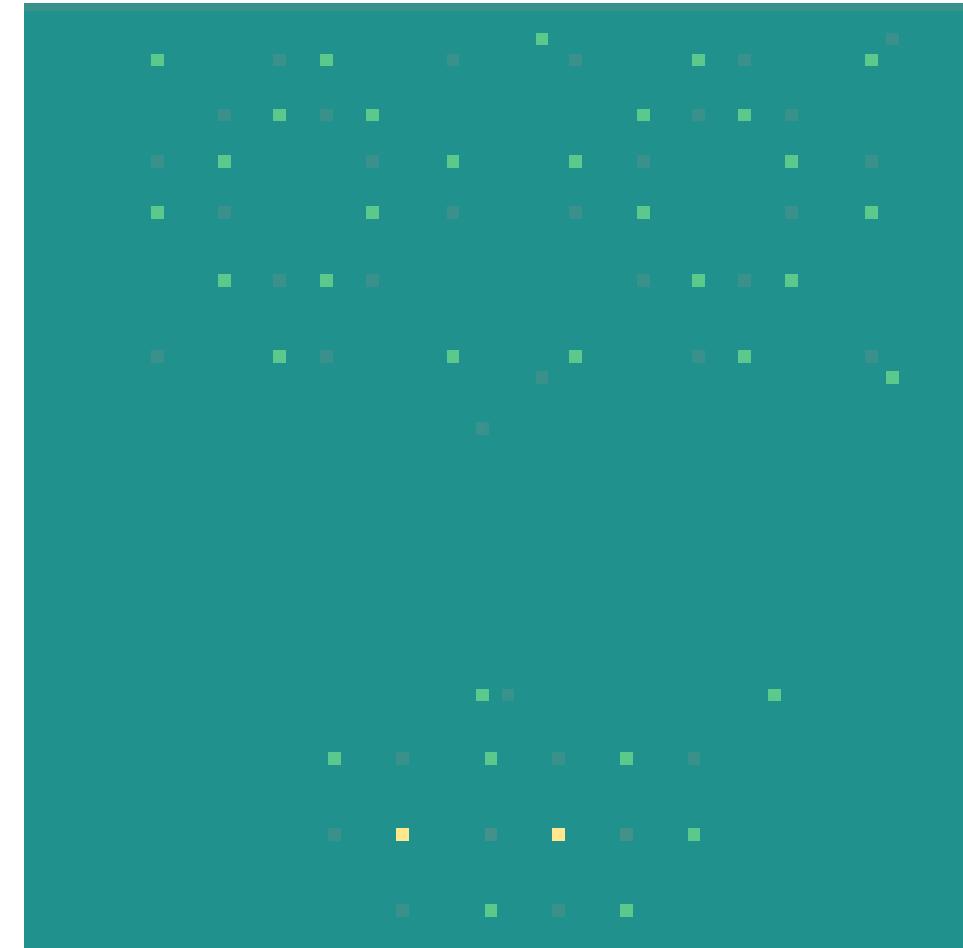


Basic operations

Convolution

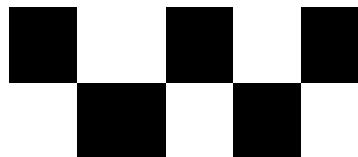
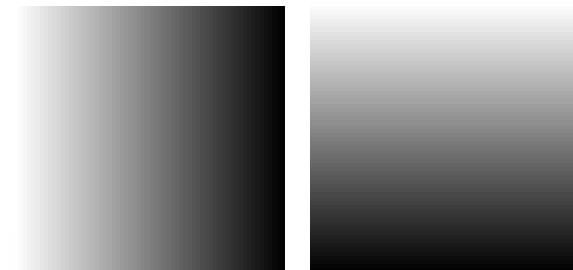
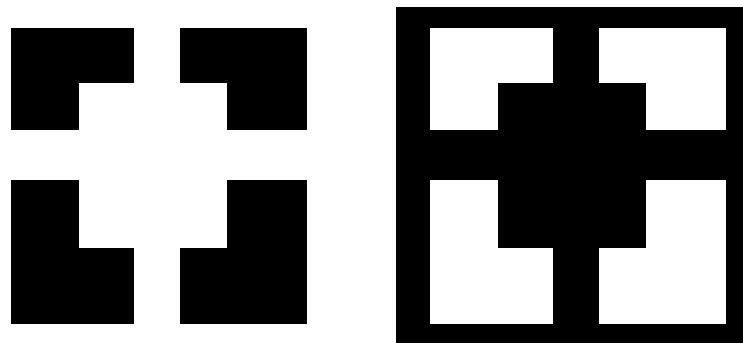


$$\otimes \begin{matrix} \text{dark gray} & \text{white} \\ \text{white} & \text{dark gray} \end{matrix} =$$

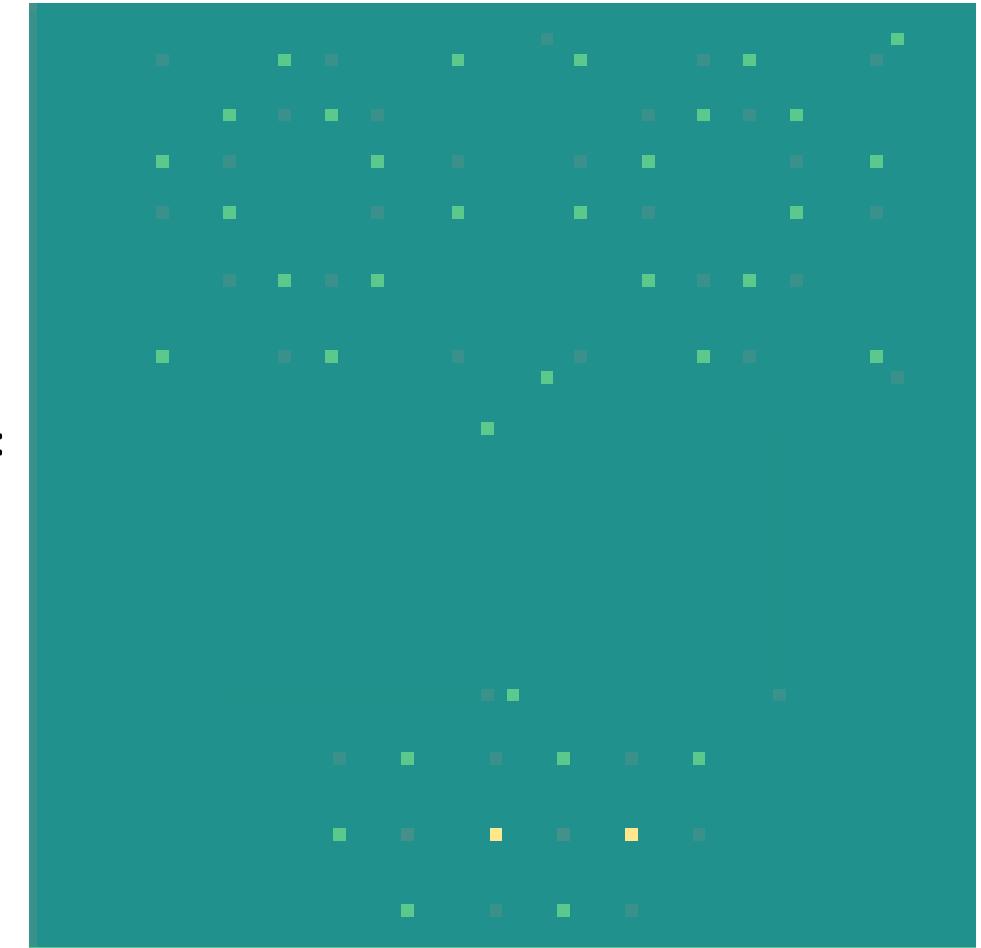


Basic operations

Convolution

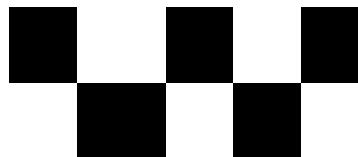
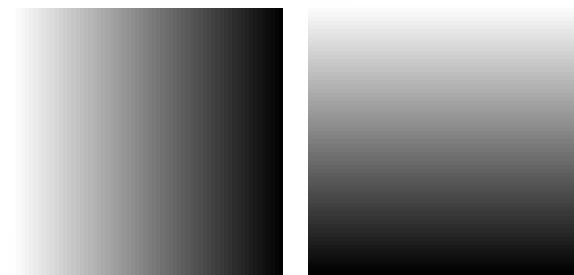
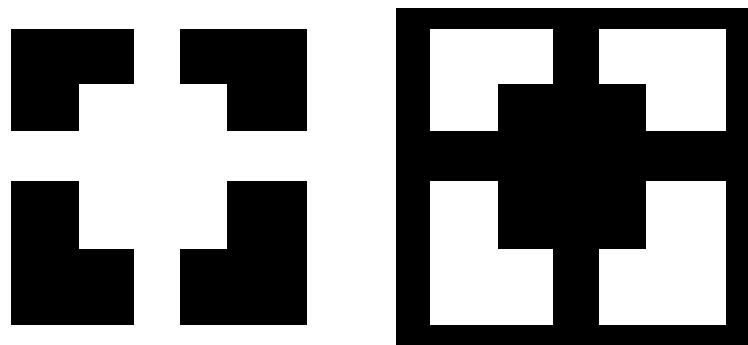


=

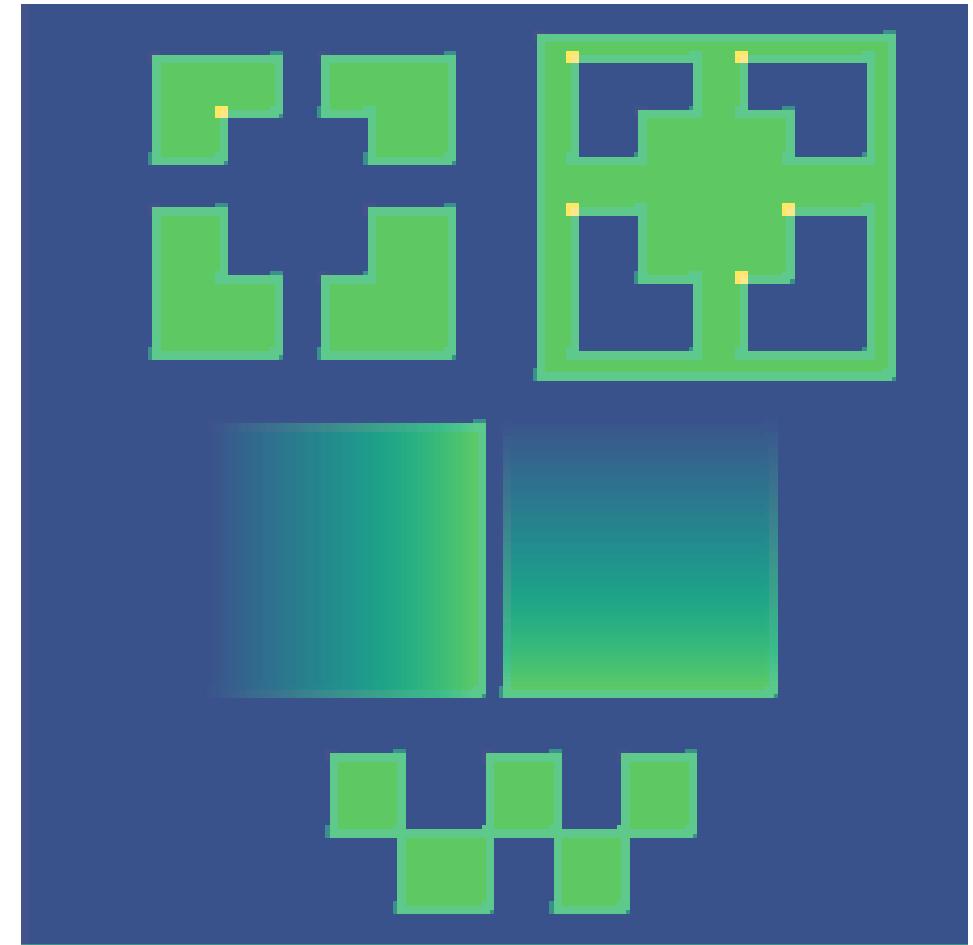


Basic operations

Convolution

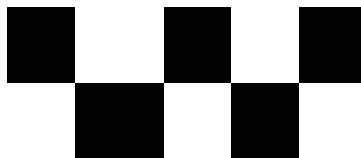
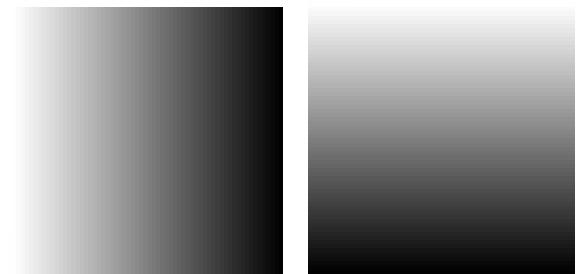
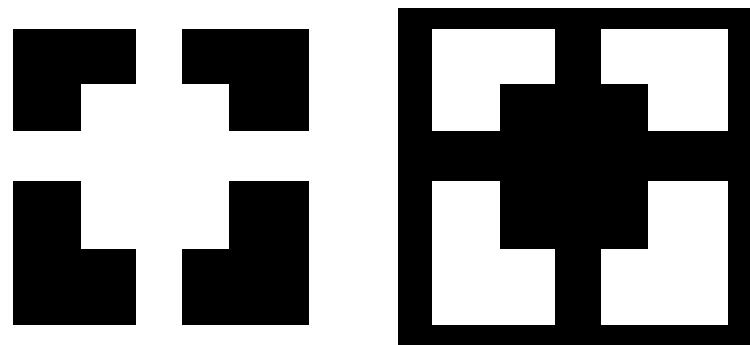


$$\otimes \begin{matrix} & & \\ & \text{---} & \\ & & \end{matrix} =$$

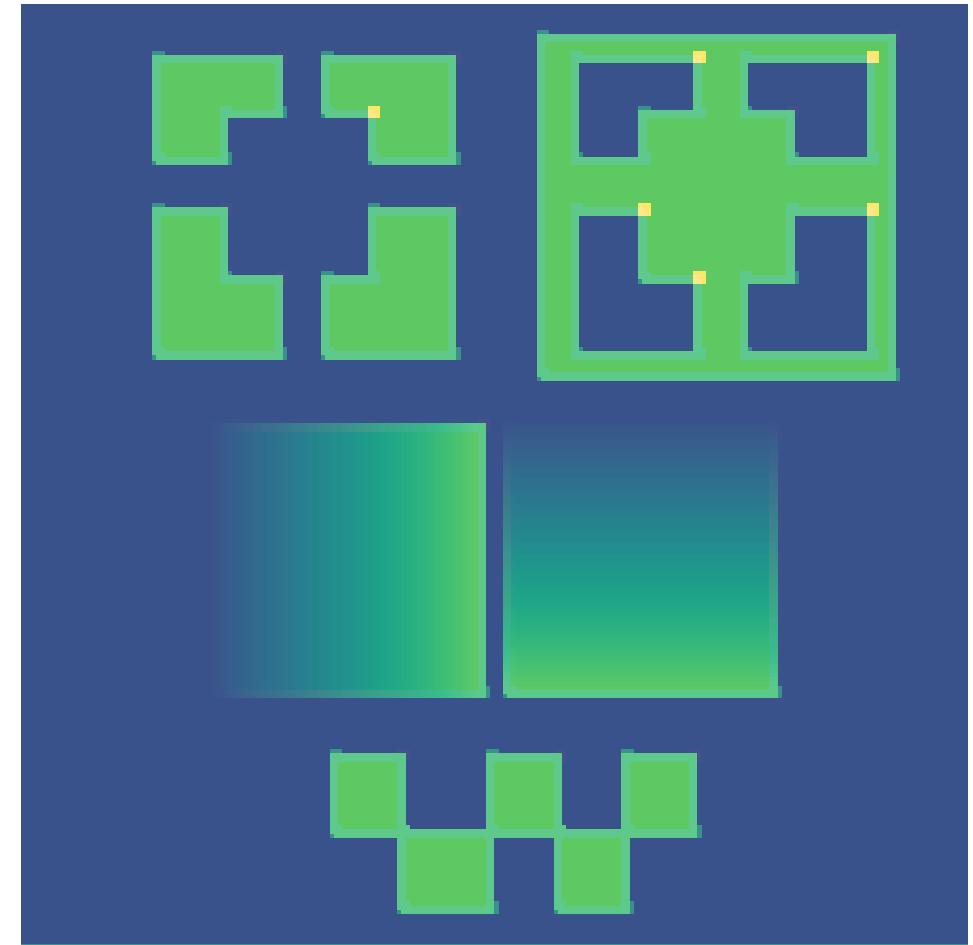


Basic operations

Convolution

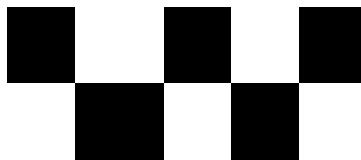
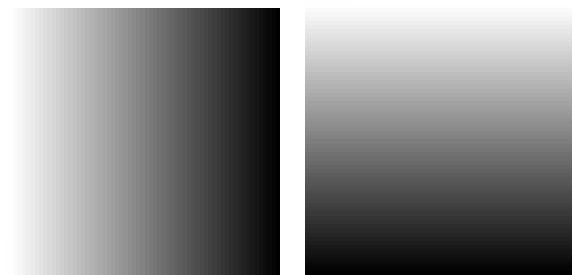
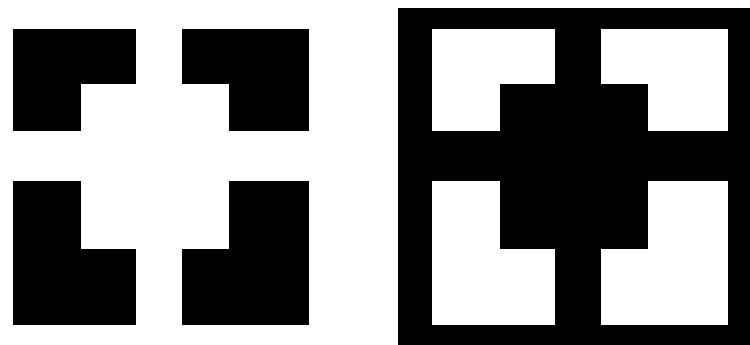


$$\otimes \begin{matrix} & & \\ & \text{dark gray} & \\ & & \end{matrix} =$$

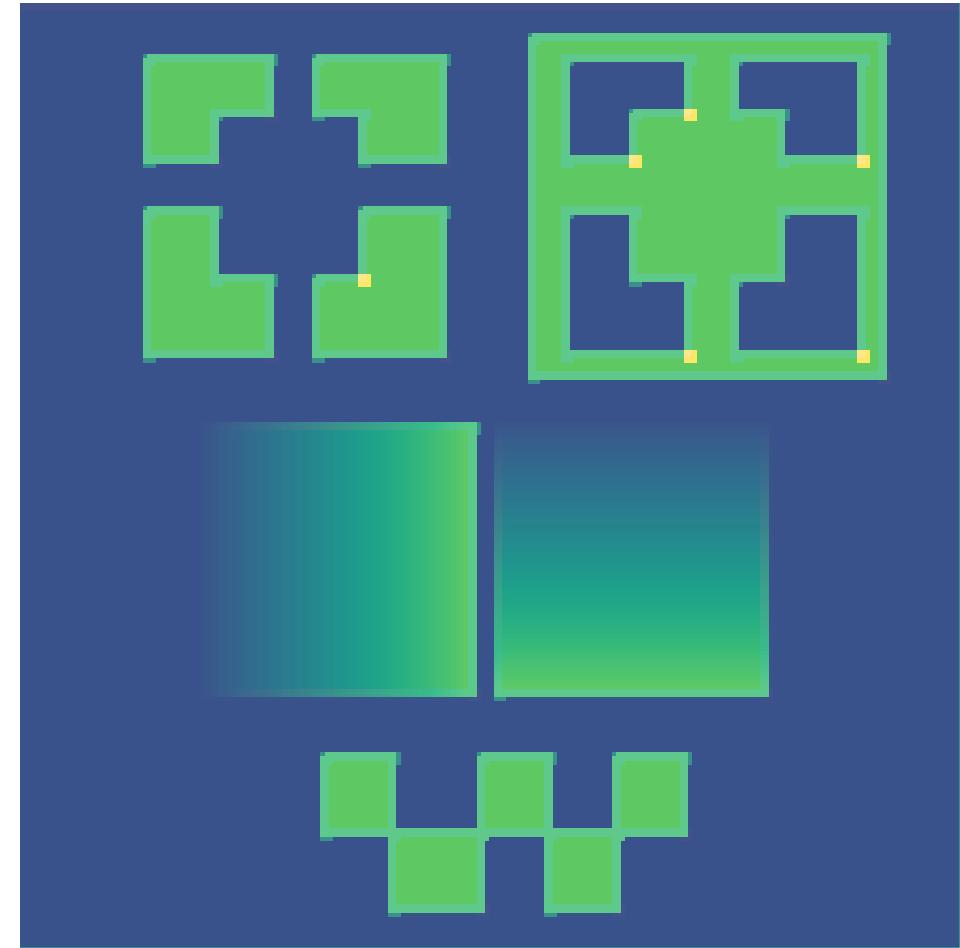


Basic operations

Convolution

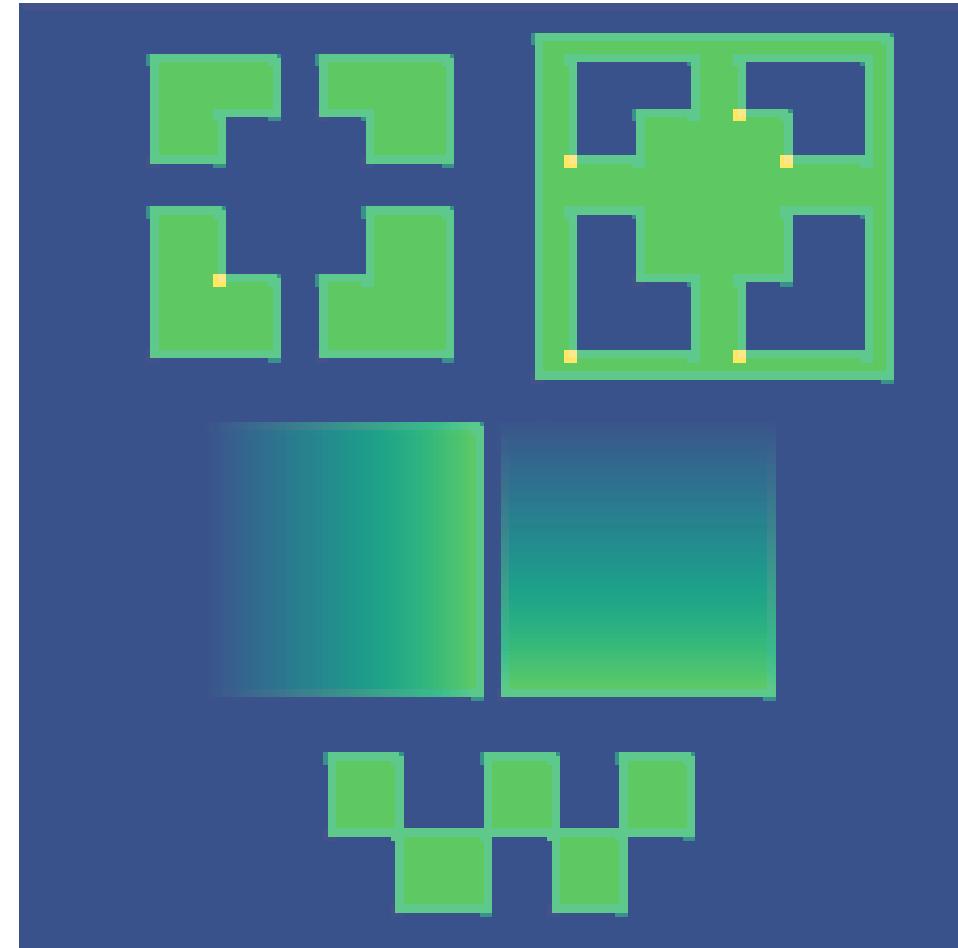
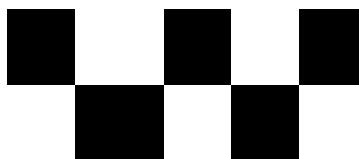
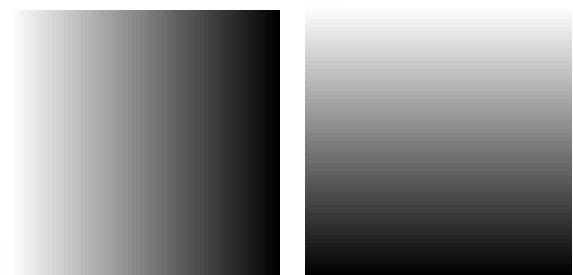
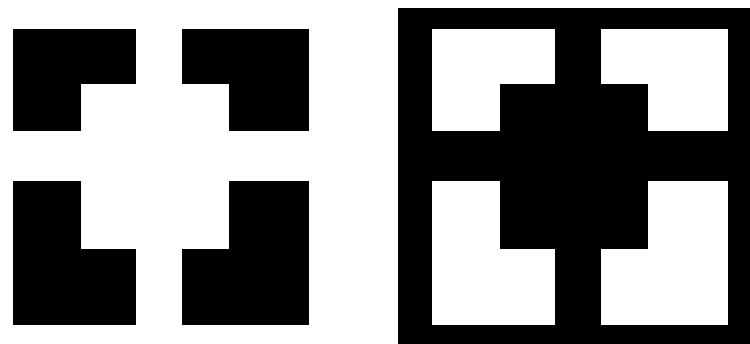


$$\otimes \begin{matrix} & & \\ & & \\ \textcolor{gray}{\otimes} & & \\ & & \end{matrix} =$$



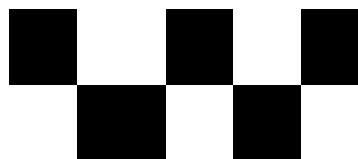
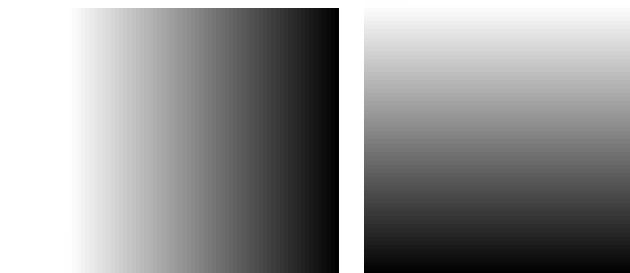
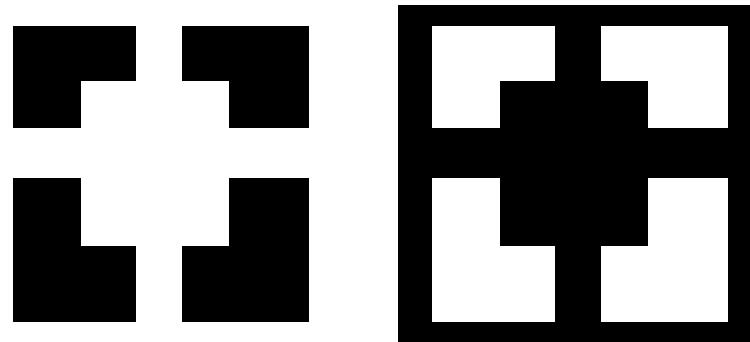
Basic operations

Convolution



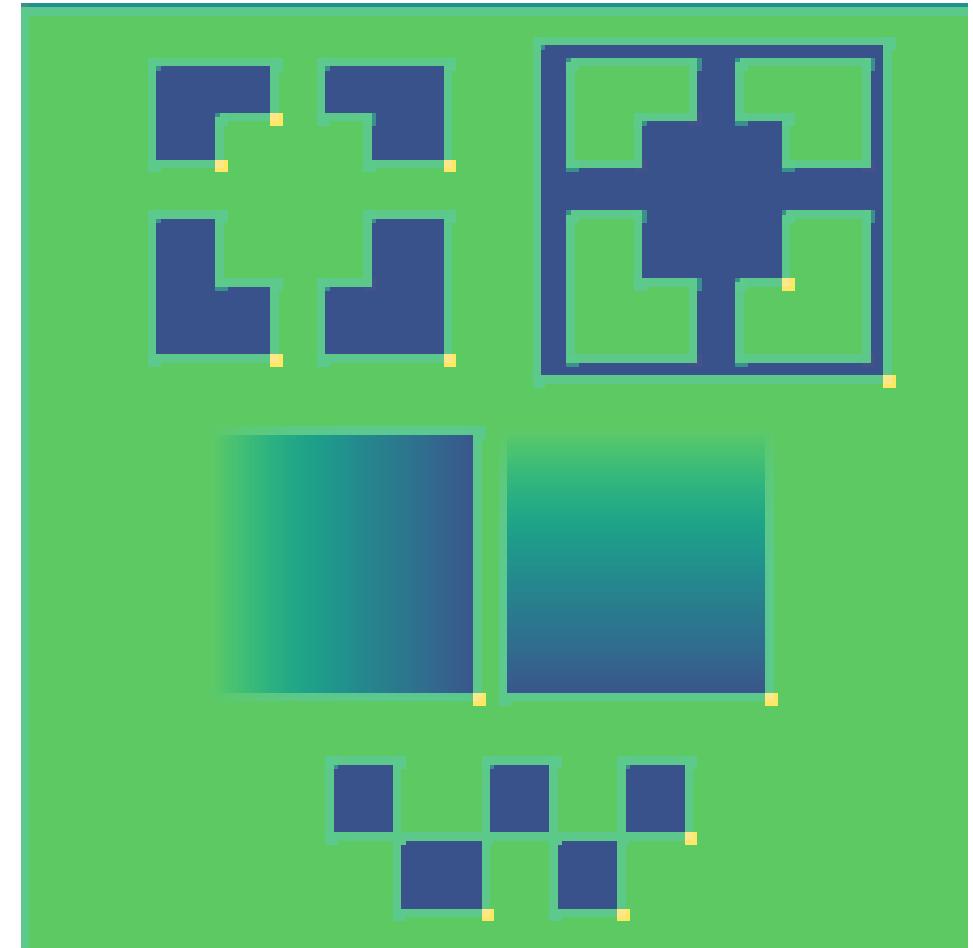
Basic operations

Convolution



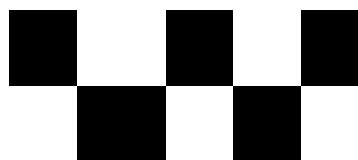
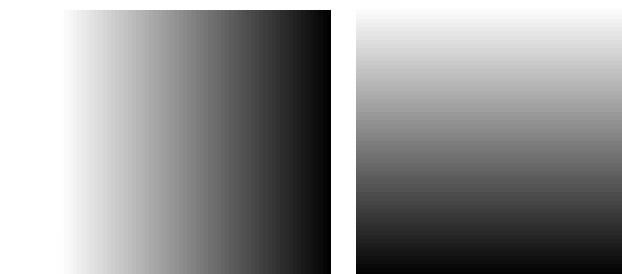
$$\otimes \begin{matrix} \text{dark gray} \\ \text{white} \end{matrix}$$

=

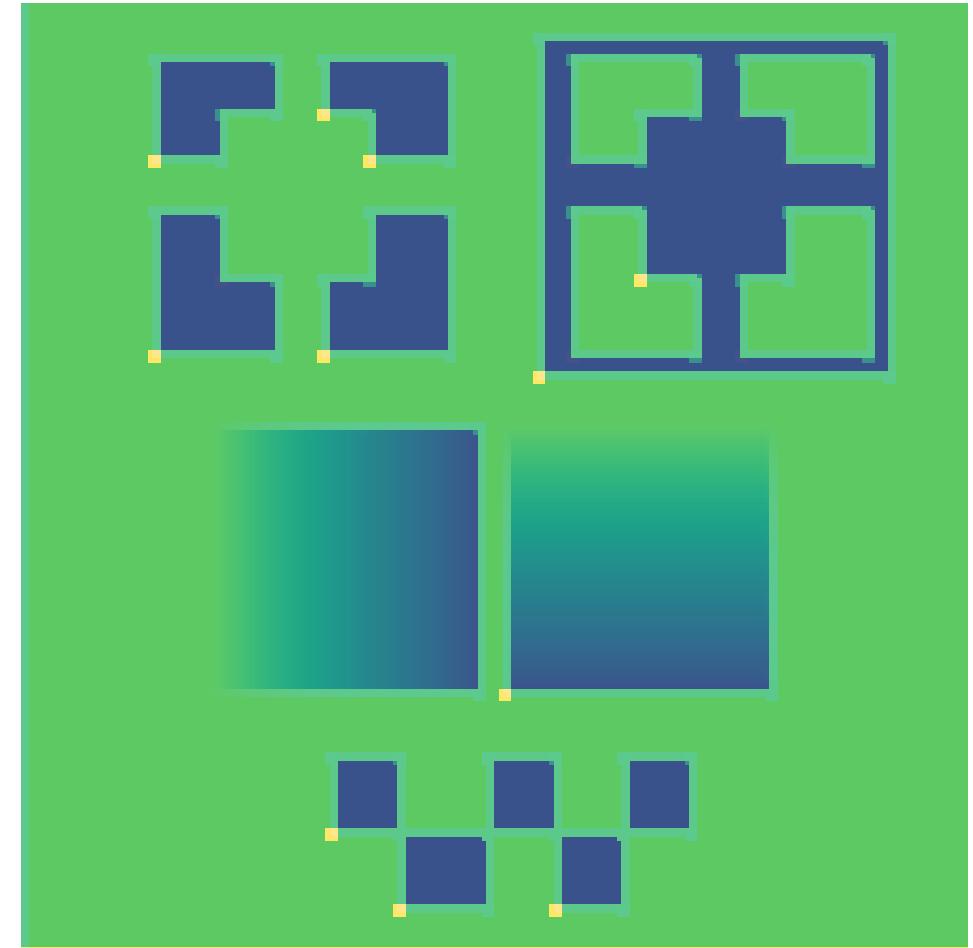


Basic operations

Convolution

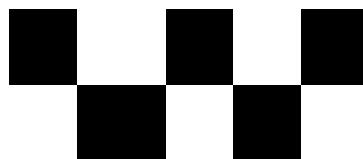
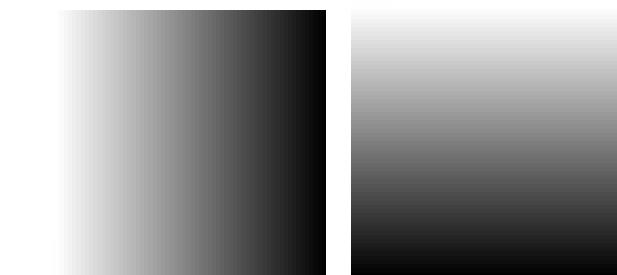
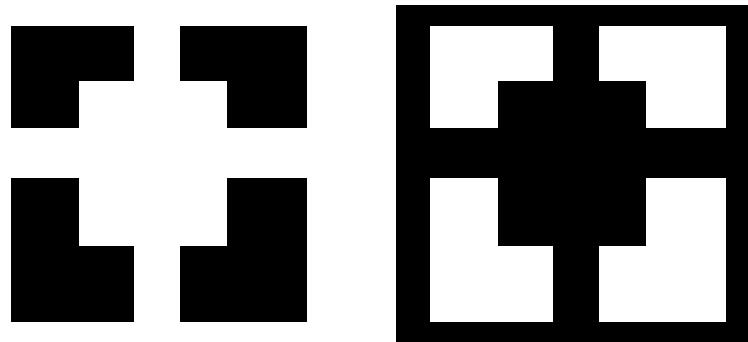


=

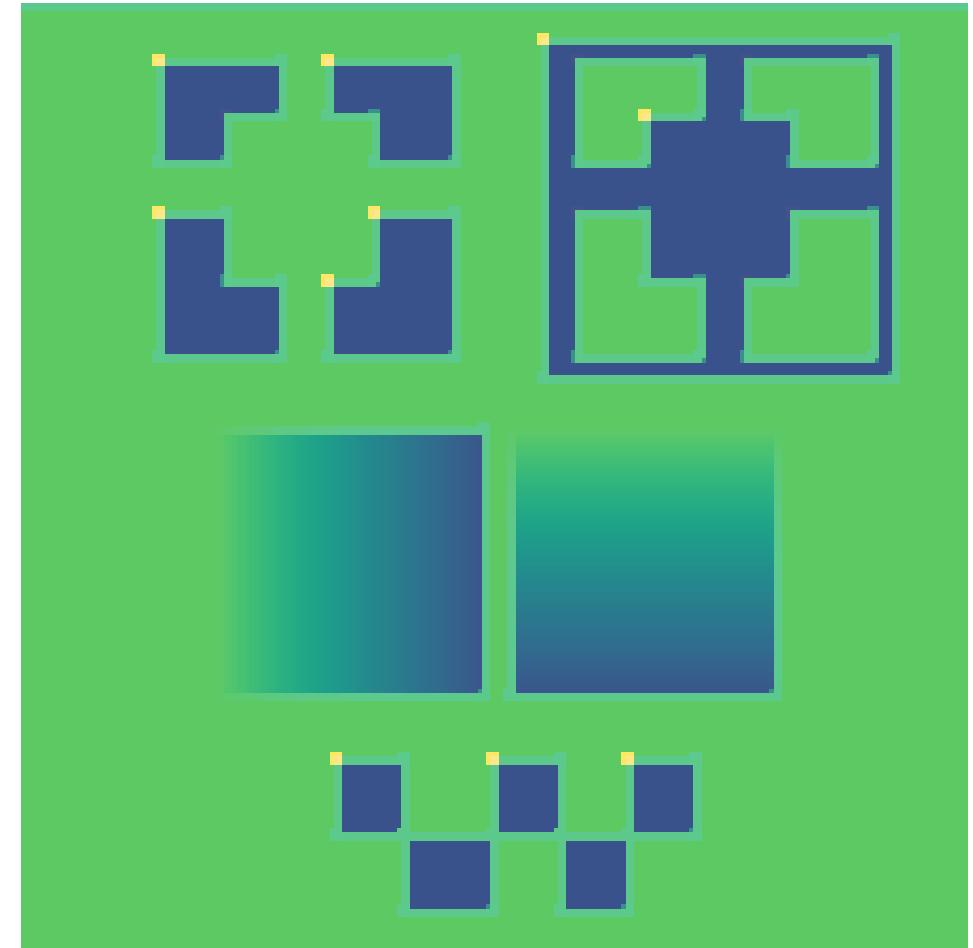


Basic operations

Convolution

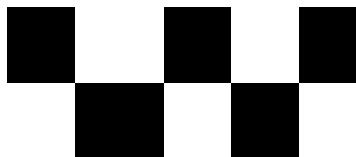
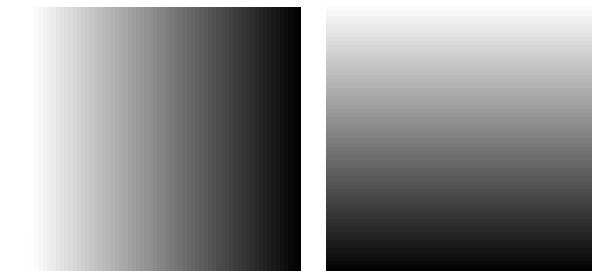


$$\otimes \quad \begin{matrix} & & \\ & & \\ & \text{---} & \\ & & \end{matrix} =$$



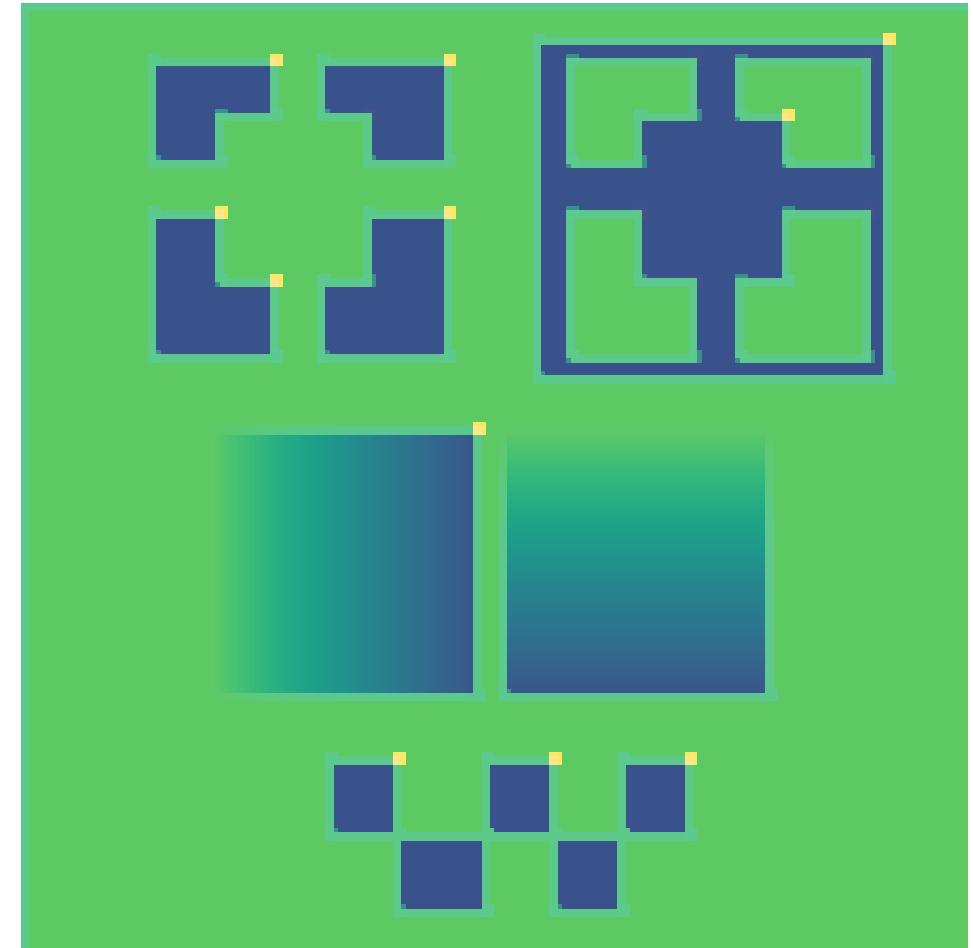
Basic operations

Convolution



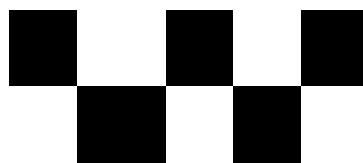
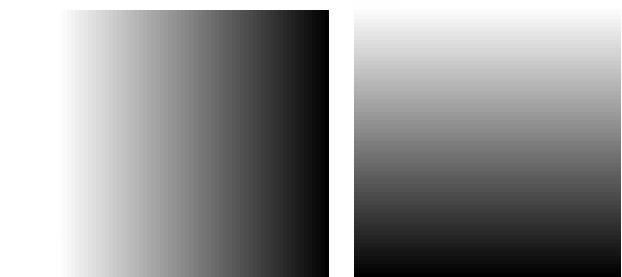
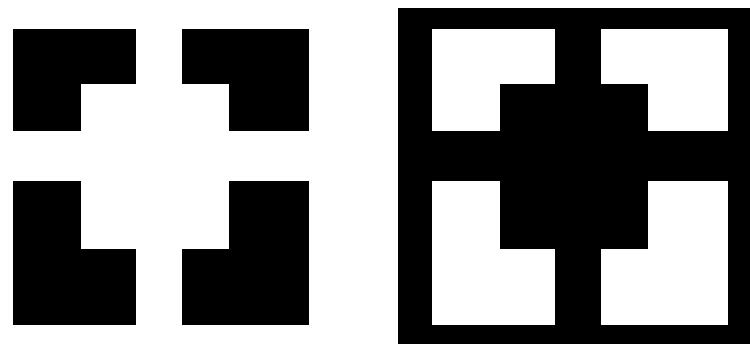
$$\otimes \begin{matrix} & & \\ & & \\ \text{---} & \text{---} & \text{---} \\ & & \end{matrix}$$

=

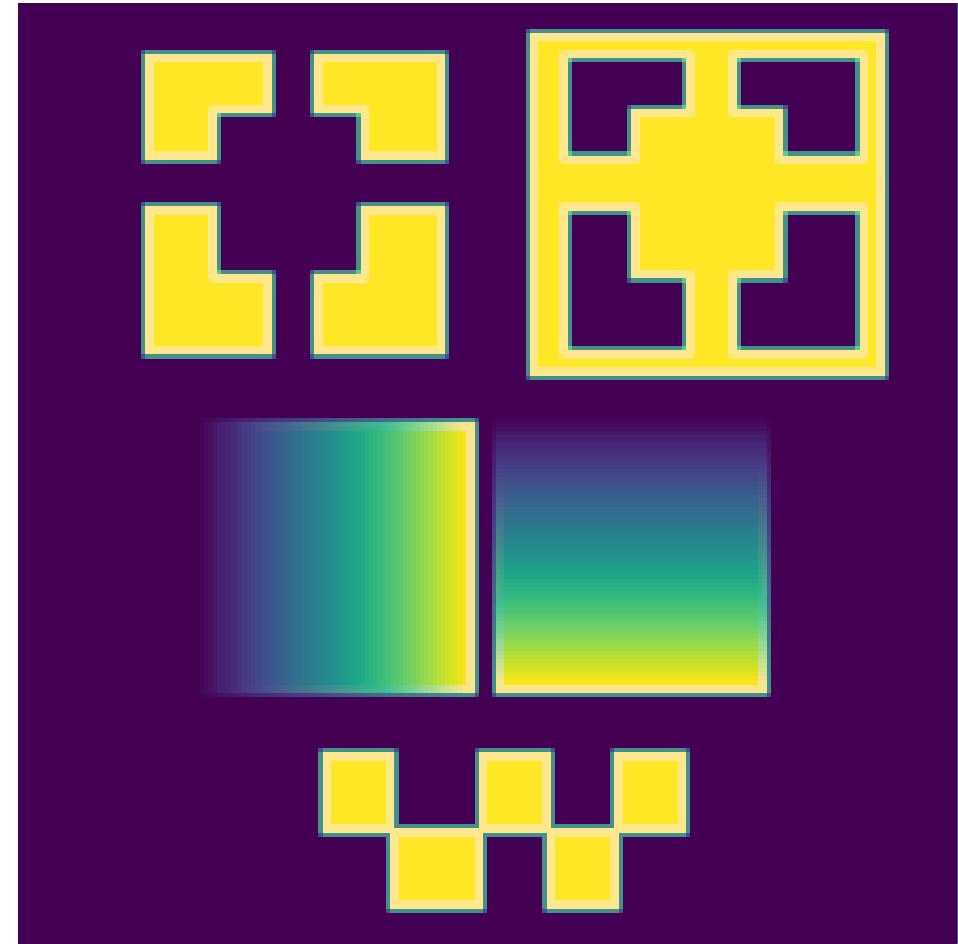


Basic operations

Convolution

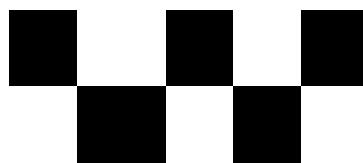
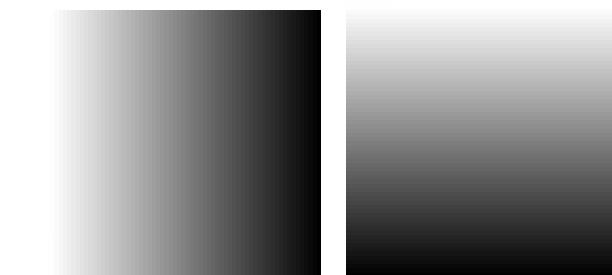
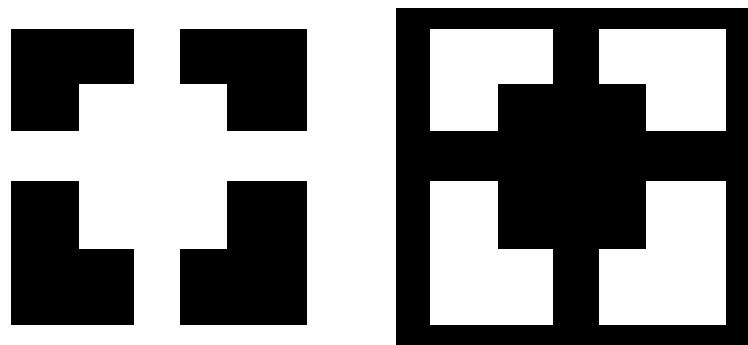


$$\otimes \begin{array}{|c|c|}\hline & & \\ \hline & & \\ \hline & & \\ \hline\end{array} =$$

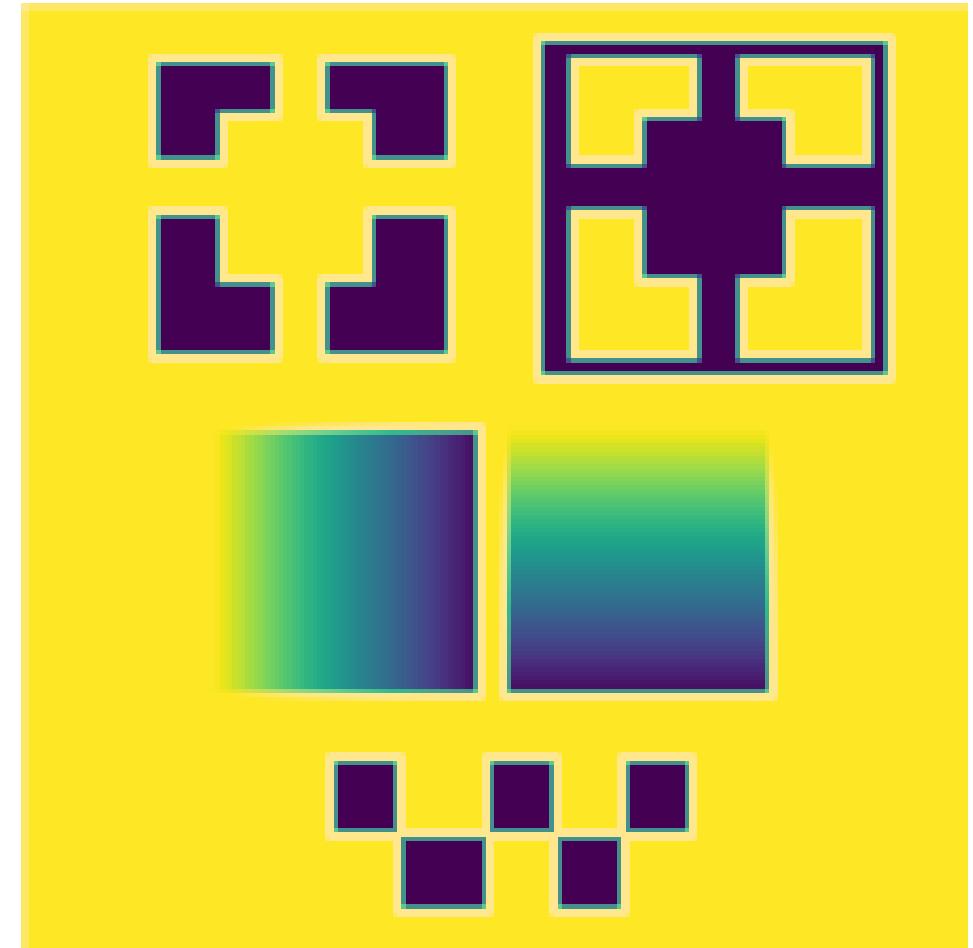


Basic operations

Convolution

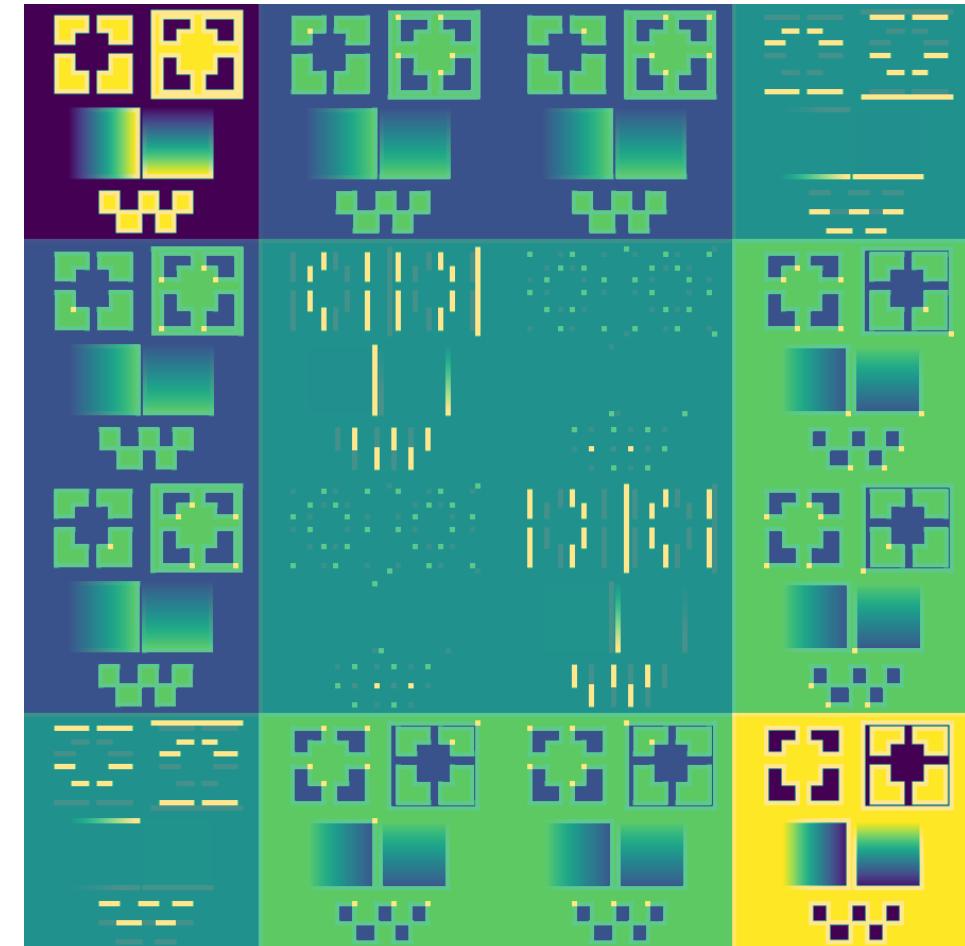
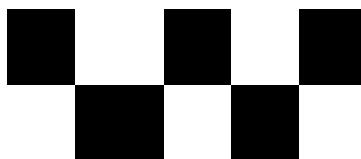
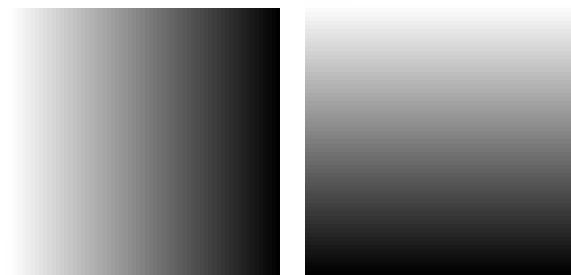
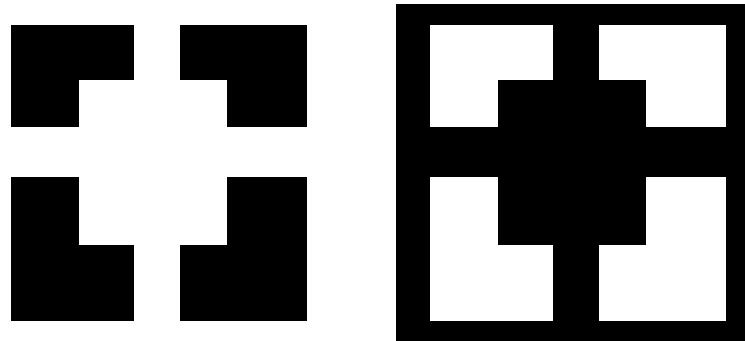


=



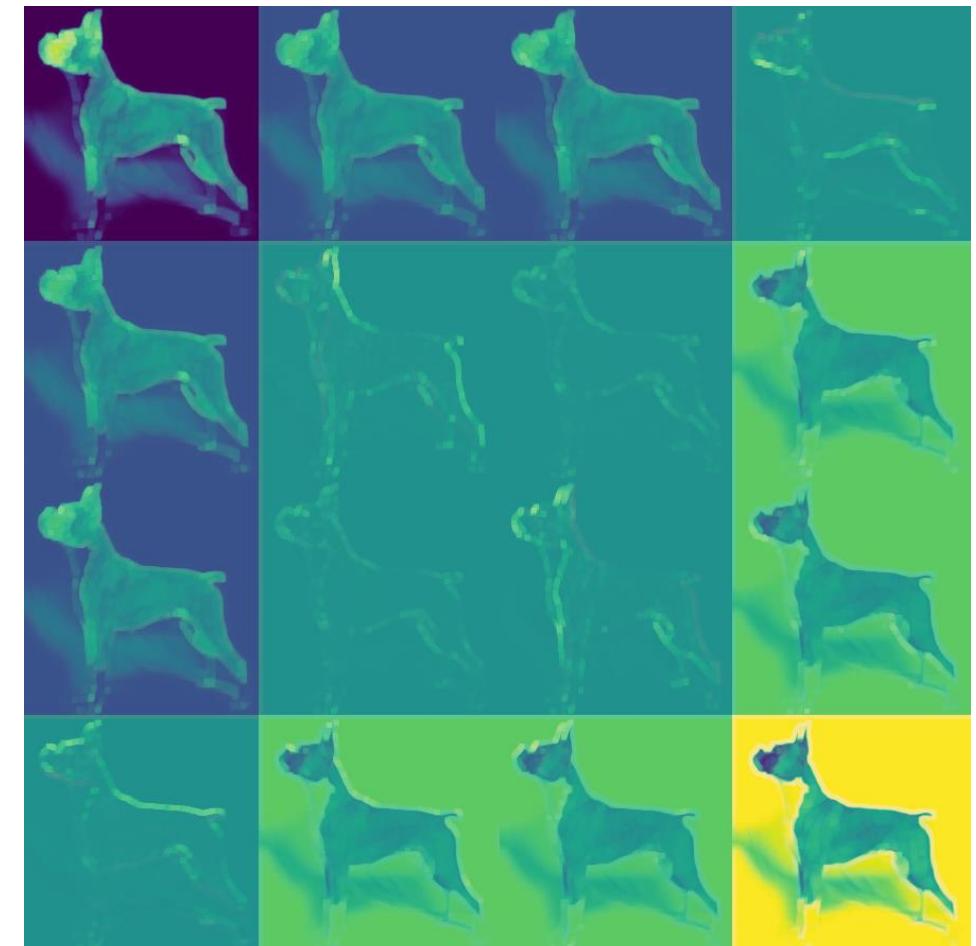
Basic operations

Convolution



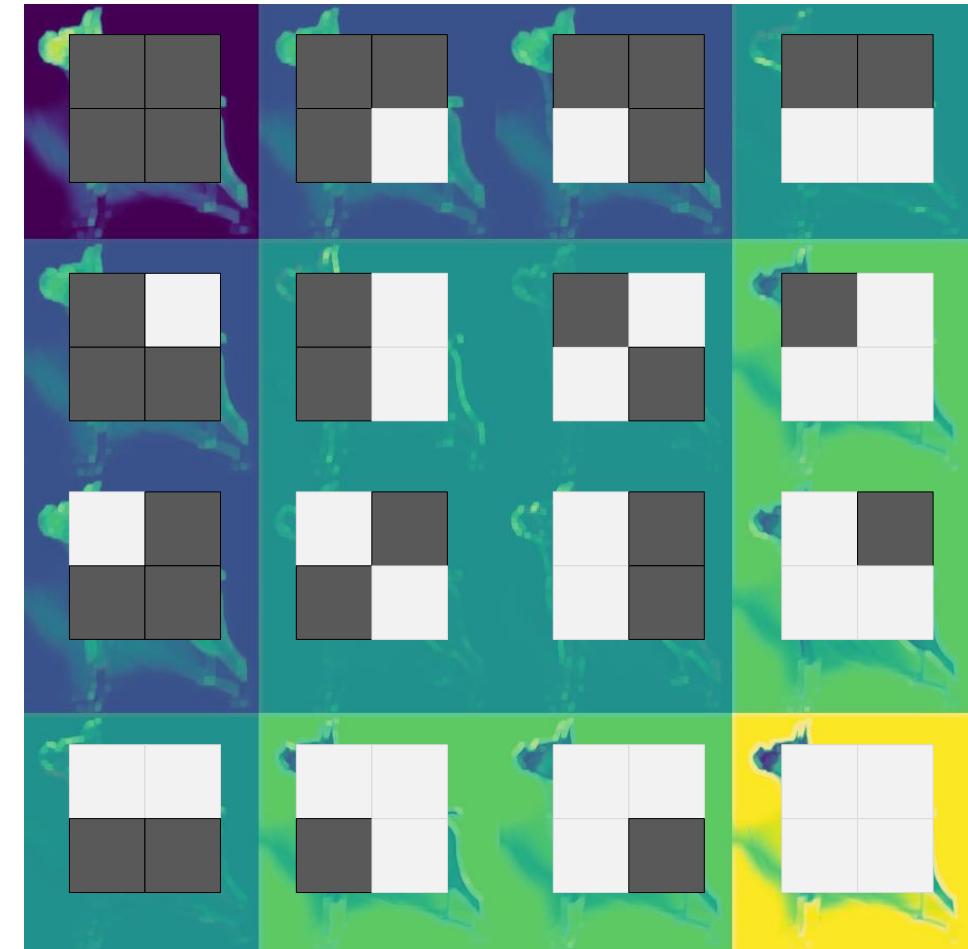
Basic operations

Convolution



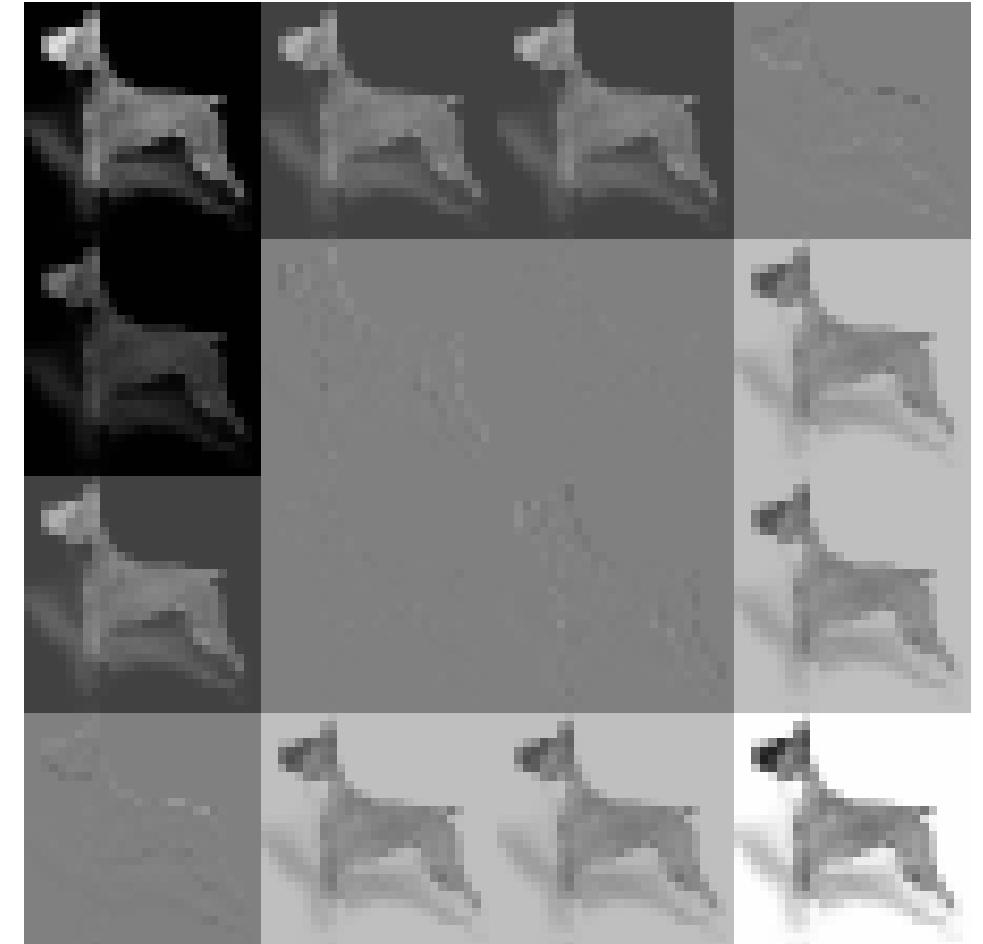
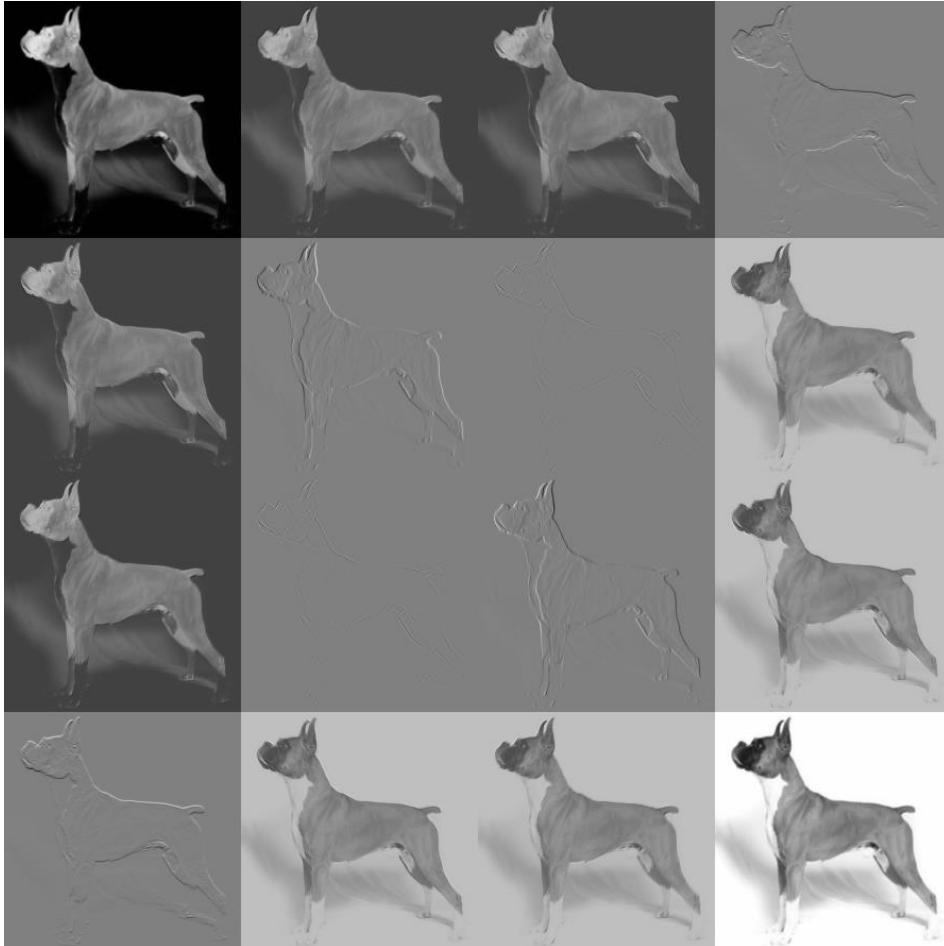
Basic operations

De-Convolution



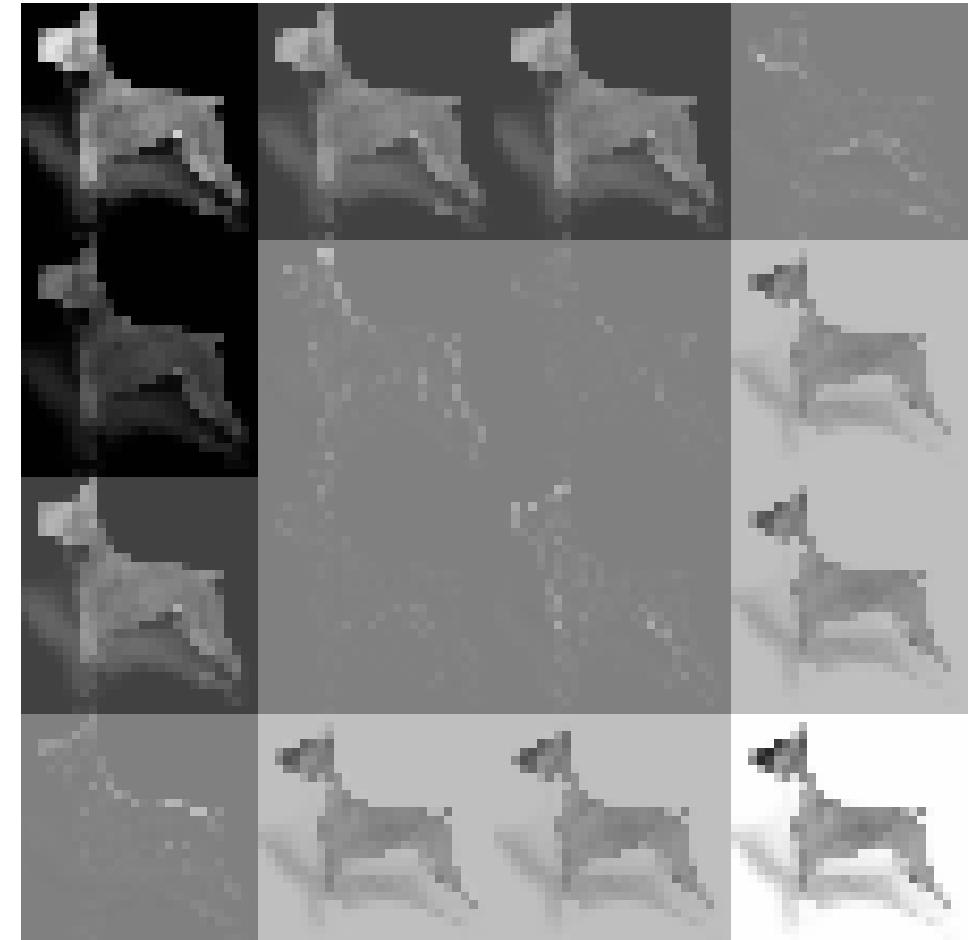
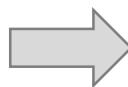
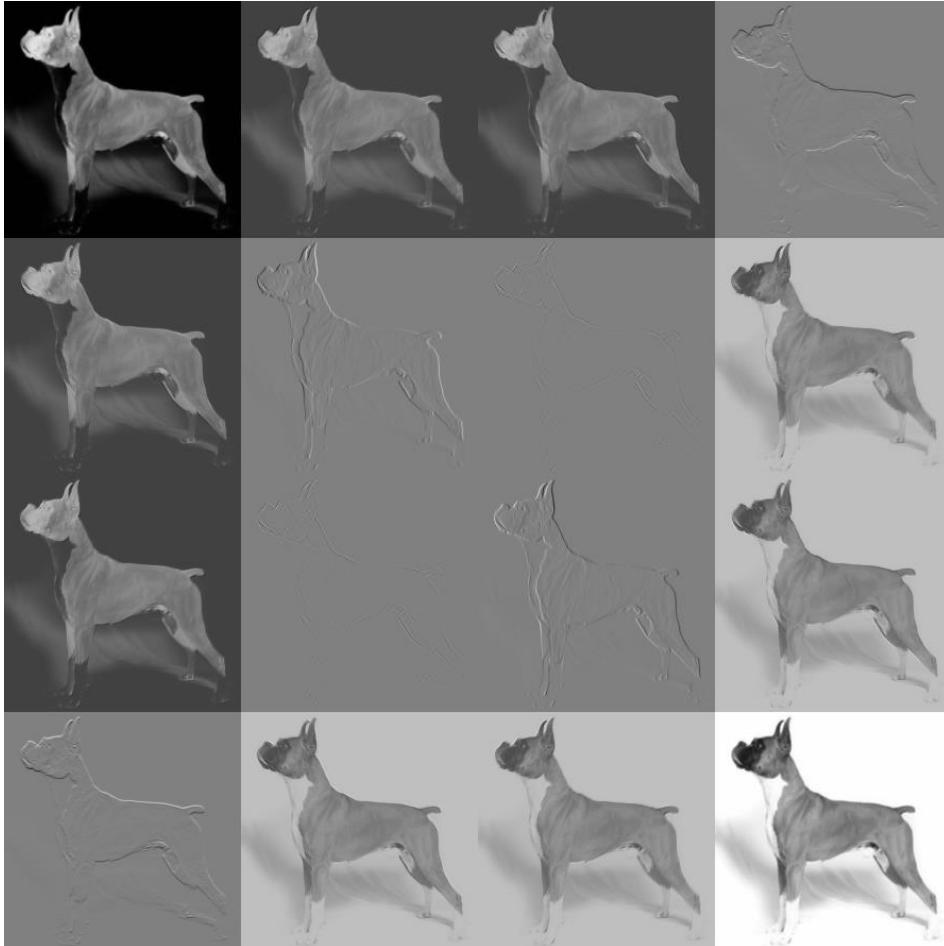
Basic operations

Pool: avg



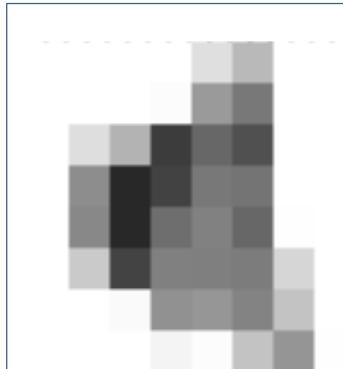
Basic operations

Pool: max



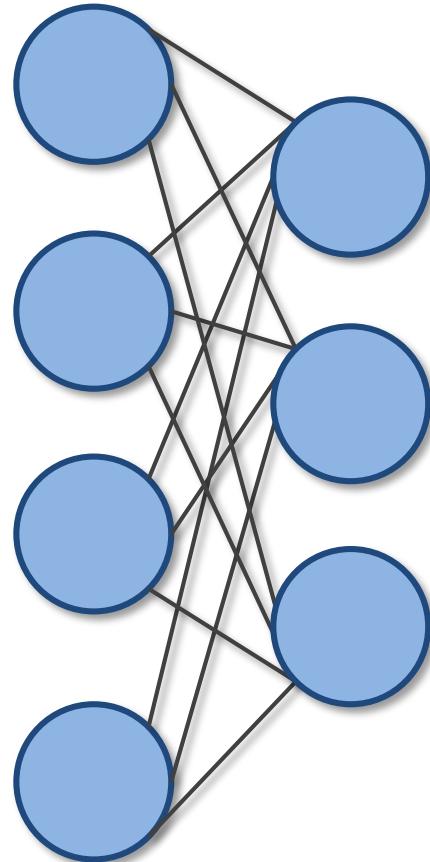
Basic operations

Flatten



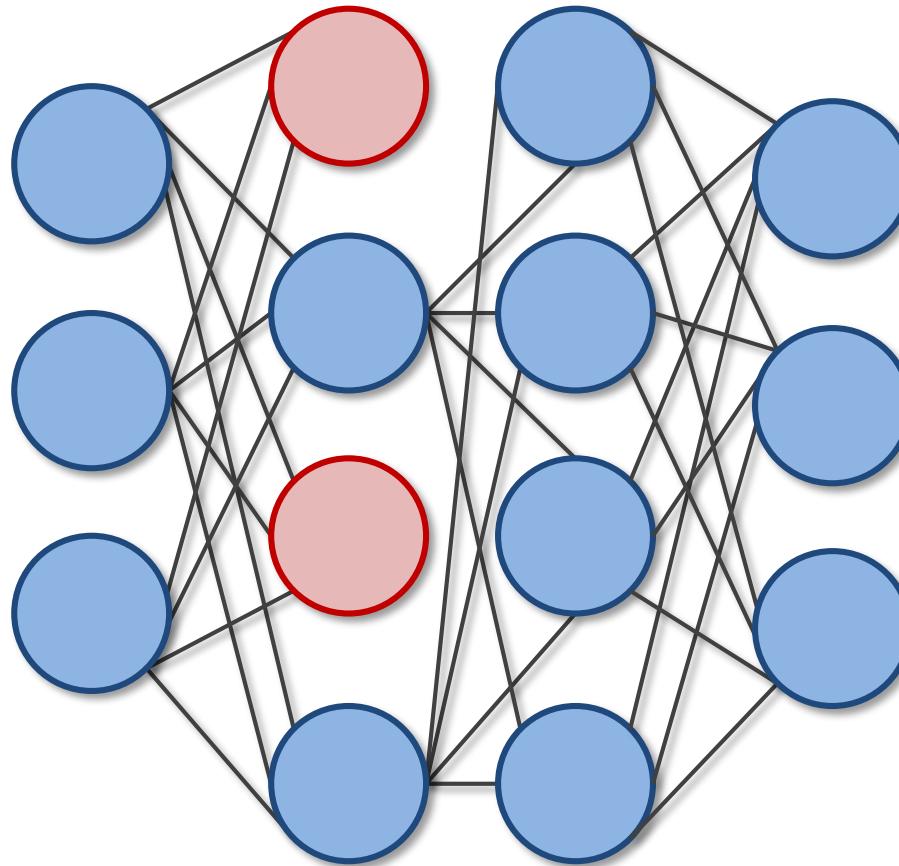
Basic operations

Fully Connected



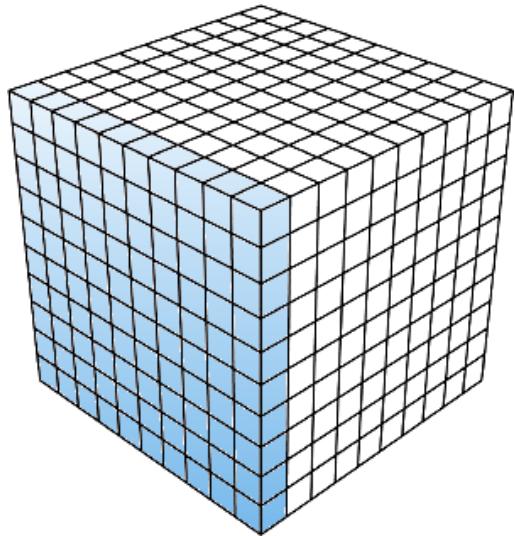
Basic operations

Dropout

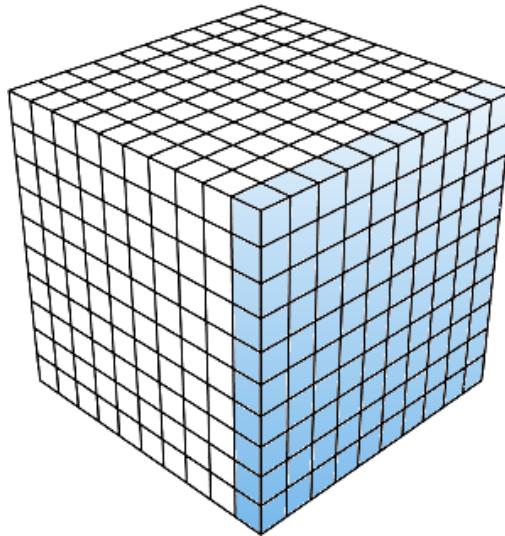


Basic operations

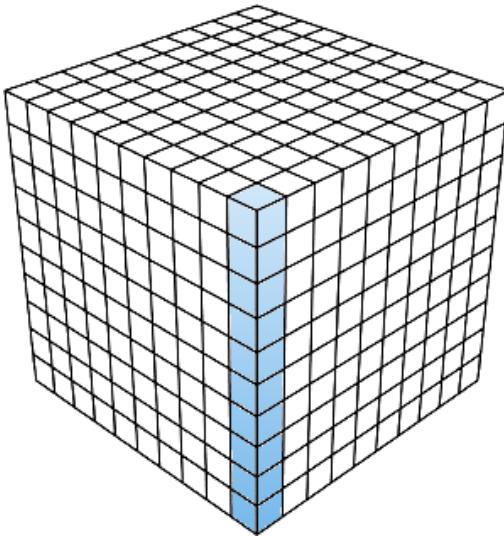
Normalization



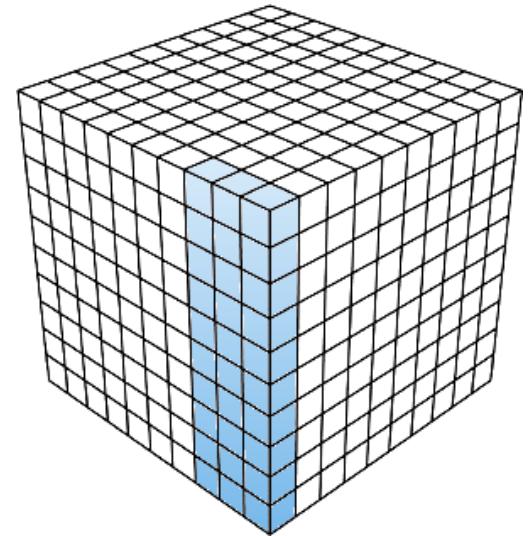
layer



batch



instance

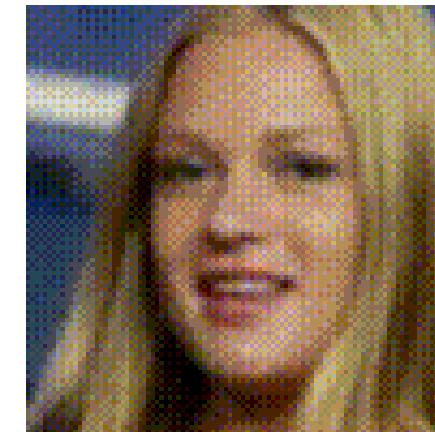
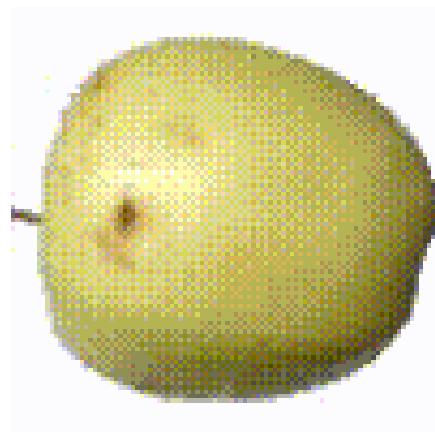
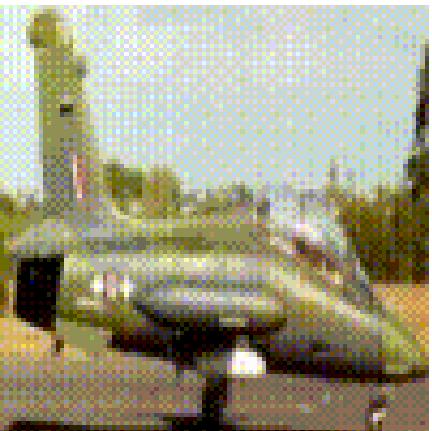


group

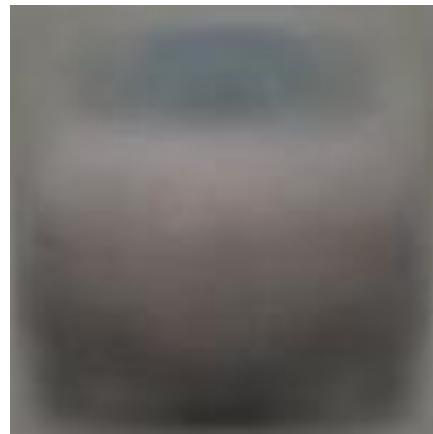
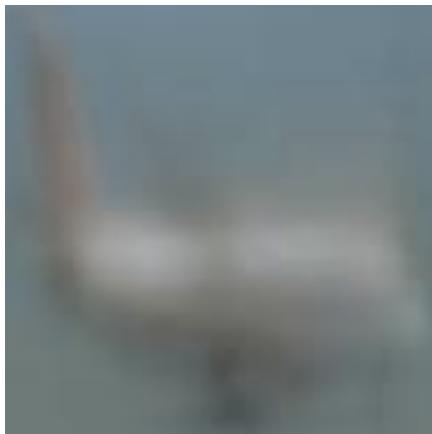
Feature extraction



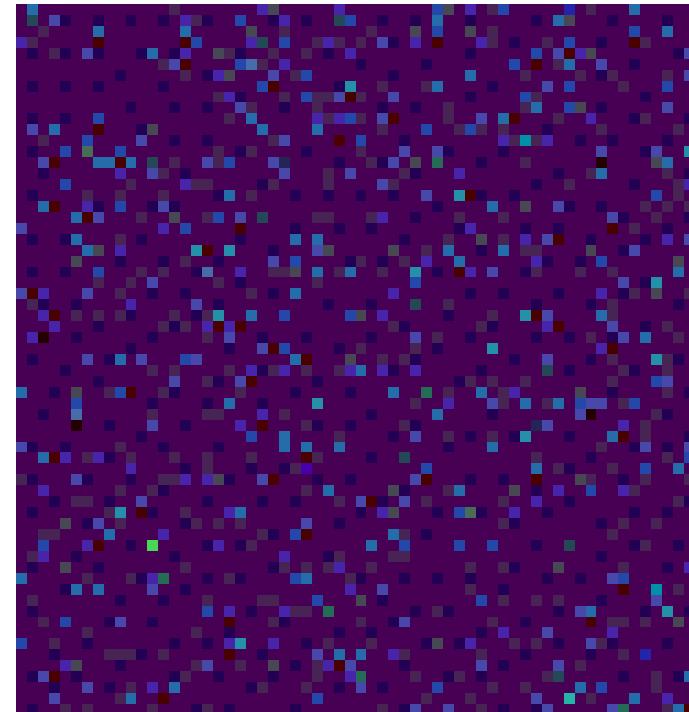
Feature extraction



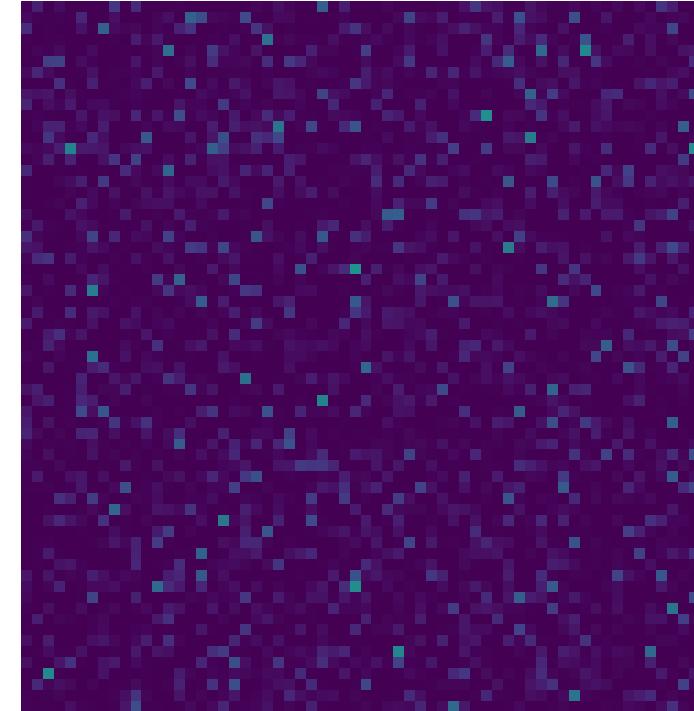
Feature extraction



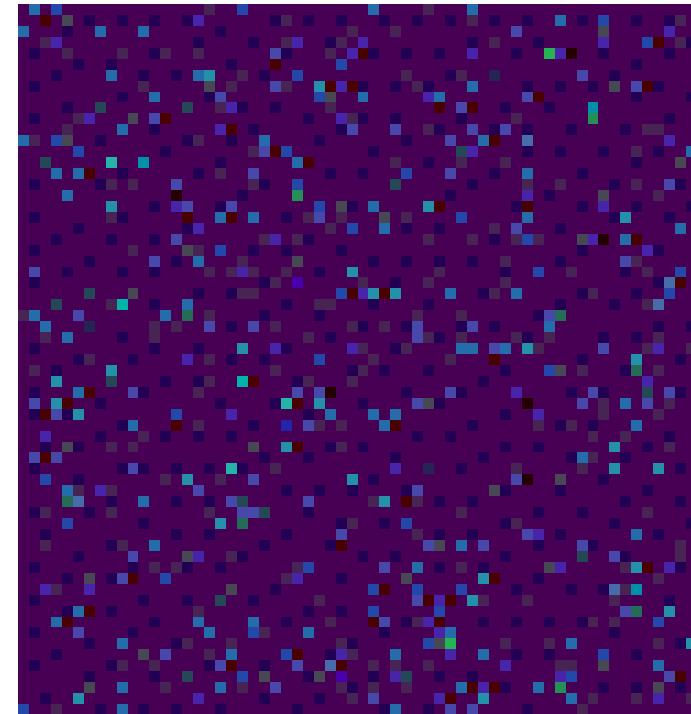
Extract the features



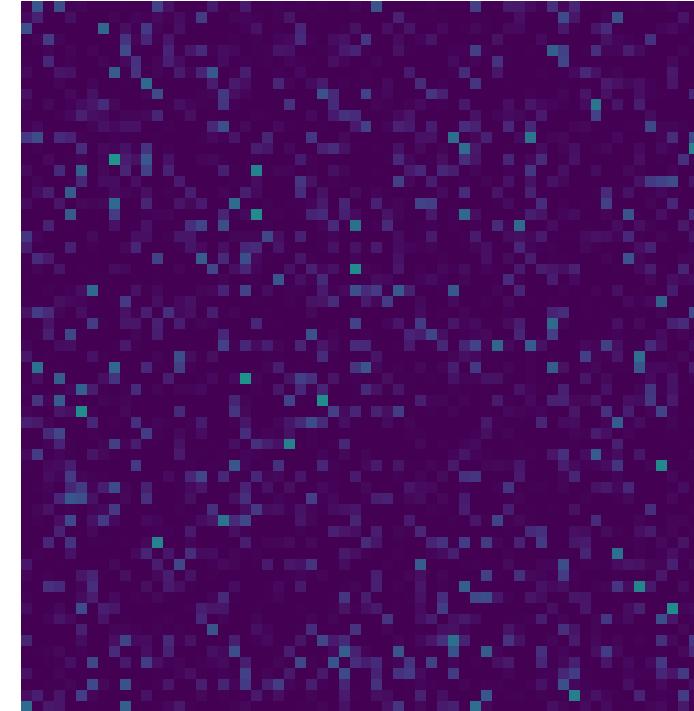
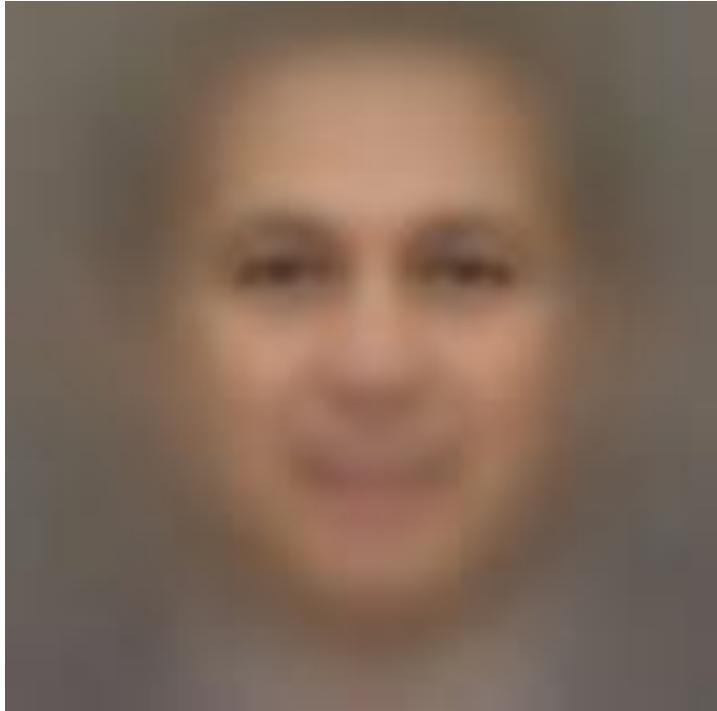
Feature extraction



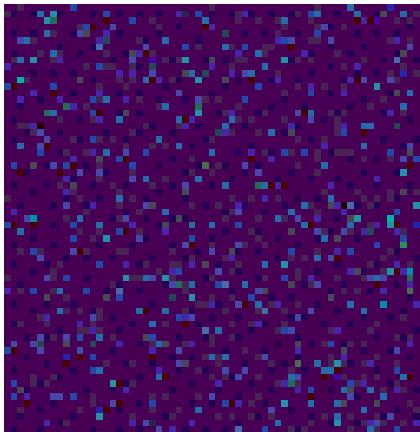
Feature extraction



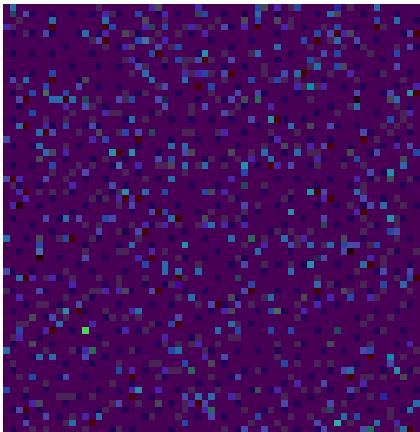
Feature extraction



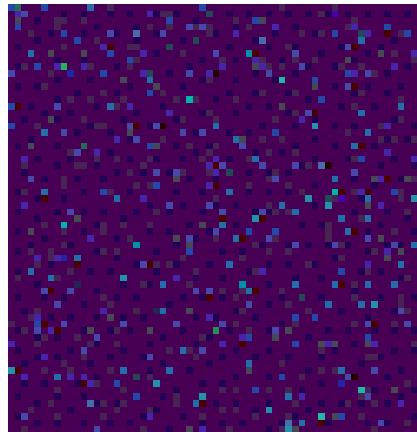
Feature extraction



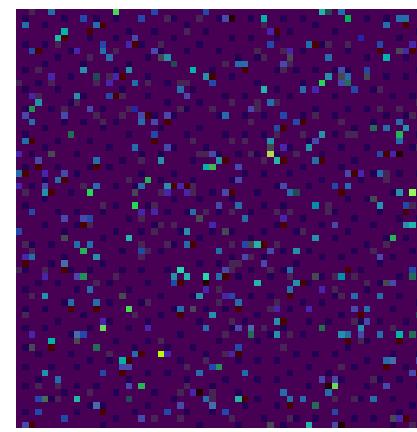
airplane



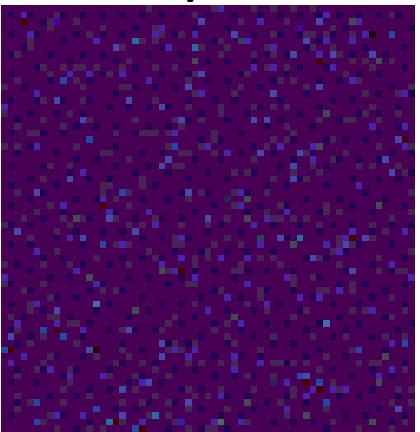
car



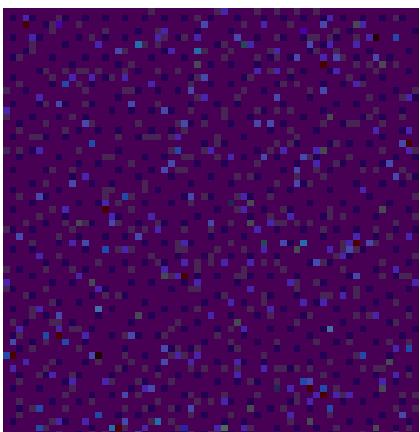
cat



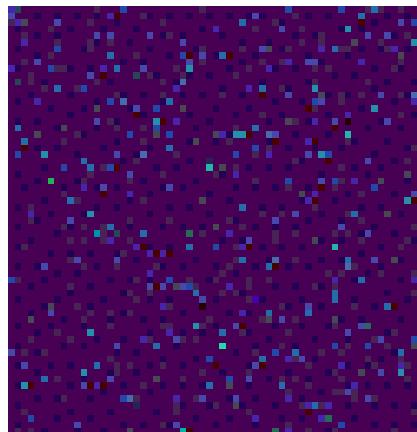
dog



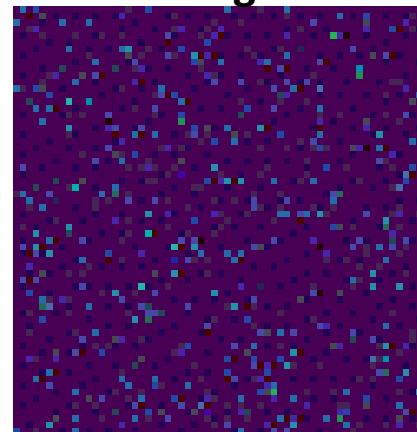
flower



fruit

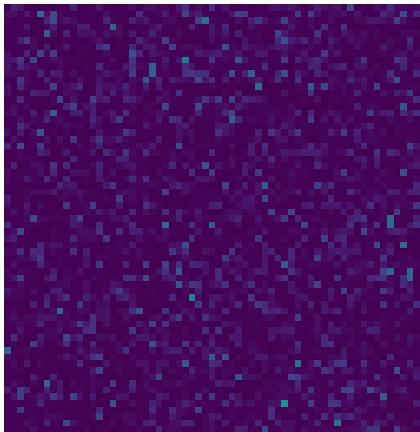


motorbike

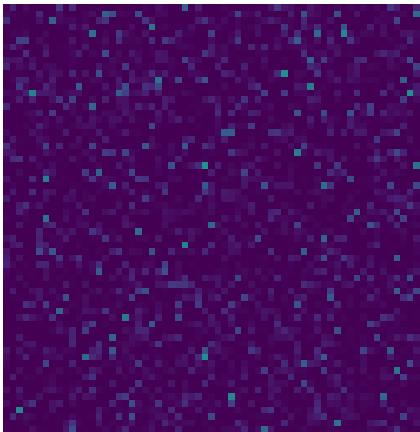


person

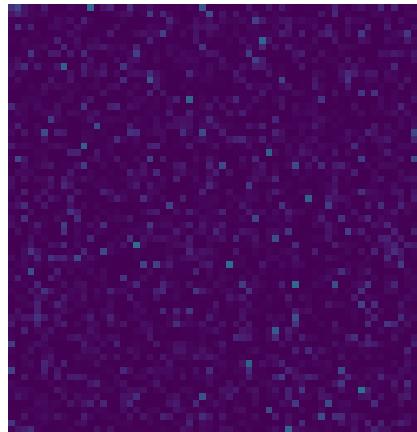
Feature extraction



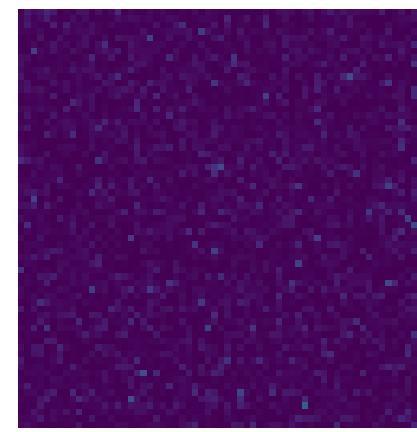
airplane



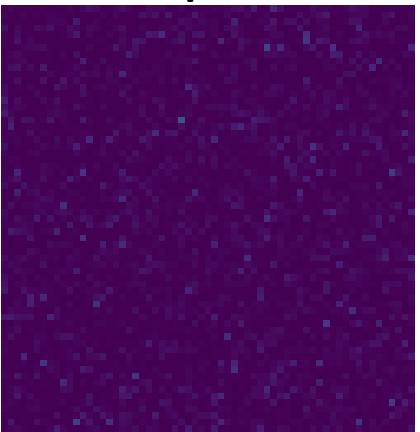
car



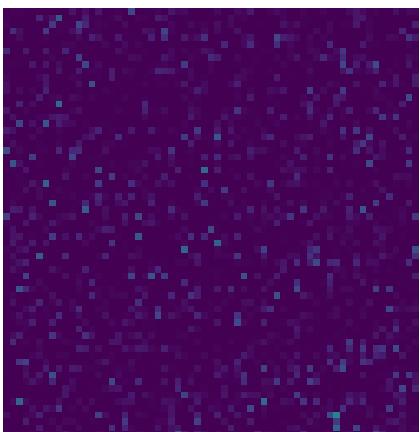
cat



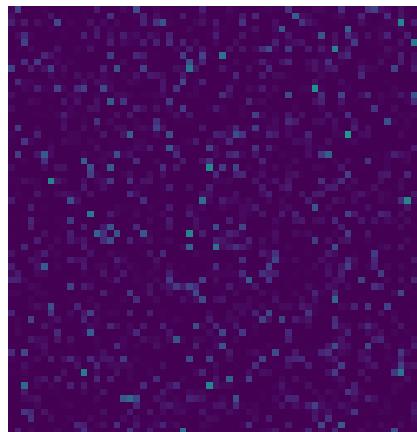
dog



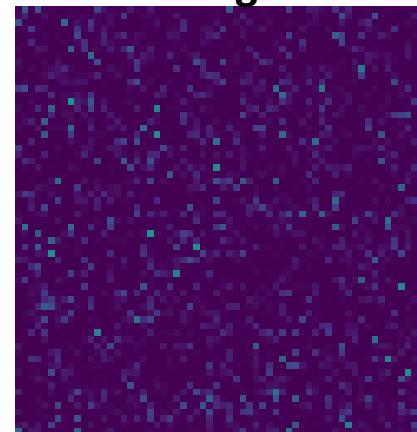
flower



fruit



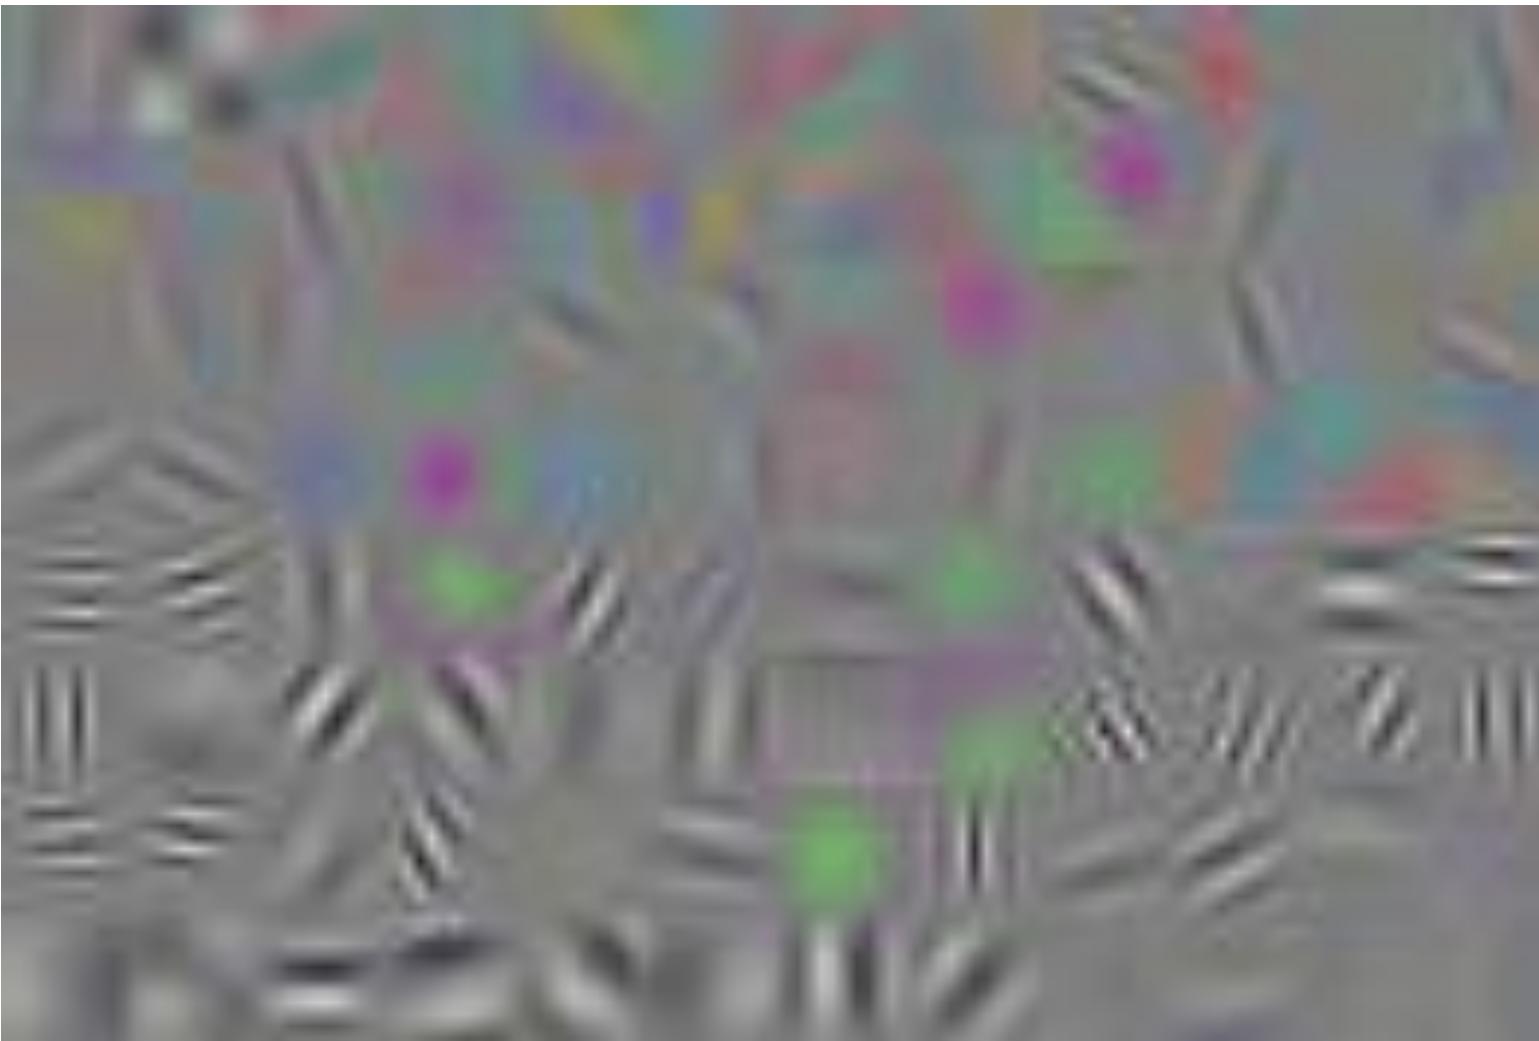
motorbike



person

AlexNet

Feature extraction



AlexNet



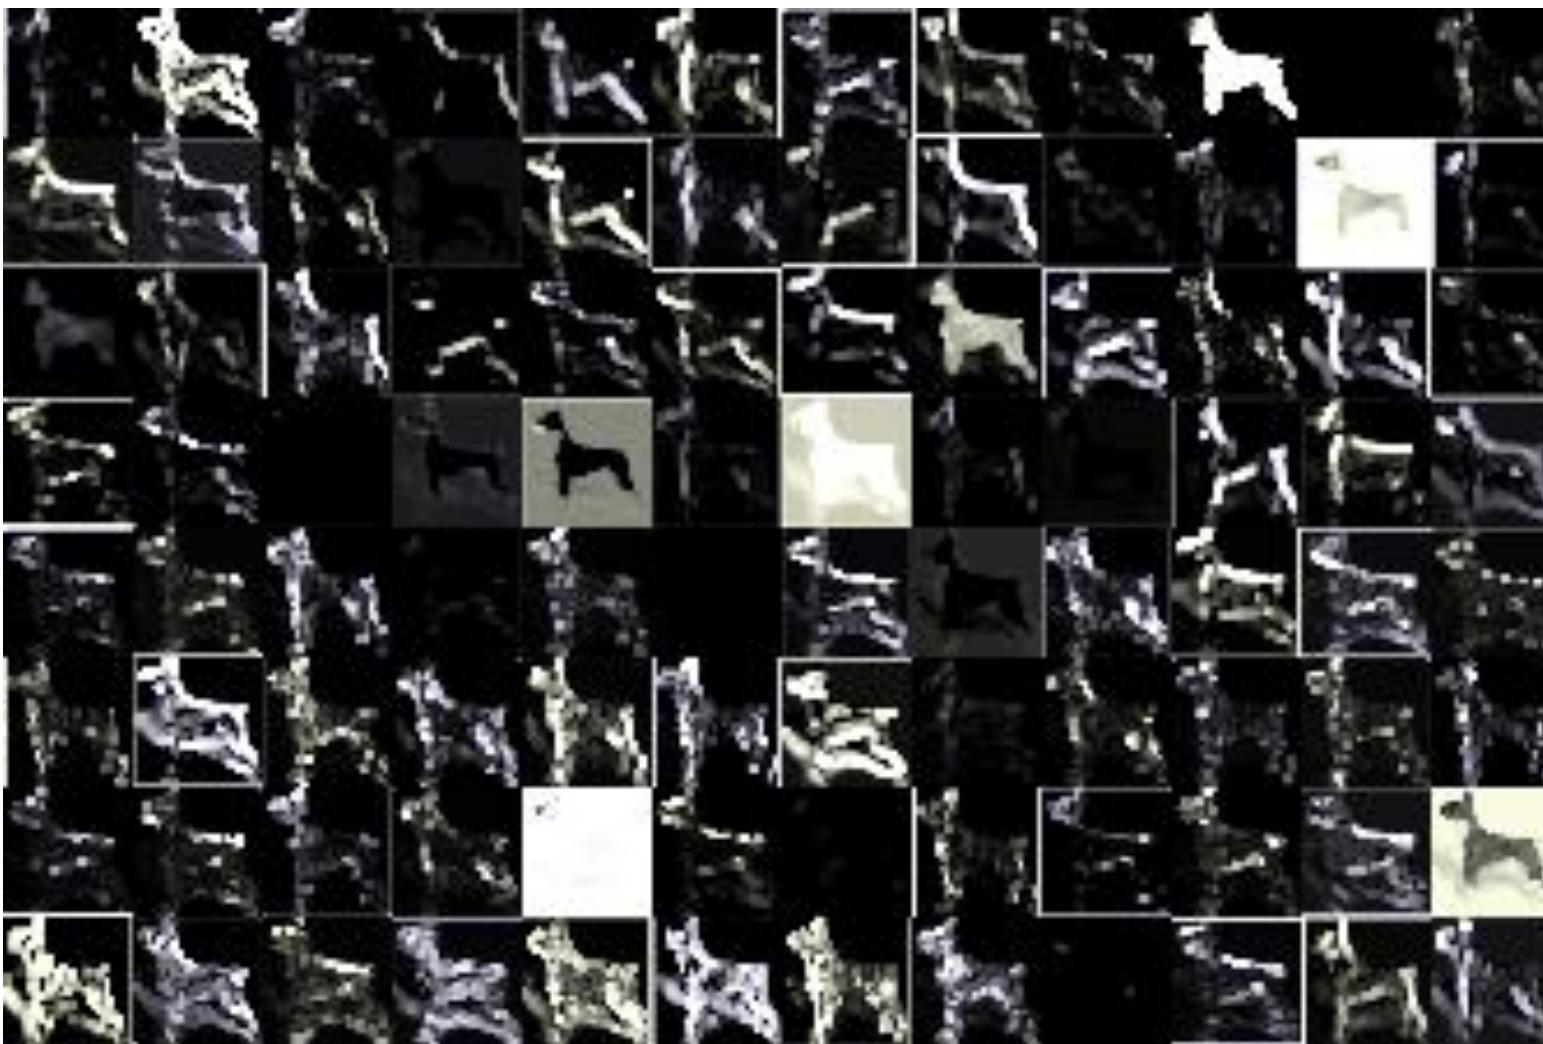
Feature extraction



AlexNet



Feature extraction



AlexNet



Feature extraction



AlexNet

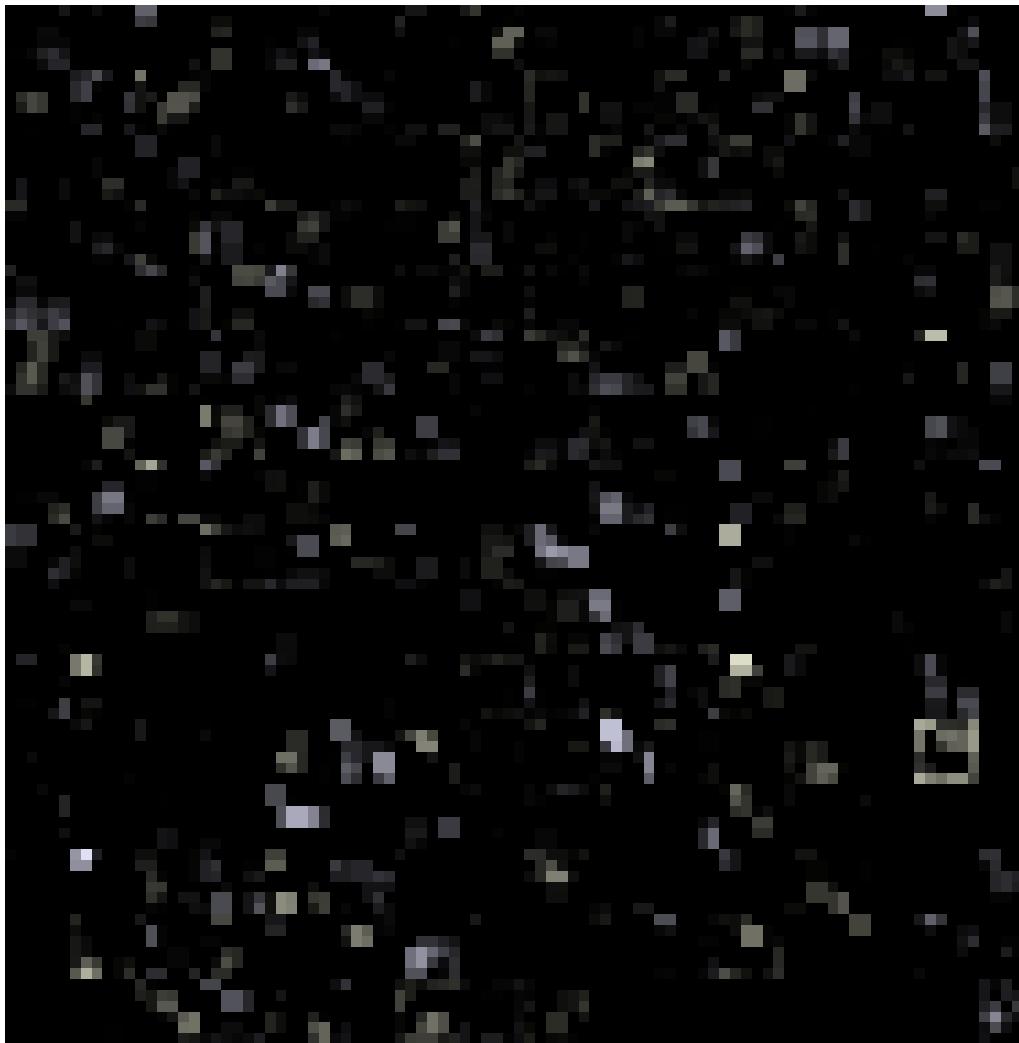


Feature extraction



AlexNet

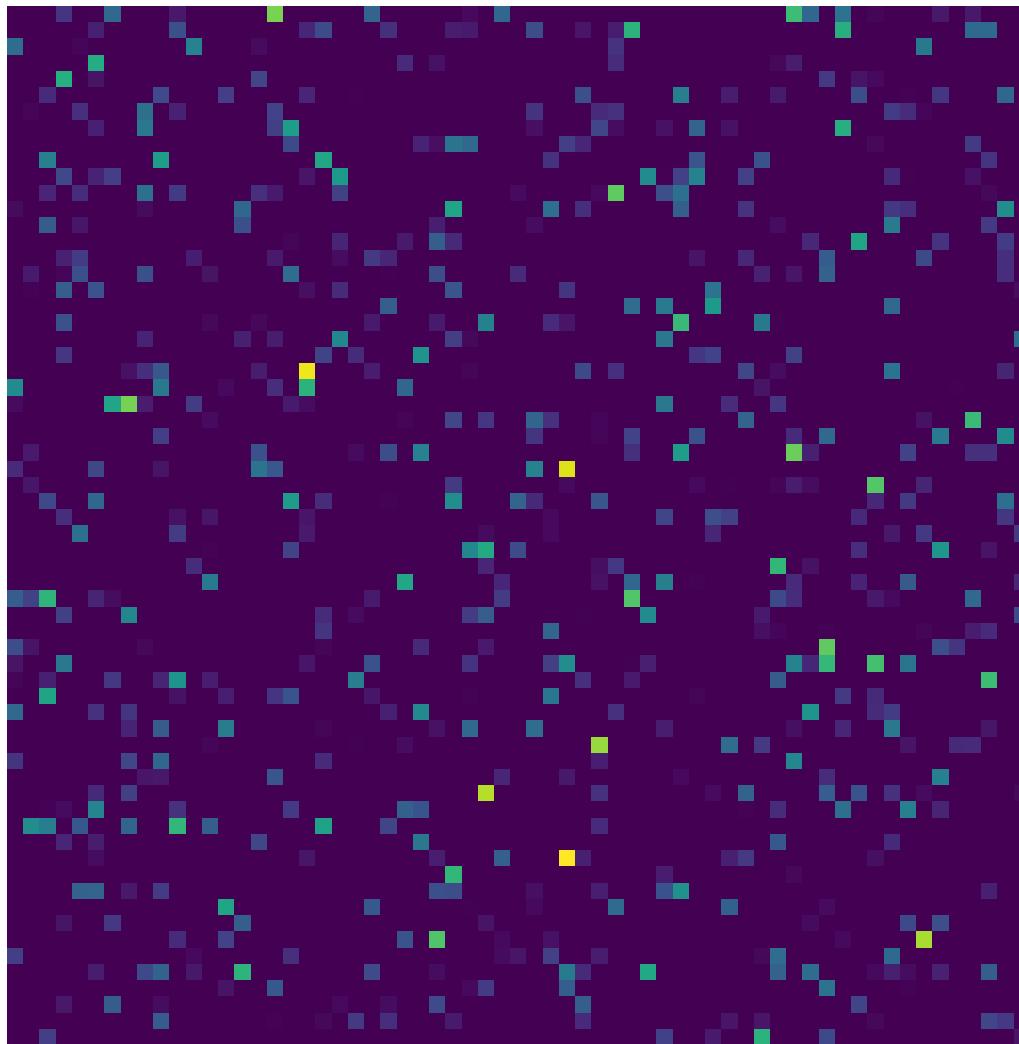
Feature extraction



AlexNet

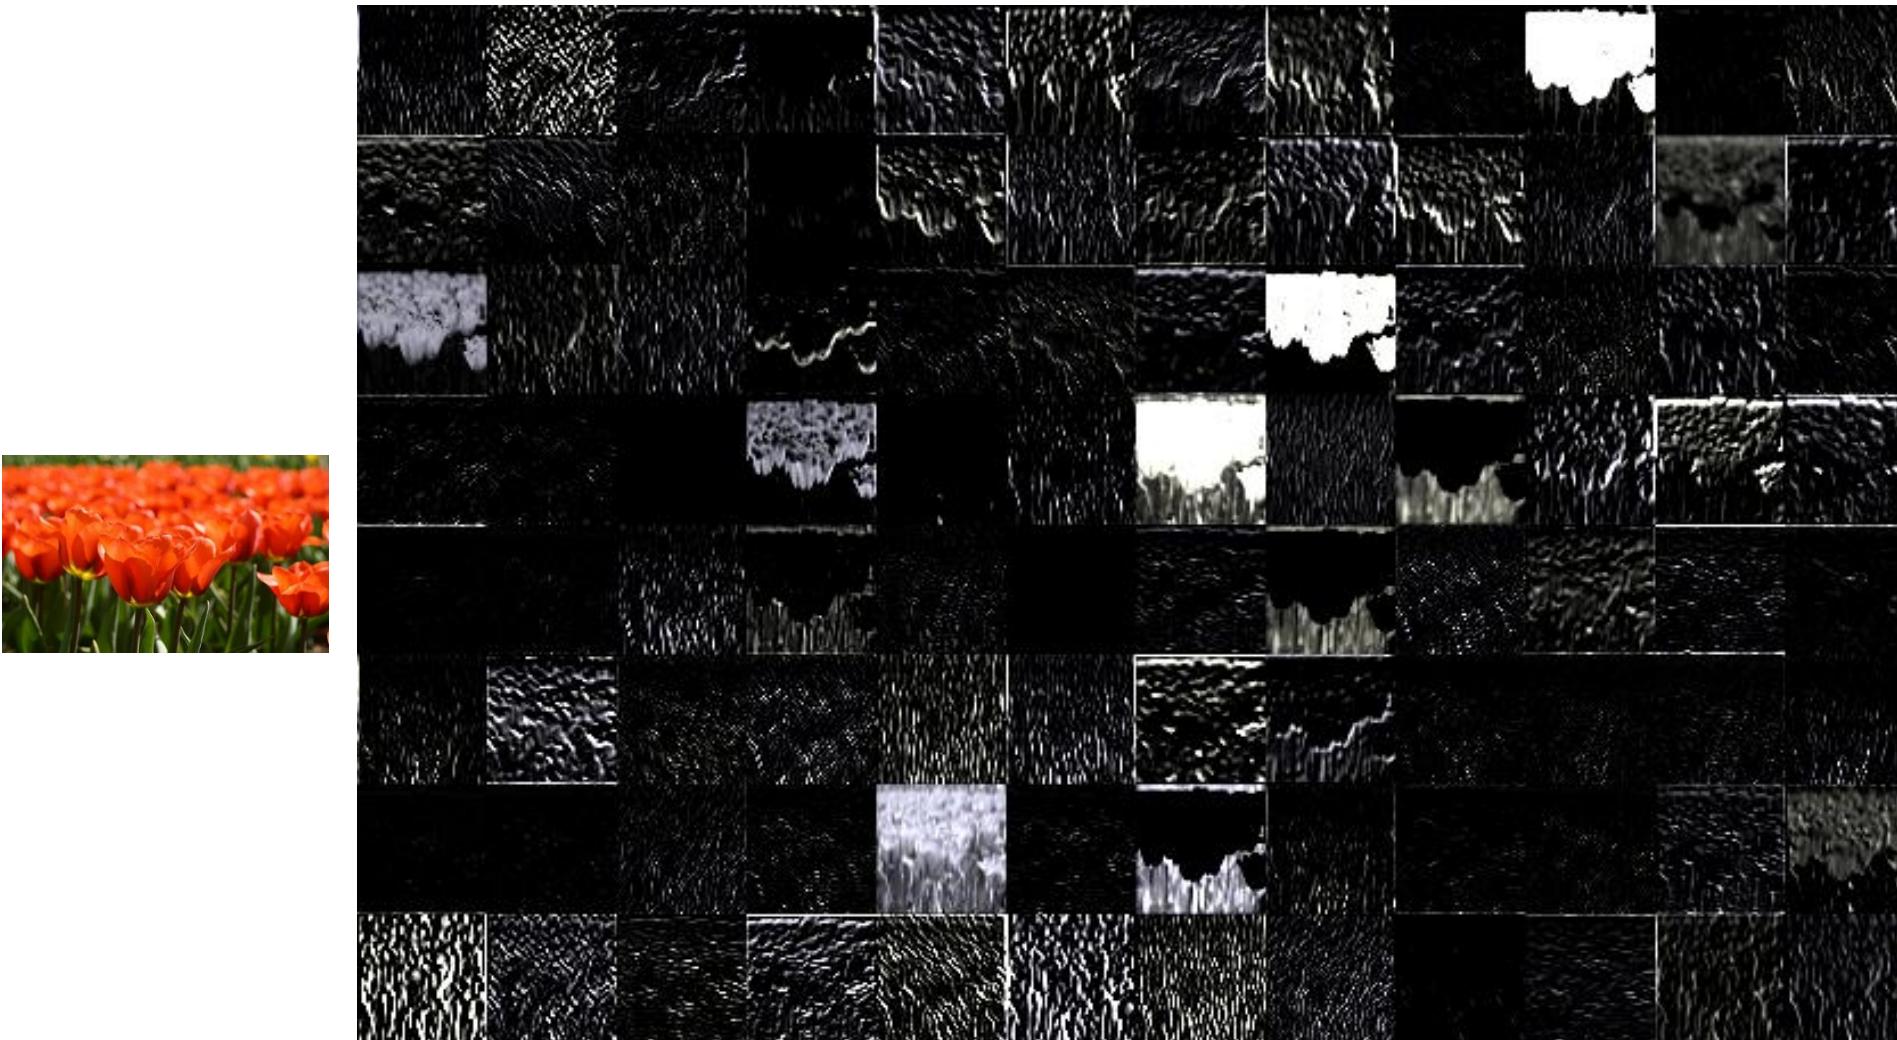


Feature extraction



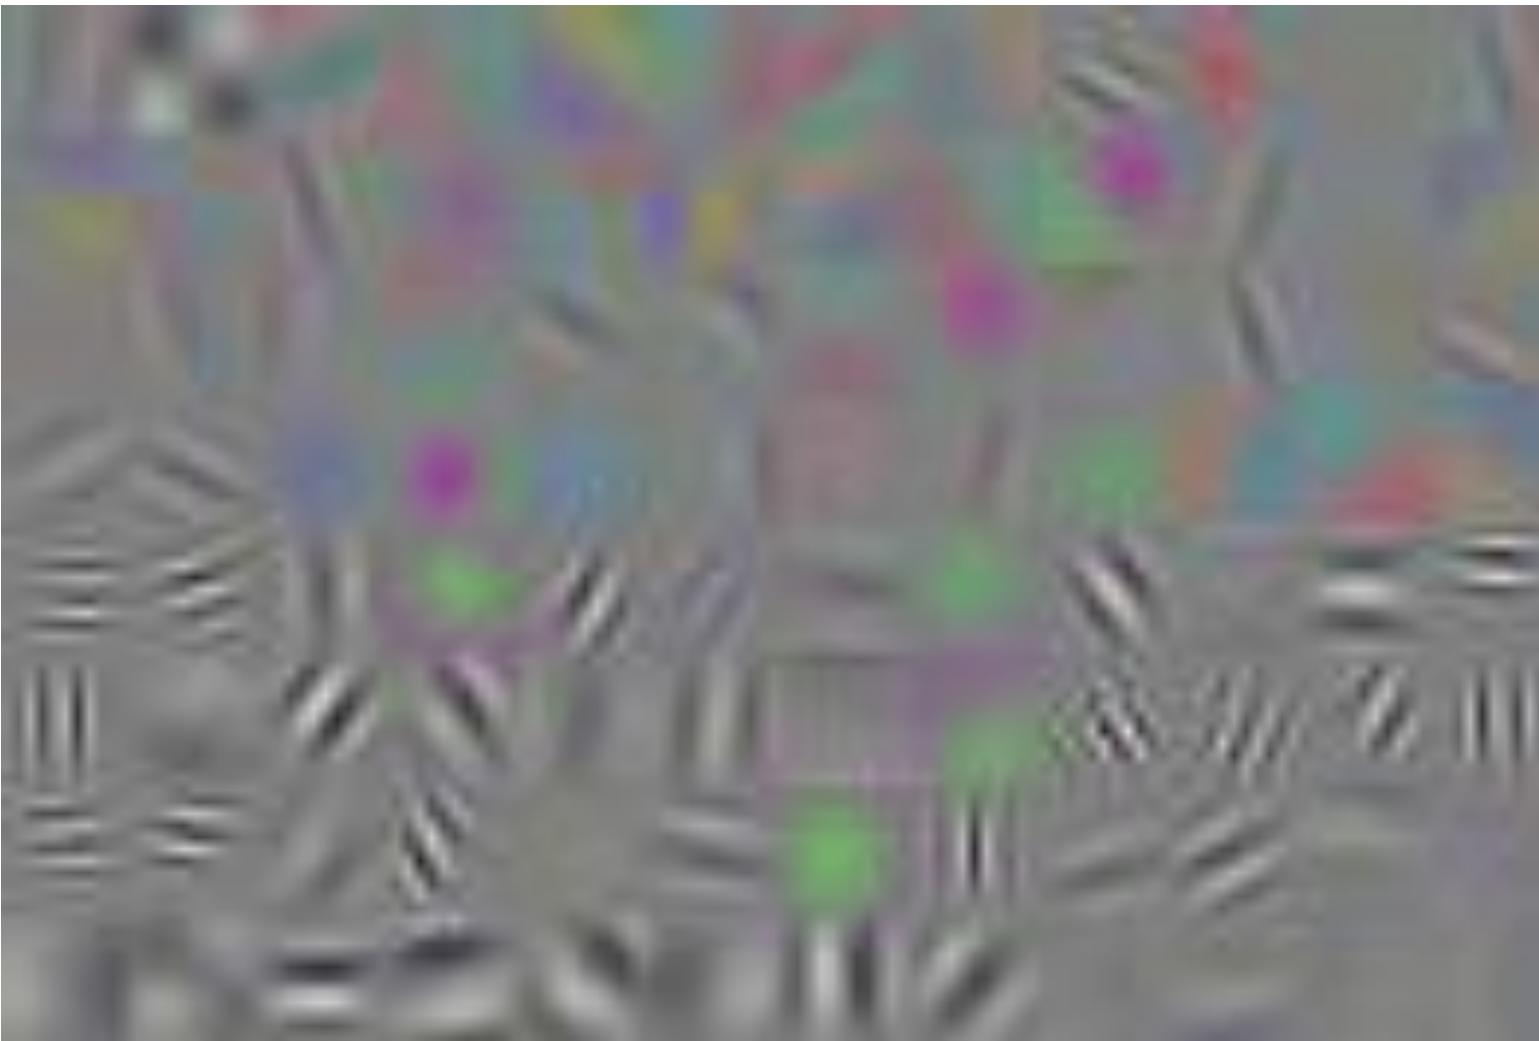
AlexNet

Feature extraction



AlexNet

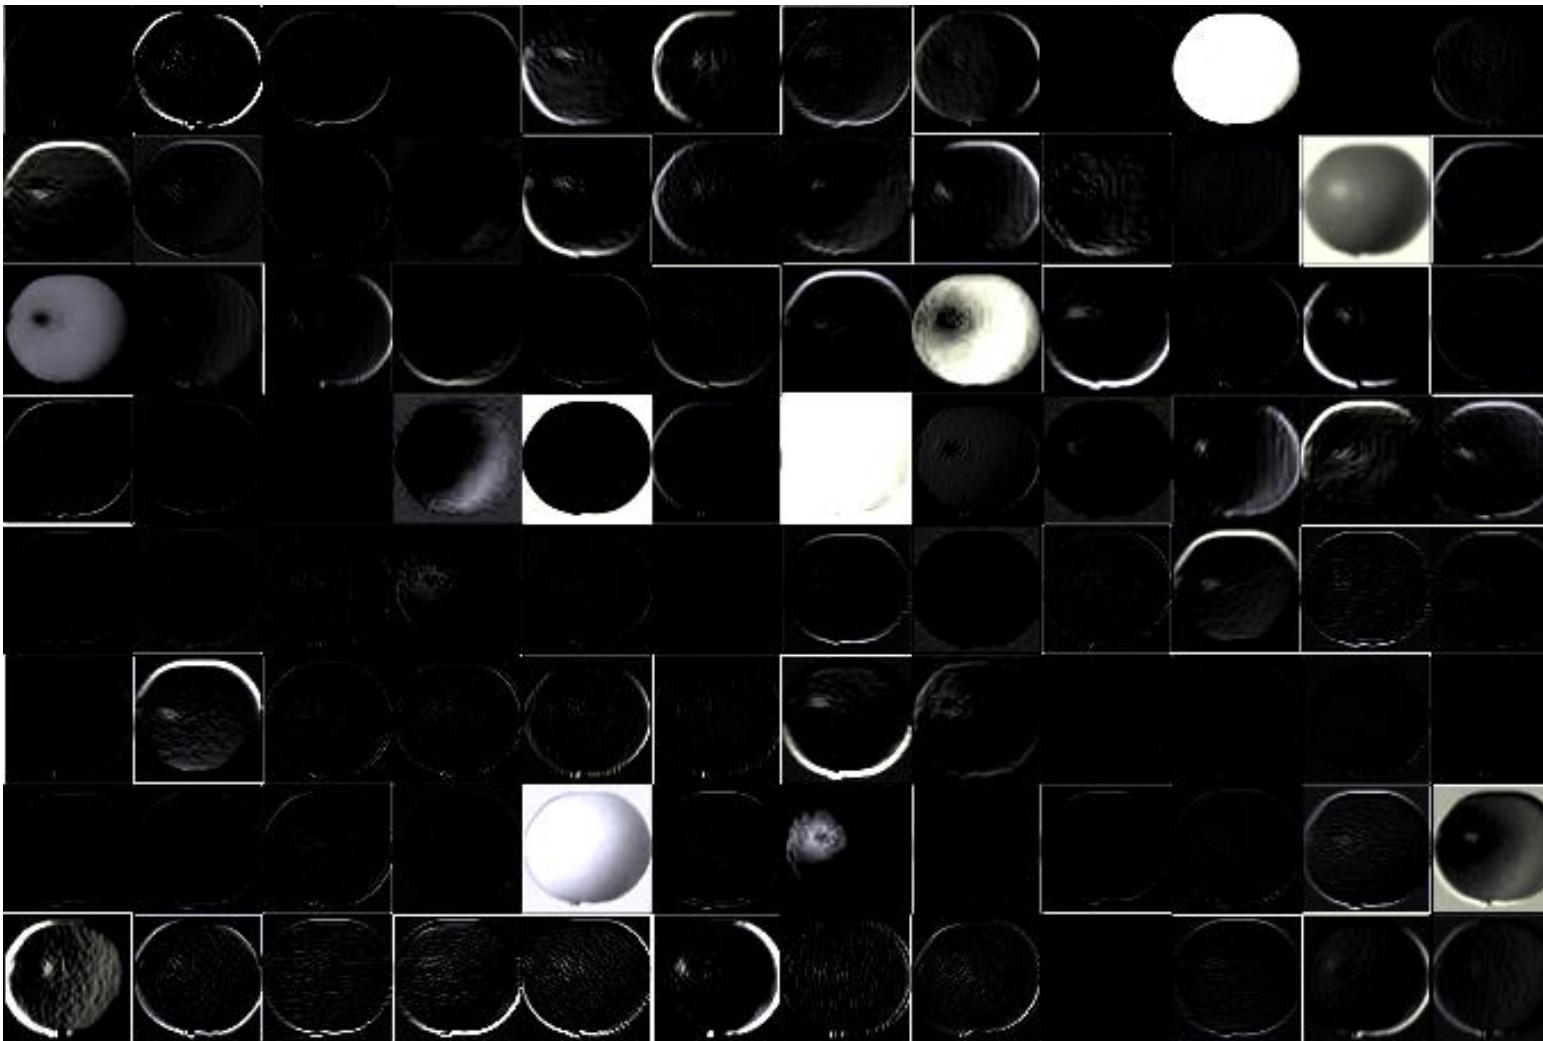
Feature extraction



AlexNet



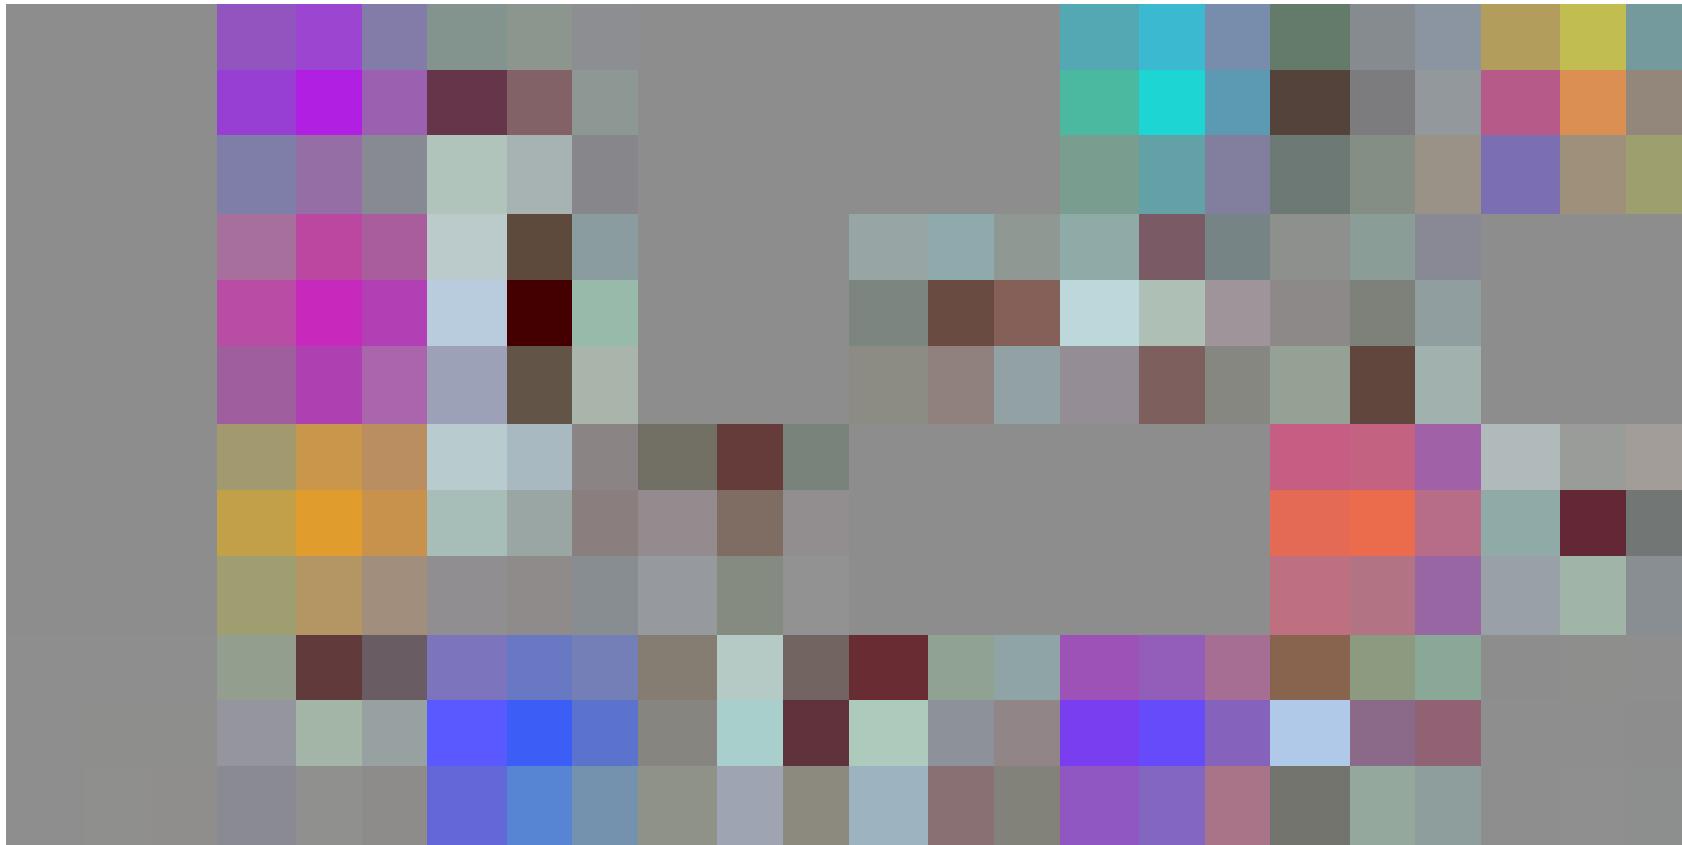
Feature extraction



MobileNet

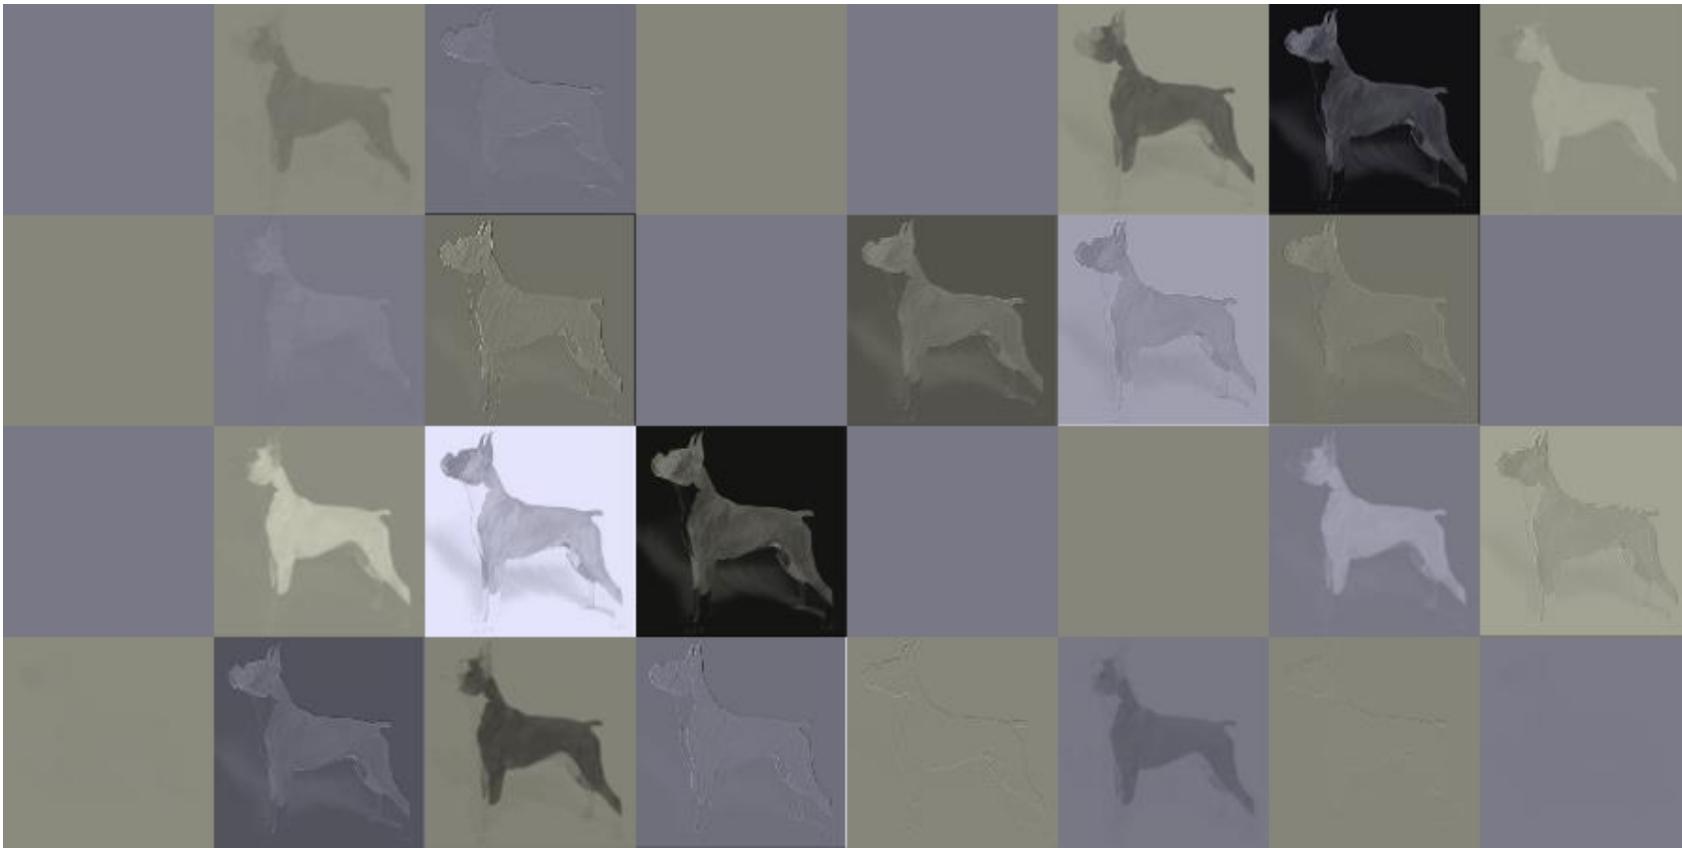


Feature extraction



MobileNet

Feature extraction



Feature extraction

Xception



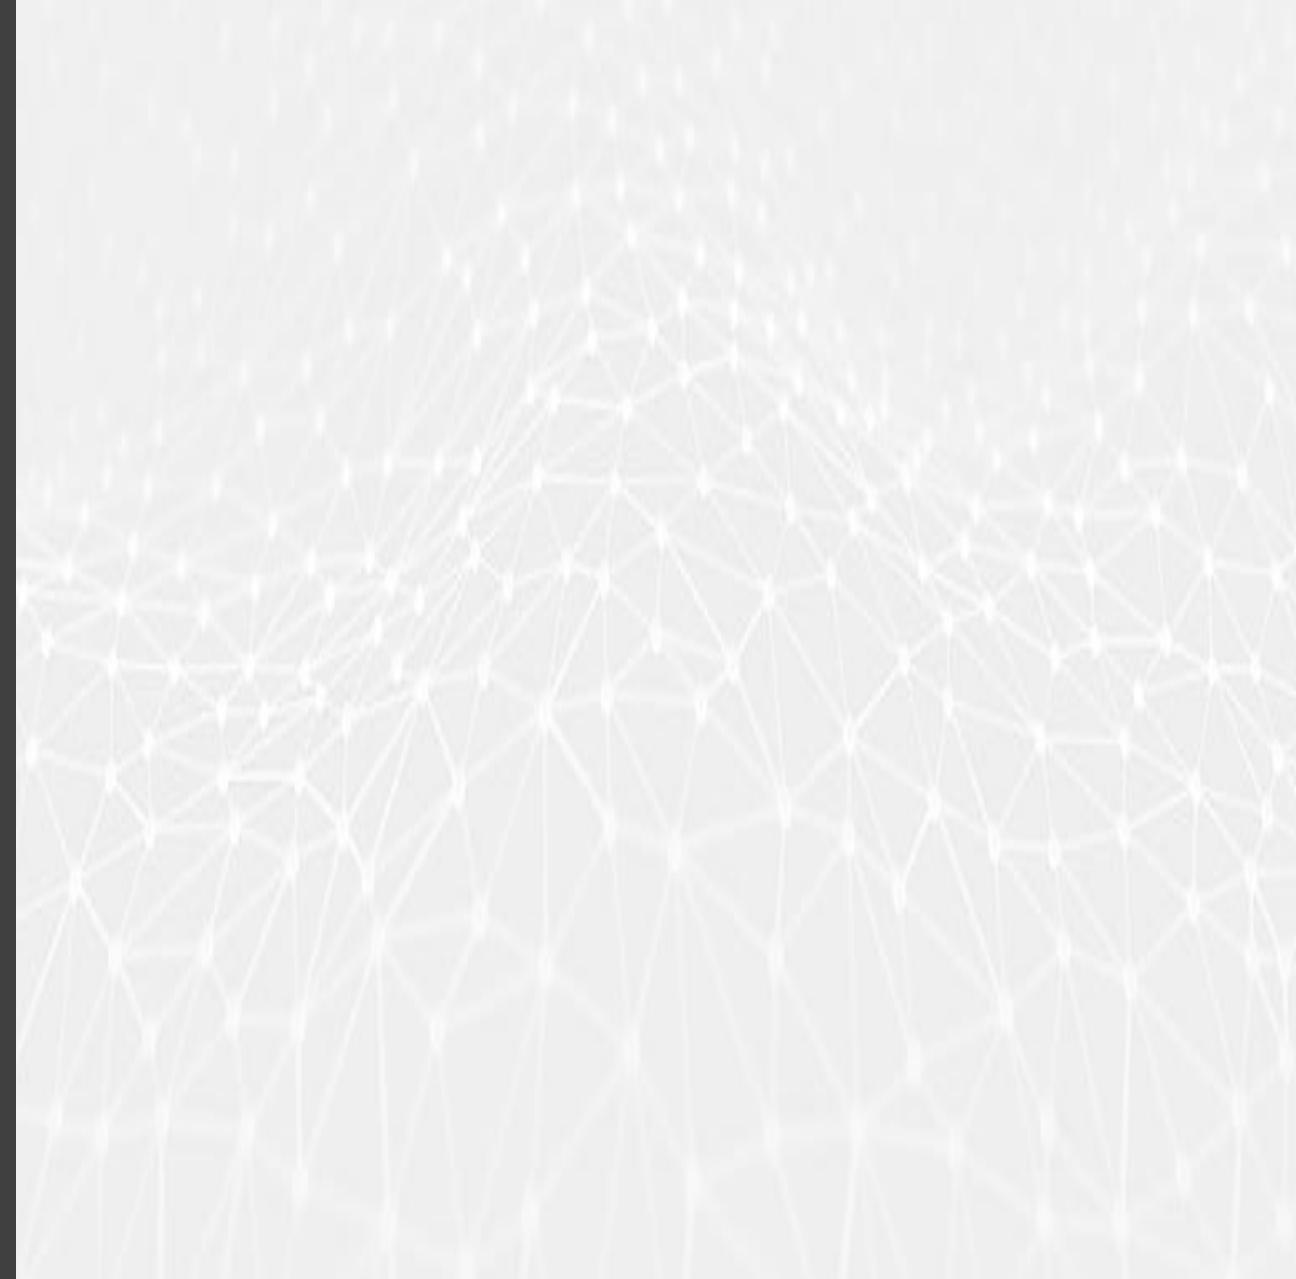
MobileNet



Feature extraction

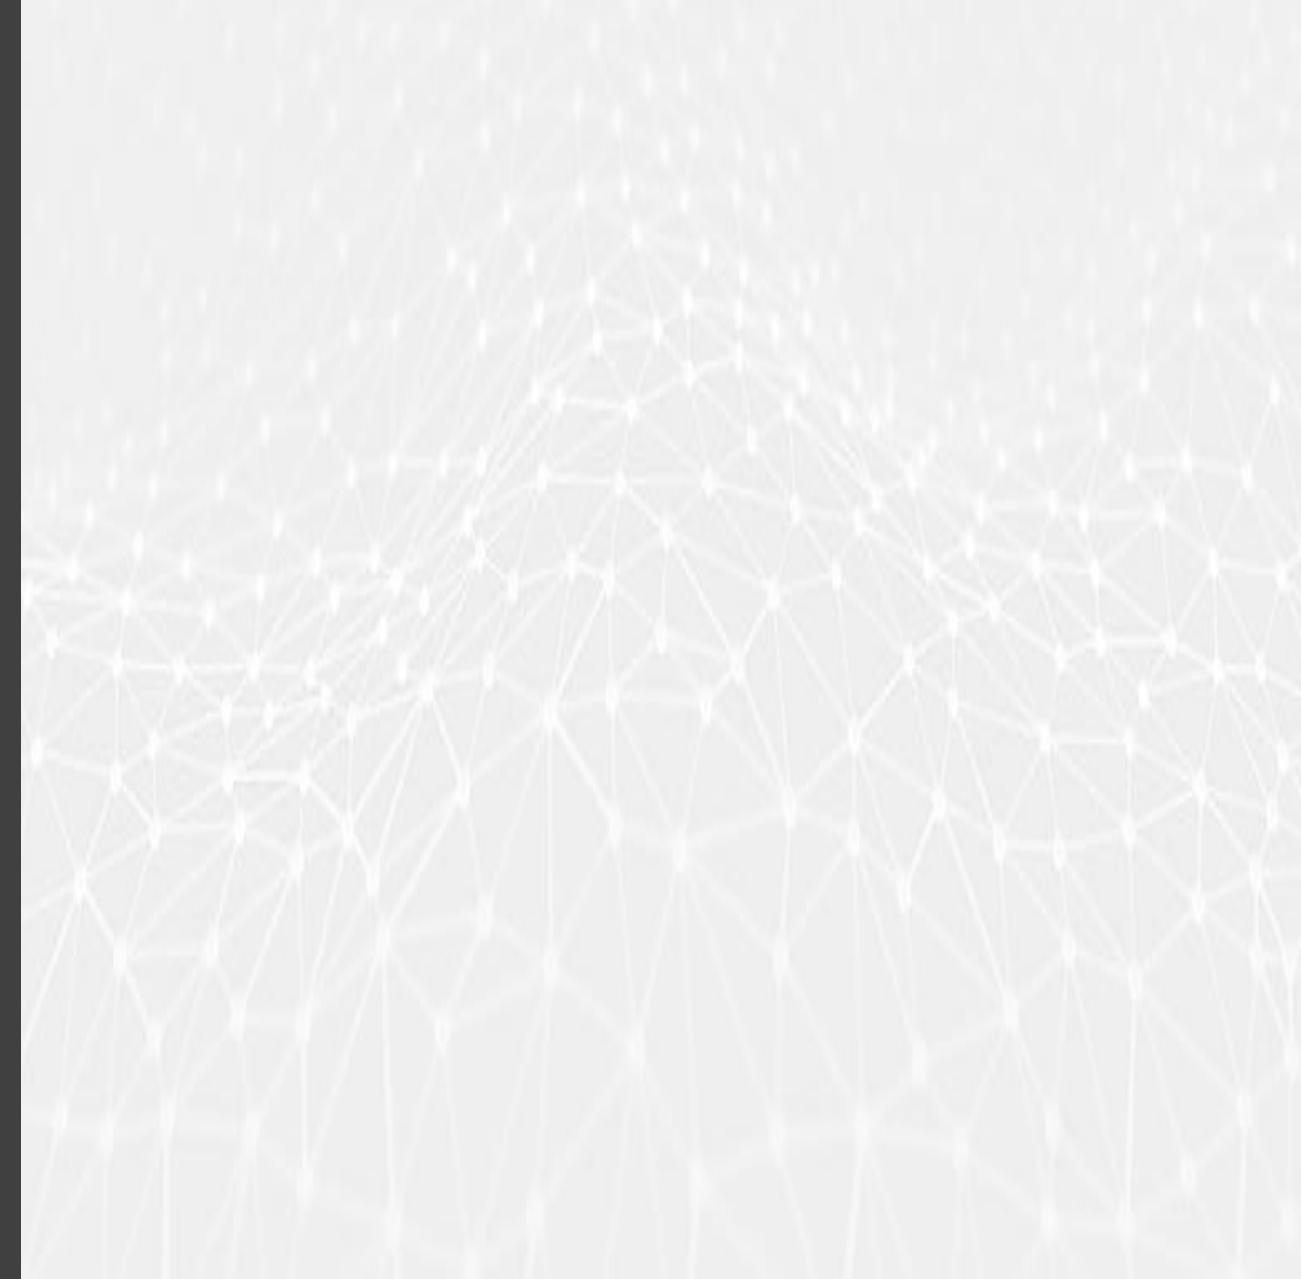


Transfer learning



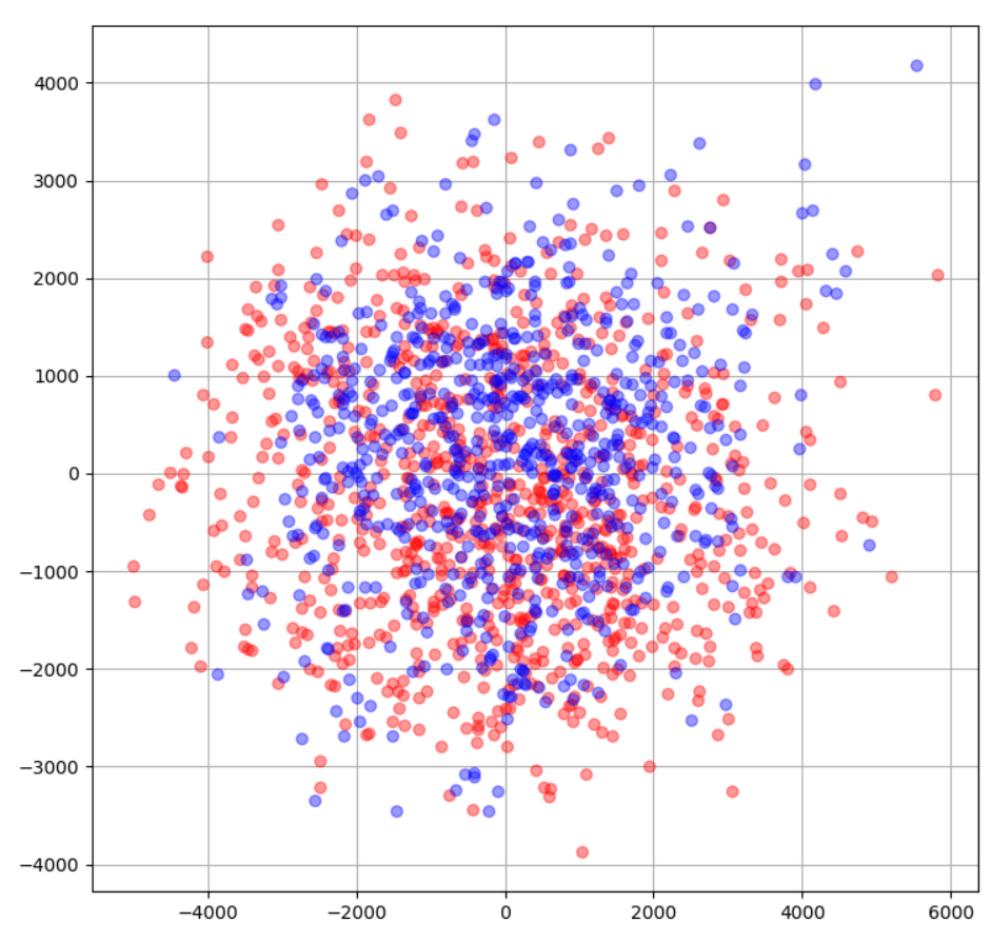
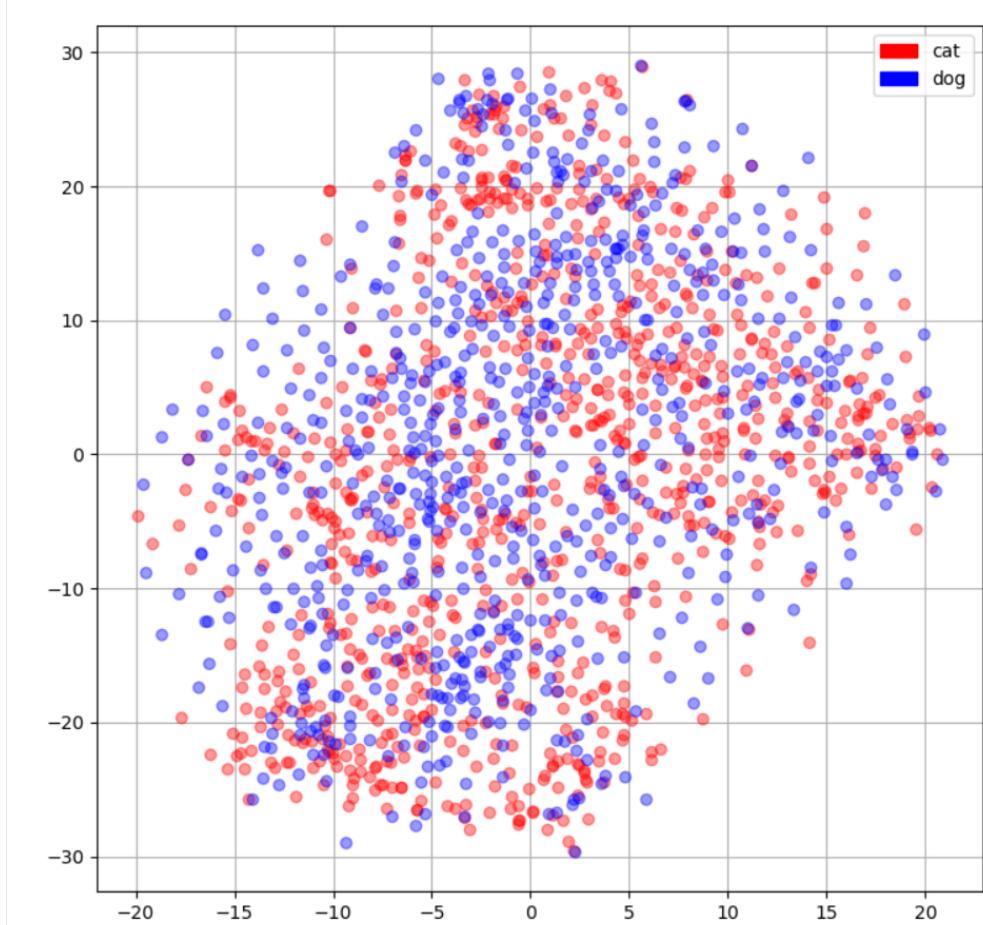
Transfer learning

Fine tuning

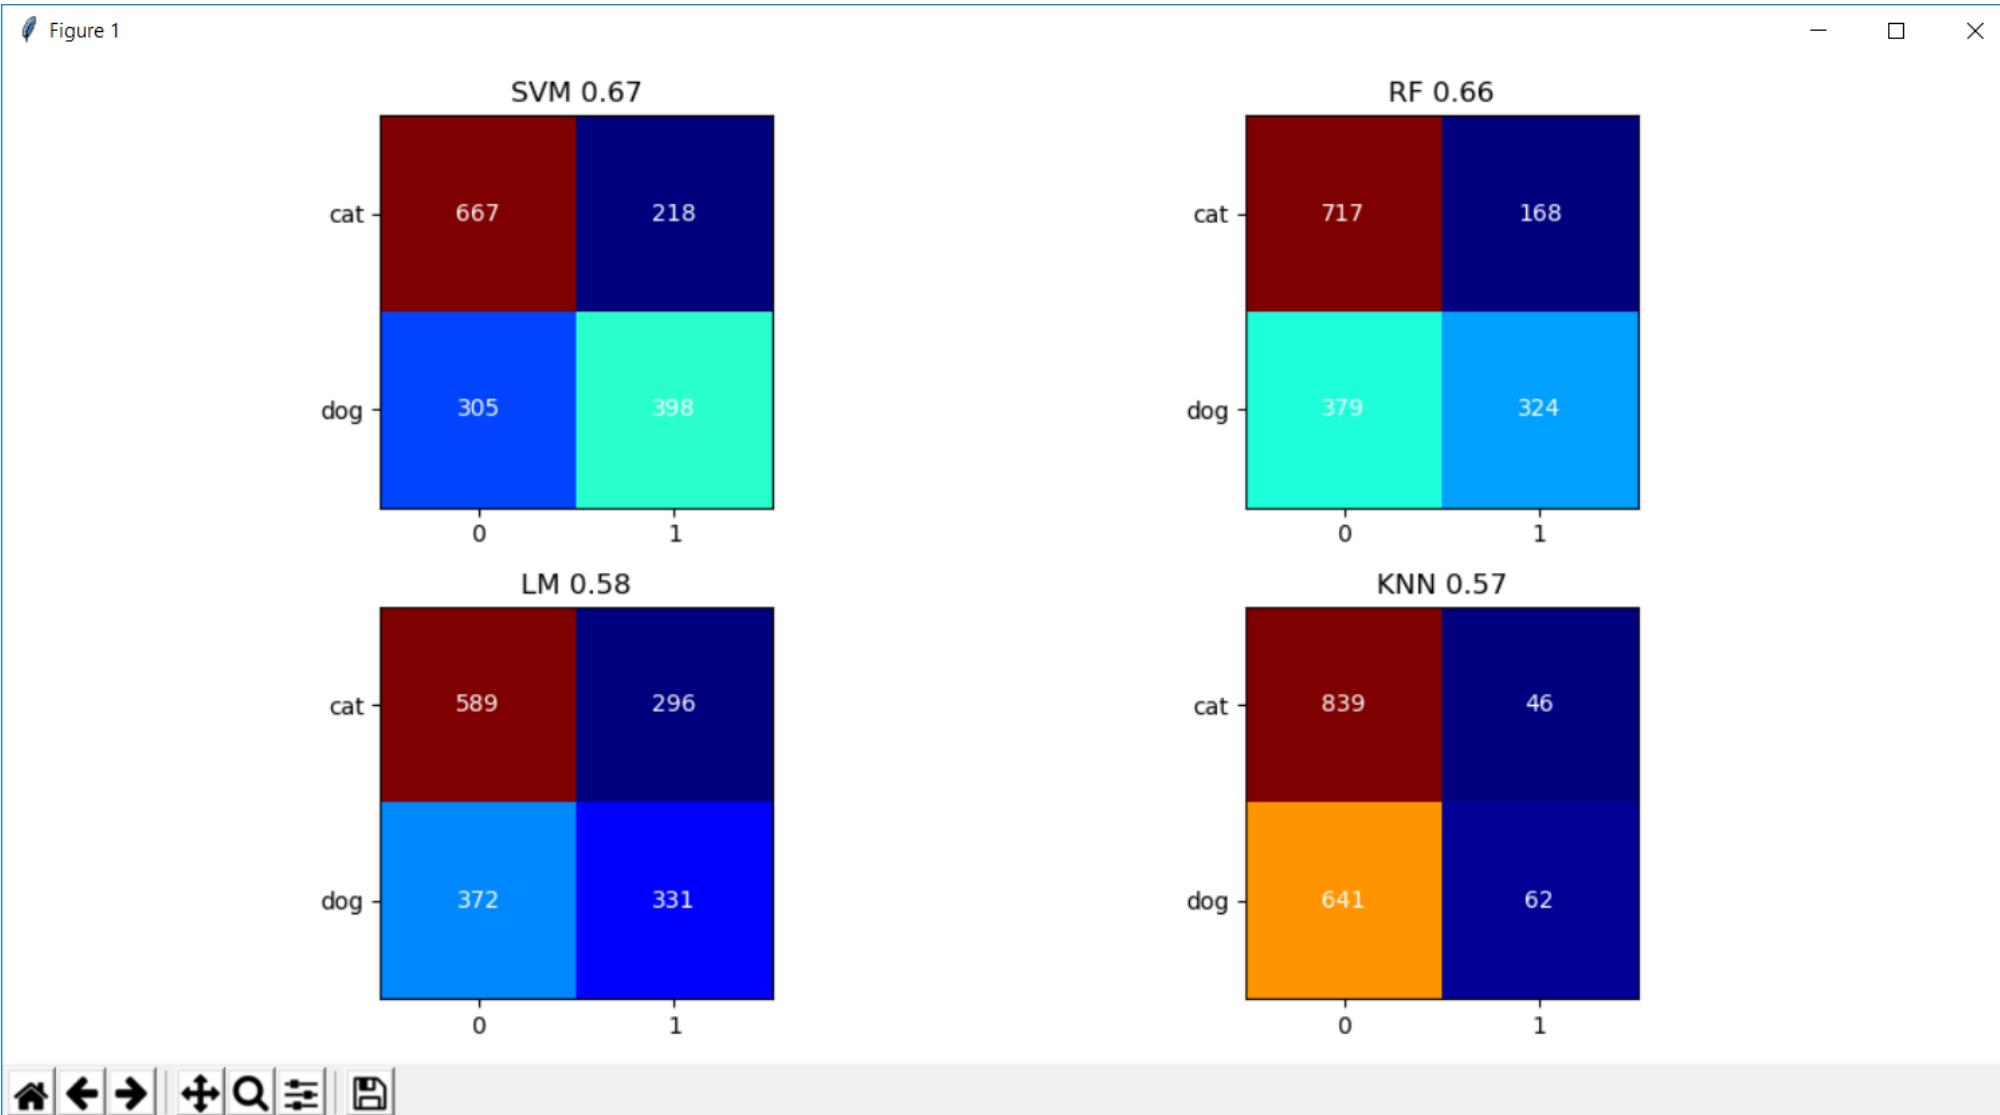


Finetuning

Figure 1

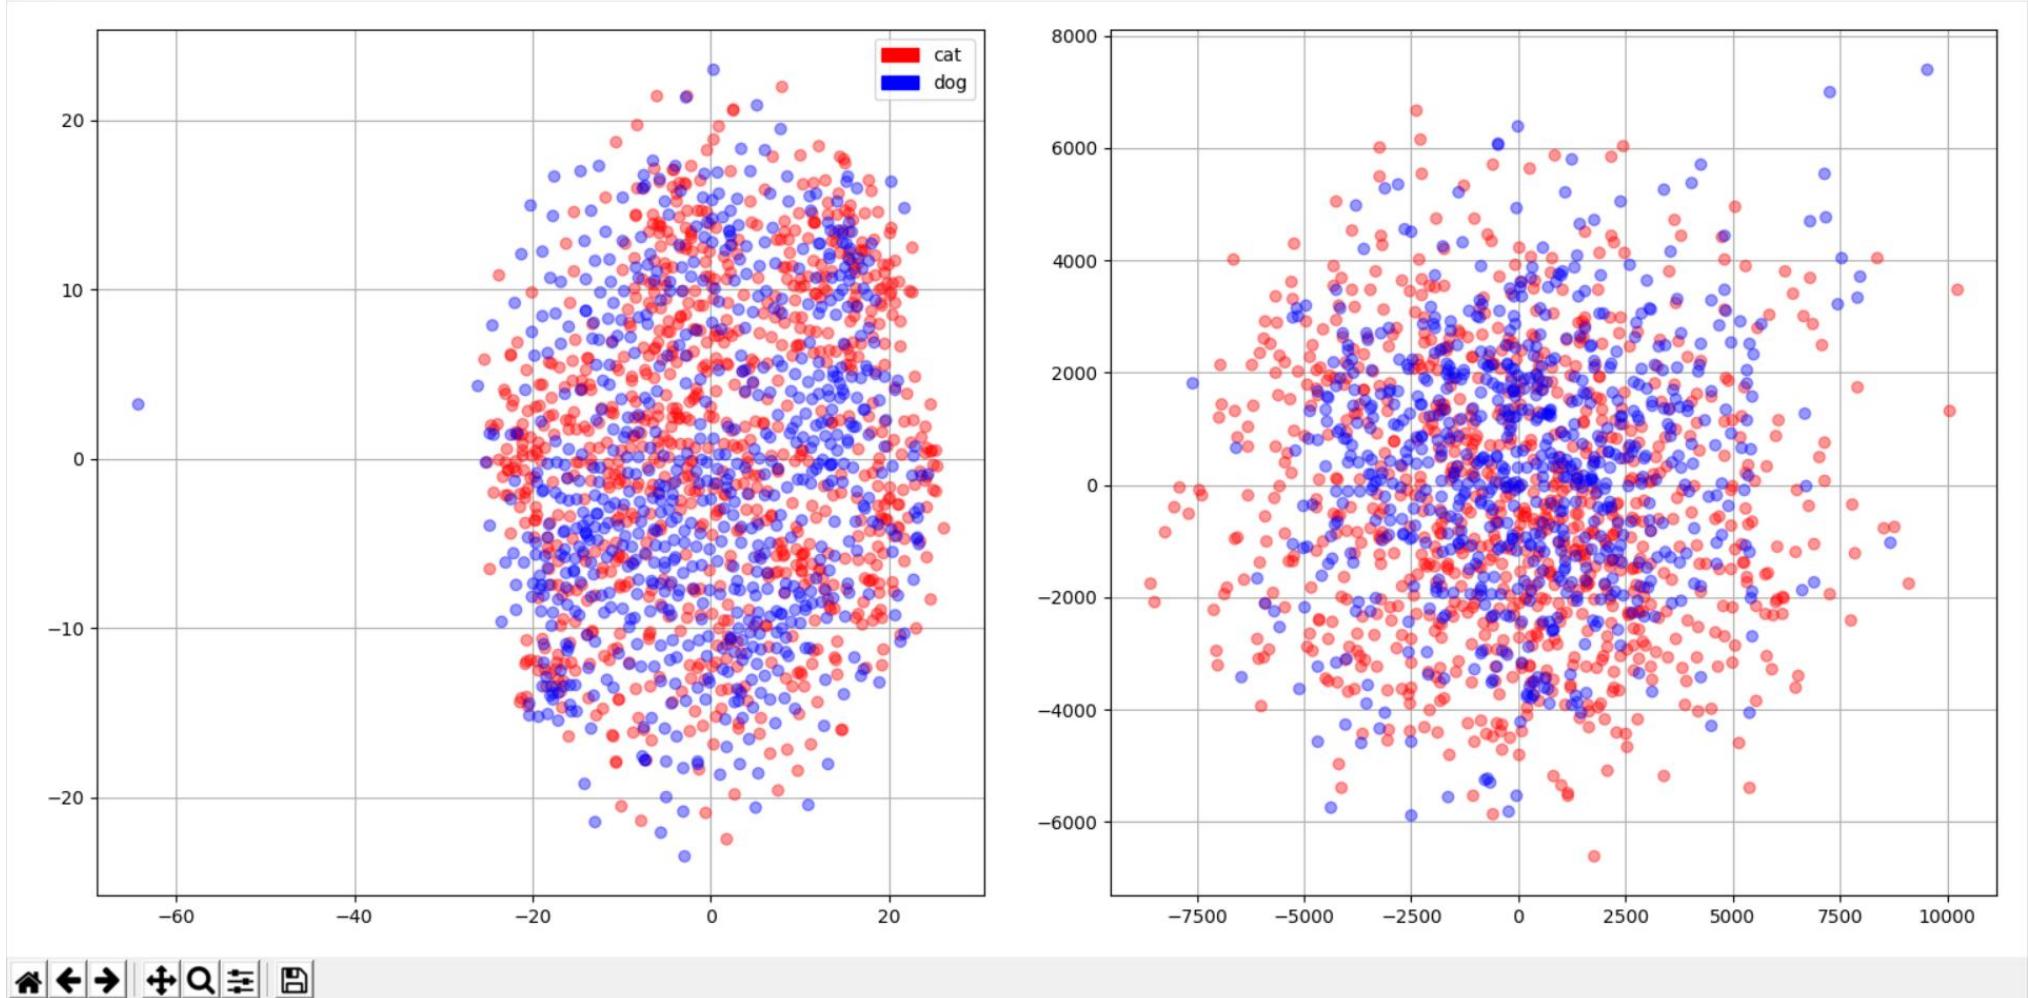


Finetuning

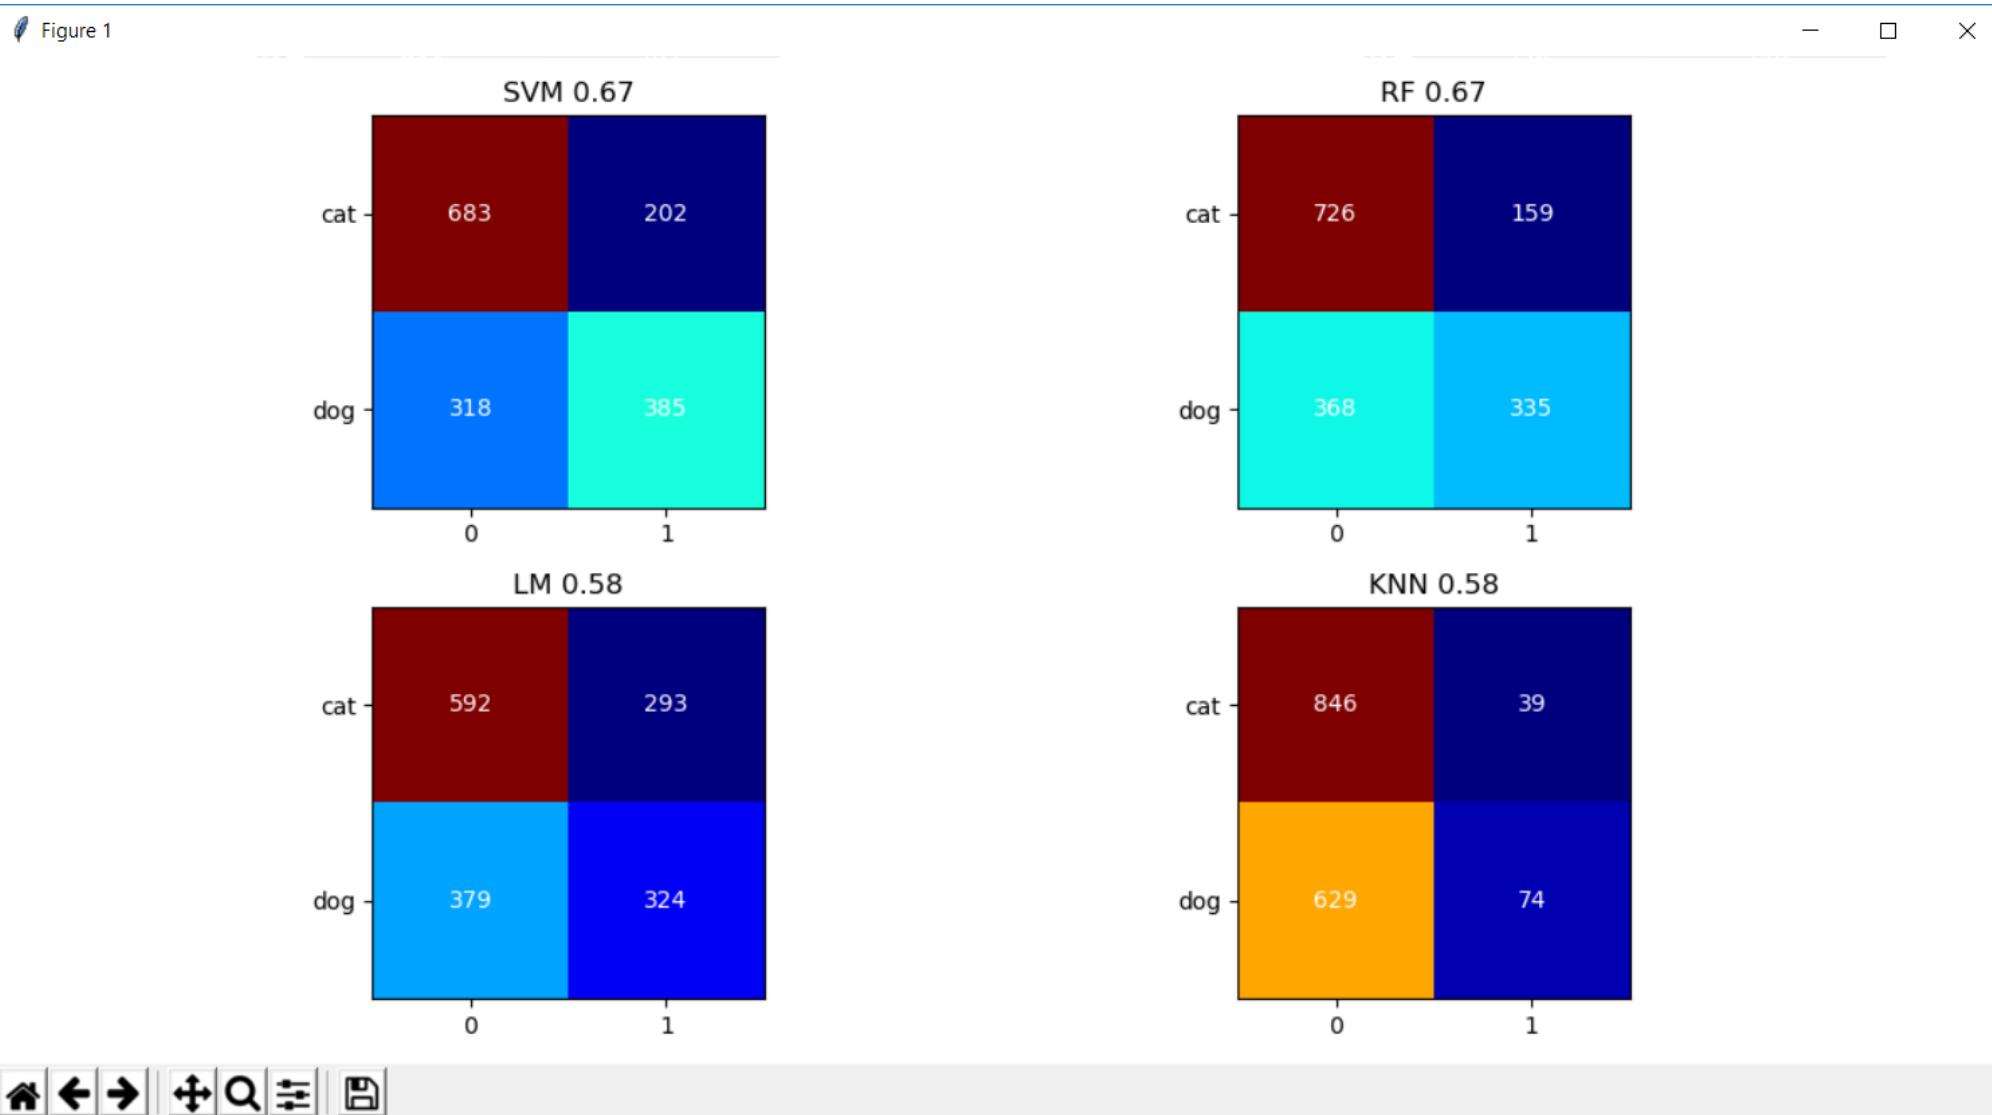


Finetuning

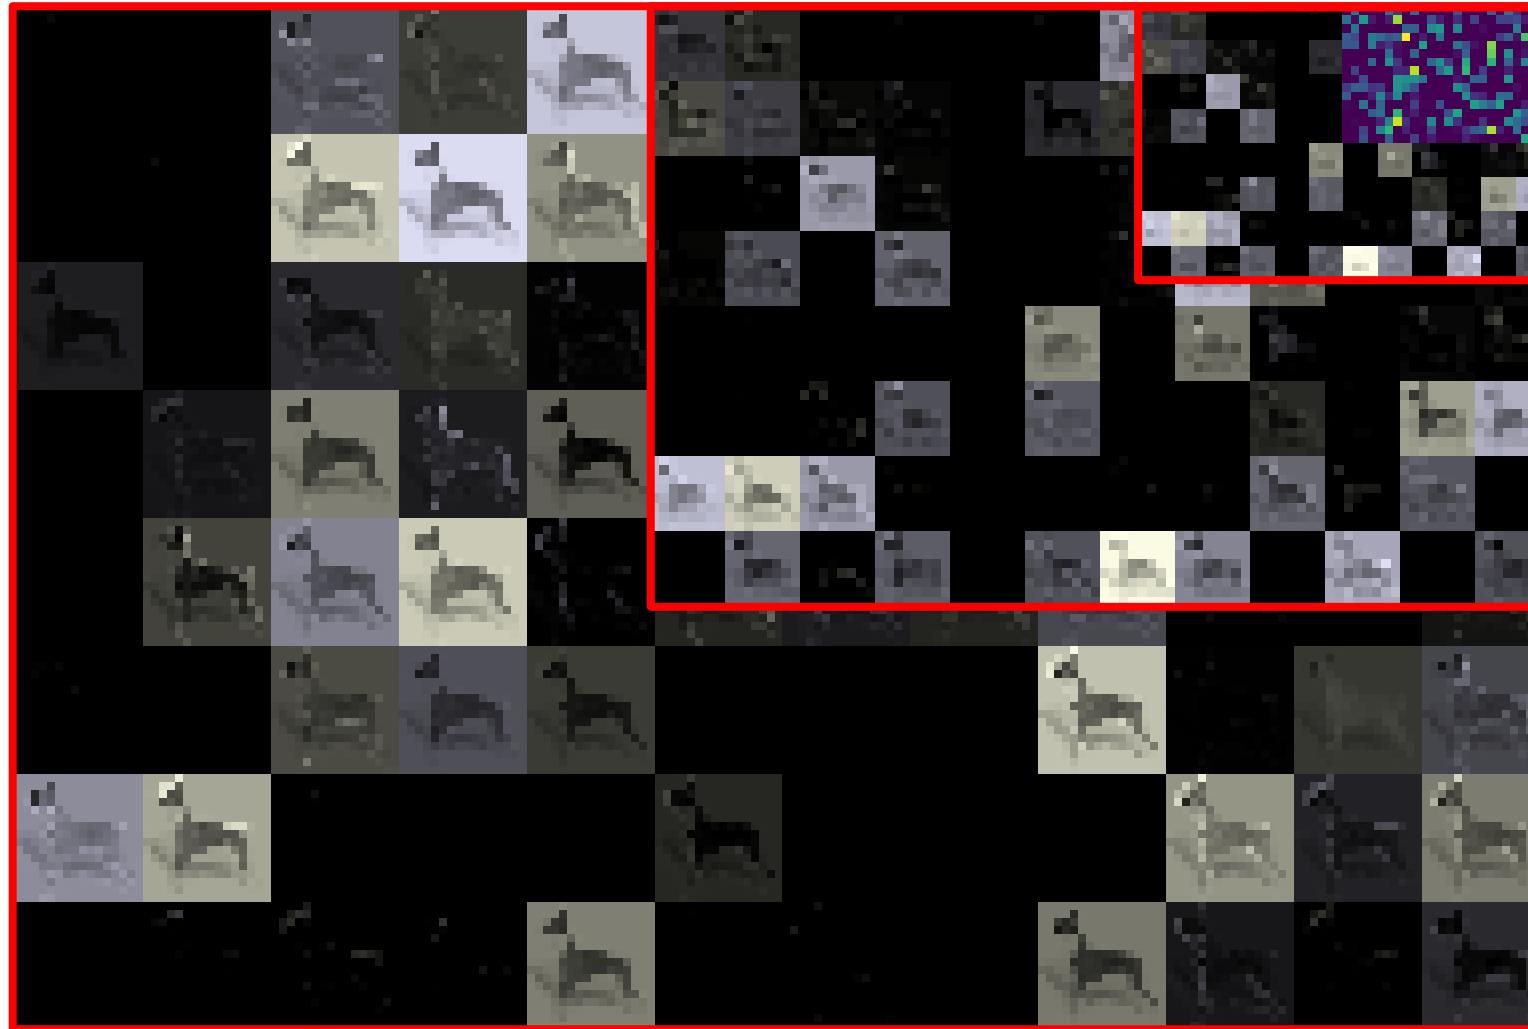
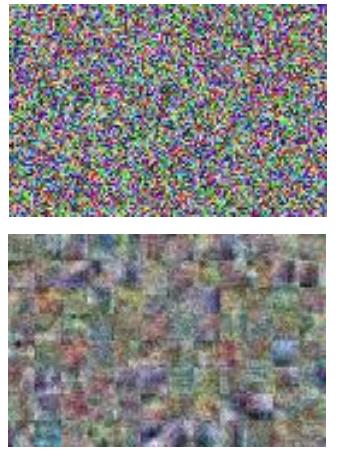
Figure 1



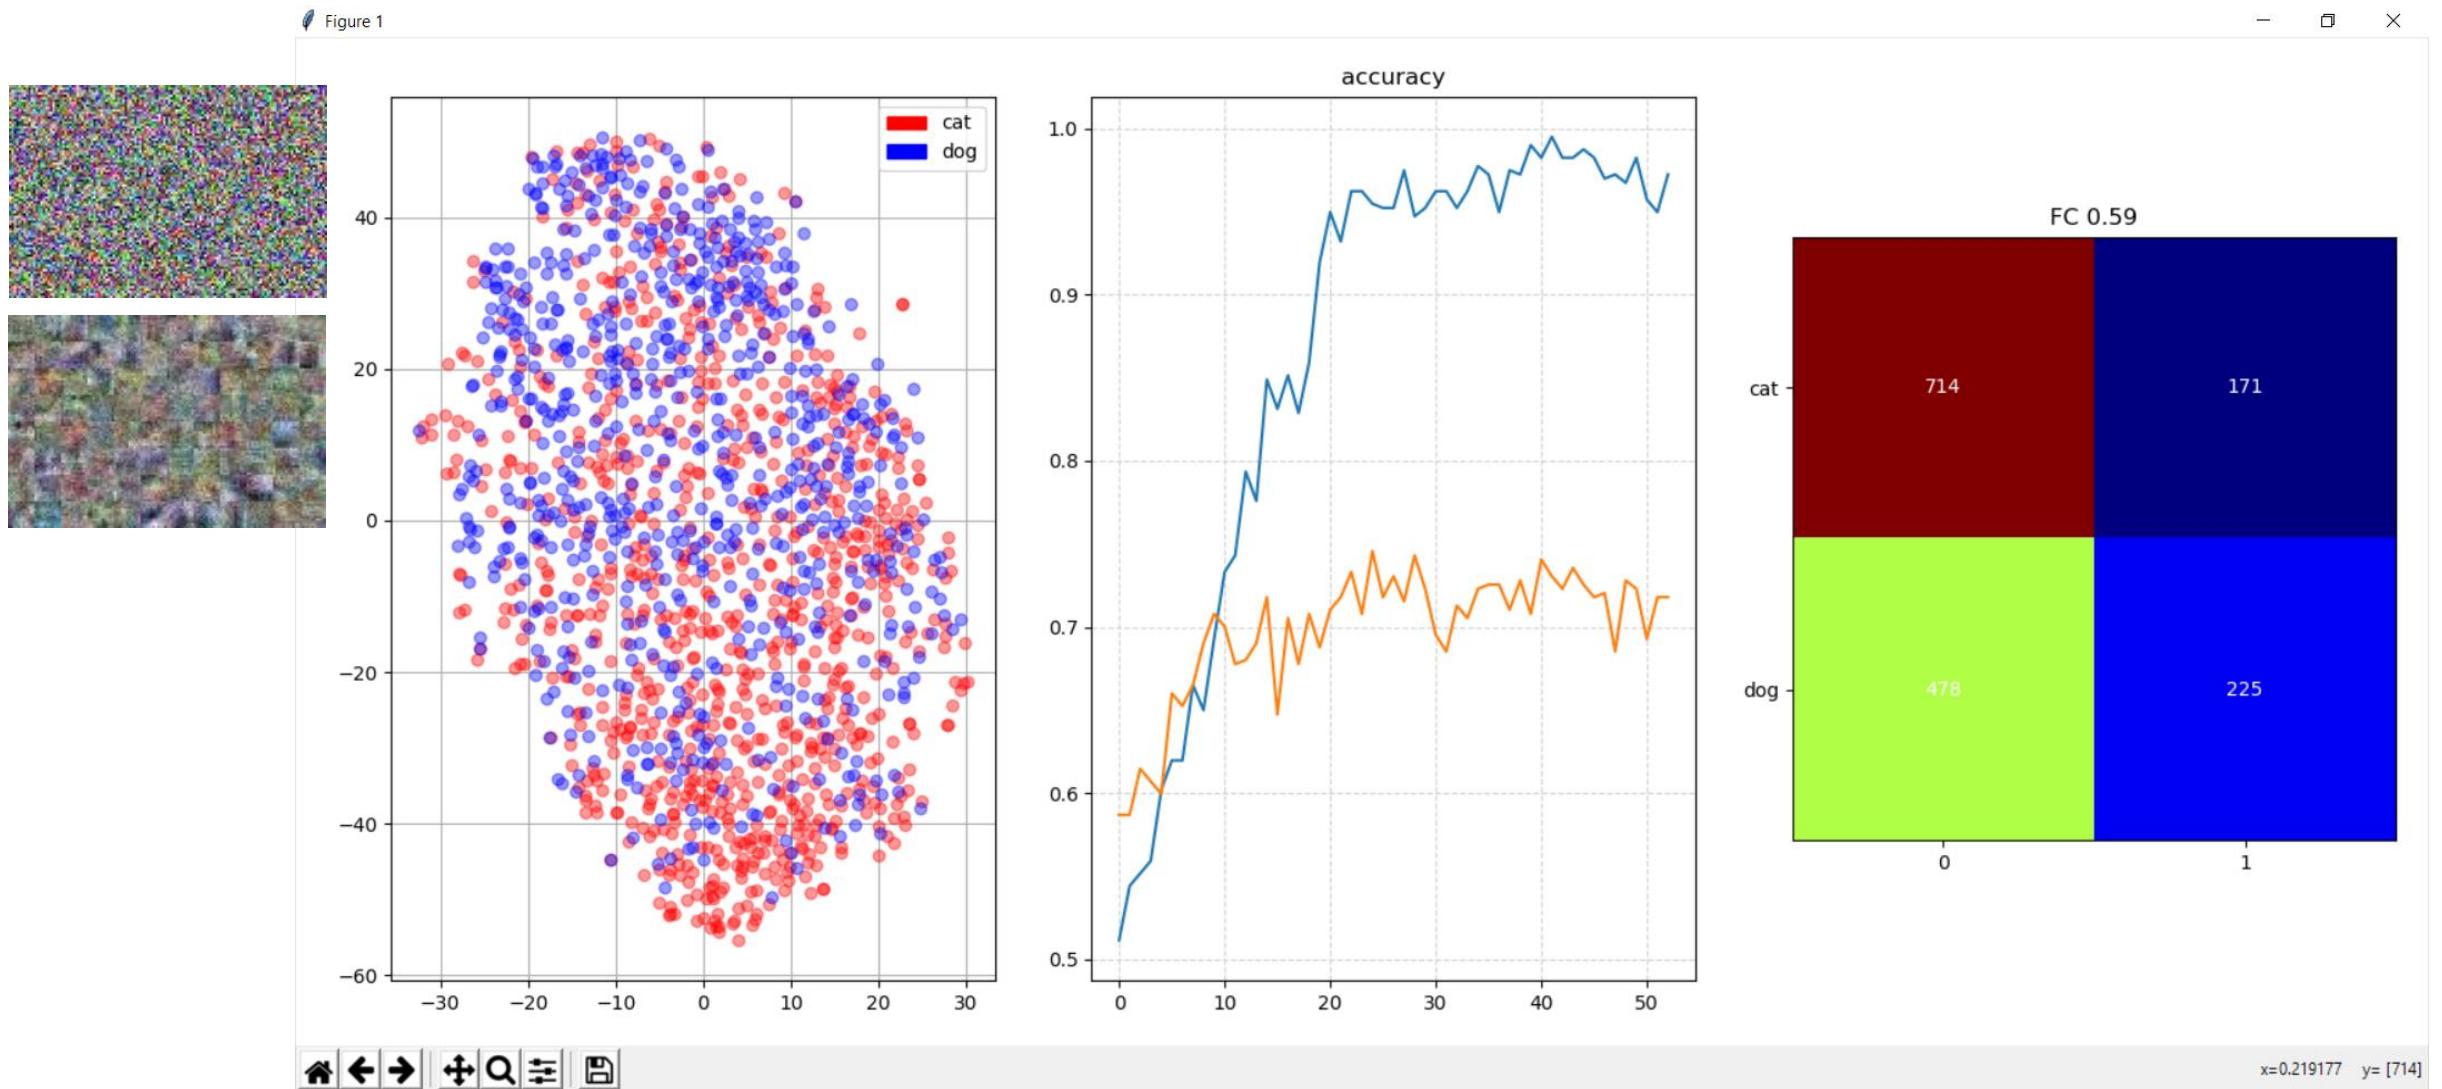
Finetuning



Finetuning



Finetuning



Finetuning

Random before learning

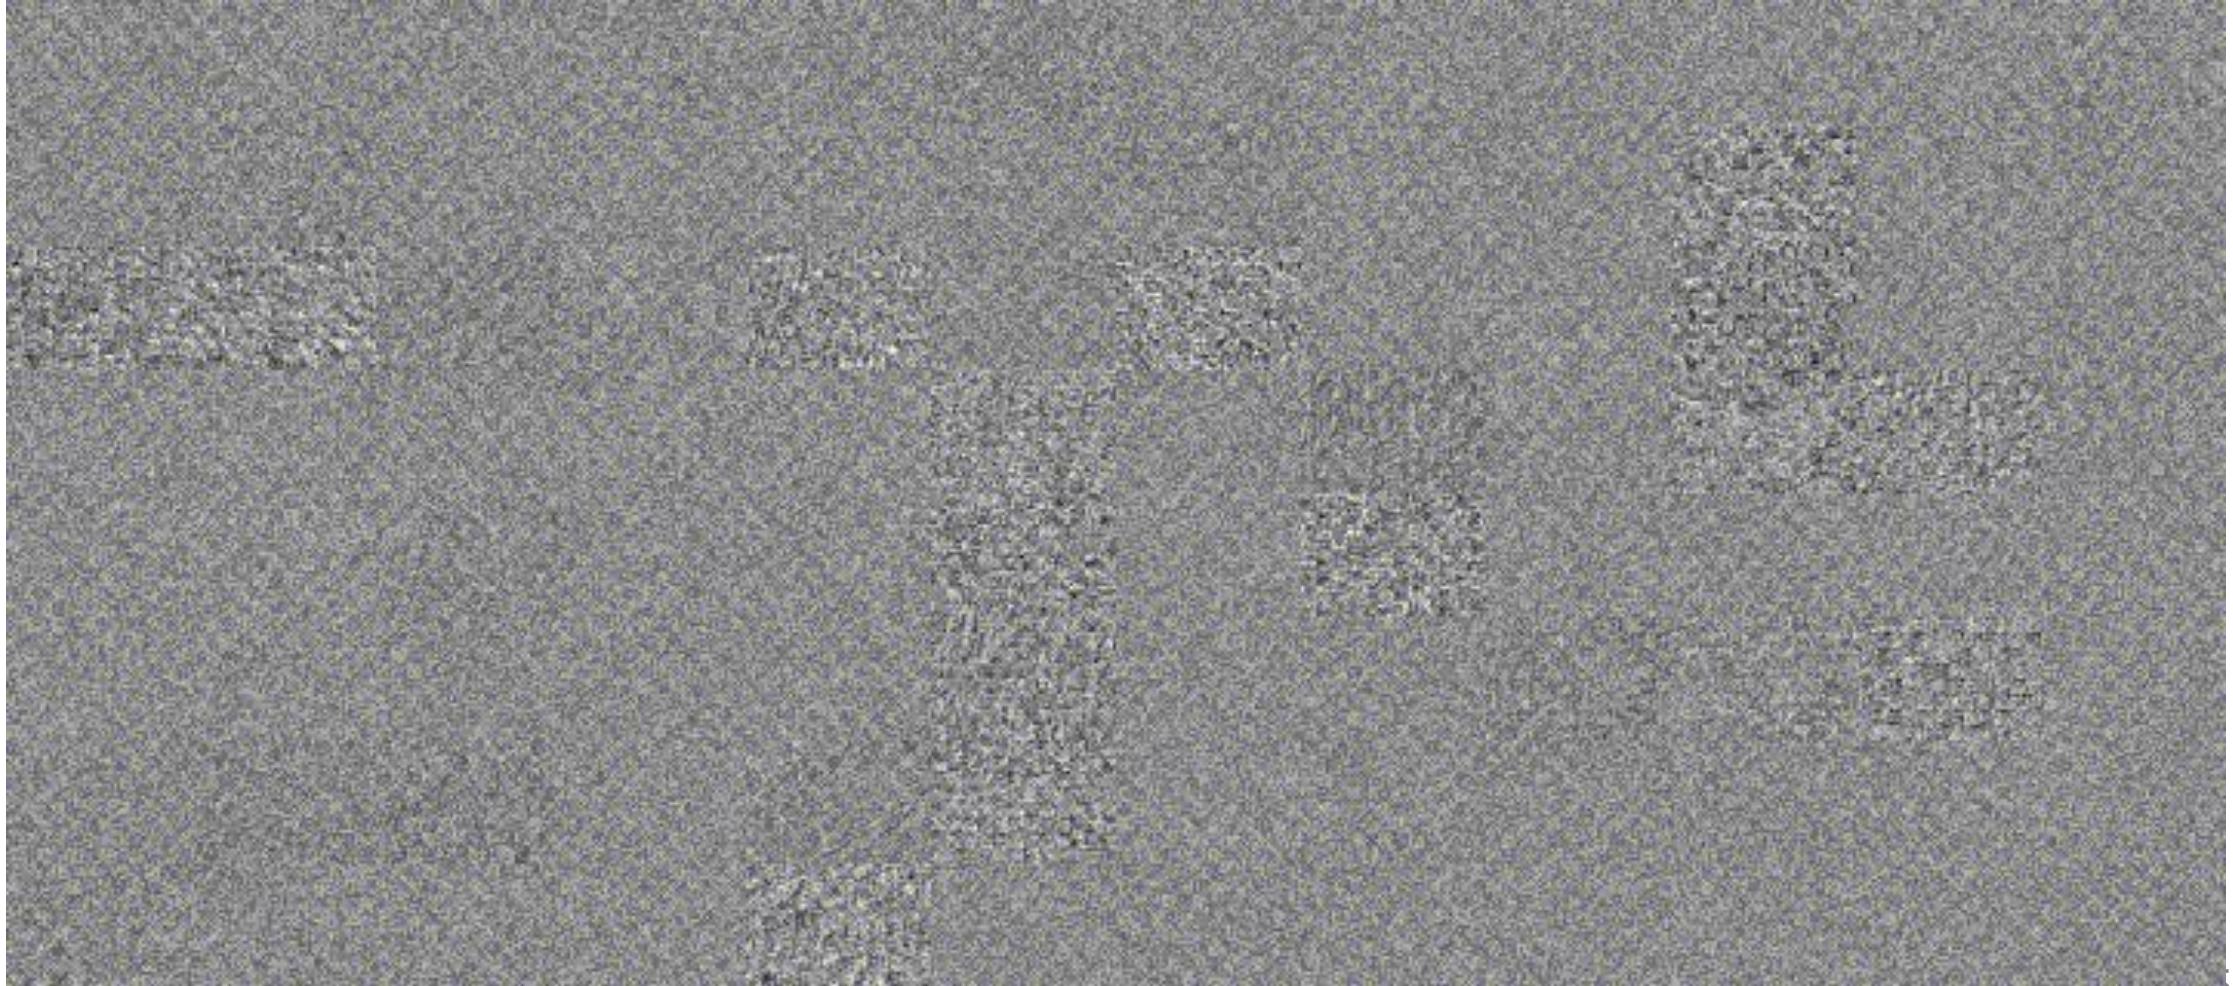


After learning

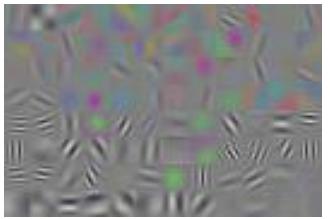


Finetuning

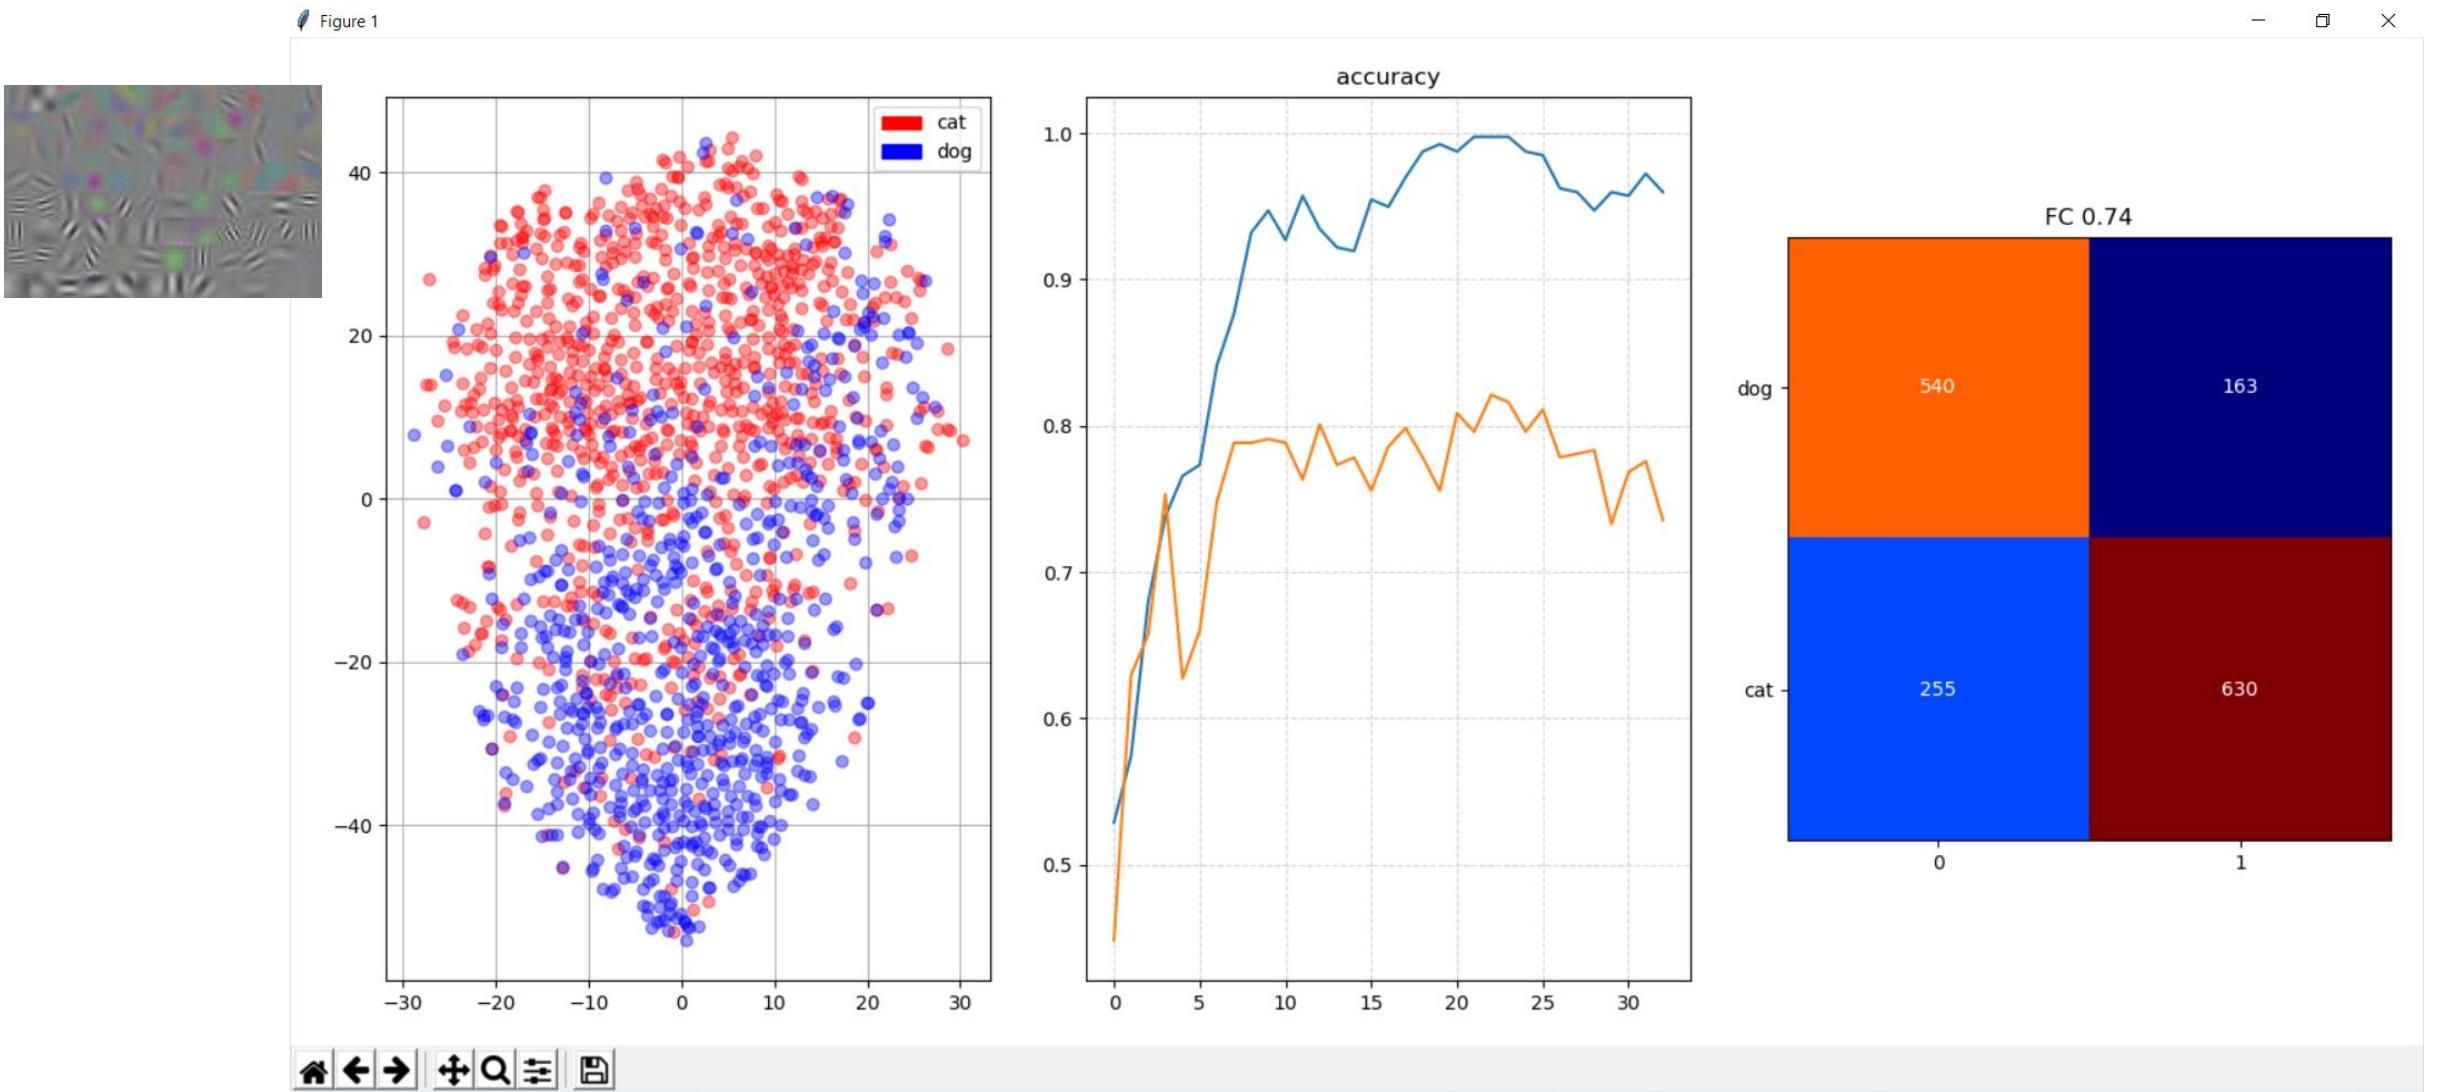
Filters after learning



Finetuning



Finetuning

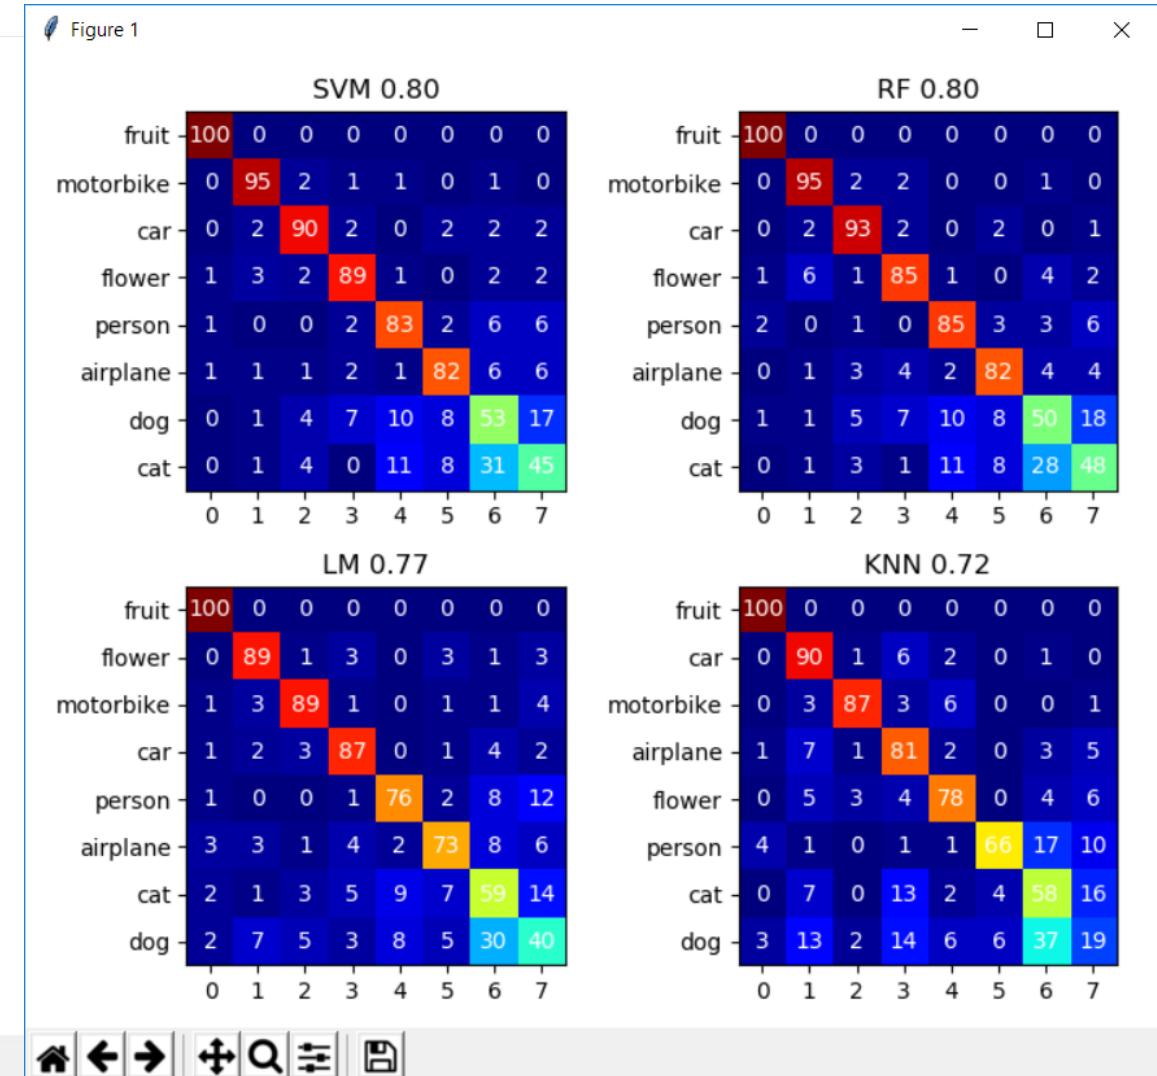
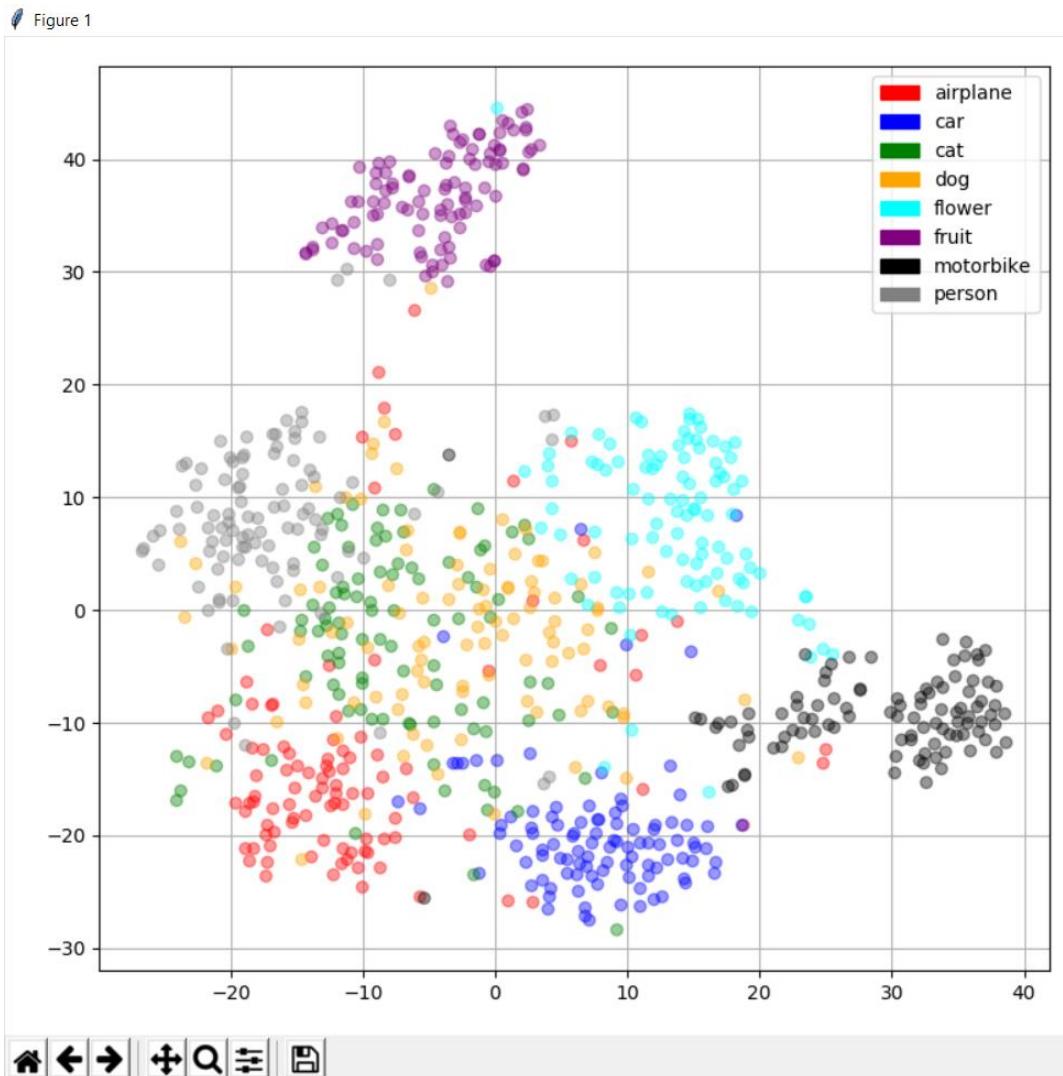


Benchmarking the CNNs



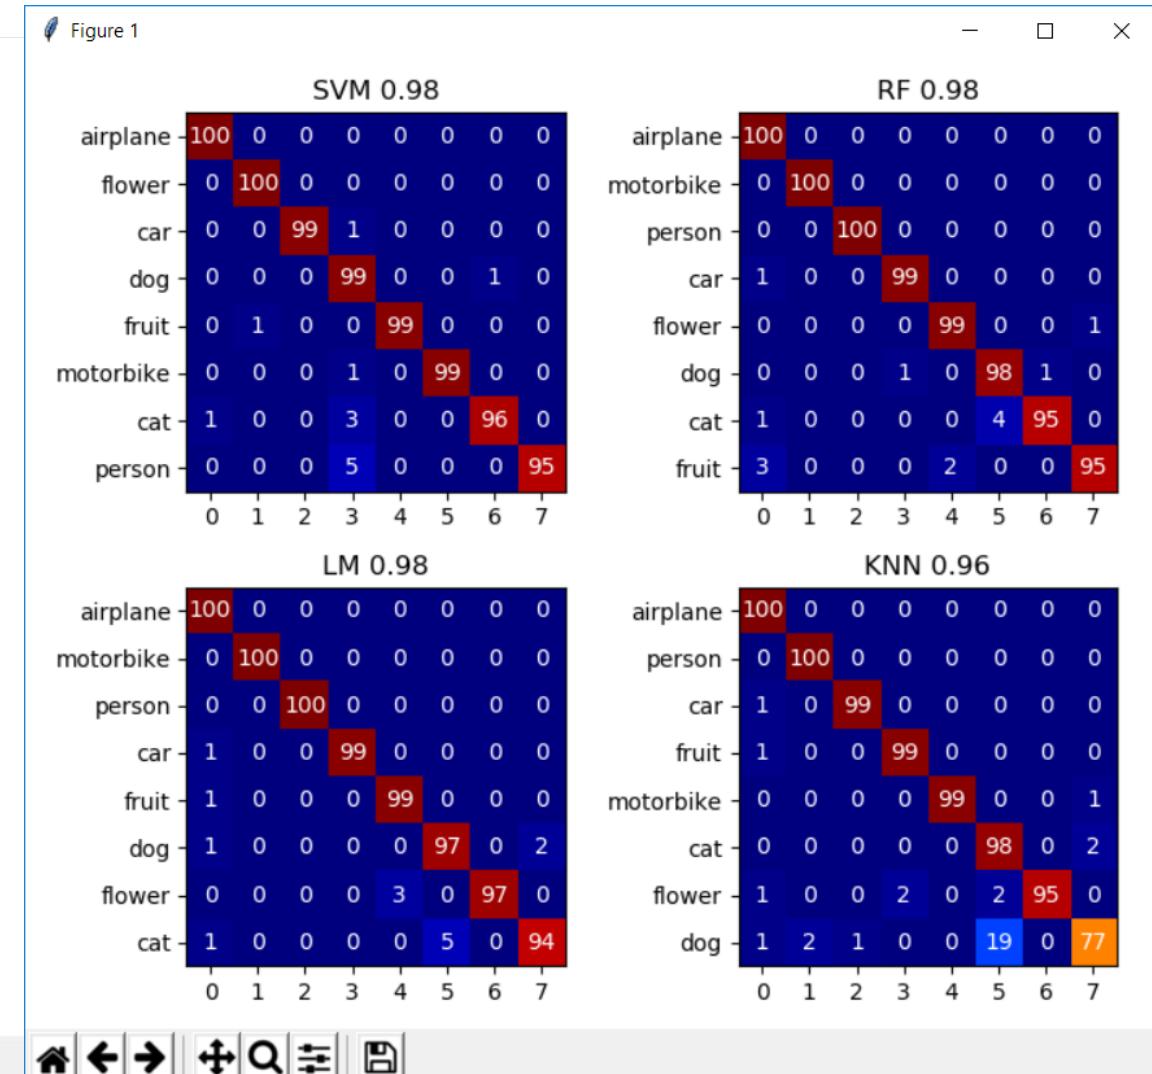
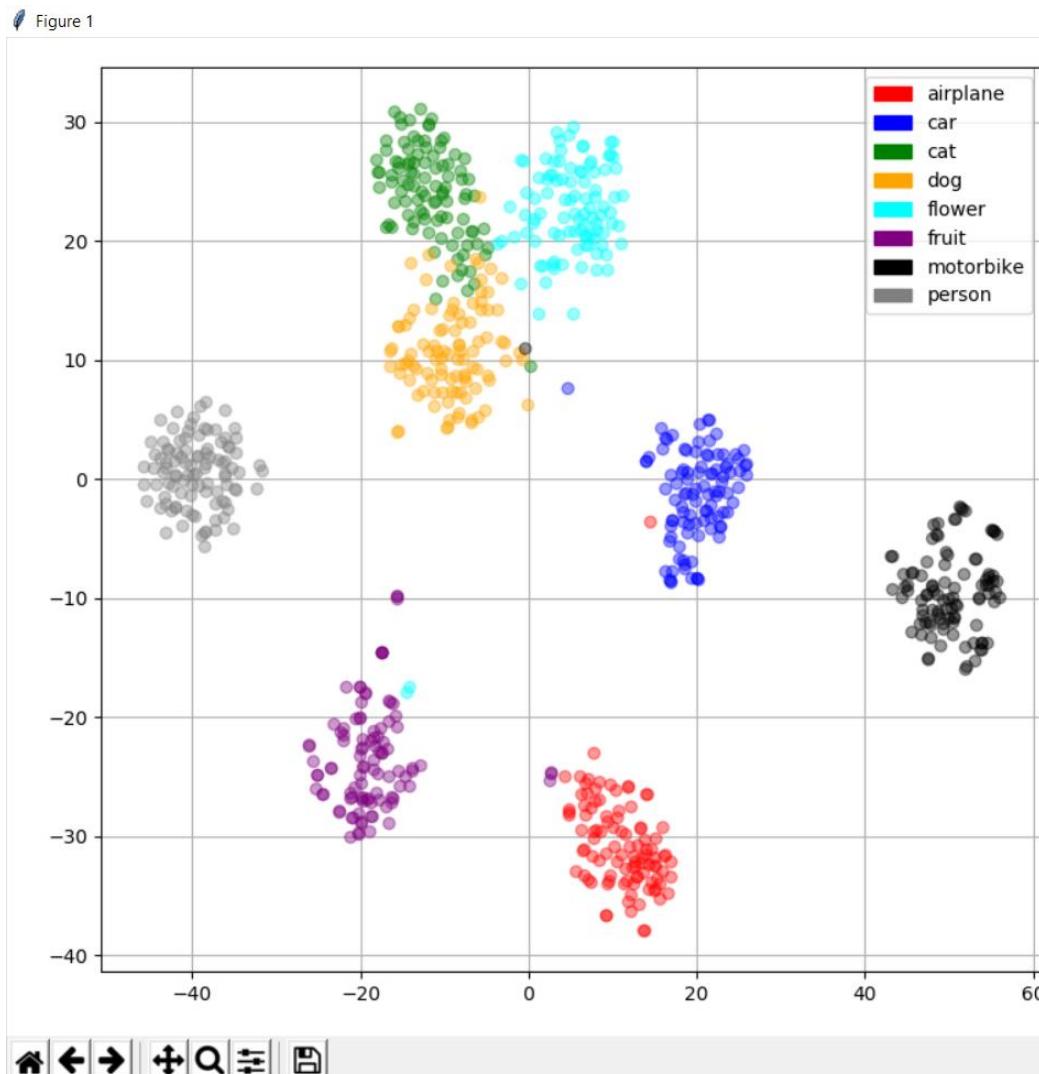
2 conv layers

Benchmarking the CNNs



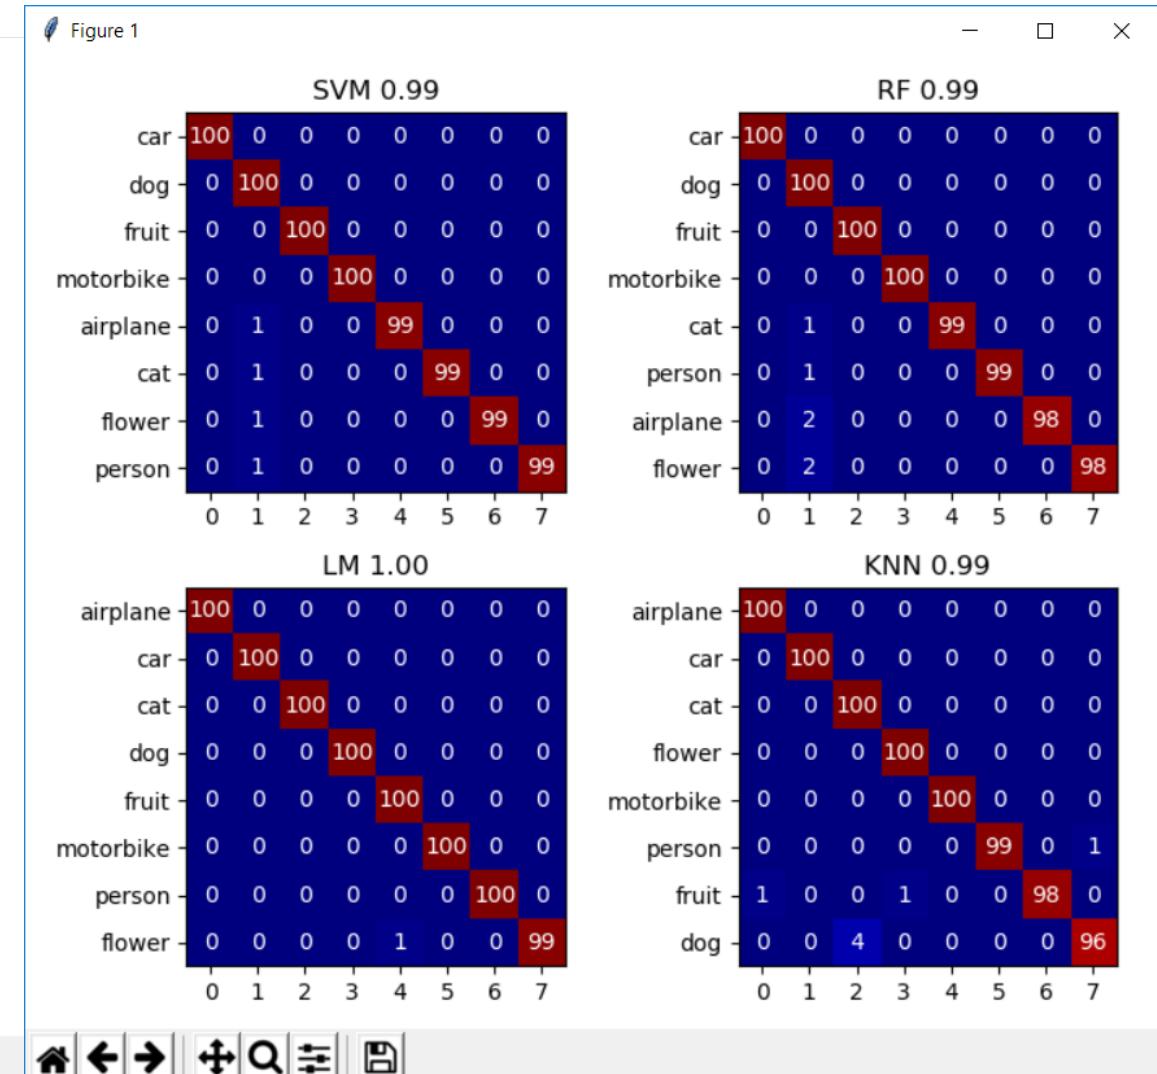
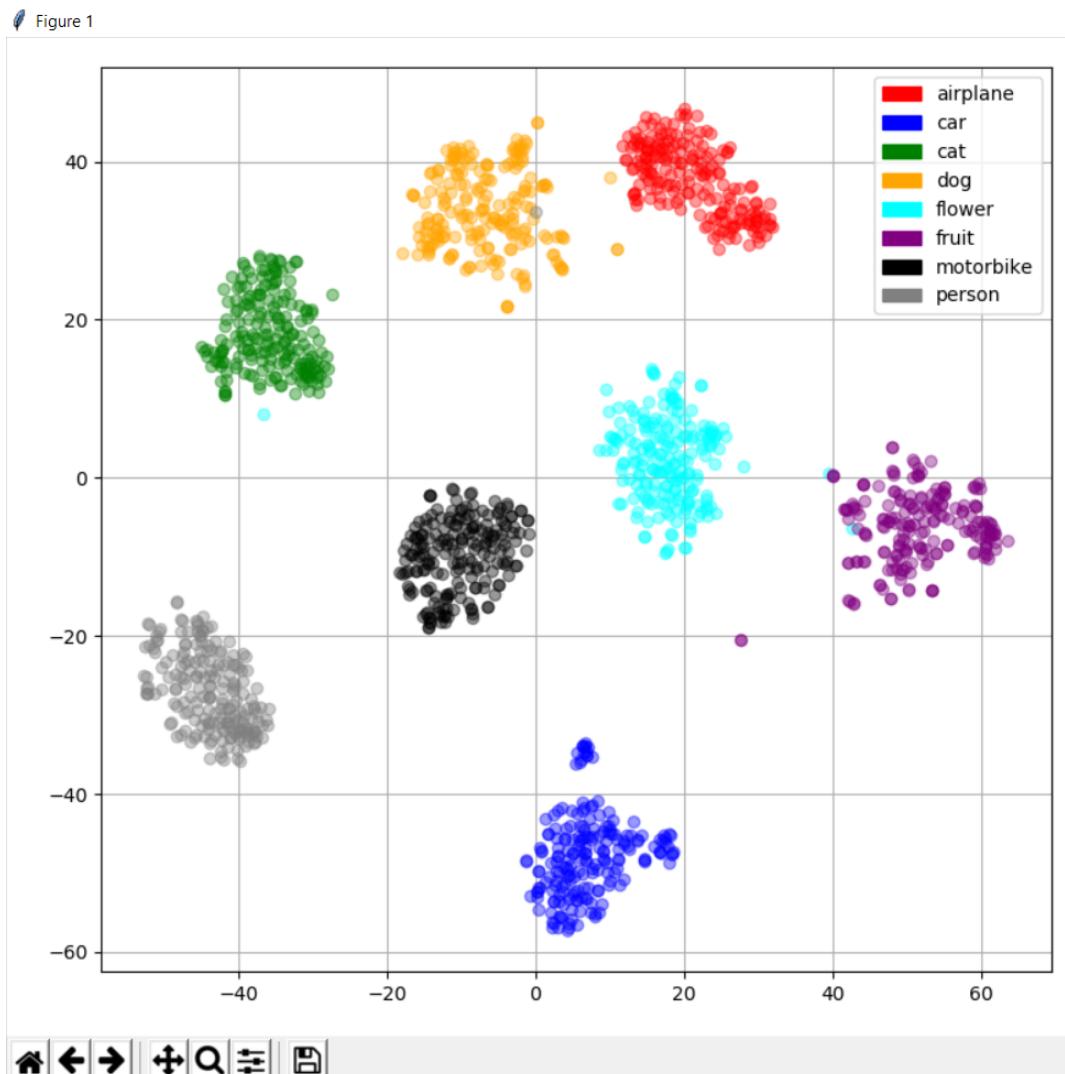
AlexNet

Benchmarking the CNNs



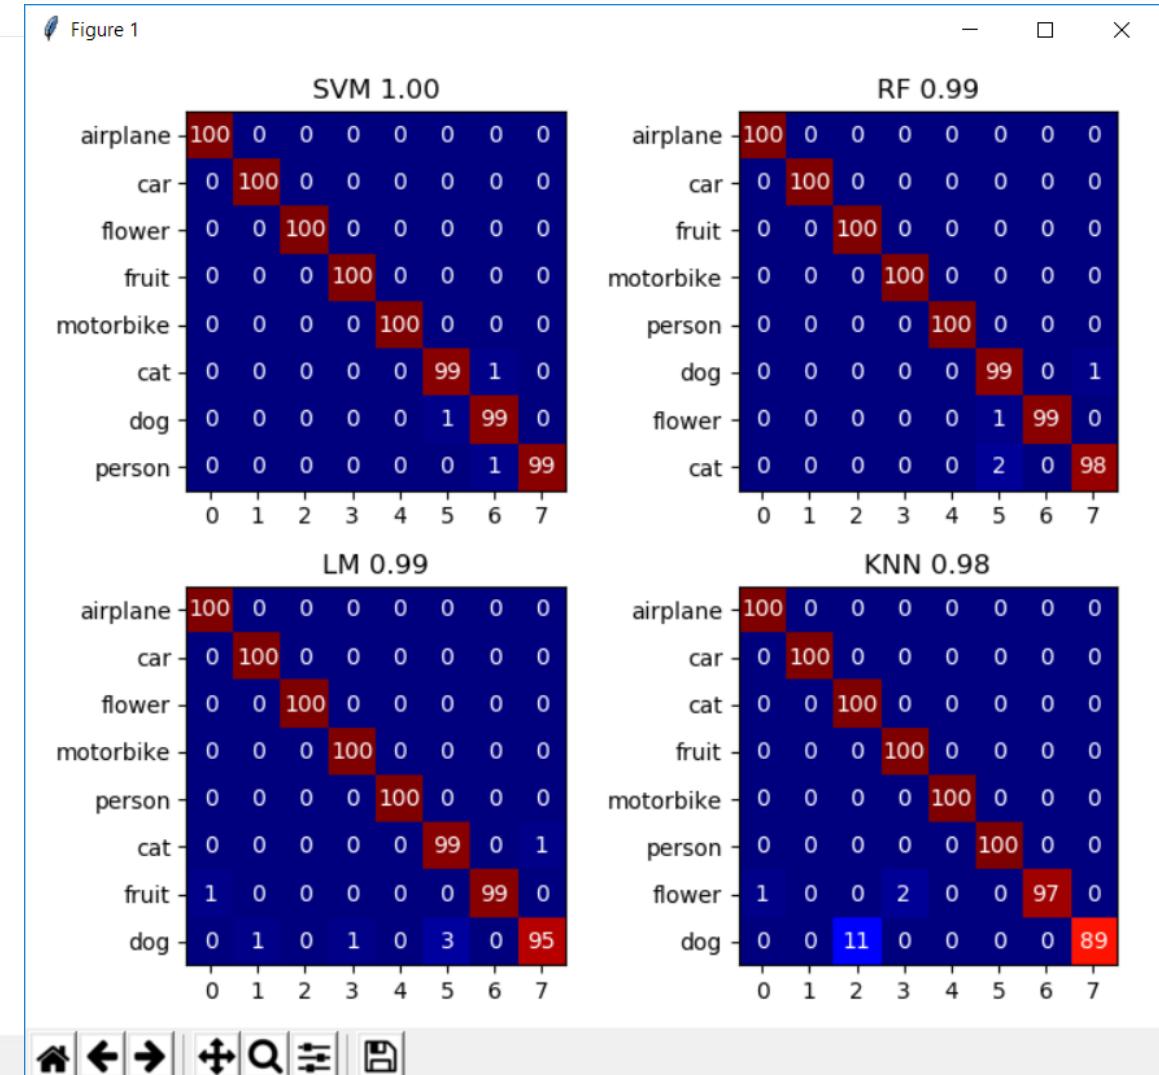
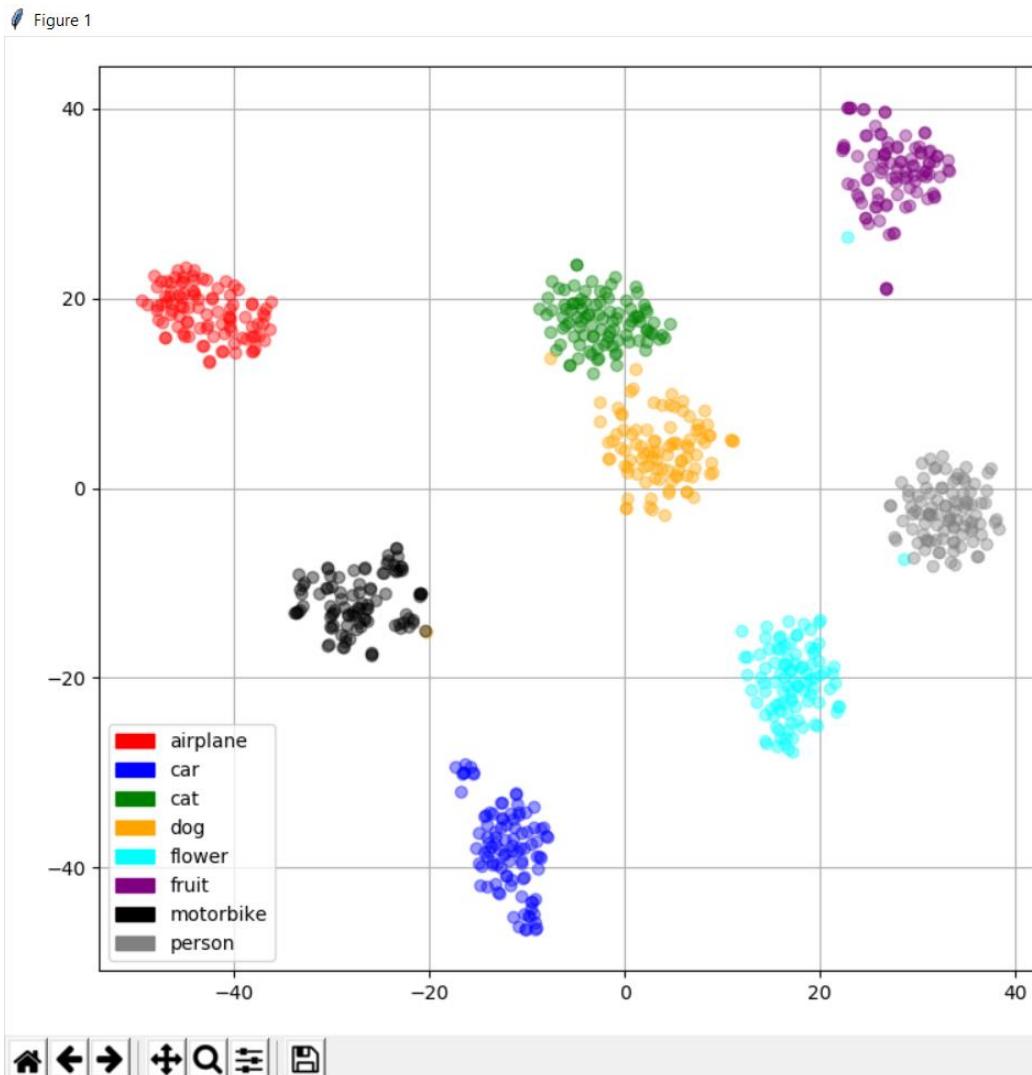
Inception

Benchmarking the CNNs

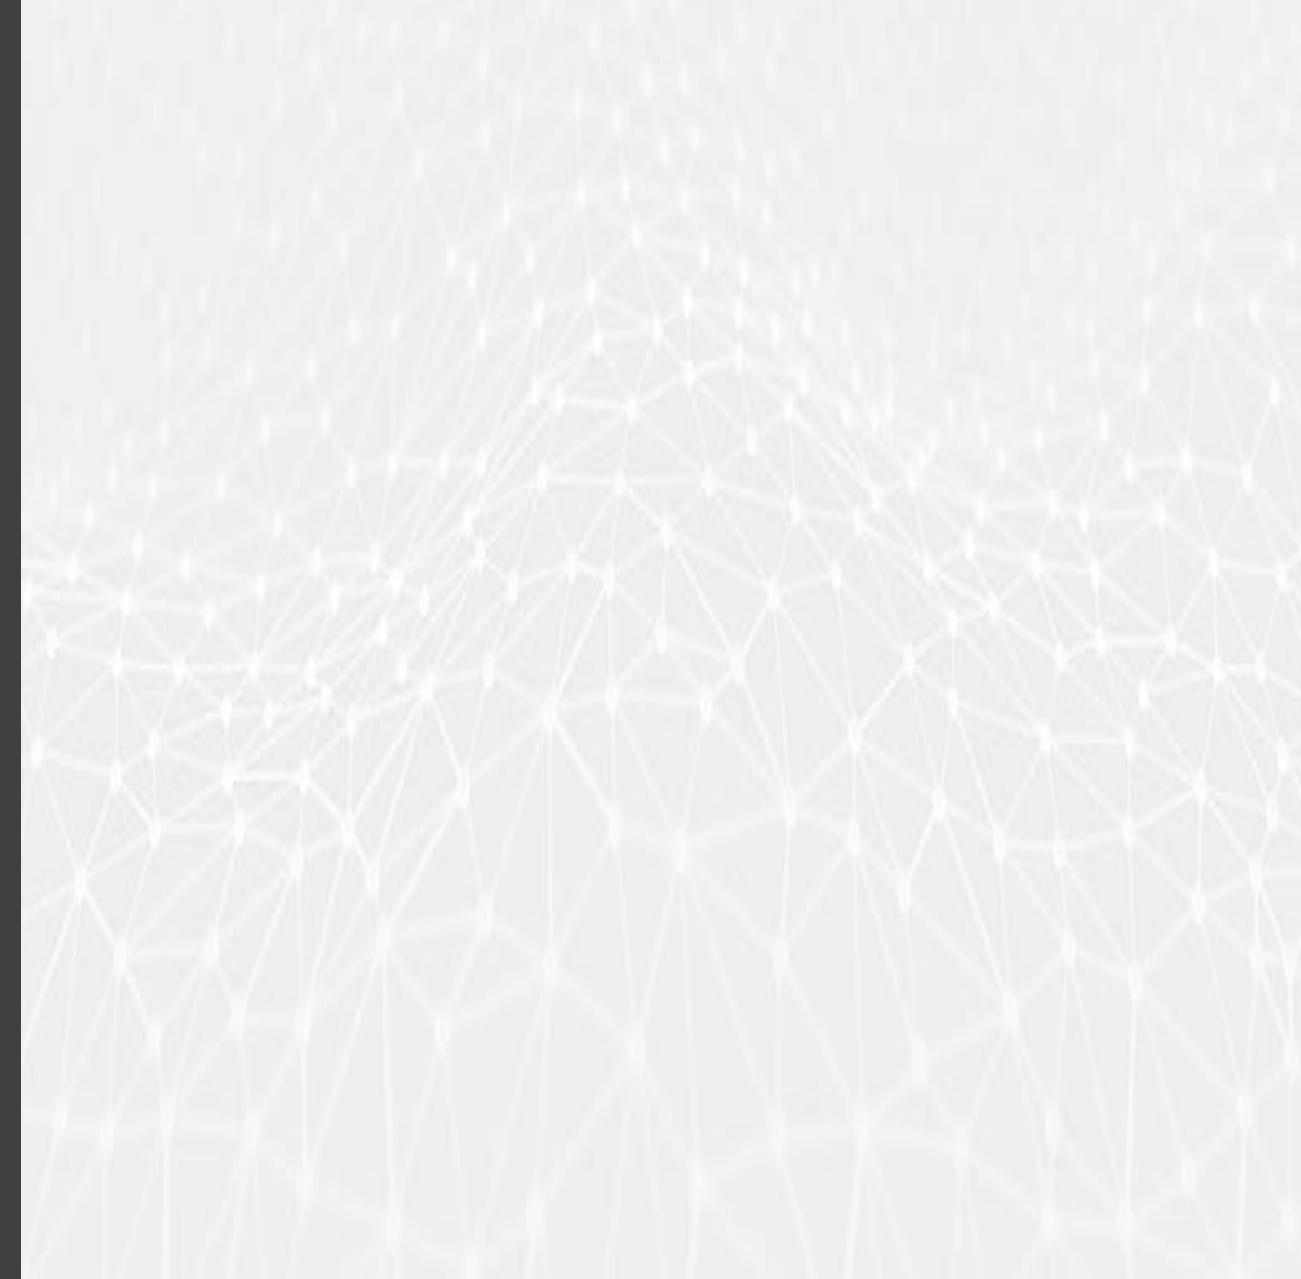


MobileNet

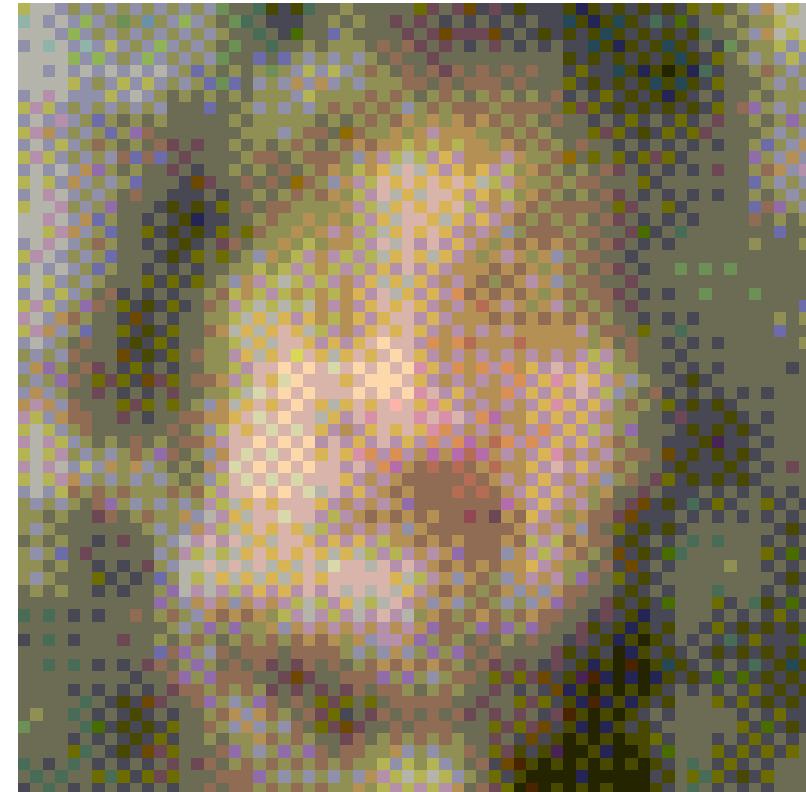
Benchmarking the CNNs



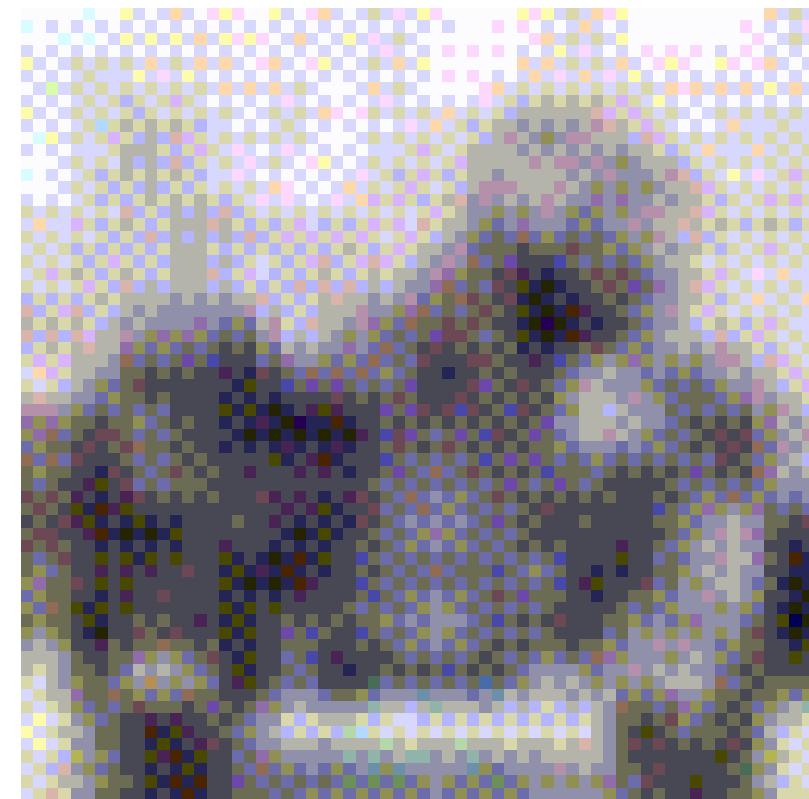
Auto-encoders



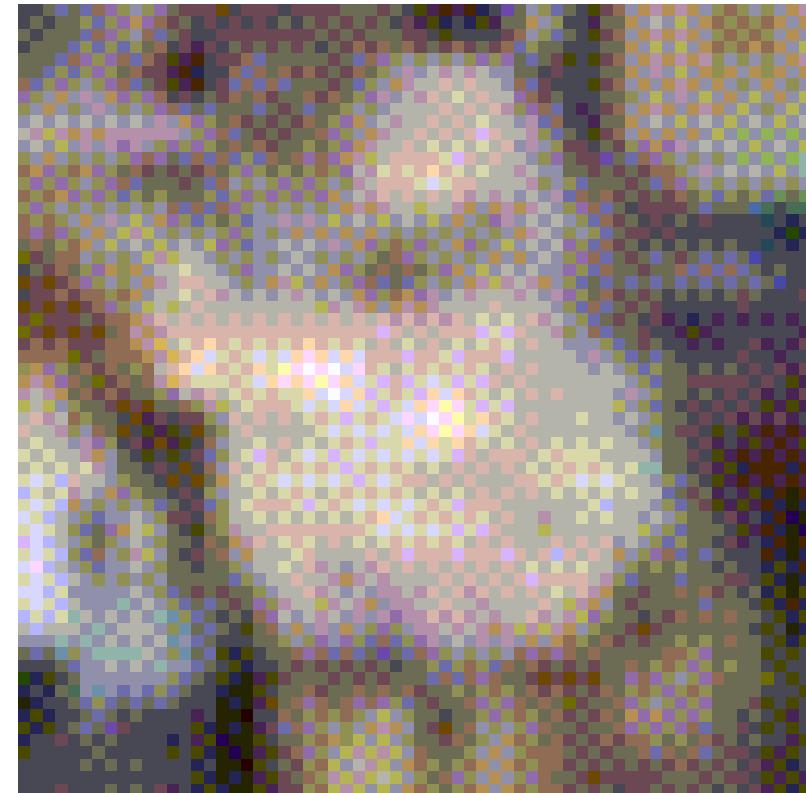
Reconstruct features back into images



Reconstruct features back into images

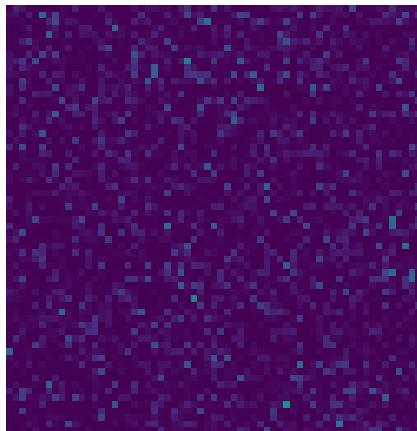


Reconstruct features back into images

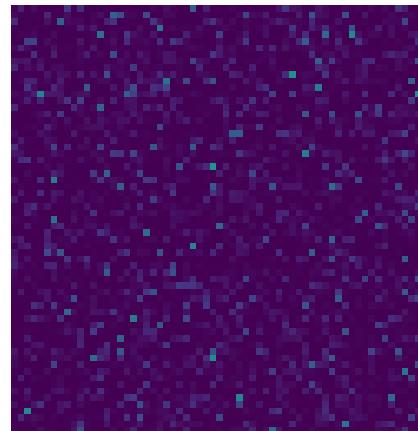


Reconstruct features back into images

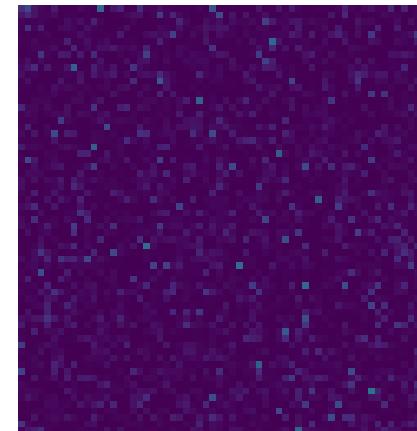
Features



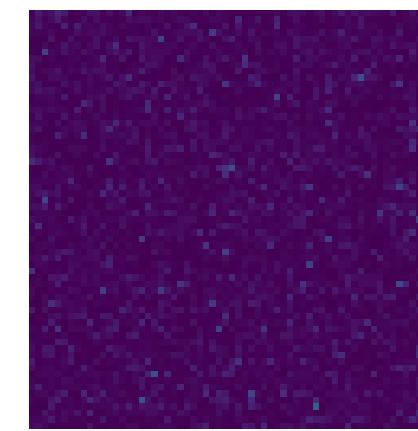
airplane



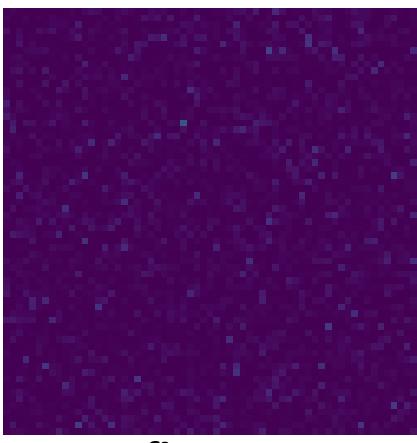
car



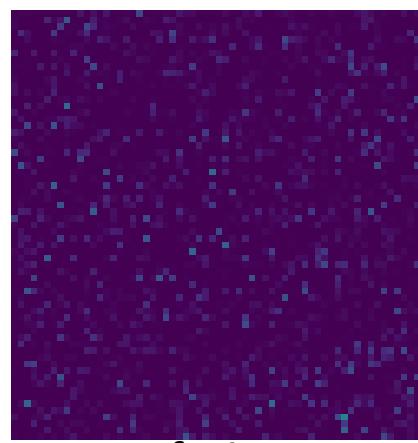
cat



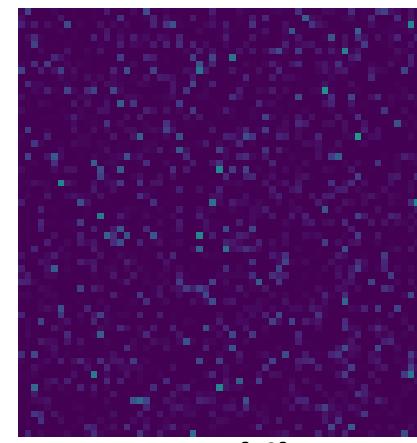
dog



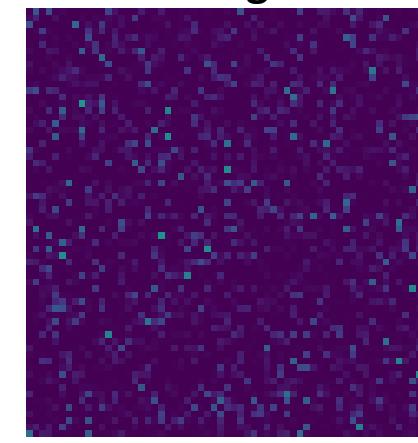
flower



fruit



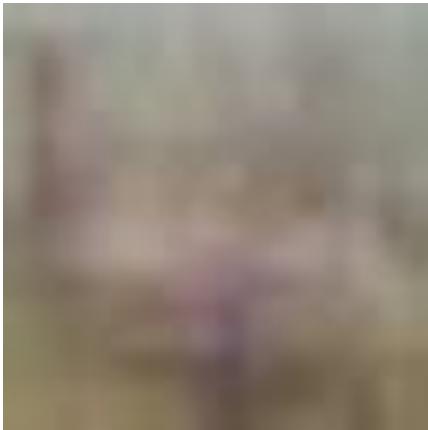
motorbike



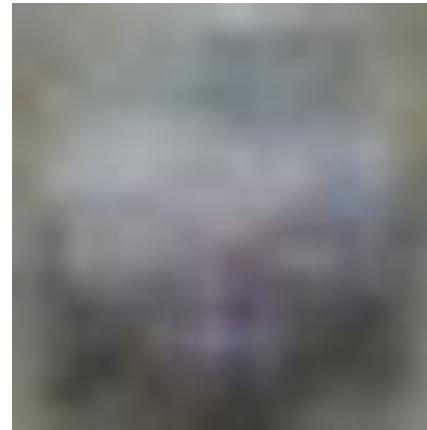
person

Reconstruct features back into images

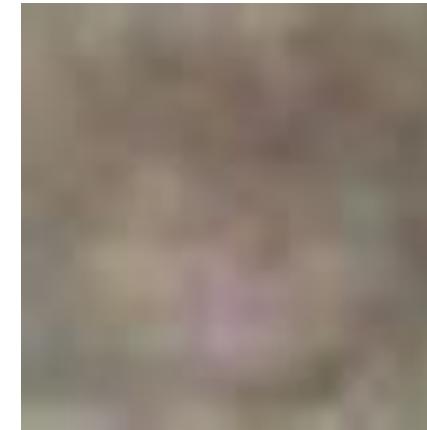
Reconstructed



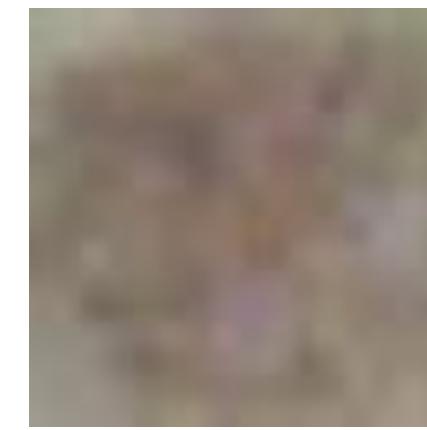
airplane



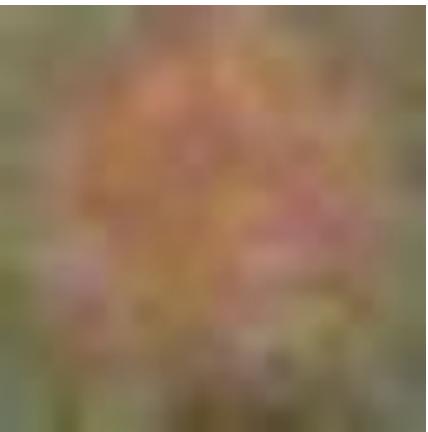
car



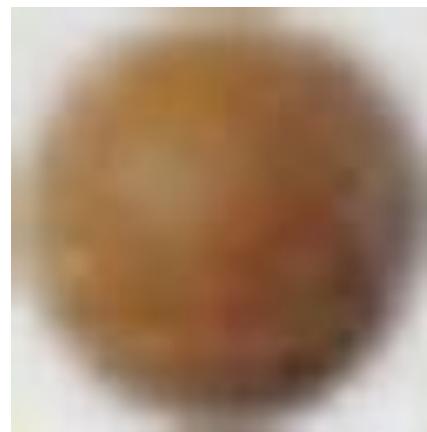
cat



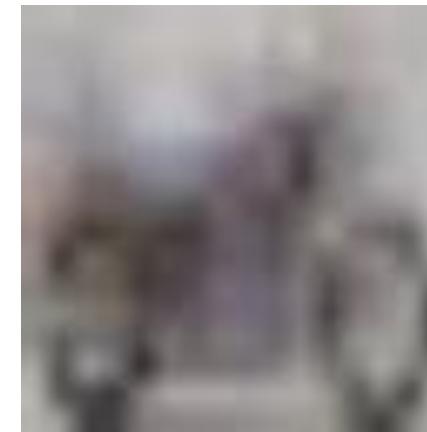
dog



flower



fruit



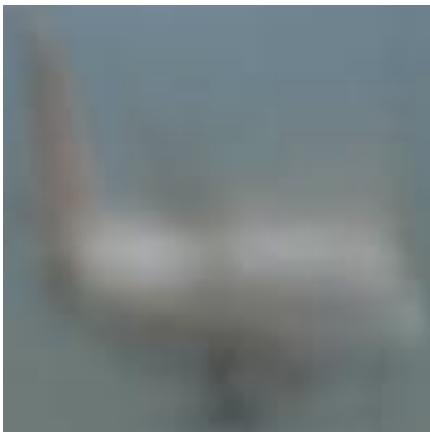
motorbike



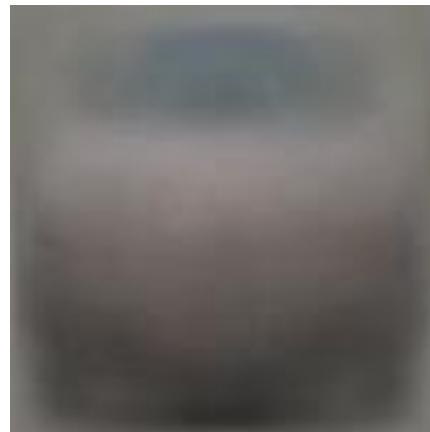
person

Reconstruct features back into images

Original



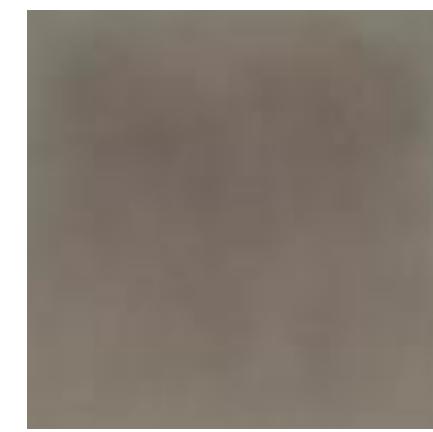
airplane



car



cat



dog



flower



fruit



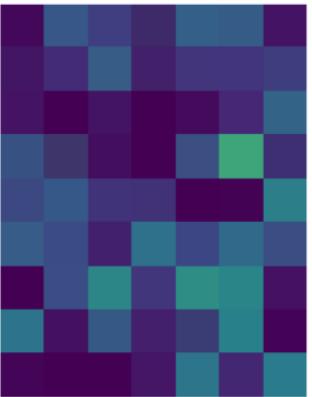
motorbike



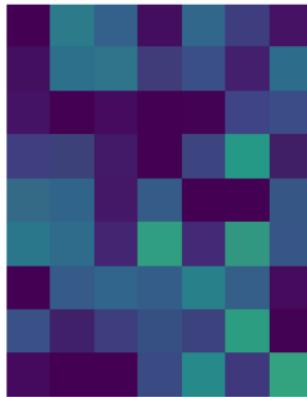
person

Reconstruct features back into images

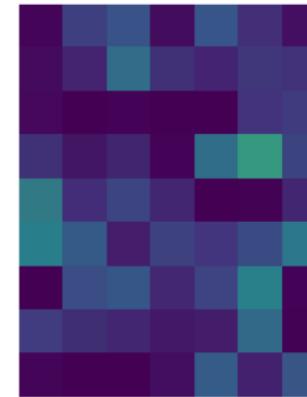
Features



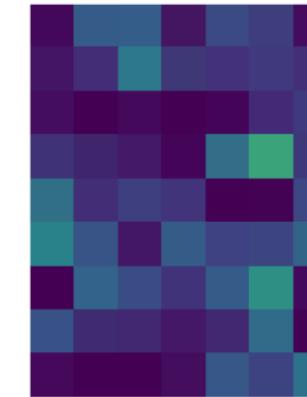
airplane



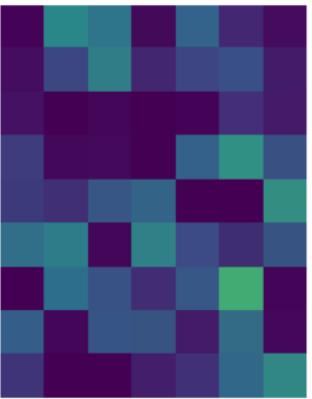
car



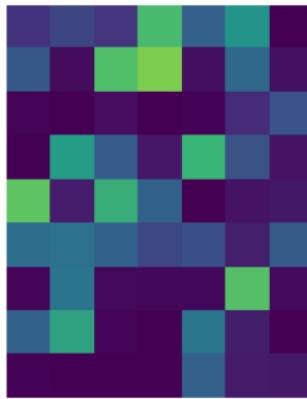
cat



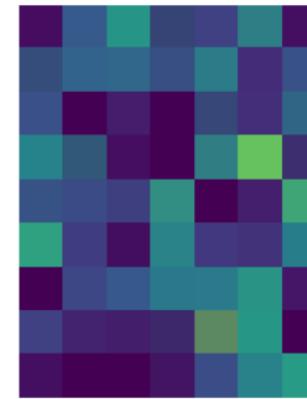
dog



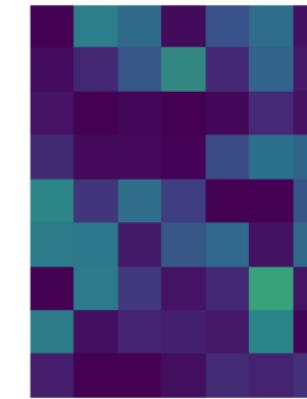
flower



fruit



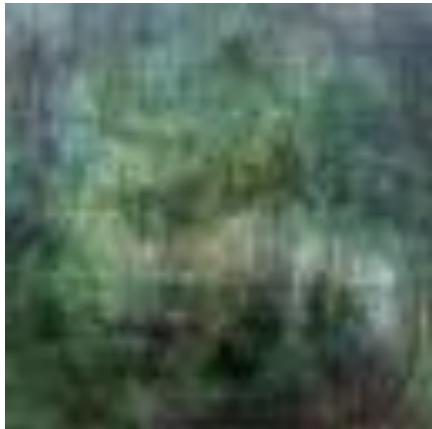
motorbike



person

Reconstruct features back into images

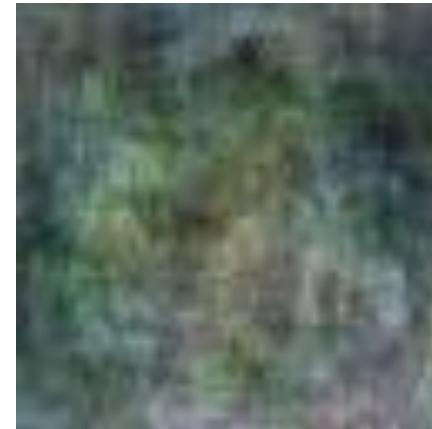
Reconstructed



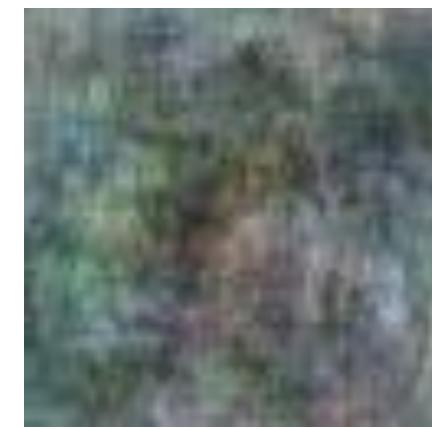
airplane



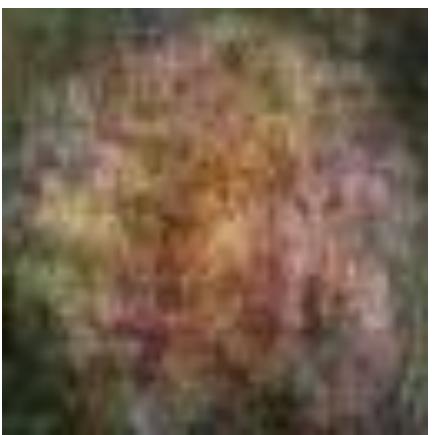
car



cat



dog



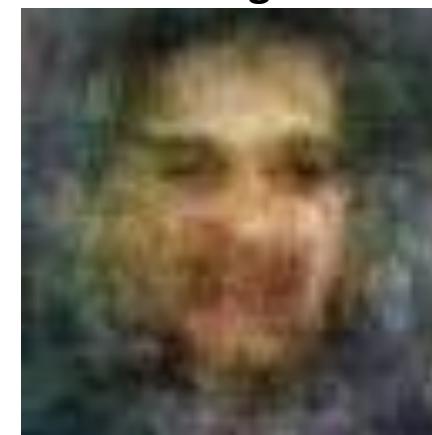
flower



fruit

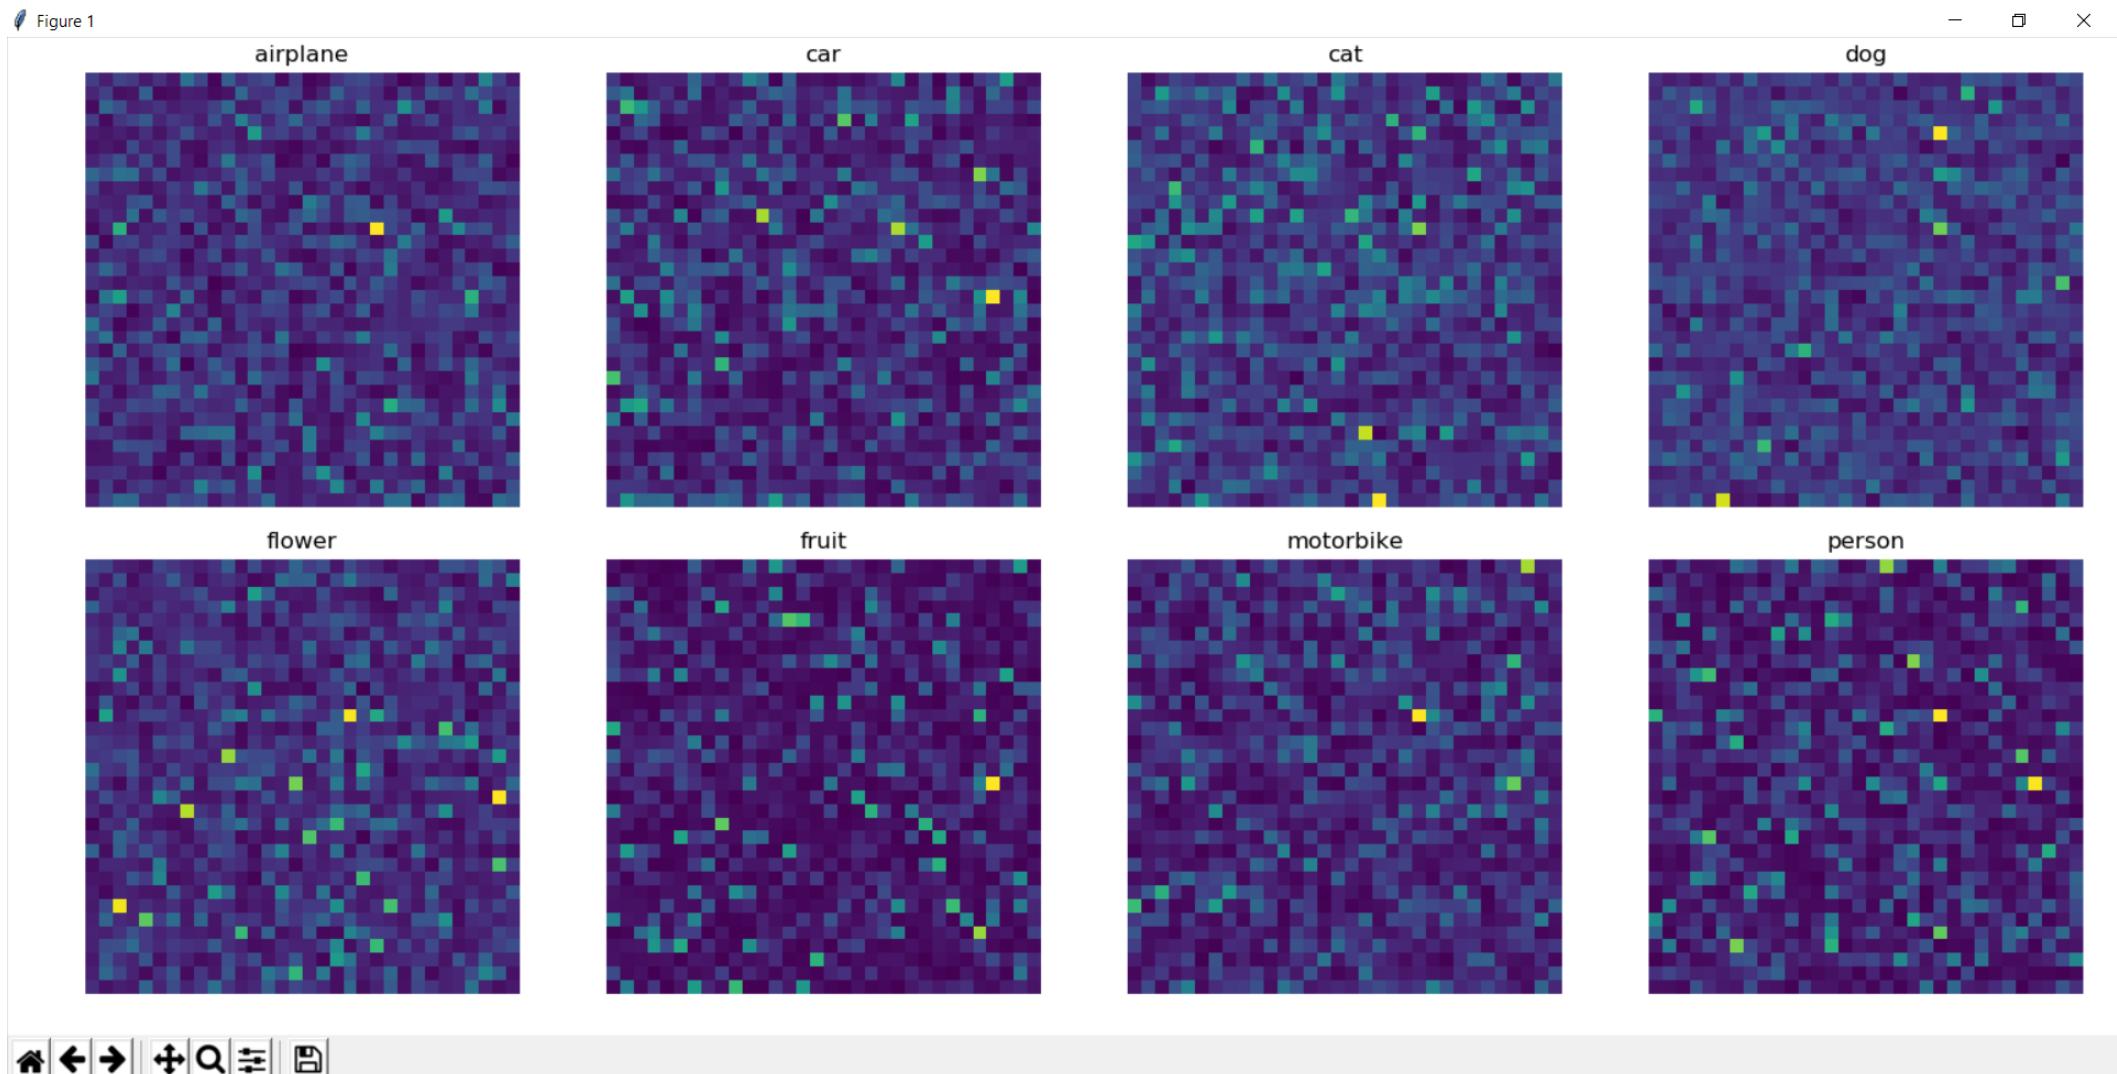


motorbike

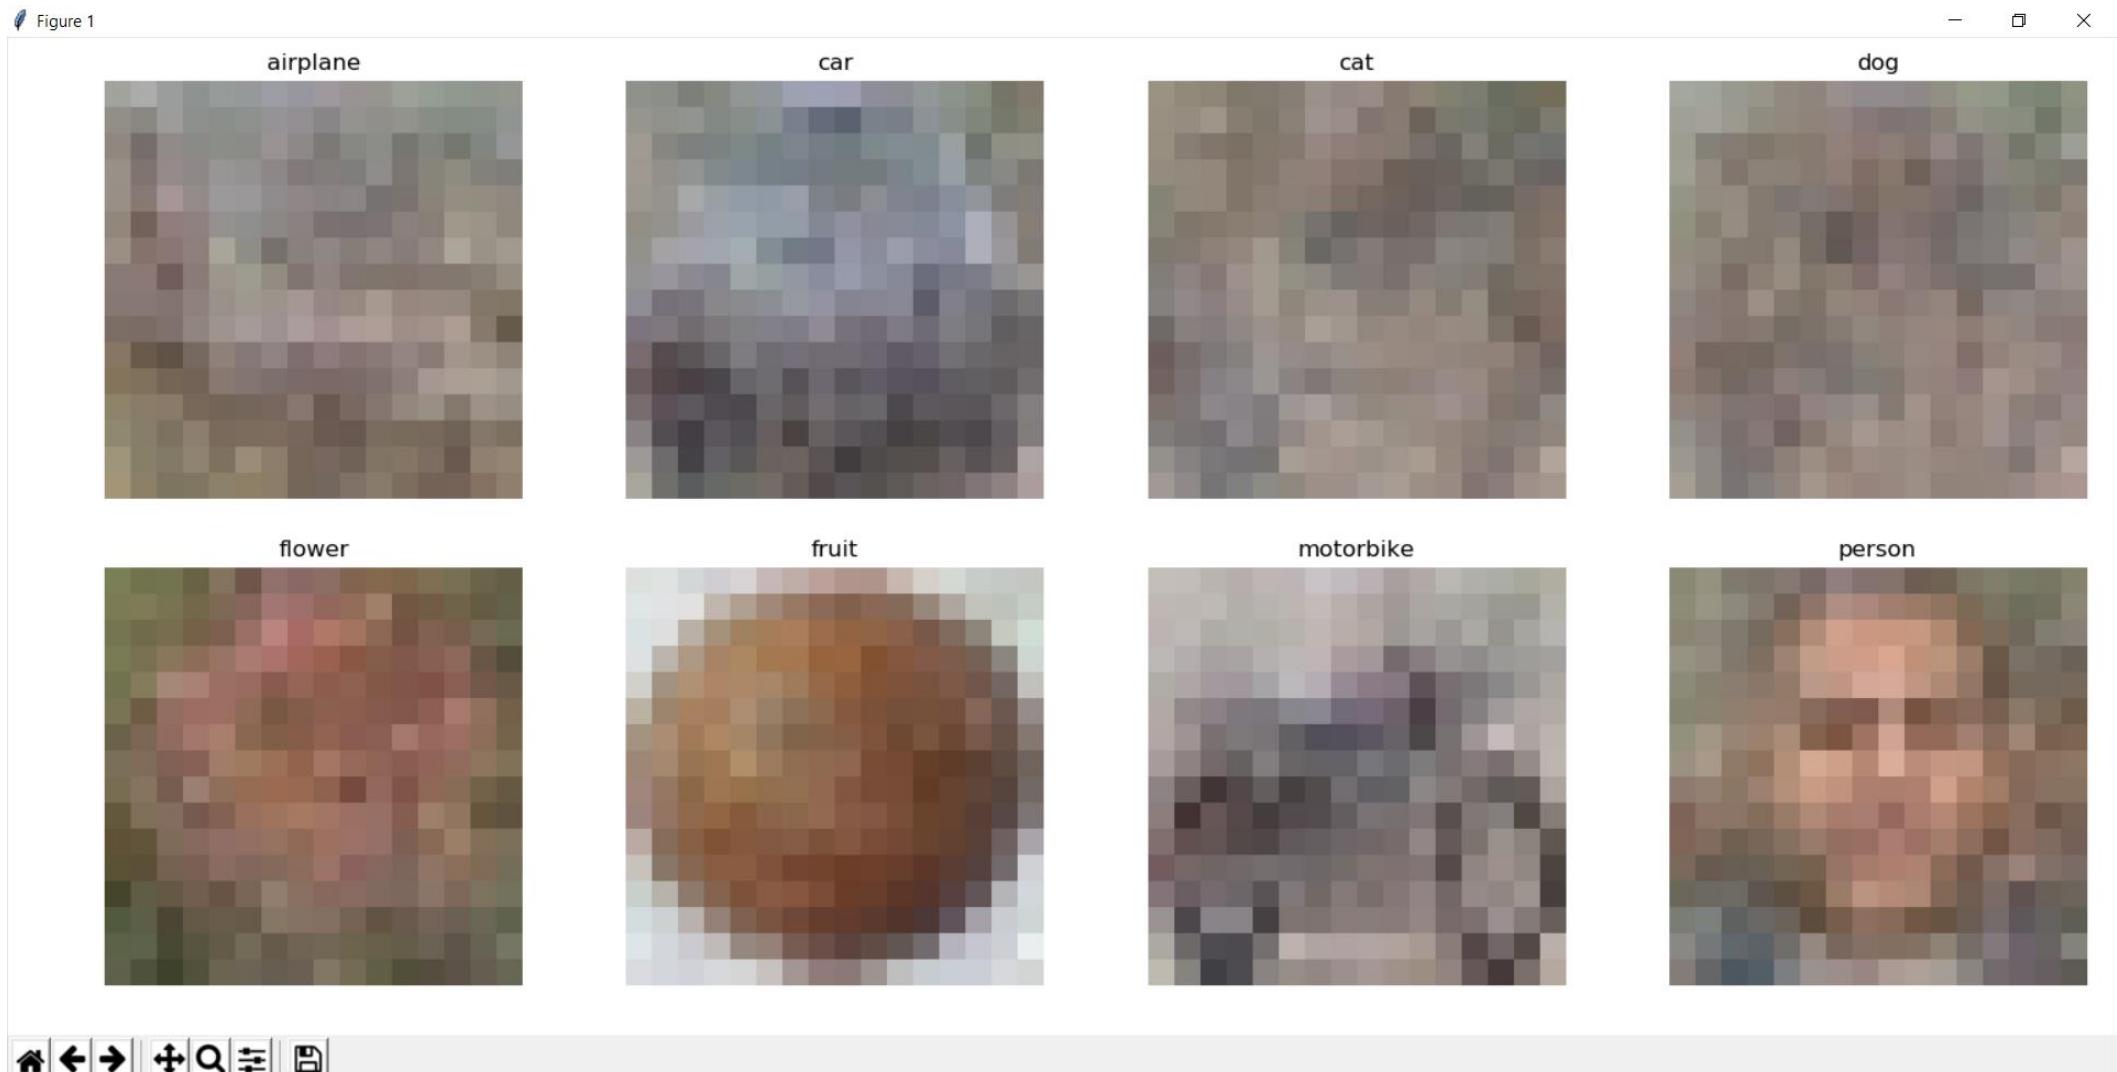


person

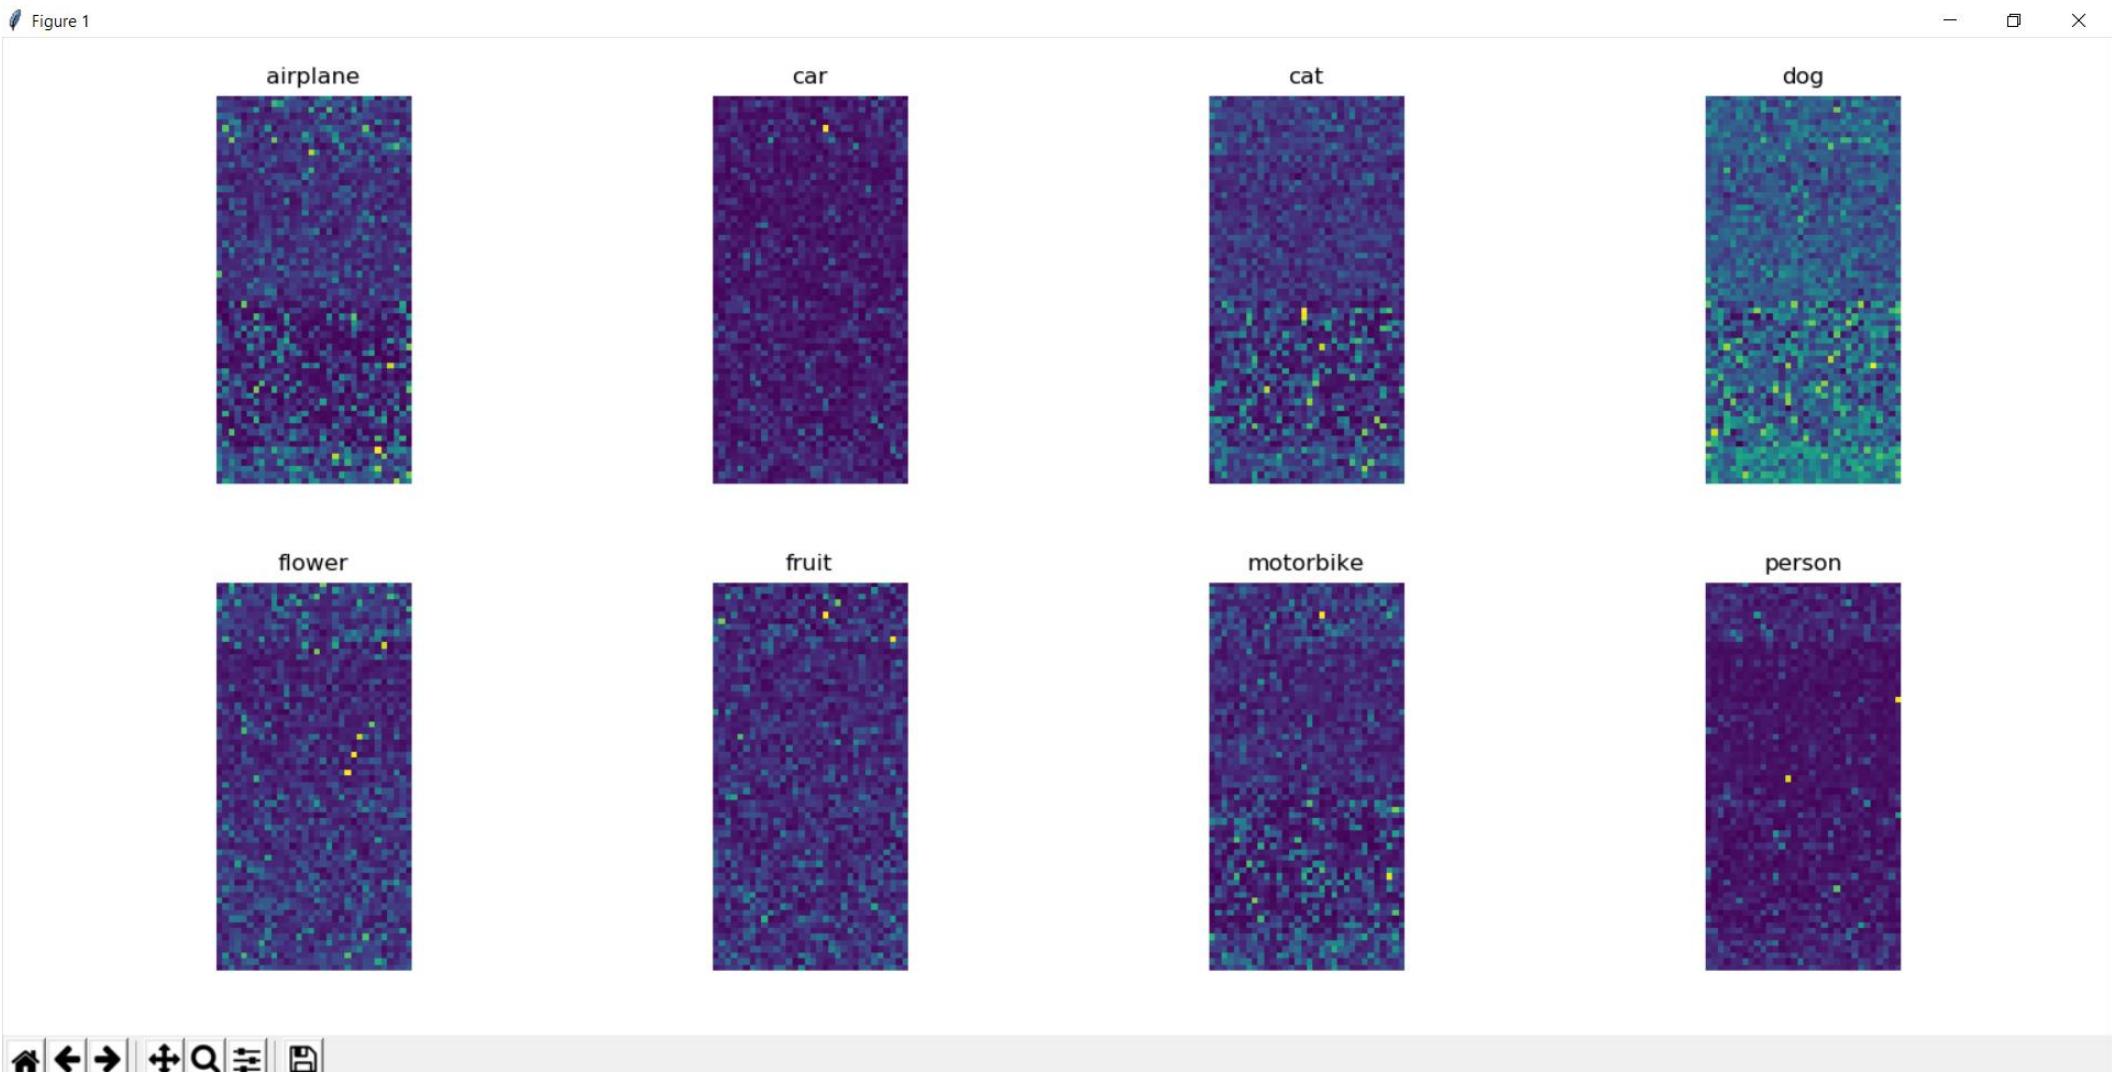
Reconstruct features back into images



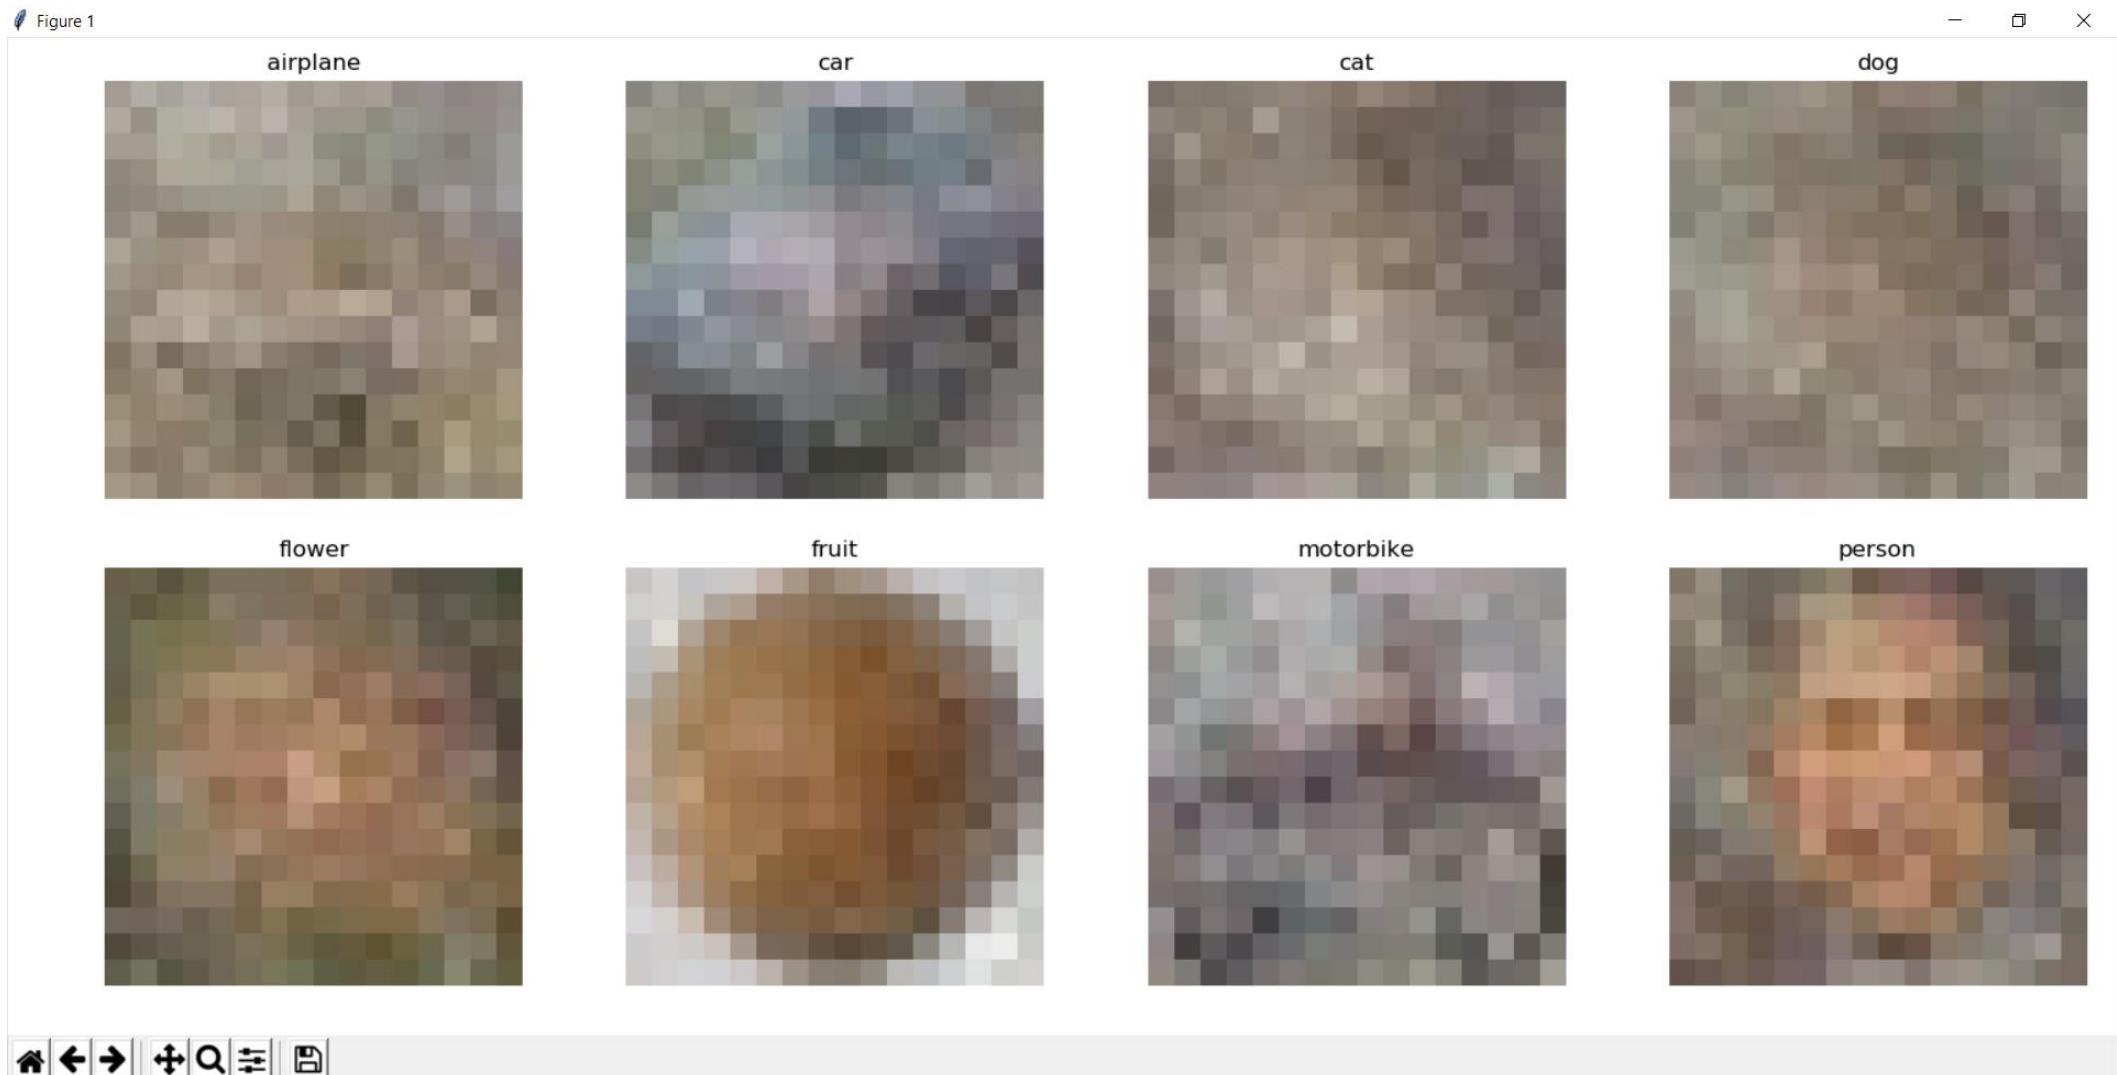
Reconstruct features back into images



Reconstruct features back into images



Reconstruct features back into images



Object Detection

by Cascades

<https://github.com/opencv/opencv/wiki/TensorFlow-Object-Detection-API>

Object detection by Cascades

```
def demo_cascade():

    if USE_CAMERA:
        cap = cv2.VideoCapture(0)
    else:
        cap = []
        frame = cv2.imread(filename_in)

    while (True):
        if USE_CAMERA:
            ret, frame = cap.read()

        gray_rgb = tools_image.desaturate(frame)
        faces = faceCascade.detectMultiScale(frame,scaleFactor=1.1,minNeighbors=5,minSize=(30, 30))

        for (x, y, w, h) in faces:
            cv2.rectangle(gray_rgb, (x, y), (x + w, y + h), (0, 255, 0), 2)

        cv2.imshow('frame', gray_rgb)
        key = cv2.waitKey(1)
        if key & 0xFF == 27:
            break

        if (key & 0xFF == 13) or (key & 0xFF == 32):
            cv2.imwrite(filename_out, gray_rgb)

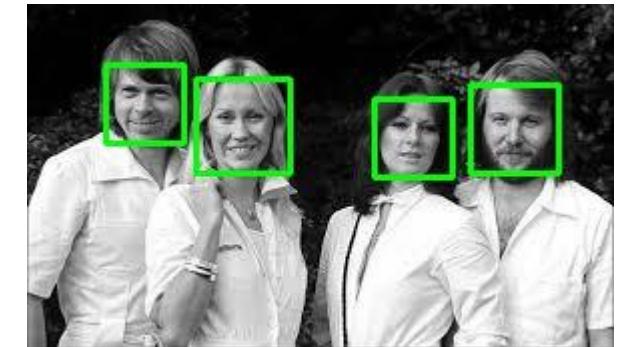
    if USE_CAMERA:
        cap.release()

    cv2.destroyAllWindows()
    if not USE_CAMERA:
        cv2.imwrite(filename_out, gray_rgb)

return
```

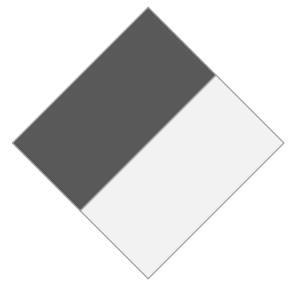
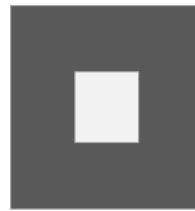
Object detection by Cascades

haarcascade_eye.xml
haarcascade_eye_tree_eyeglasses.xml
haarcascade_frontalcatface.xml
haarcascade_frontalcatface_extended.xml
haarcascade_frontalface_alt.xml
haarcascade_frontalface_alt2.xml
haarcascade_frontalface_alt_tree.xml
haarcascade_frontalface_default.xml
haarcascade_fullbody.xml
haarcascade_lefteye_2splits.xml
haarcascade_licence_plate_rus_16stages.xml
haarcascade_lowerbody.xml
haarcascade_profileface.xml
haarcascade_righteye_2splits.xml
haarcascade_russian_plate_number.xml
haarcascade_smile.xml
haarcascade_upperbody.xml



Object detection: Haar cascades

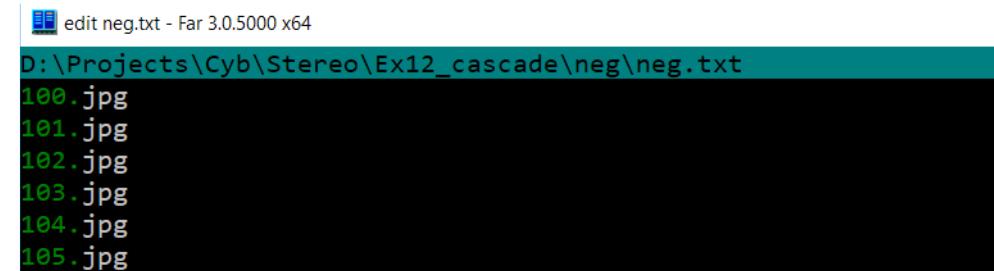
1. Create samples
2. Verify samples
3. Train the cascade



Object detection: create samples from single positive

```
opencv_createsamples -vec pos.vec  
                      -img pos_single\object.jpg  
                      -bg neg\neg.txt  
                      -num 100  
                      -bgcolor 0  
                      -bgthresh 0  
                      -w 28 -h 28  
                      -maxxangle 1.1  
                      -maxyangle 1.1  
                      -maxzangle 0.5  
                      -maxidev 40
```

result vec file
input positive sample file
list of negative samples



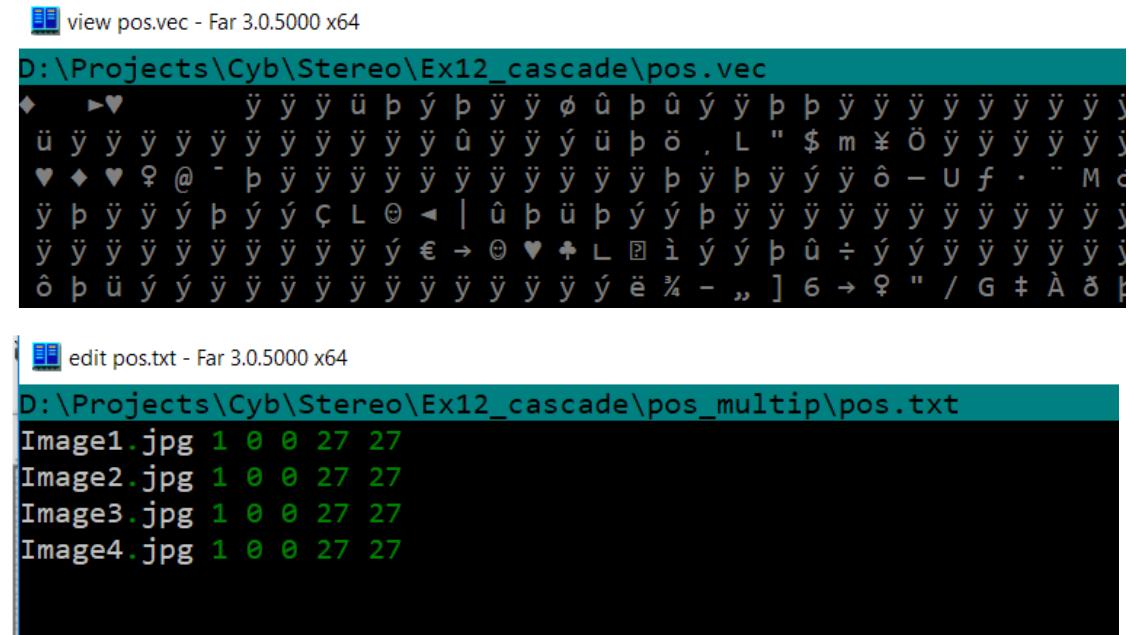
The screenshot shows a terminal window titled "edit neg.txt - Far 3.0.5000 x64". The current directory is "D:\Projects\Cyb\Stereo\Ex12_cascade\neg\". The file "neg.txt" contains the following list of image filenames:

```
D:\Projects\Cyb\Stereo\Ex12_cascade\neg\neg.txt  
100.jpg  
101.jpg  
102.jpg  
103.jpg  
104.jpg  
105.jpg
```

Object detection: create samples from many positive

```
opencv_createsamples -vec pos.vec  
                      -info pos_multip\pos.txt  
                      -num 4  
                      -w 28 -h 28
```

result file
list of pos samples



The screenshot shows two terminal windows side-by-side.

The left terminal window is titled "view pos.vec - Far 3.0.5000 x64" and displays the command:

```
D:\Projects\Cyb\Stereo\Ex12_cascade\pos.vec
```

The right terminal window is titled "edit pos.txt - Far 3.0.5000 x64" and displays the command:

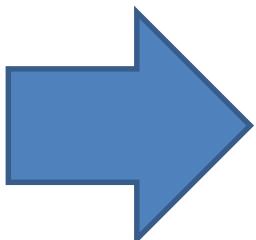
```
D:\Projects\Cyb\Stereo\Ex12_cascade\pos_multip\pos.txt
```

Both windows show the following output:

```
Image1.jpg 1 0 0 27 27  
Image2.jpg 1 0 0 27 27  
Image3.jpg 1 0 0 27 27  
Image4.jpg 1 0 0 27 27
```

Object detection: verify vec file

```
opencv_createsamples -vec pos.vec  
-show
```



Object detection: training

```
opencv_traincascade -data output_folder
                    -vec ../positive/positives.vex
                    -bg ../negative/bg_train.txt
                    -numPos 800
                    -numNeg 400
                    -numStages 10
                    -w 20
                    -h 20
```

edit cascade.xml - Far 3.0.5000 x64

```
C:\Users\dryabokon\source\Digits\ML\data\ex12_single\output_s
<?xml version="1.0"?>
<opencv_storage>
<cascade>
    <stageType>BOOST</stageType>
    <featureType>HAAR</featureType>
    <height>40</height>
    <width>40</width>
    <stageParams>
        <boostType>GAB</boostType>
        <minHitRate>9.9500000476837158e-01</minHitRate>
        <maxFalseAlarm>5.000000000000000e-01</maxFalseAlarm>
        <weightTrimRate>9.499999999999996e-01</weightTrimRate>
        <maxDepth>1</maxDepth>
        <maxWeakCount>100</maxWeakCount></stageParams>
    <featureParams>
        <maxCatCount>0</maxCatCount>
        <featSize>1</featSize>
        <mode>BASIC</mode></featureParams>
```

existing folder
made by opencv_createsamples
list of negative samples

edit bg_train.txt - Far 3.0.5000 x64

```
D:\Projects\Cyb\Stereo\Ex11_cascade\negative\bg_train.txt
.../negative/bg/62.jpg
.../negative/bg/22.jpg
.../negative/bg/372.jpg
.../negative/bg/712.jpg
.../negative/bg/156.jpg
.../negative/bg/416.jpg
.../negative/bg/275.jpg
.../negative/bg/811.jpg
.../negative/bg/152.jpg
.../negative/bg/451.jpg
.../negative/bg/606.jpg
.../negative/bg/136.jpg
.../negative/bg/286.jpg
.../negative/bg/116.jpg
.../negative/bg/398.jpg
.../negative/bg/707.jpg
```

view positives.vex - Far 3.0.5000 x64

```
D:\Projects\Cyb\Stereo\Ex11_cascade\positive\positives.vex
00000000: 33 03 00 00 90 01 00 00 00 00 00 00 98 00 9D 3▼ ⑧ ~
000000010: 00 93 00 7B 00 84 00 93 00 C0 00 B7 00 AF 00 9B " { „ “ Á - >
000000020: 00 91 00 8D 00 79 00 75 00 80 00 9E 00 DA 00 D7 ‘ Ÿ y u € ž Ú ×
000000030: 00 75 00 72 00 9A 00 95 00 8F 00 8B 00 9E 00 98 u r š • Ÿ < ž ~
000000040: 00 A6 00 A2 00 A3 00 99 00 A3 00 91 00 8F 00 90 i ¢ £ ™ € Ÿ Ÿ
000000050: 00 9C 00 AE 00 CB 00 BC 00 66 00 7E 00 60 00 68 ø ® £ % f ~ ` h
000000060: 00 97 00 CF 00 9B 00 80 00 68 00 69 00 66 00 77 – İ > € h i f w
000000070: 00 8B 00 78 00 5E 00 72 00 87 00 8B 00 B0 00 A9 < x ^ r ũ < ° @
000000080: 00 72 00 81 00 62 00 56 00 A8 00 A3 00 95 00 9D r Ÿ b V “ £ • Ÿ
000000090: 00 A3 00 88 00 8C 00 B7 00 B2 00 9F 00 A7 00 A0 £ ^ € . ° Ÿ Ÿ
0000000A0: 00 A2 00 BF 00 D7 00 F2 00 ED 00 72 00 88 00 71 ¢ ū x ò í r < q
0000000B0: 00 CF 00 87 00 87 00 A4 00 C8 00 C9 00 C7 00 C5 i ũ ũ x È É C Á
0000000C0: 00 C4 00 C3 00 C1 00 C6 00 C5 00 C6 00 C6 Á Á Á È Á È Á È
0000000D0: 00 DE 00 8C 00 9B 00 84 00 CC 00 A9 00 95 00 82 p € > „ ī Ø • ,
0000000E0: 00 C2 00 BB 00 C6 00 B1 00 A9 00 97 00 A1 00 A6 Á » £ ± Ø - i ū
0000000F0: 00 A9 00 AE 00 C0 00 CE 00 DD 00 99 00 83 00 93 Ø ® Á ī Ÿ ™ f “
```

Train the Cascades

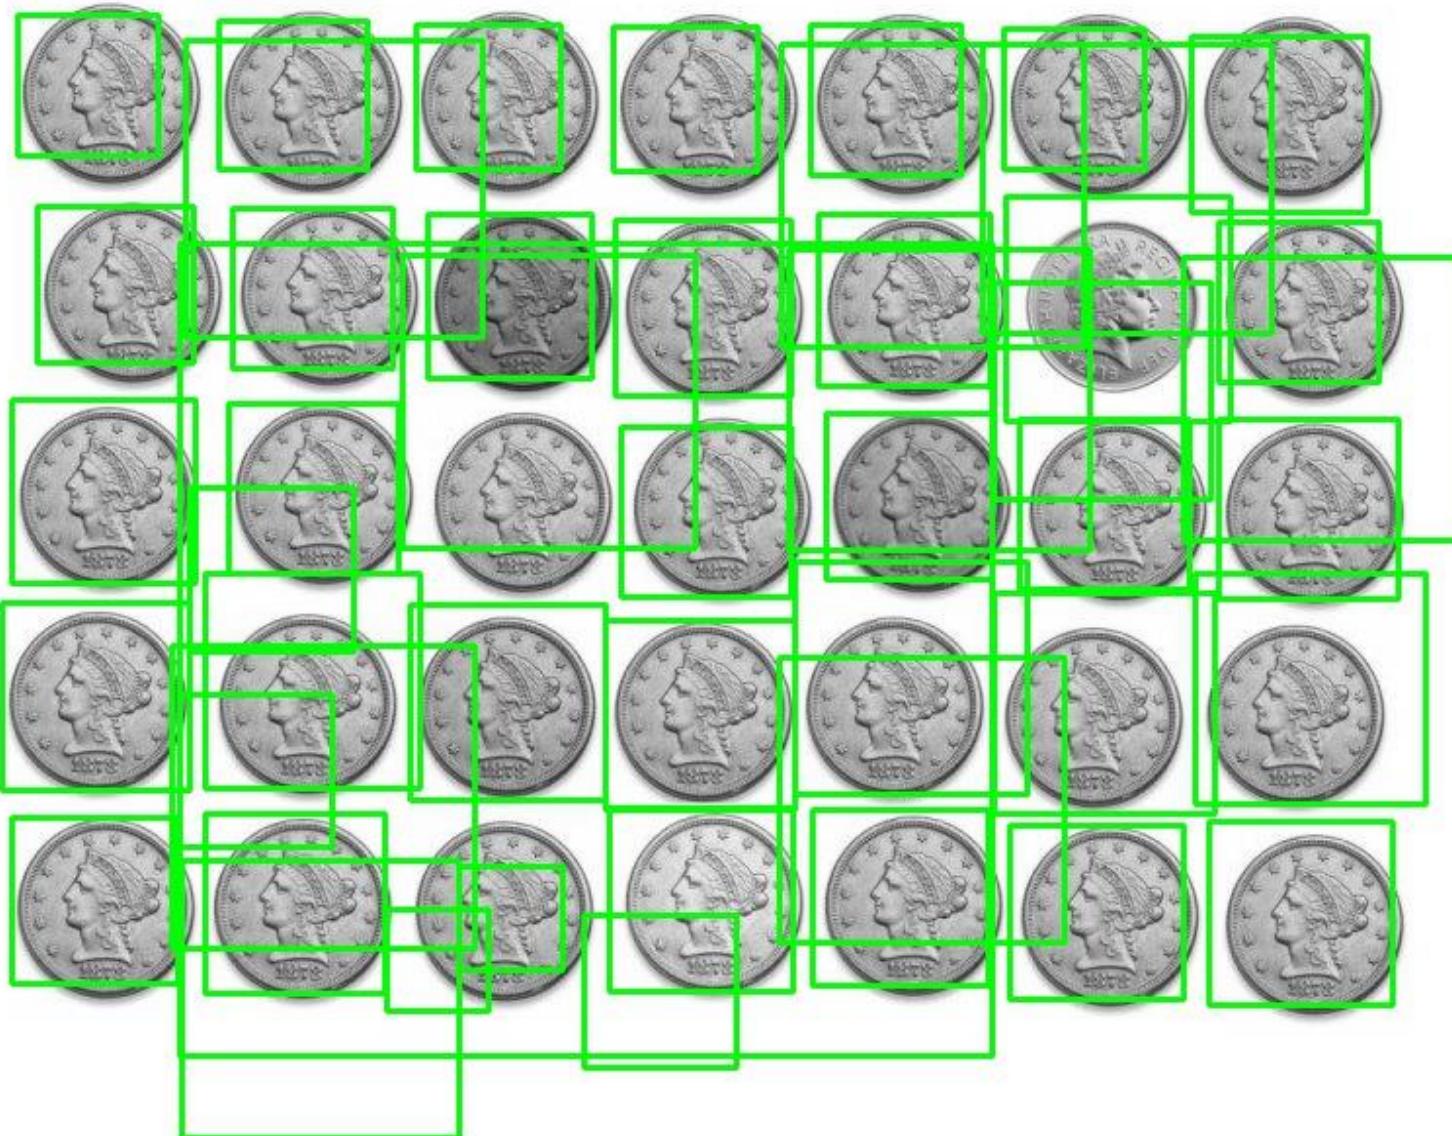
```
import cv2
import os
import tools_image
# -----
path = 'data/ex12_single/'
filename_in = path + 'pos_single/object_to_detect.jpg'
filename_out = 'data/output/detect_result.jpg'
object_detector= cv2.CascadeClassifier(path + 'output_single/cascade.xml')

# -----
def train_cascade():
    os.chdir(path)
    os.system('1-create_vec_from_single.bat')#os.system('2-verify_vec_single.bat')
    os.system('3-train_from_single.bat')
    return
# -----


if __name__ == '__main__':
    #train_cascade()

    image = tools_image.desaturate(cv2.imread(filename_in))
    objects, rejectLevels, levelWeights =
        object_detector.detectMultiScale3(image, scaleFactor=1.05, minSize=(20, 20), outputRejectLevels=True)
    for (x, y, w, h) in objects:cv2.rectangle(image, (x, y), (x + w, y + h), (0, 255, 0), 2)
    cv2.imwrite(filename_out,image)
```

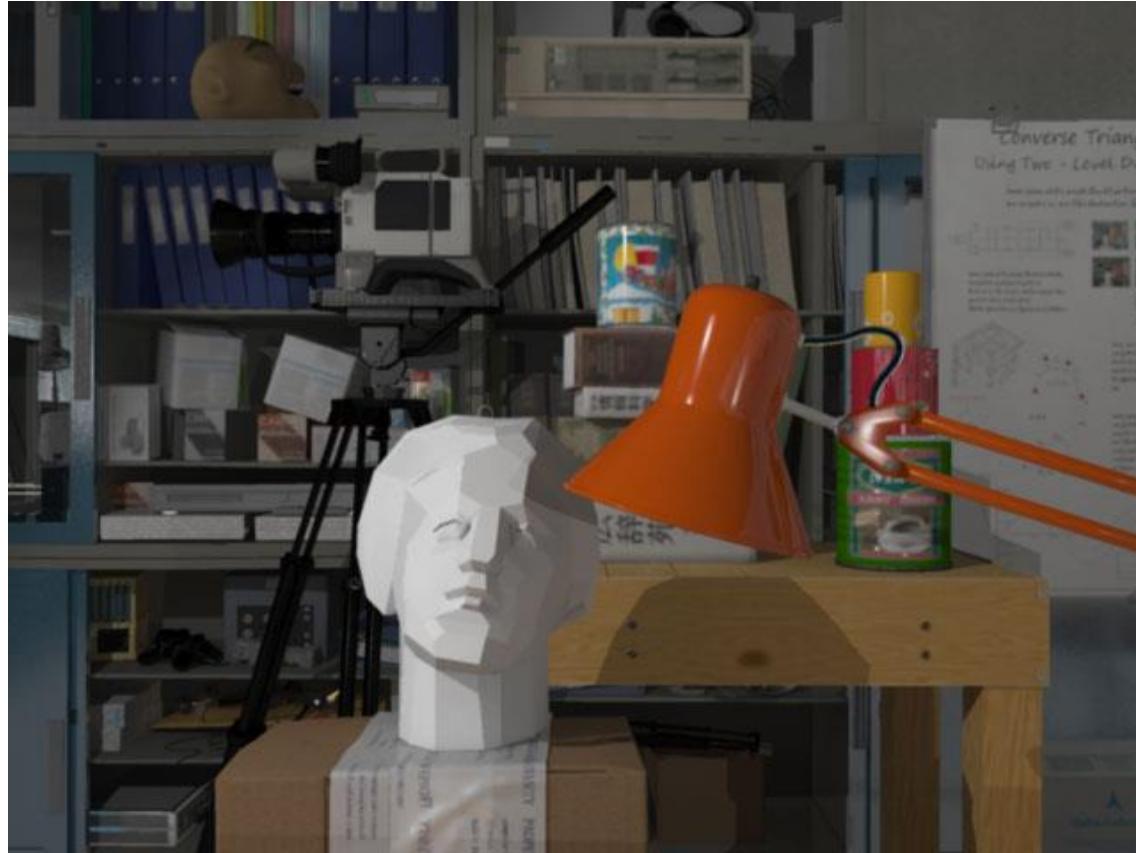
Object detection: example of training



Stereo Vision



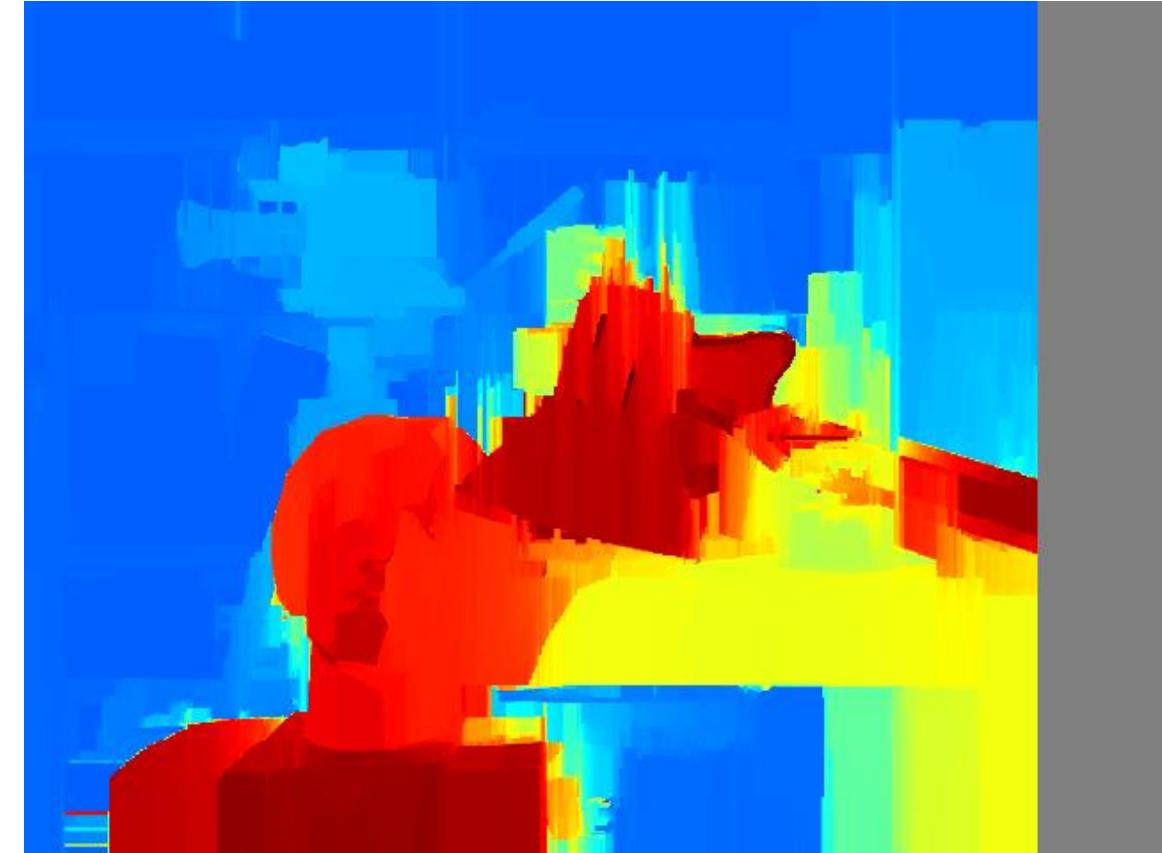
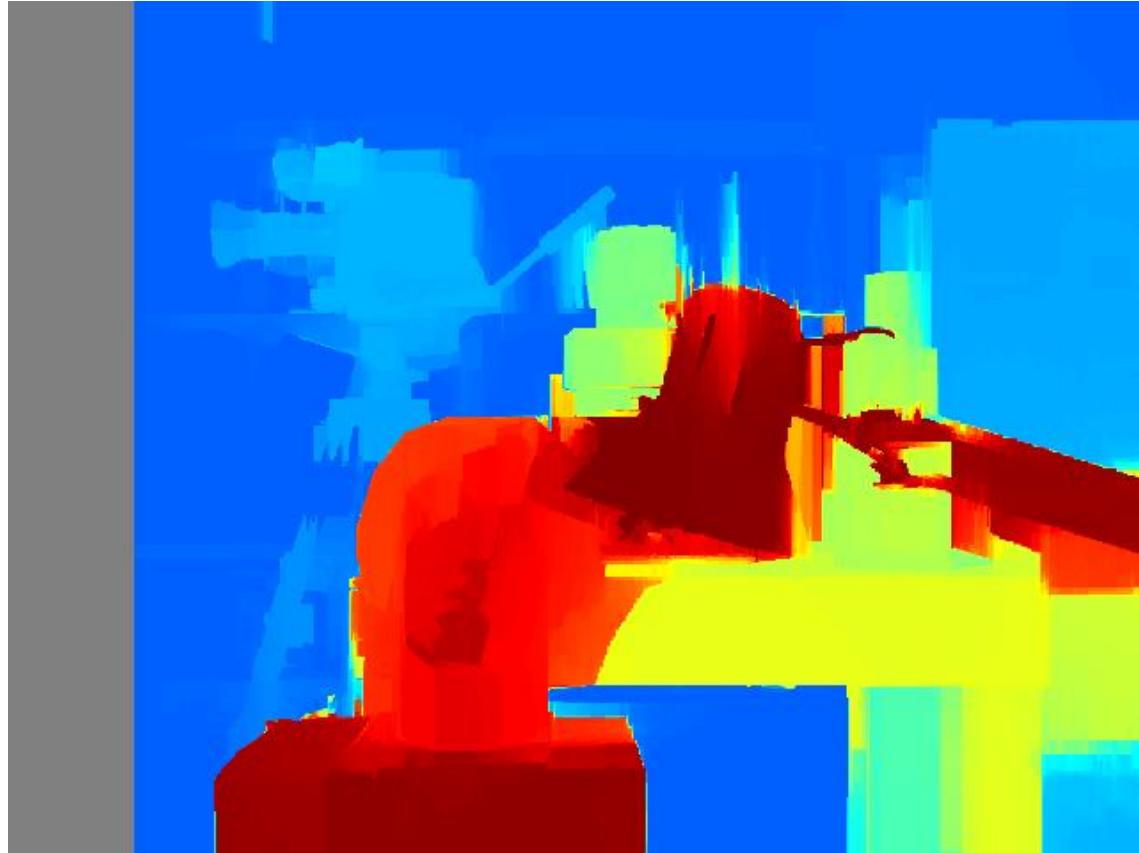
Stereo vision



Stereo vision



Stereo vision



Stereo vision

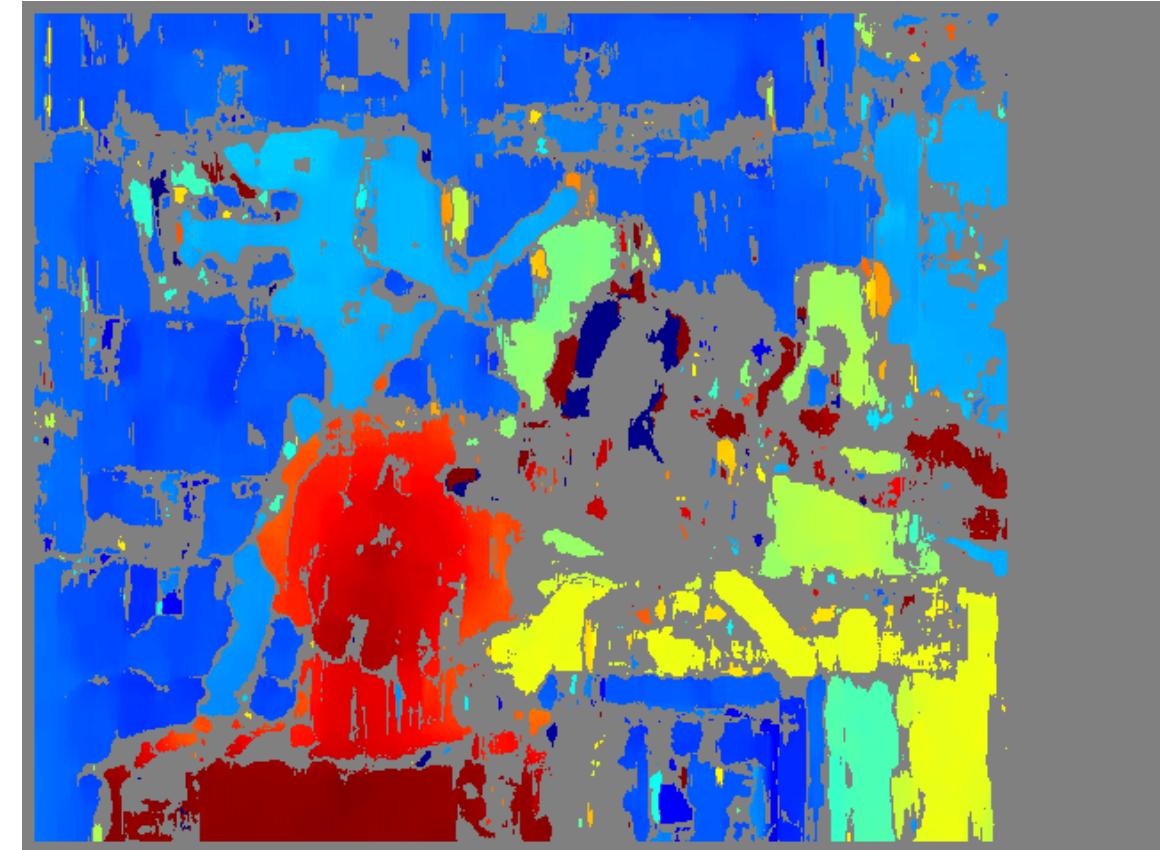
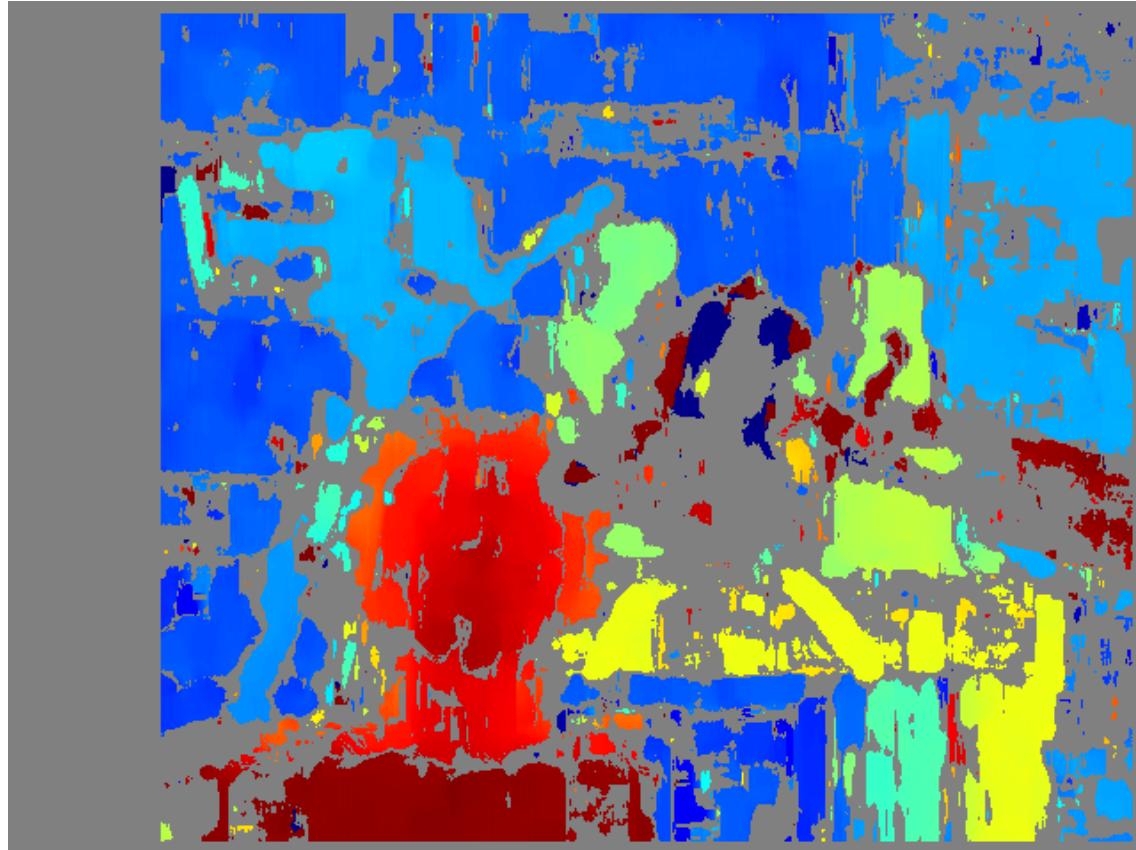
```
def get_disparity_v_01(imgL, imgR, disp_v1, disp_v2, disp_h1, disp_h2):
    window_size = 7
    left_matcher = cv2.StereoSGBM_create(
        minDisparity=disp_h1,
        numDisparities=int(0 + (disp_h2 - disp_h1) / 16) * 16,
        blockSize=5,
        P1=8 * 3 * window_size ** 2,
        P2=32 * 3 * window_size ** 2,
        disp12MaxDiff=disp_h2,
        uniquenessRatio=15,
        speckleWindowSize=0,
        speckleRange=2,
        preFilterCap=63,
        mode=cv2.STEREO_SGBM_MODE_SGBM_3WAY
    )

    right_matcher = cv2.ximgproc.createRightMatcher(left_matcher)
    dispr = right_matcher.compute(imgR, imgL)
    displ = left_matcher.compute(imgL, imgR)

    wls_filter = cv2.ximgproc.createDisparityWLSFilter(matcher_left=left_matcher)
    wls_filter.setLambda(80000)
    wls_filter.setSigmaColor(1.2)

    filteredImg_L = wls_filter.filter(displ, imgL, None, dispr)
    wls_filter = cv2.ximgproc.createDisparityWLSFilter(matcher_left=right_matcher)
    filteredImg_R = wls_filter.filter(dispr, imgR, None, displ)
```

Stereo vision

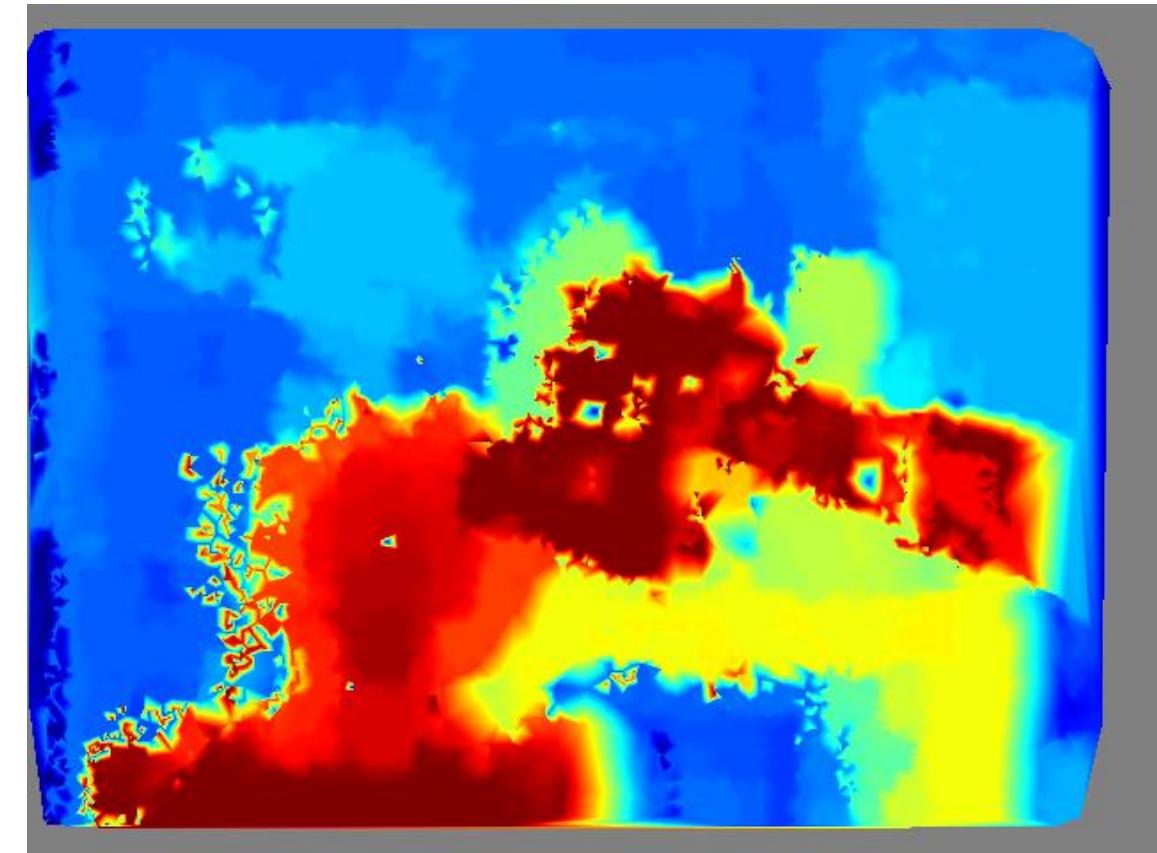
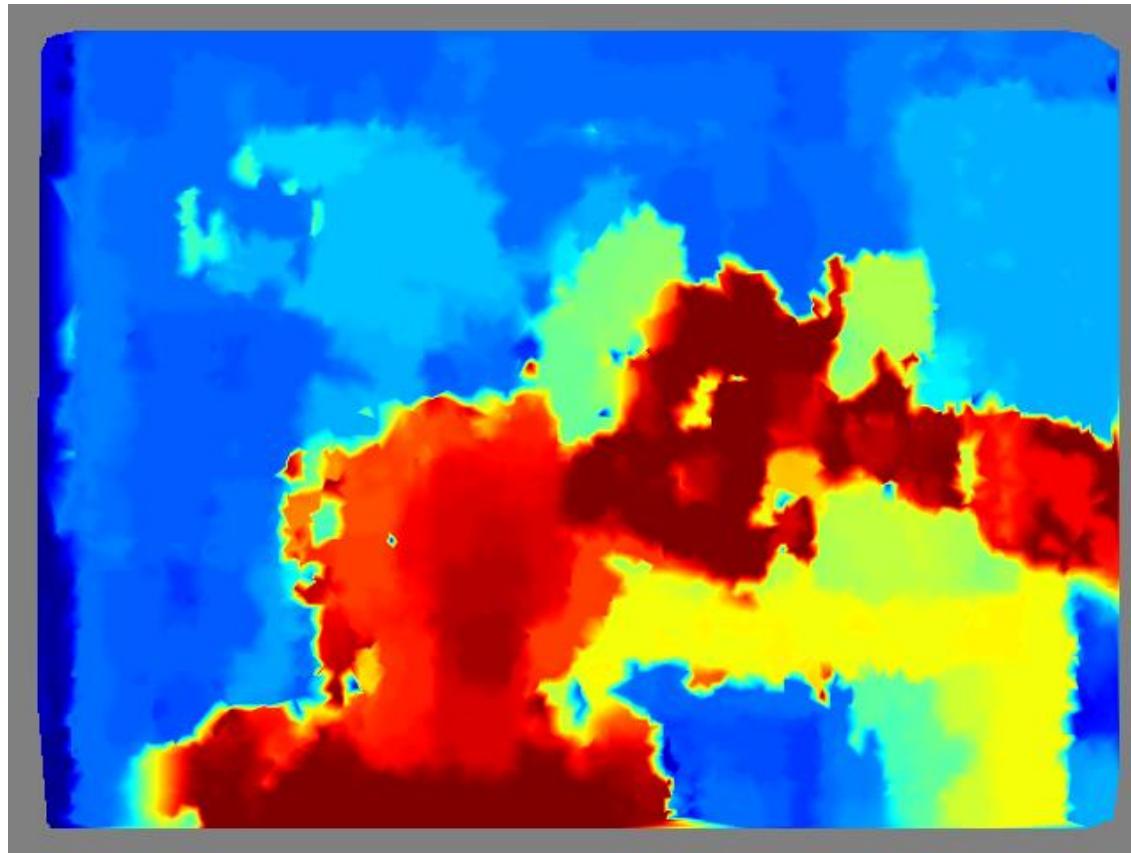


Stereo vision

```
def get_disparity_v_02(imgL, imgR, disp_v1, disp_v2, disp_h1, disp_h2):
    max = numpy.maximum(math.fabs(disp_h1), math.fabs(disp_h2))
    levels = int(1 + (max) / 16) * 16
    stereo = cv2.StereoBM_create(numDisparities=levels, blockSize=15)
    displ = stereo.compute(imgL, imgR)

    dispR = numpy.flip(stereo.compute(numpy.flip(imgR, axis=1), numpy.flip(imgL, axis=1)), axis=1)
    return -displ / 16, -dispR / 16
```

Stereo vision



Stereo vision

```
def get_best_matches(image1,image2,disp_v1, disp_v2, disp_h1, disp_h2, window_size=15,step=10):

    N = int(image1.shape[0] * image1.shape[1] / step / step)
    rand_points = numpy.random.rand(N,2)
    rand_points[:, 0] = window_size + (rand_points[:, 0]*(image1.shape[0]-2*window_size))
    rand_points[:, 1] = window_size + (rand_points[:, 1]*(image1.shape[1]-2*window_size))

    coord1,coord2,quality = [],[],[]

    for each in rand_points:
        row,col = int(each[0]),int(each[1])

        template = tools_image.crop_image(image1,row-window_size,col-window_size,row+window_size,col+window_size)
        top,left,bottom,right = row - window_size + disp_v1, col - window_size + disp_h1, row + window_size + disp_v2, col + window_size + disp_h2
        field = tools_image.crop_image(image2,top,left,bottom,right)

        q = cv2.matchTemplate(field, template, method=cv2.TM_CCOEFF_NORMED)
        q = q[1:, 1:]
        q = (q + 1) * 128

        idx = numpy.argmax(q.flatten())
        q_best = numpy.max(q.flatten())
        qq=q[int(idx/q.shape[1]), idx % q.shape[1]]

        if q_best>0:
            dr = int(idx/q.shape[1])+disp_v1
            dc = idx % q.shape[1]+disp_h1

            if col + dc>=0 and col + dc<image1.shape[1] and row + dr>=0 and row + dr<image1.shape[0]:
                coord1.append([col ,row ])
                coord2.append([col+ dc, row+dr])
                quality.append(q_best)

    return numpy.array(coord1), numpy.array(coord2),numpy.array(quality)
```

Part 2: Projective Geometry



- KeyPoints detection
- Matching the KeyPoints
- Homography
- Blending
- Camera calibration
- Stereopair rectification
- Pose estimation

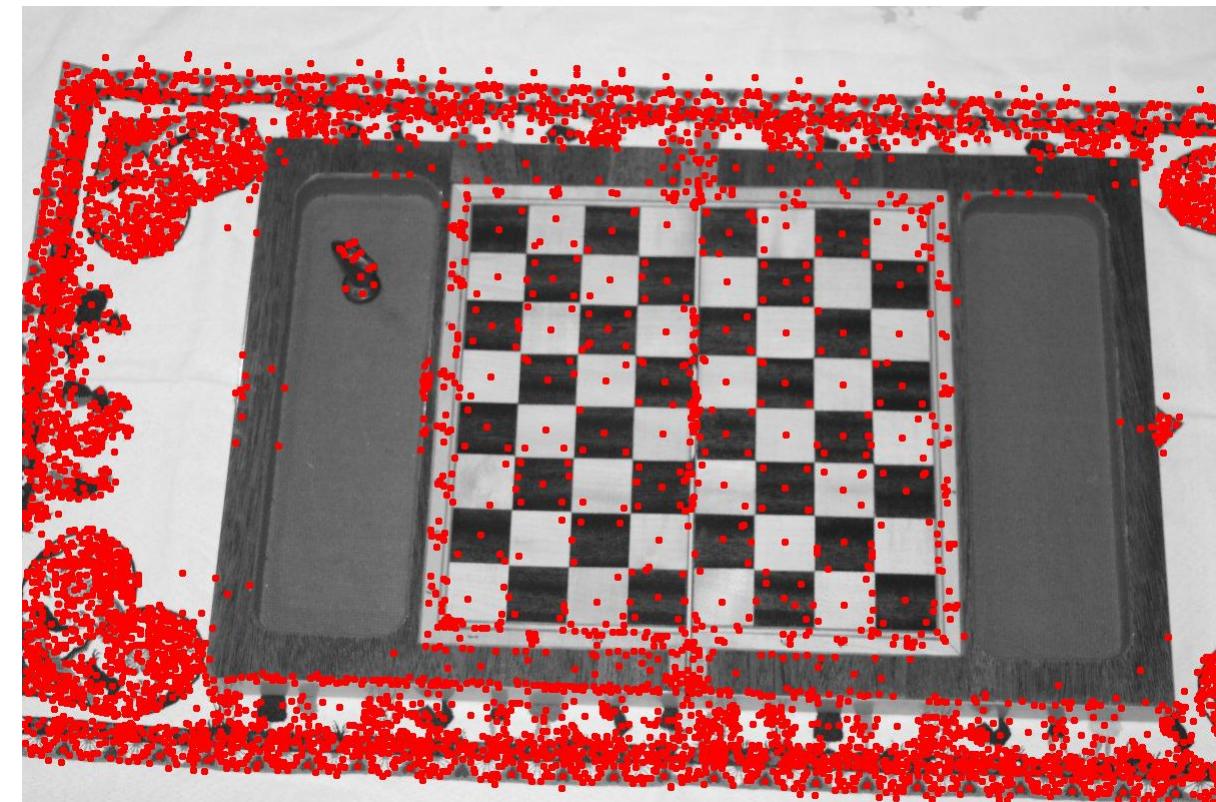
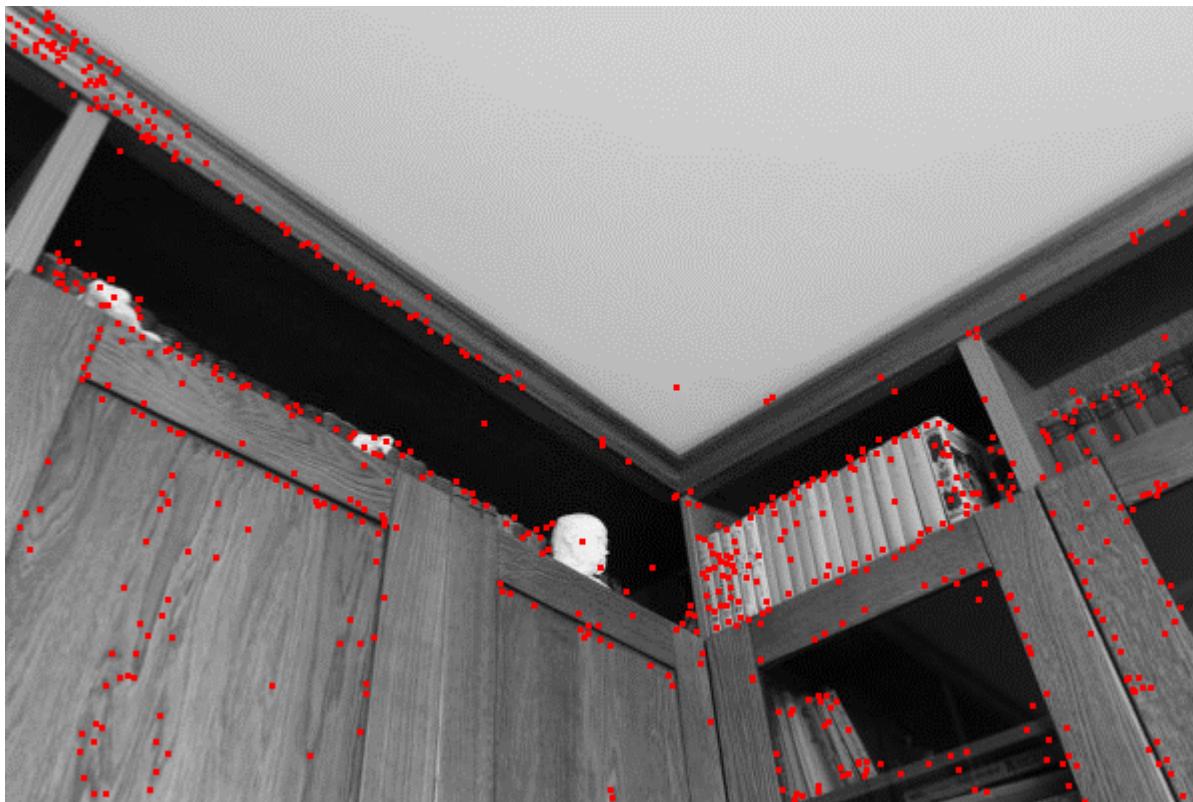
KeyPoints detection



https://github.com/dryabokon/geometry/blob/master/ex01_keypoints_detect.py

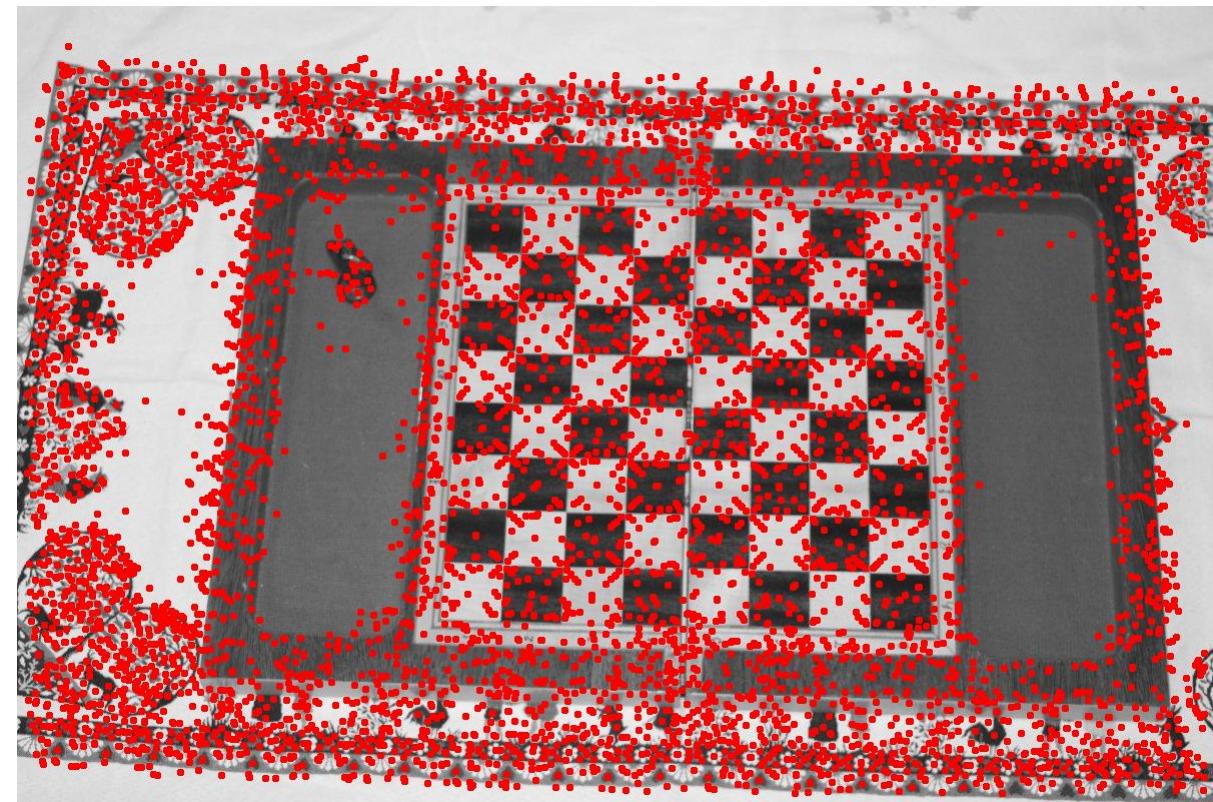
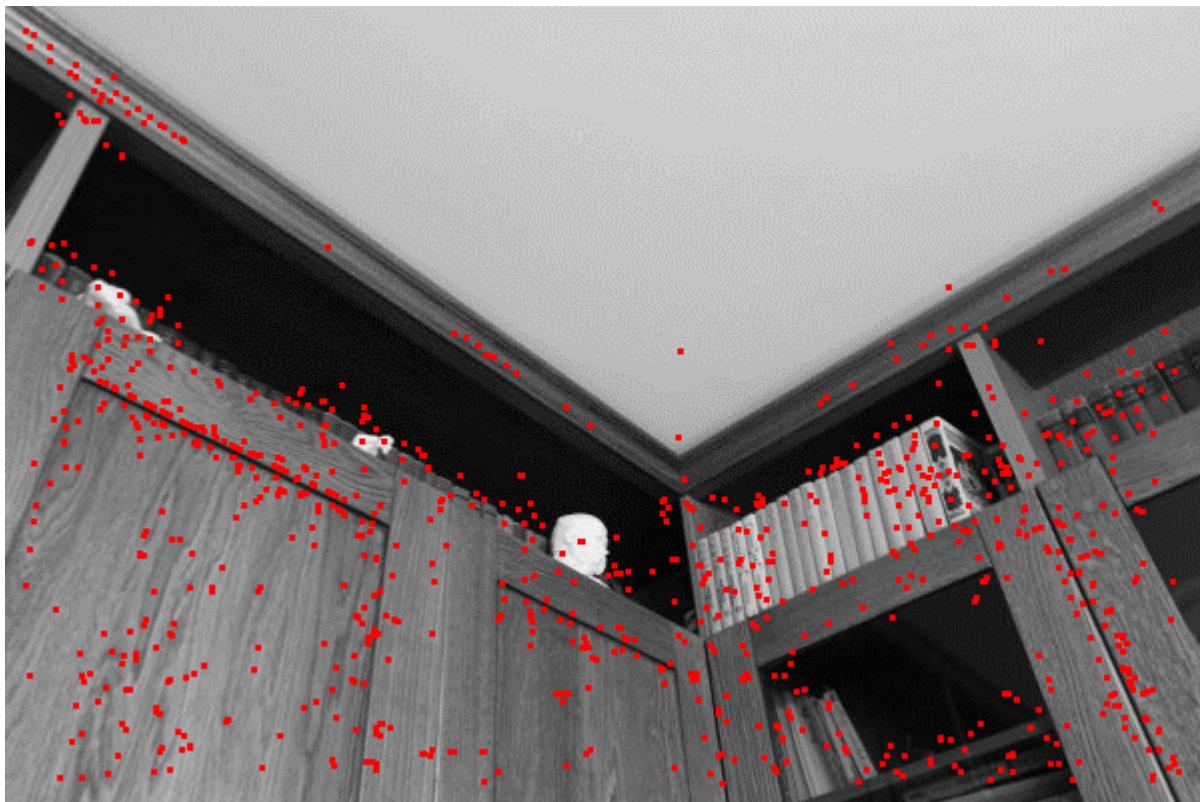
KeyPoints detection

SIFT: Scale-Invariant Feature Transform



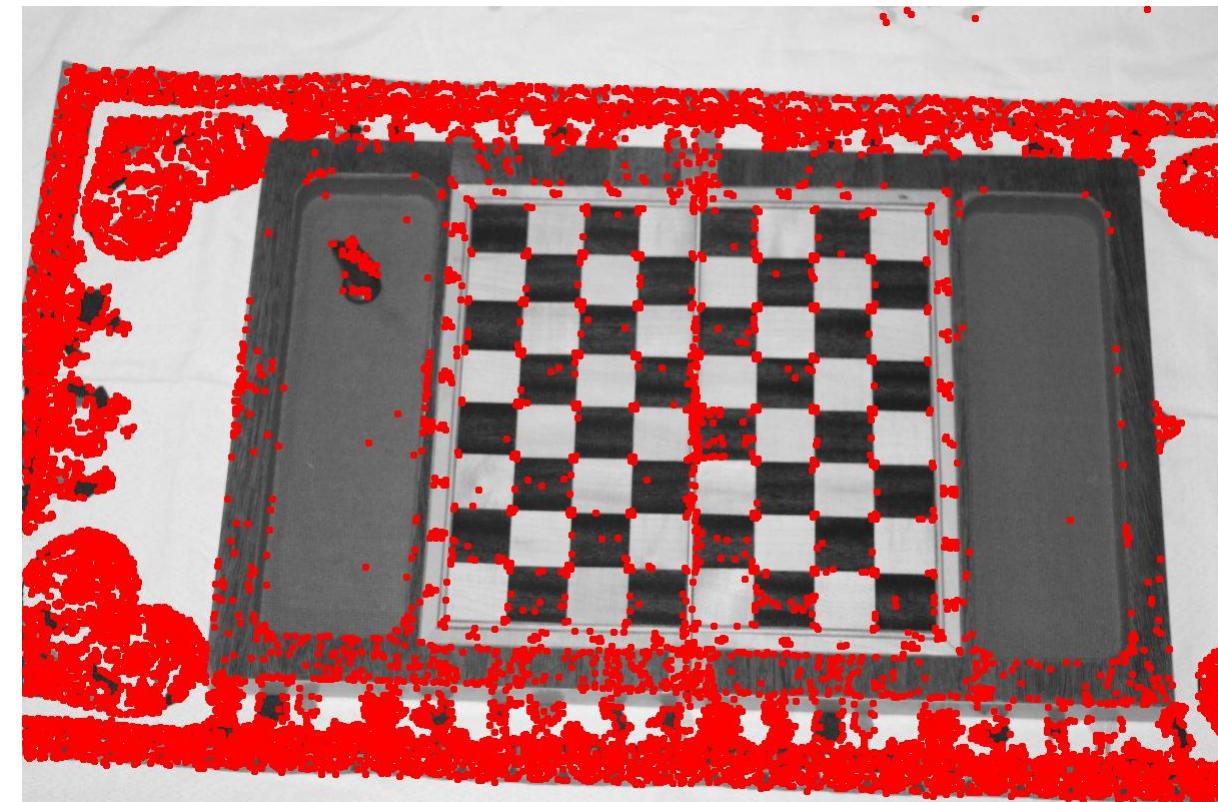
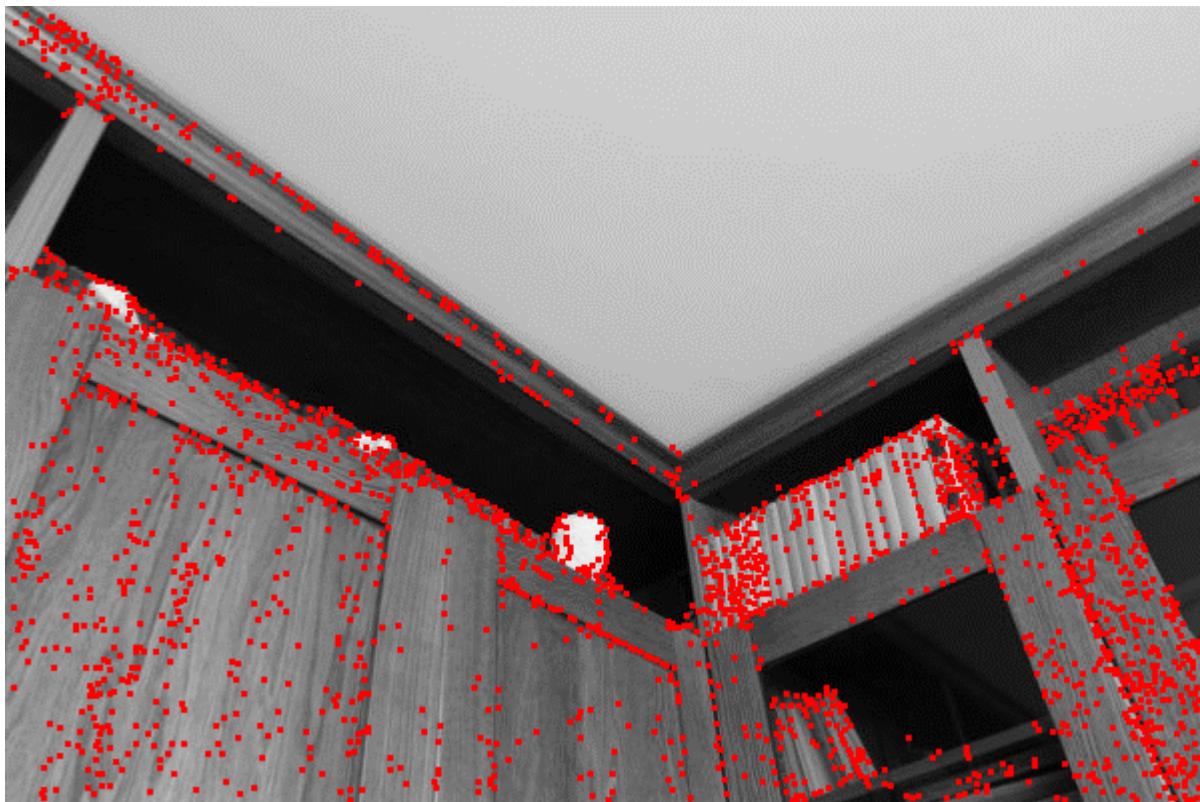
KeyPoints detection

SURF: Speeded-Up Robust Features



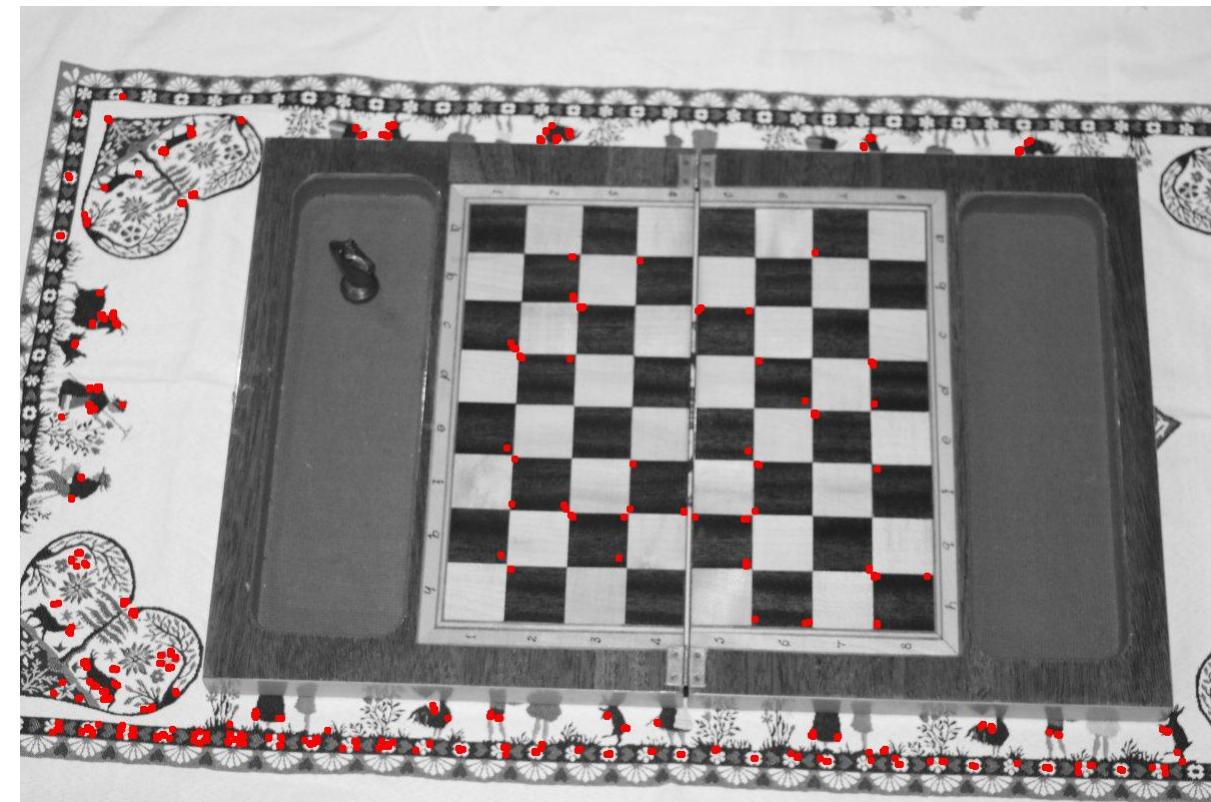
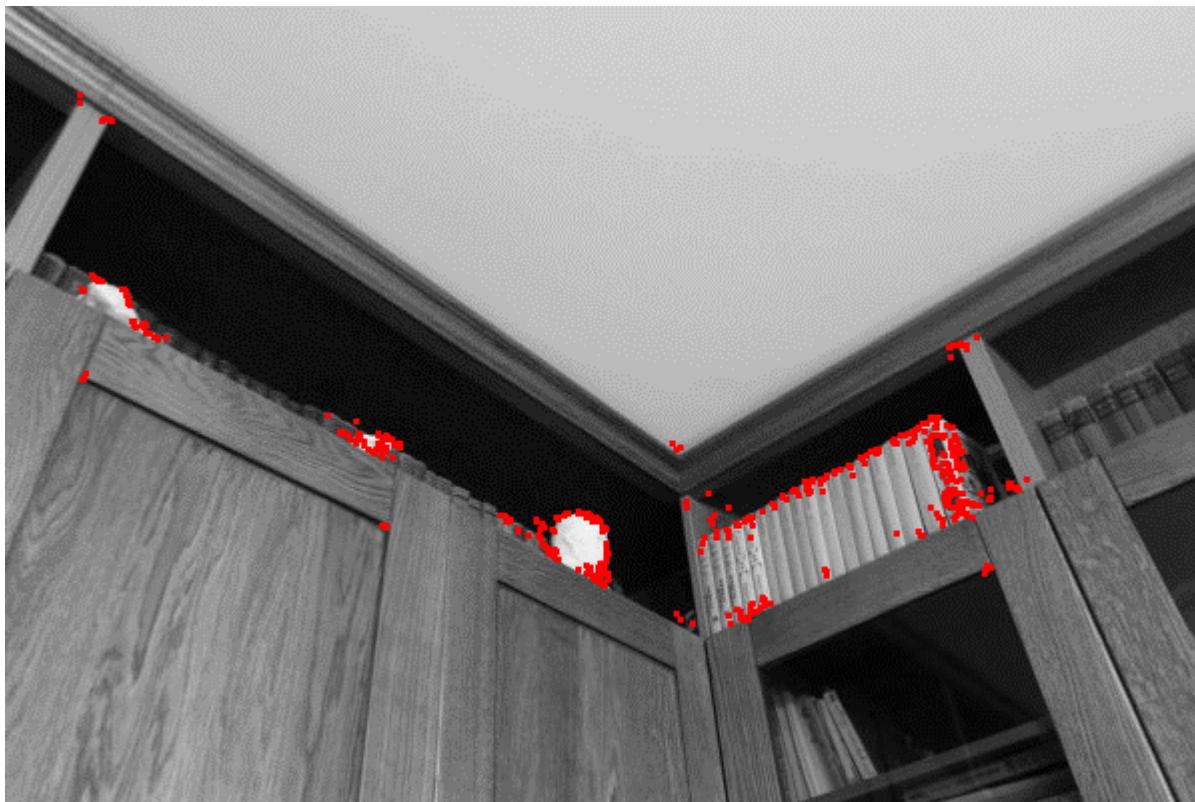
KeyPoints detection

FAST: Features from Accelerated Segment Test



KeyPoints detection

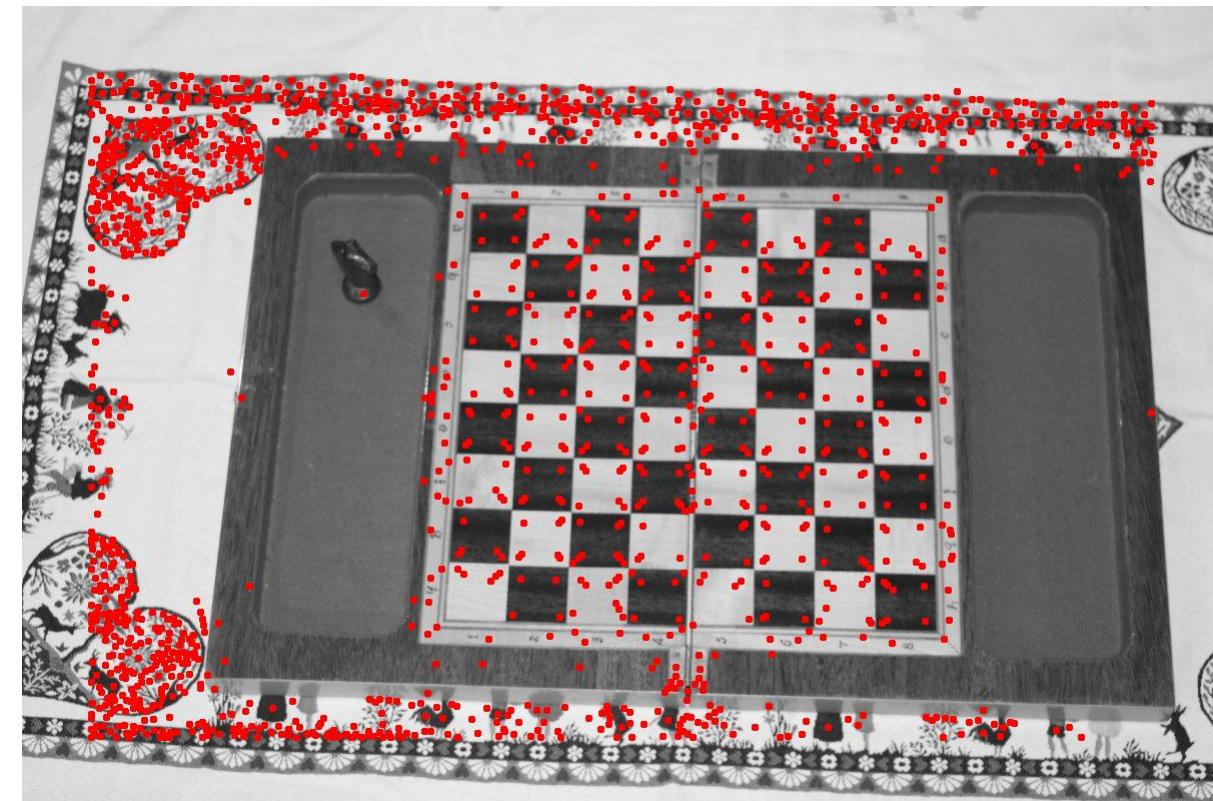
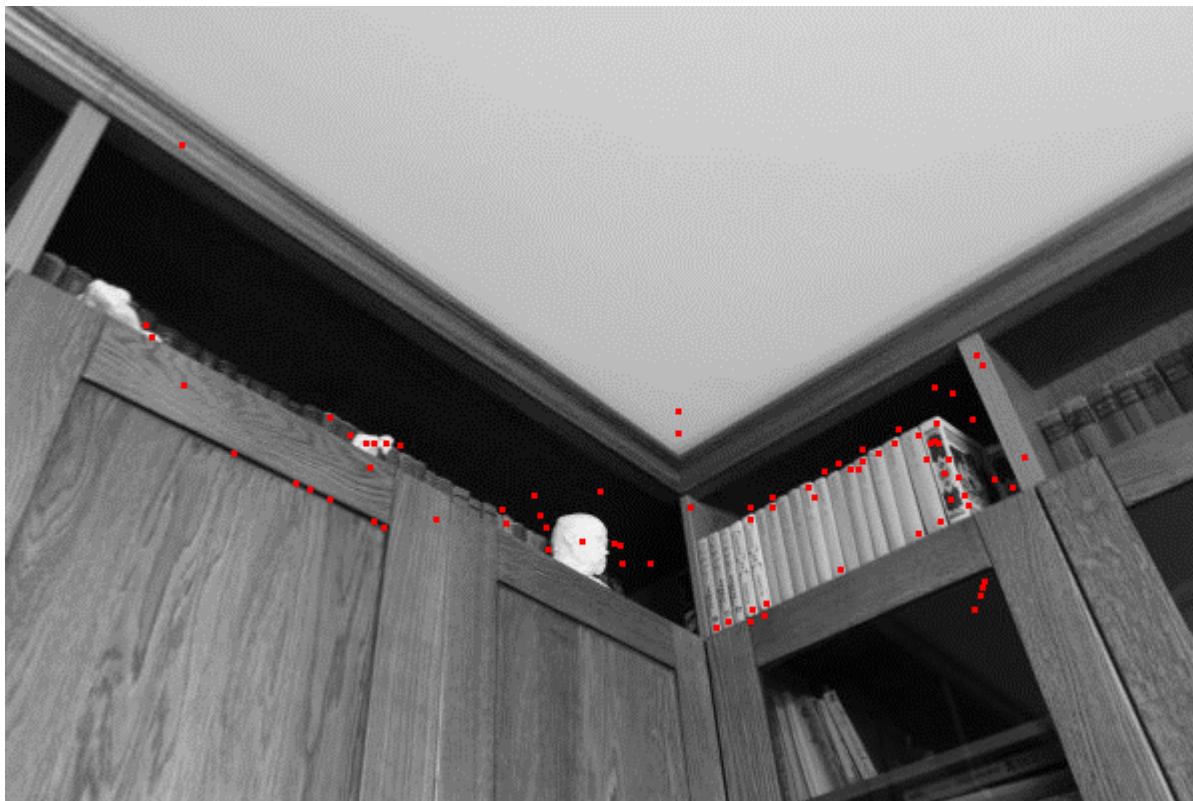
ORB: Oriented FAST and Rotated BRIEF



KeyPoints detection

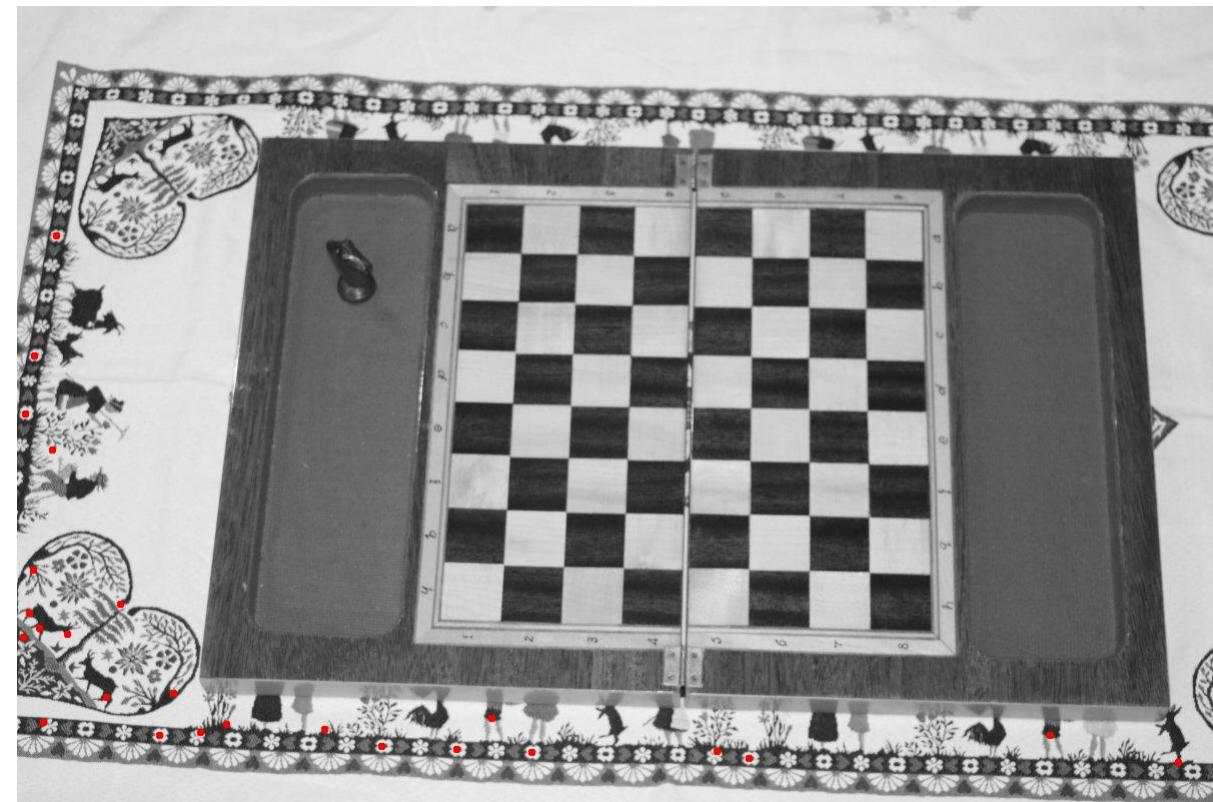
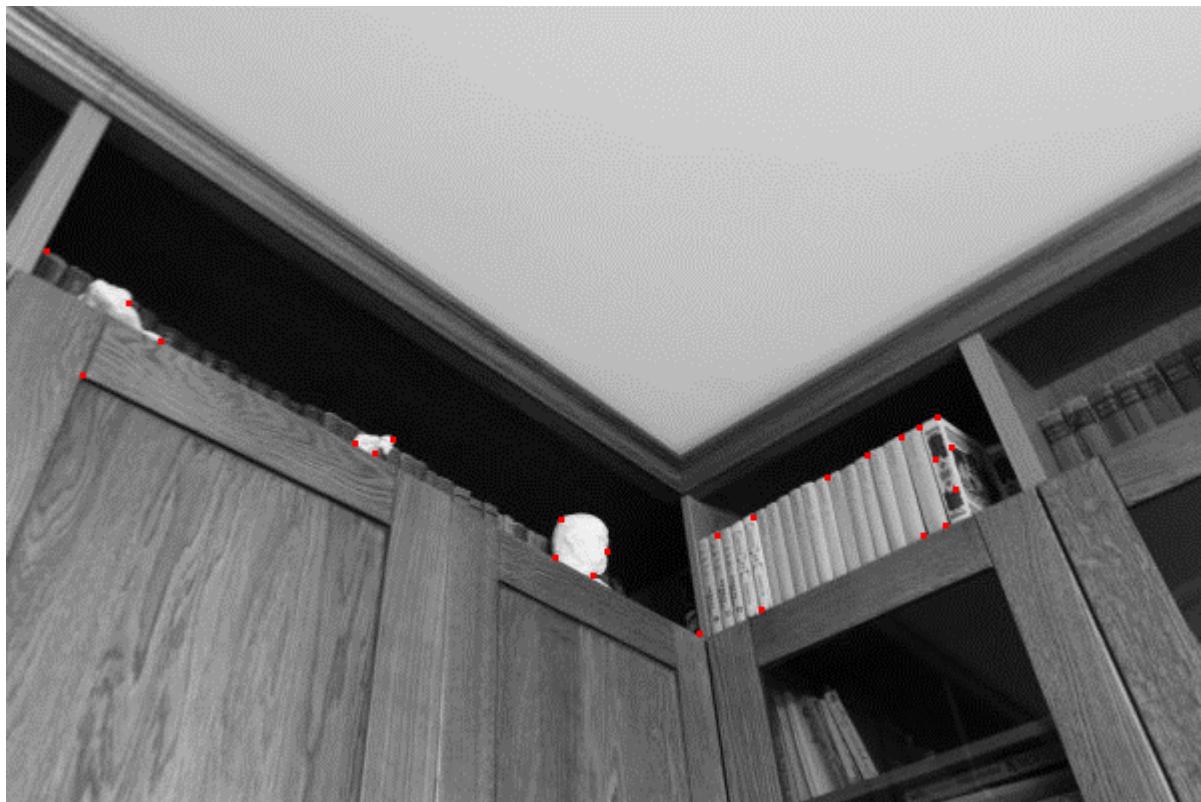
STAR: scale-invariant center-surround detector(CENSURE)

called STAR detector in OpenCV



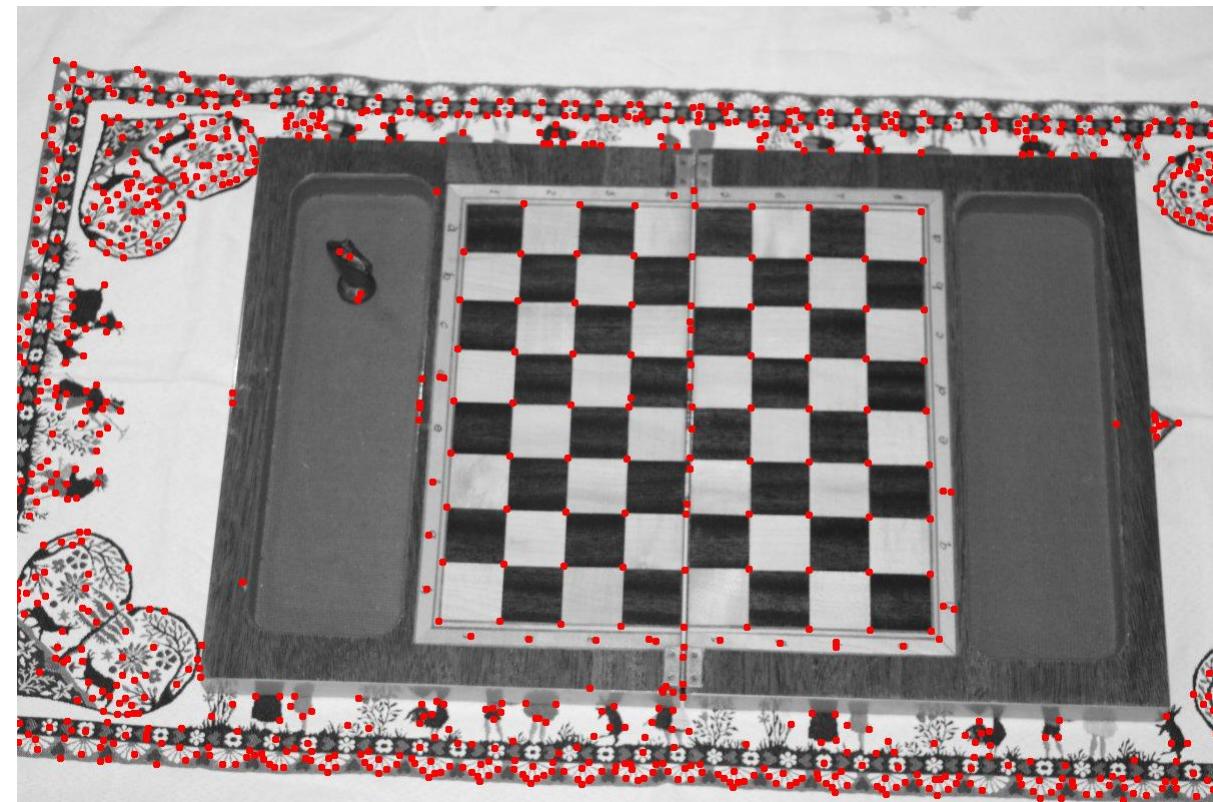
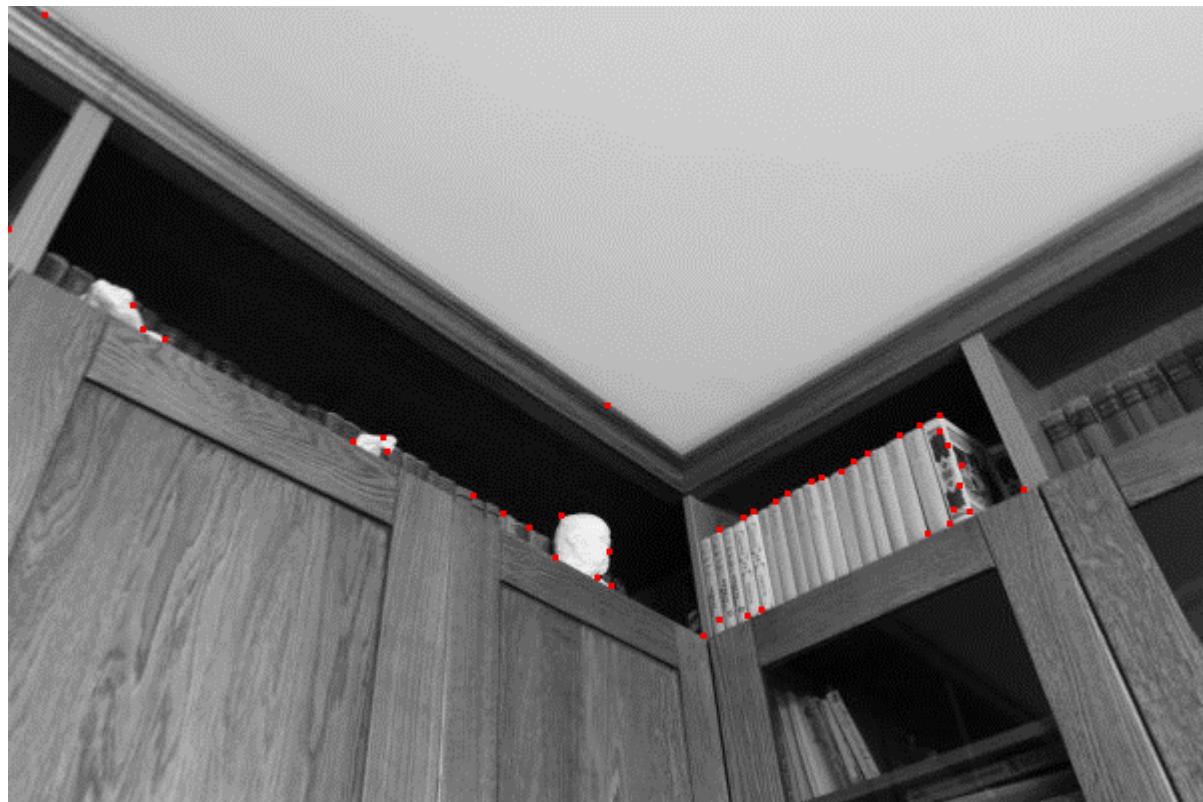
KeyPoints detection

Shi-Tomasi

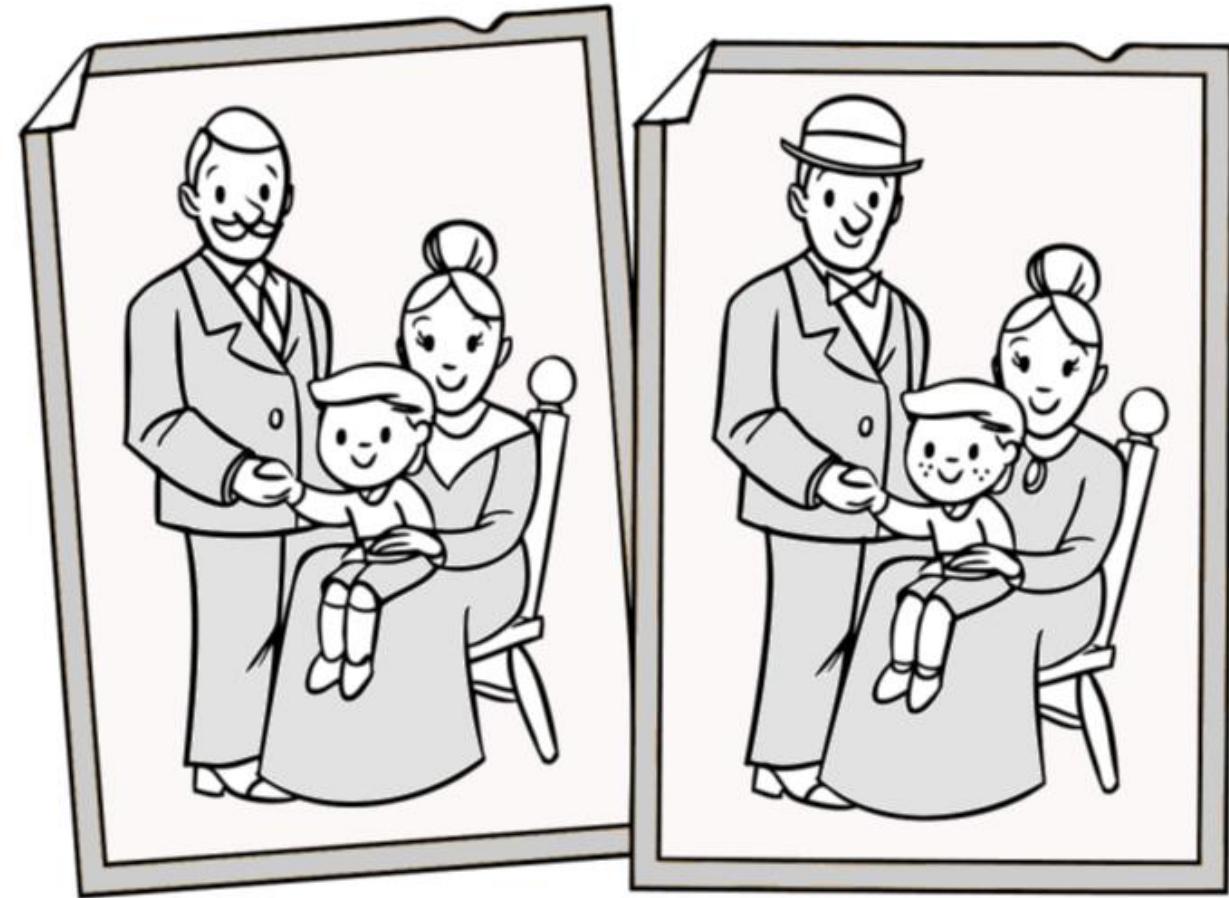


KeyPoints detection

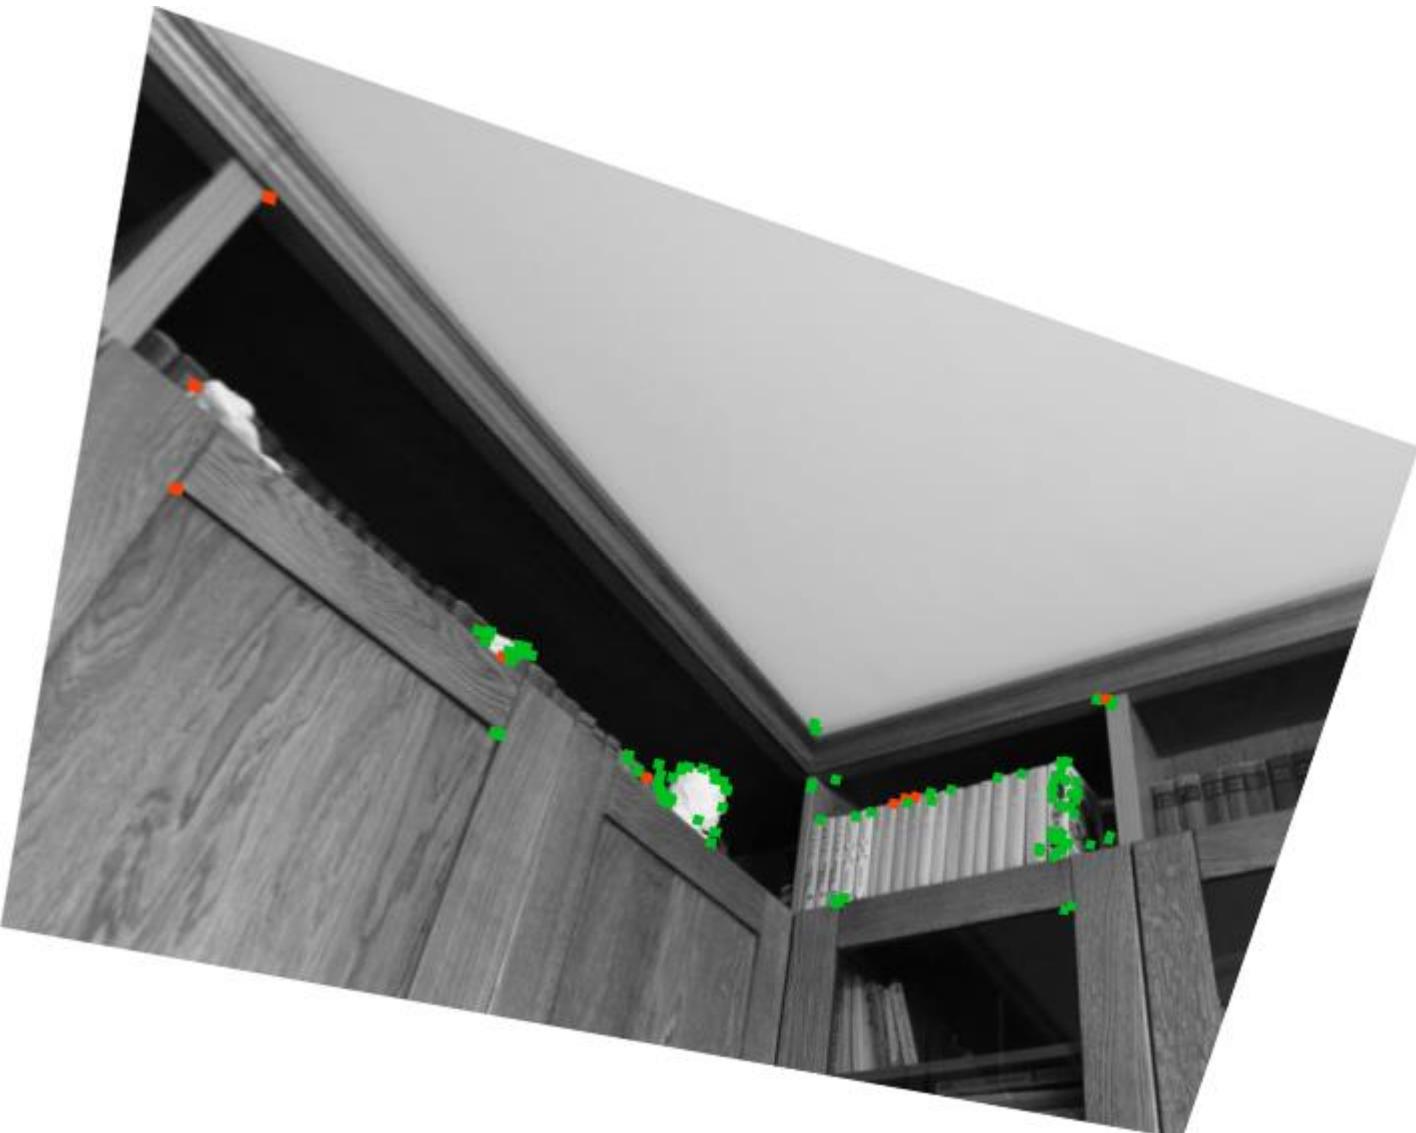
Harris



Matching the KeyPoints



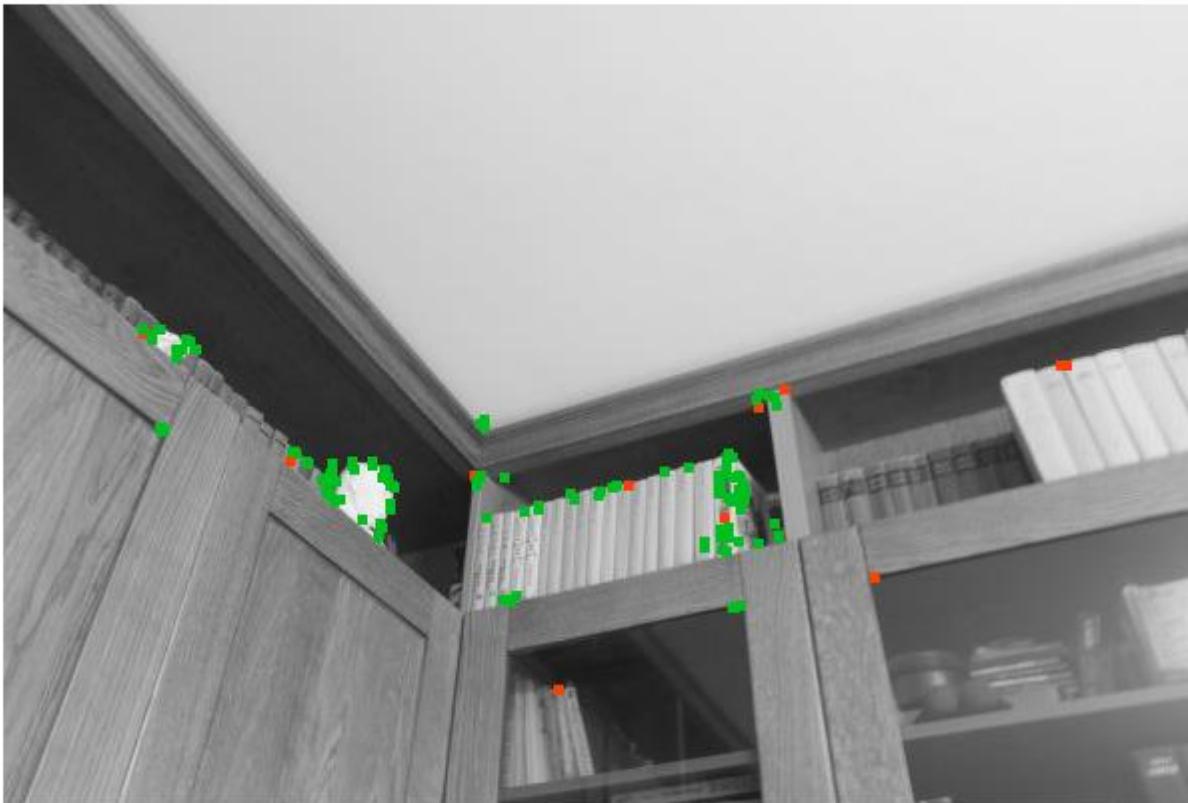
Matching the KeyPoints



*Detector: ORB
Matcher: FLANN*

True Positive
False Positive

Matching the KeyPoints



Detector: ORB

Matcher: FLANN

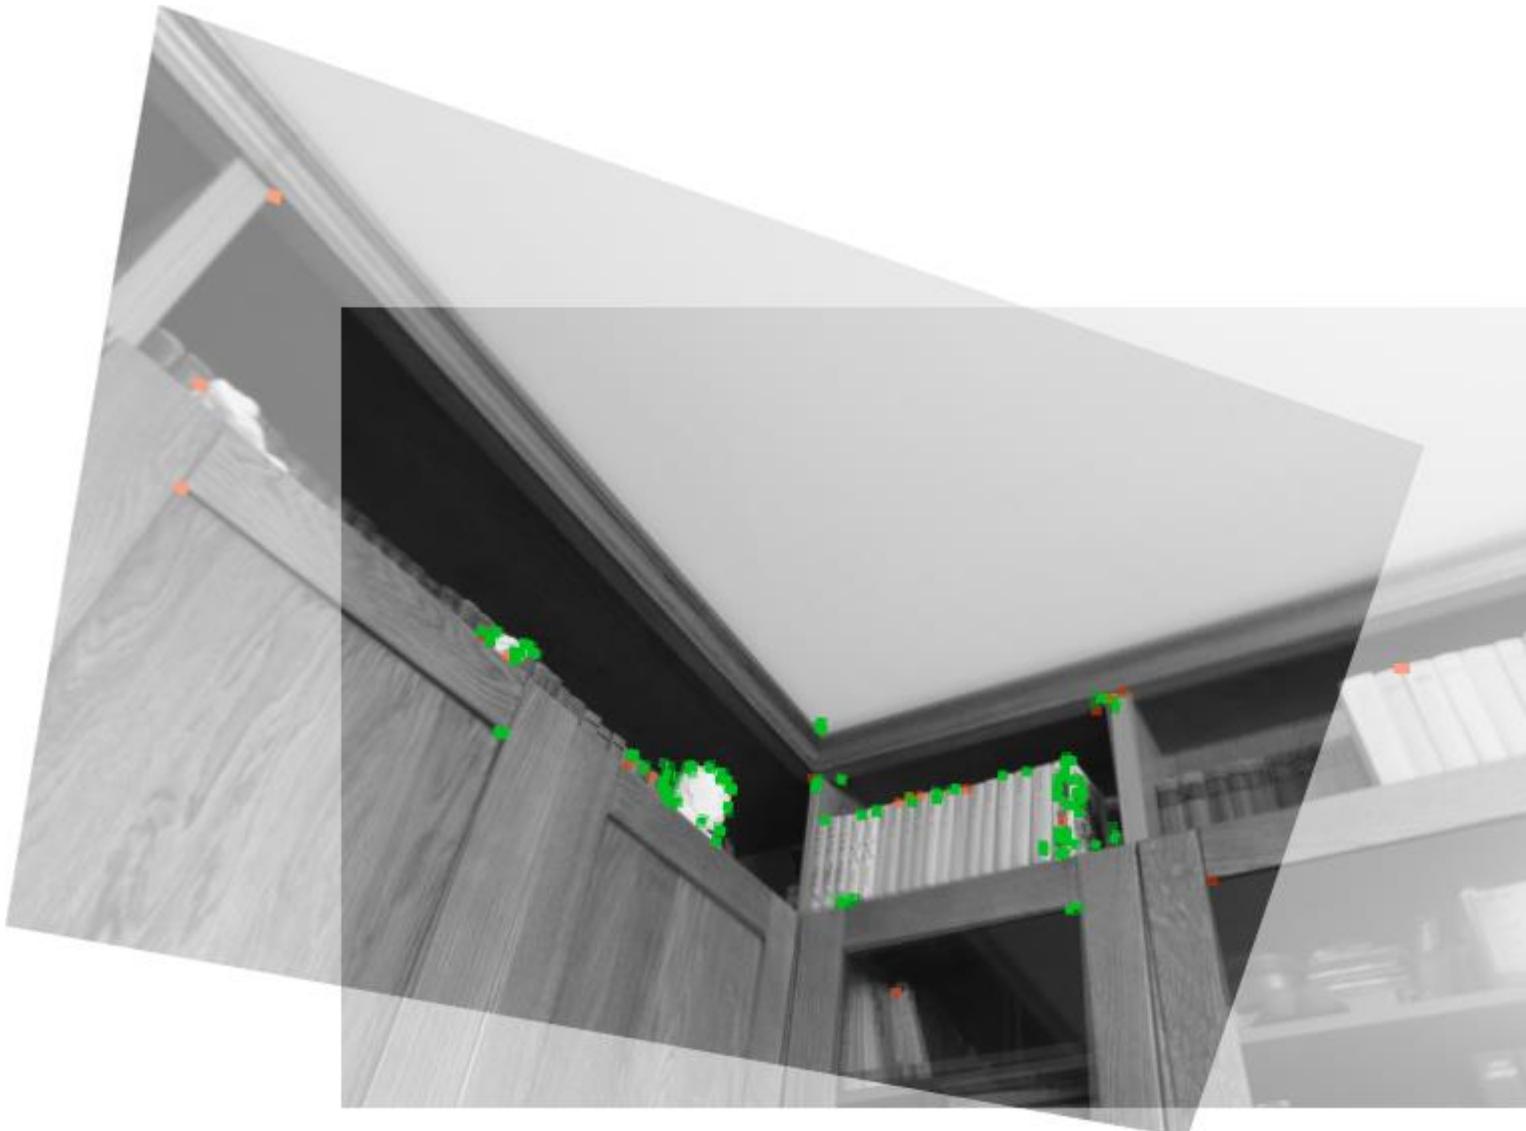
True Positive

False Positive

Matching the KeyPoints

Detector: ORB
Matcher: FLANN

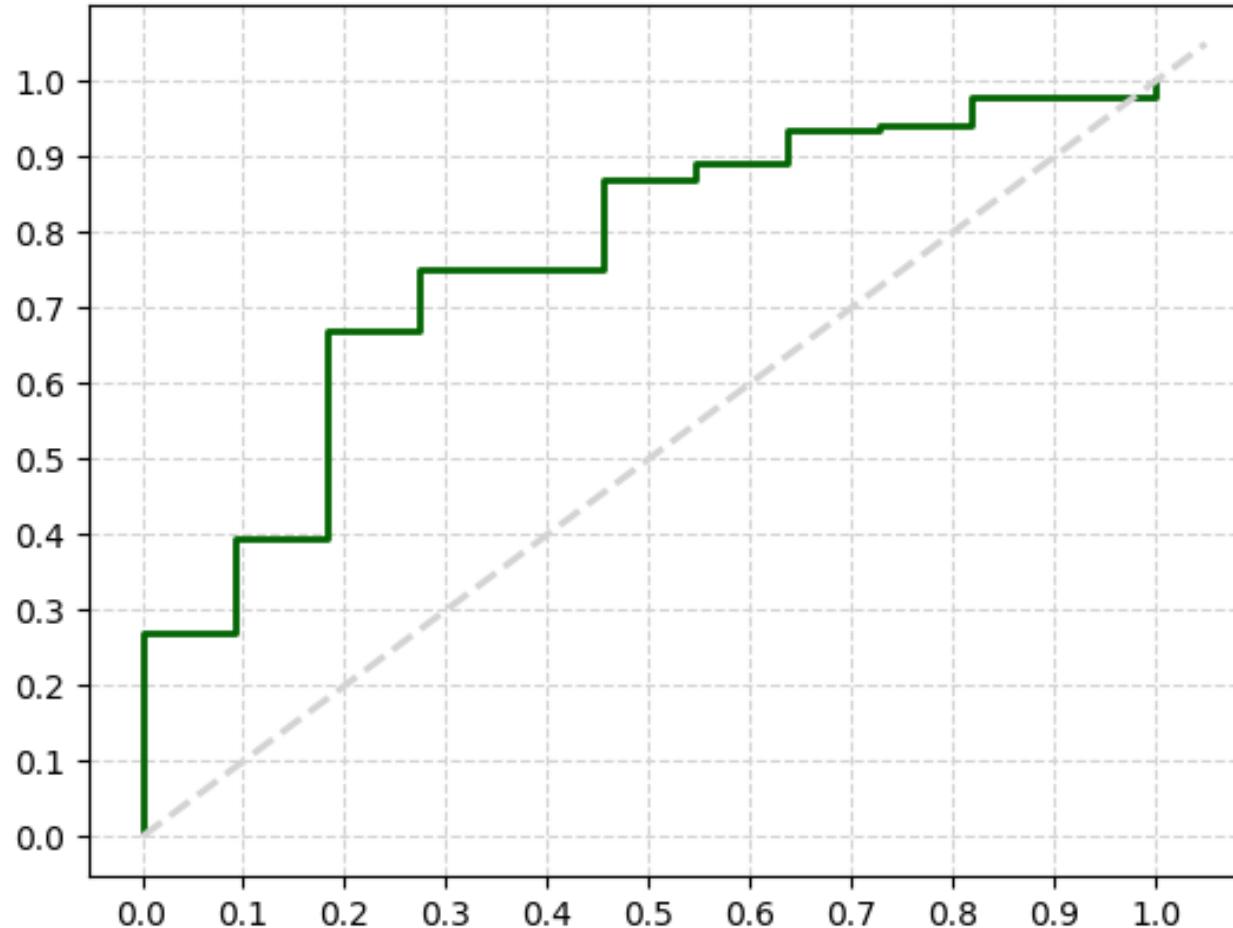
 True Positive
 False Positive



Matching the KeyPoints

*Detector: ORB
Matcher: FLANN*

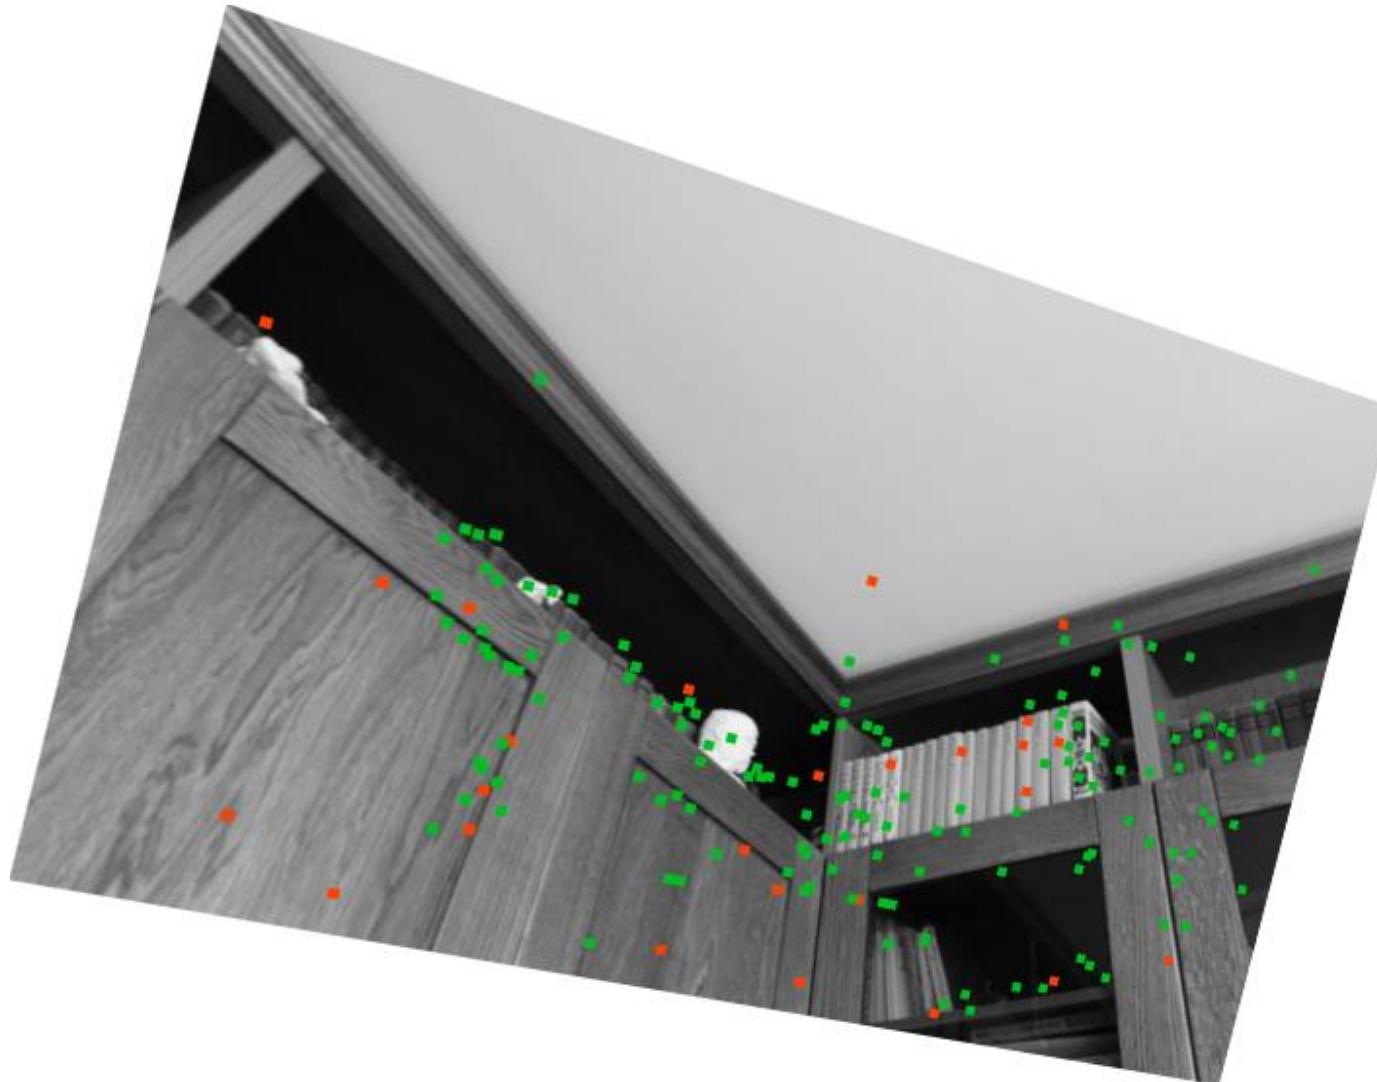
AUC = 76%



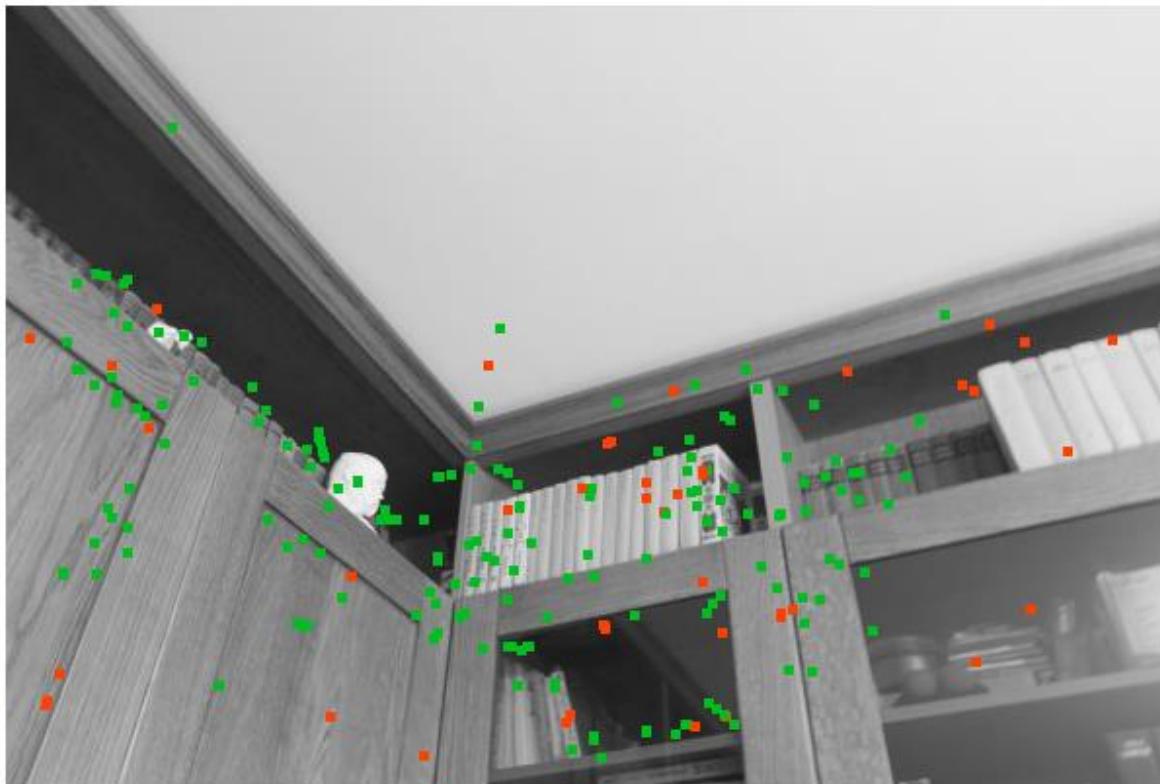
Matching the KeyPoints

*Detector: SURF
Matcher: KNN*

- *True Positive*
- *False Positive*



Matching the KeyPoints



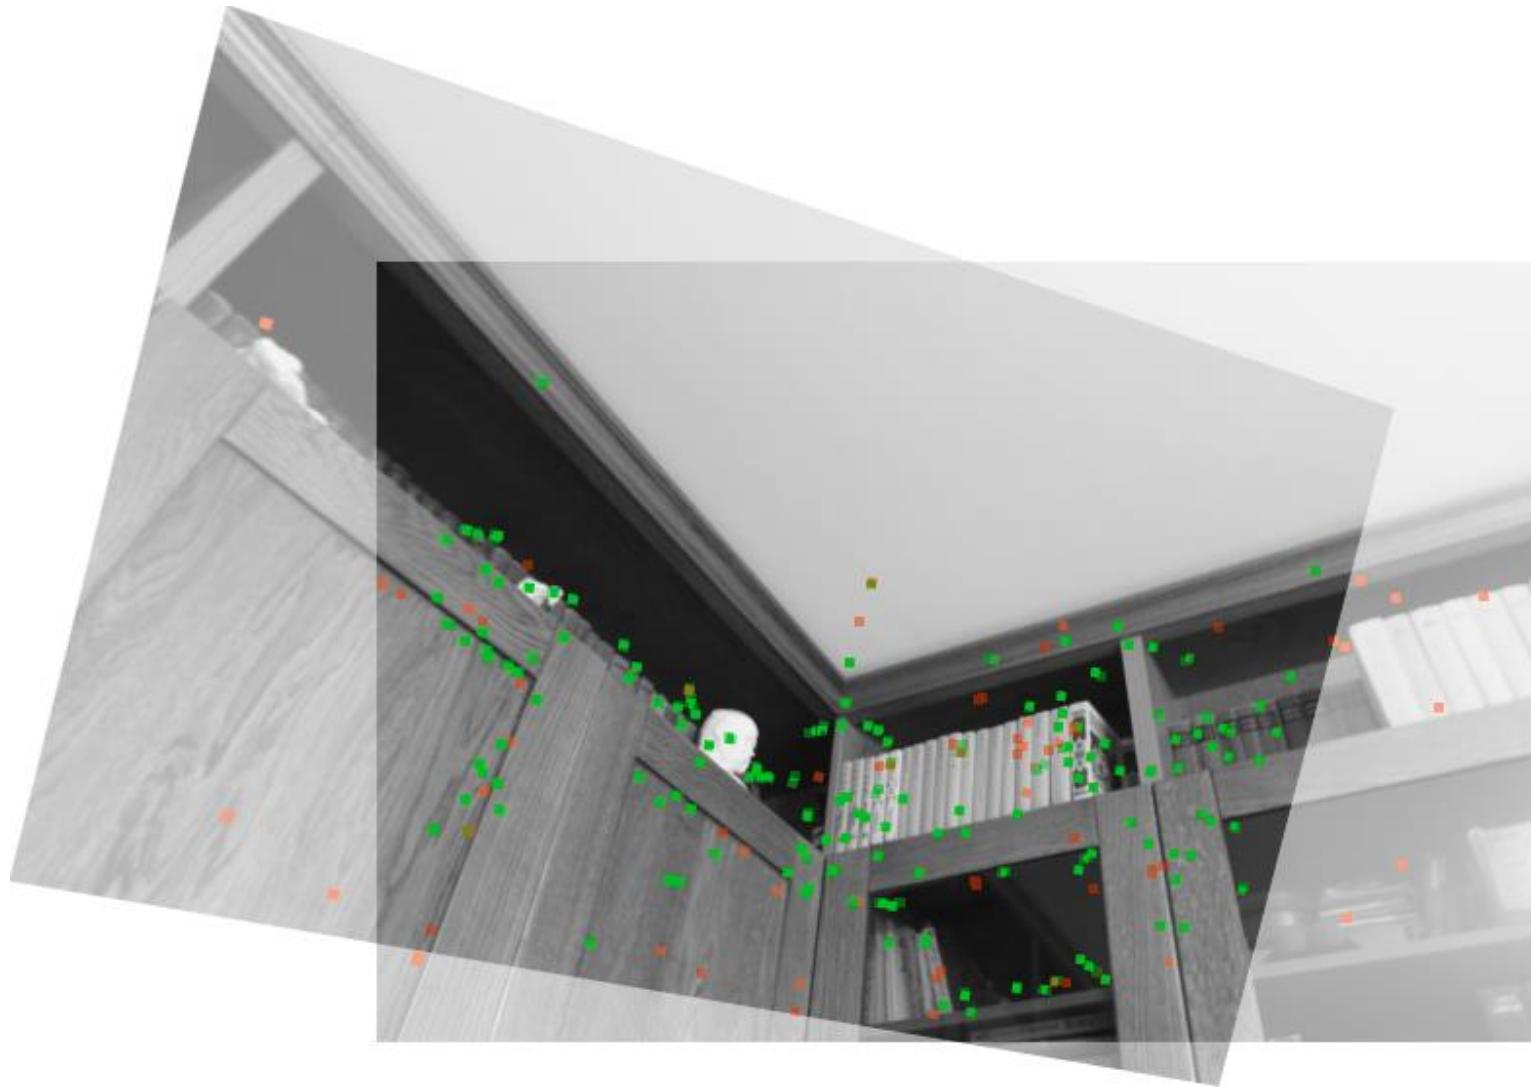
*Detector: SURF
Matcher: KNN*

- *True Positive*
- *False Positive*

Matching the KeyPoints

*Detector: SURF
Matcher: KNN*

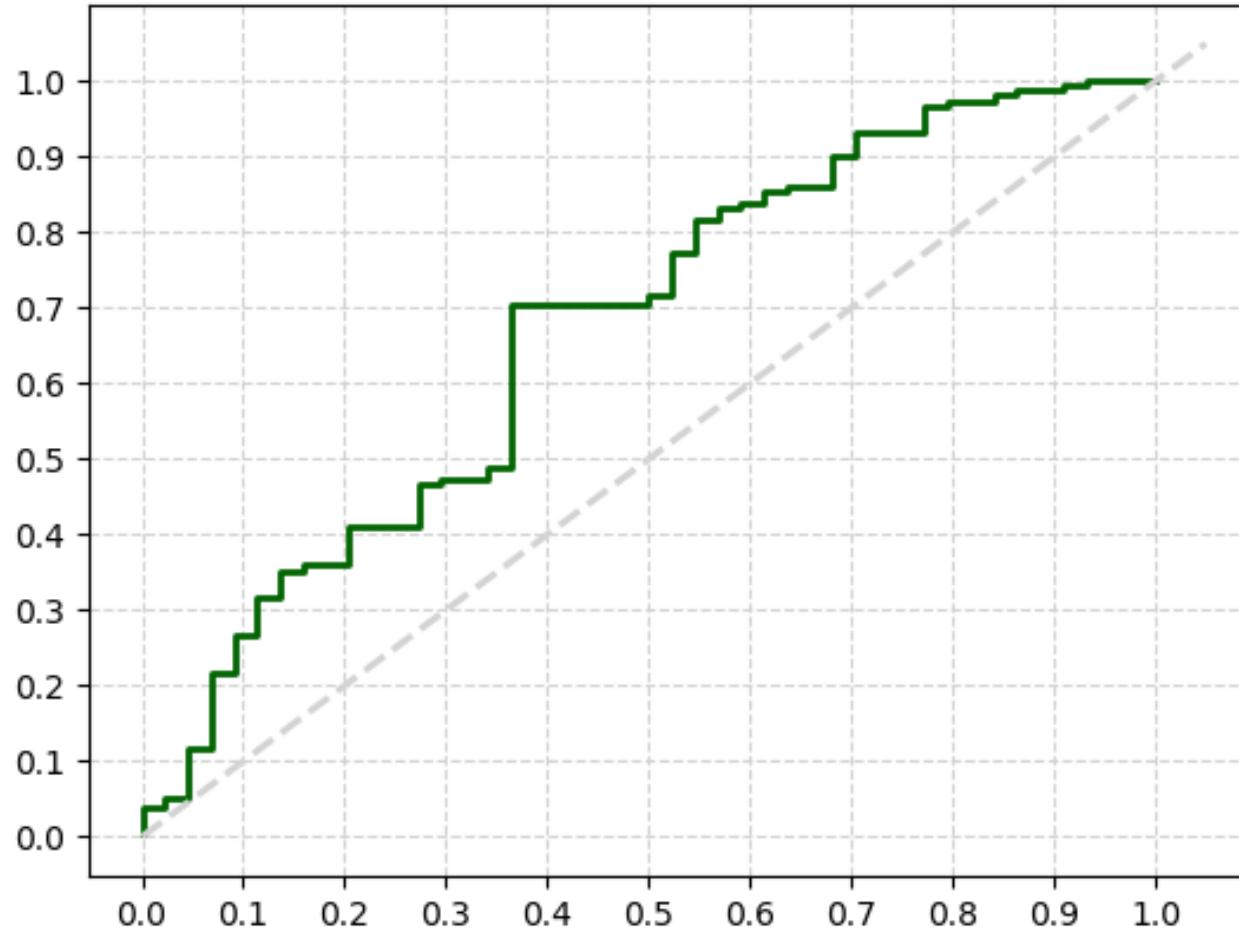
- *True Positive*
- *False Positive*



Matching the KeyPoints

*Detector: SURF
Matcher: KNN*

AUC = 67%



Matching the KeyPoints

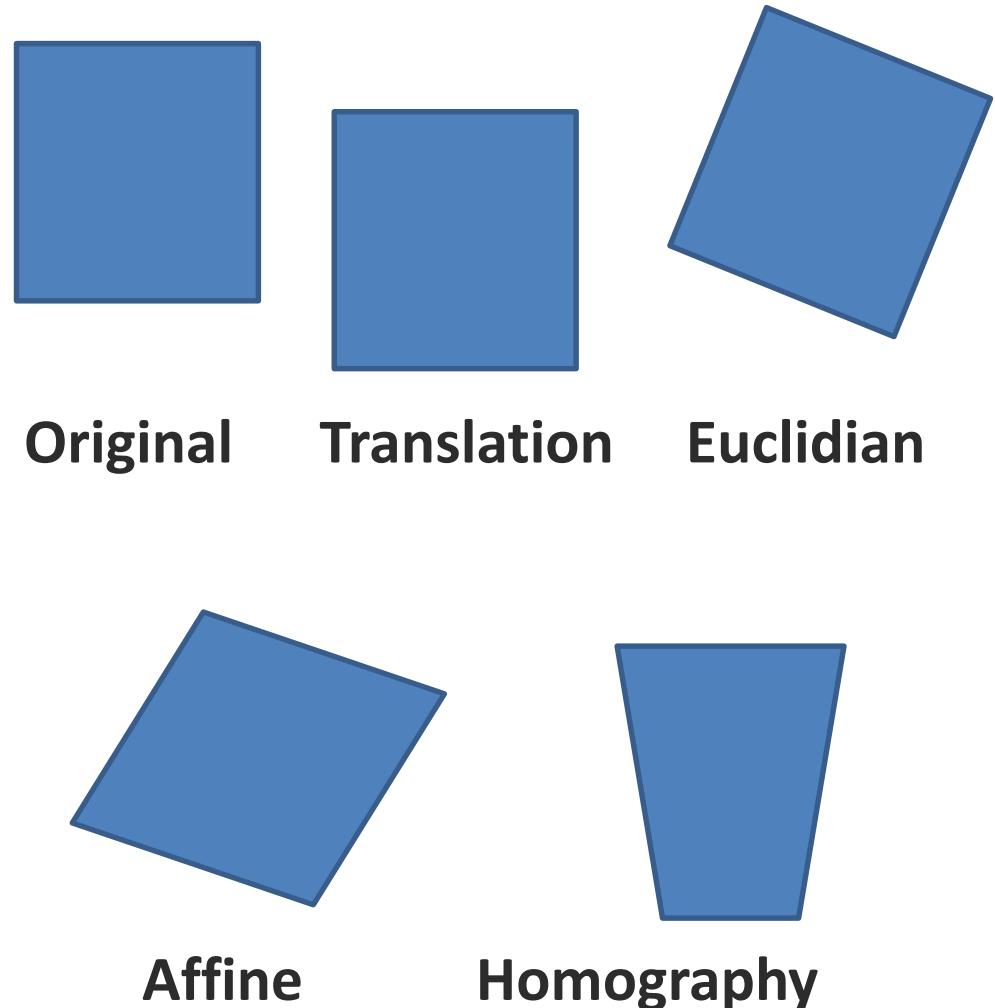
	BFMatcher.match	BFMatcher.knnMatch	FlannBasedMatcher.knnMatch
ORB	91%	75%	76%
SIFT	63%	69%	70%
SURF	-	67%	69%

3. Homography



Homography: transformations

- *findTransformECC*
- *findHomography*
- *estimateAffine2D*,
- *estimateAffinePartial2D*
- *transform*
- *perspectiveTransform*
- *getPerspectiveTransform*
- *warpAffine*
- *warpPerspective*



Homography: manual match



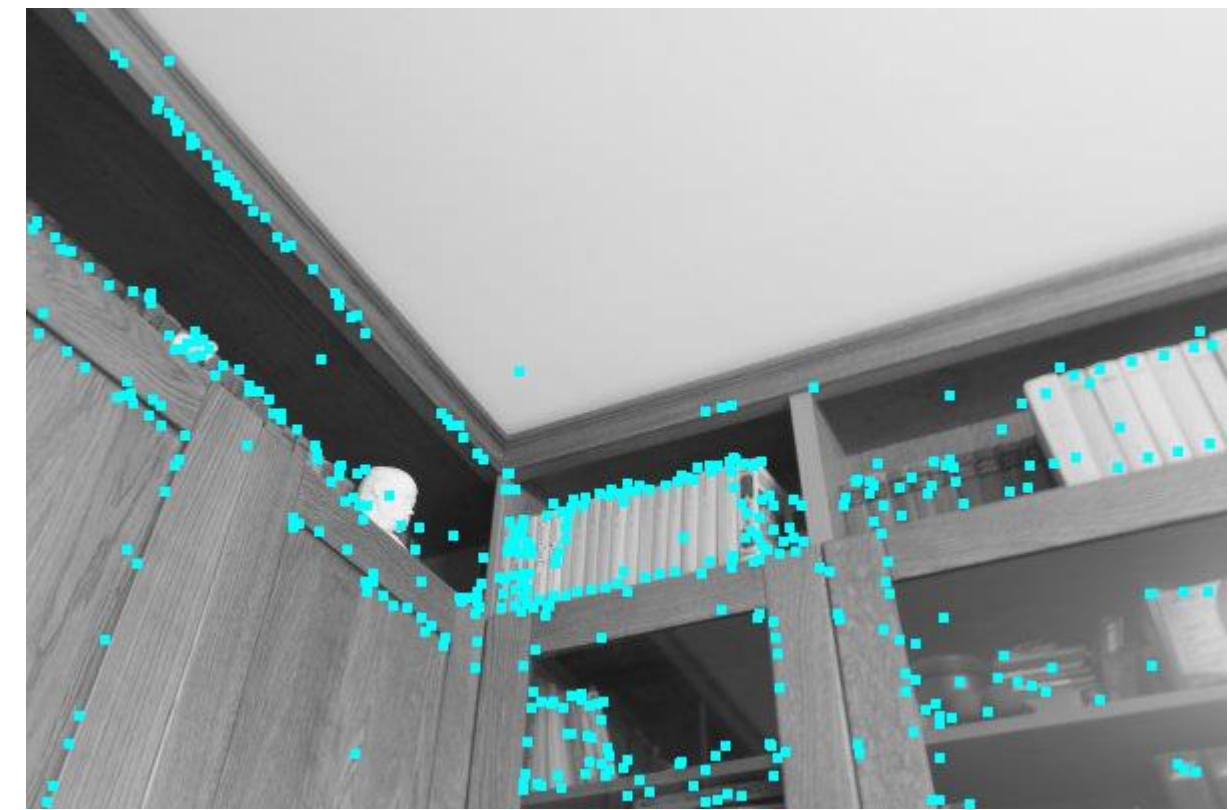
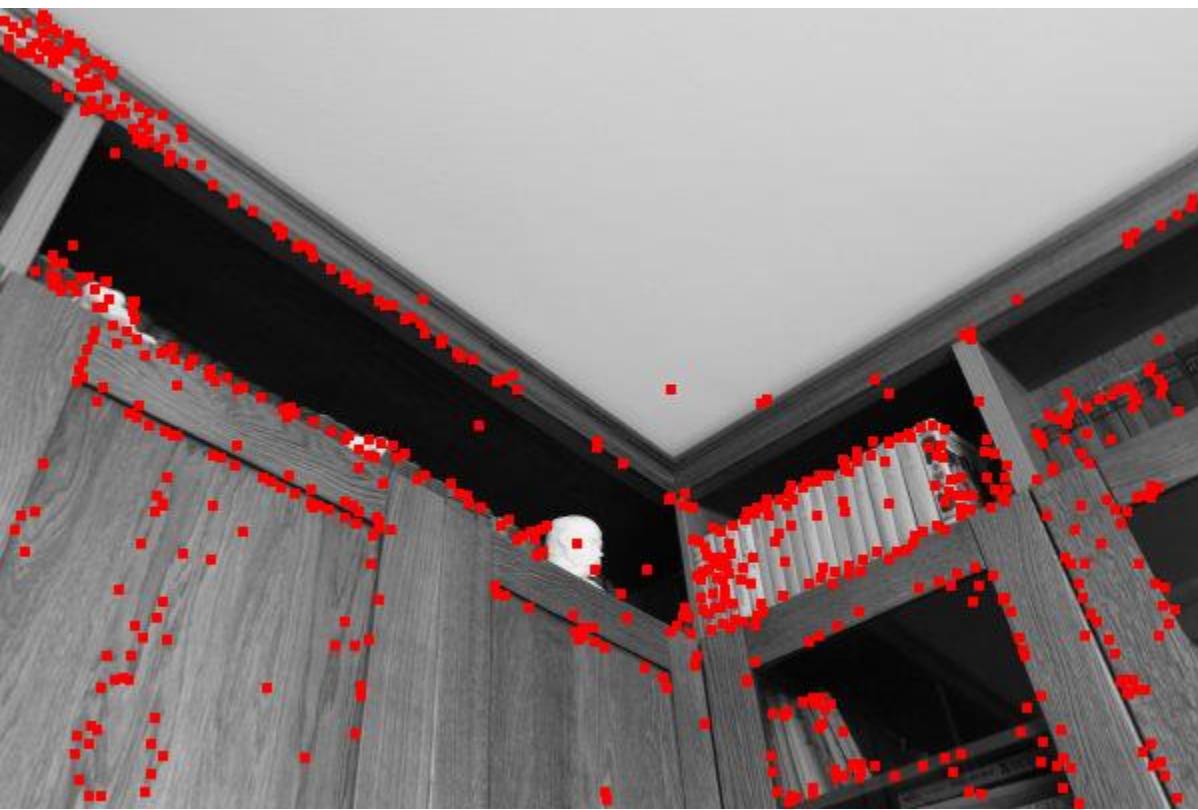
+



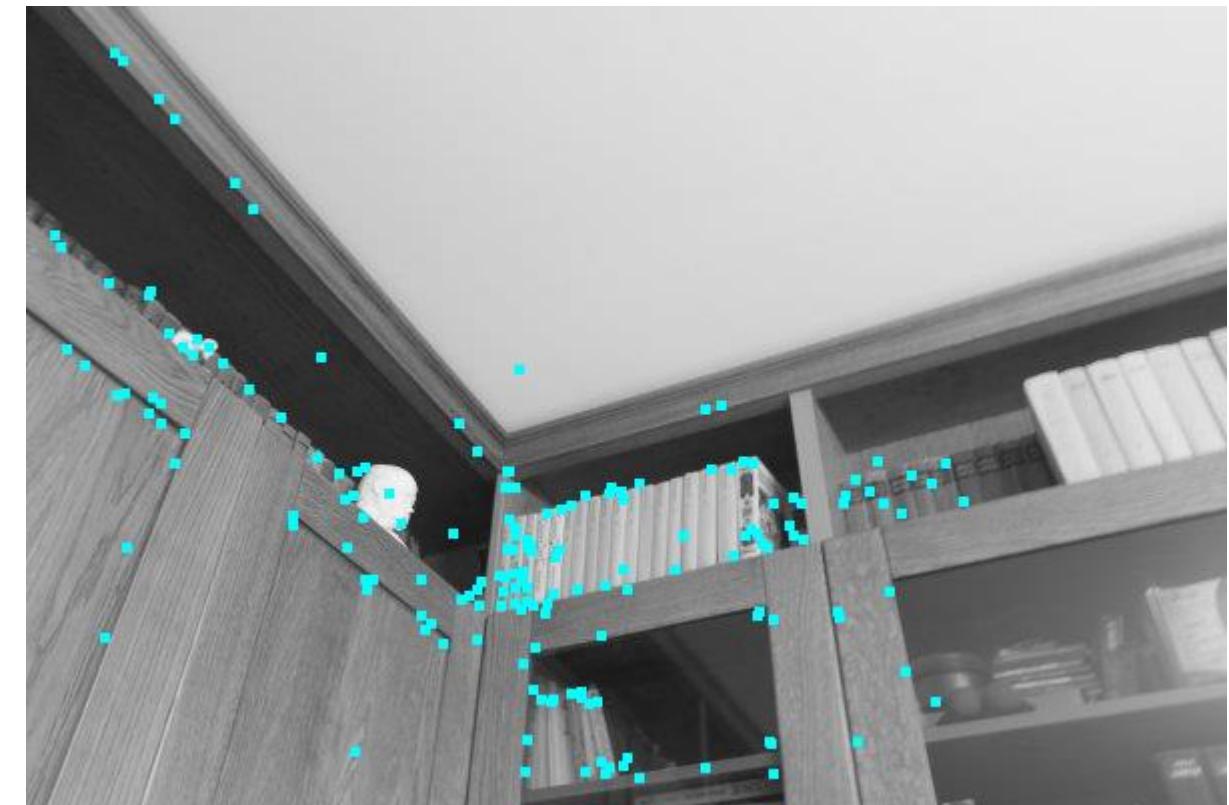
=



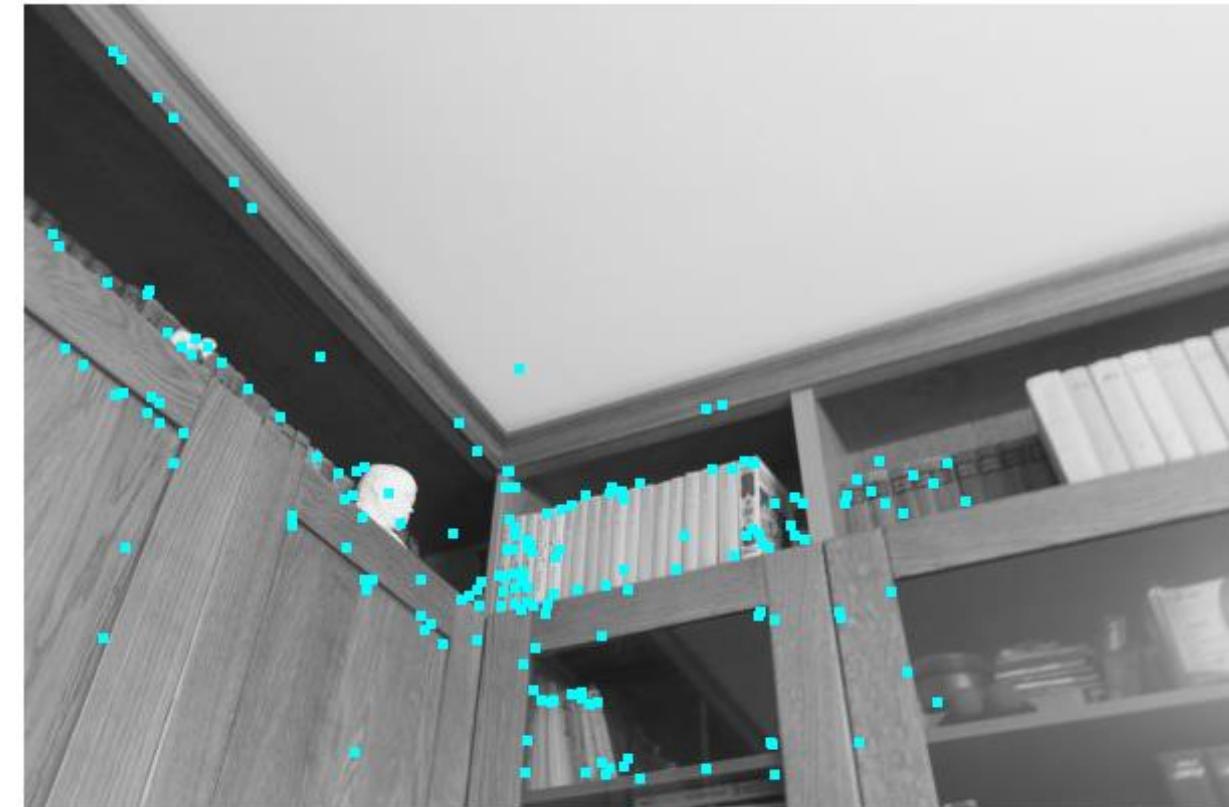
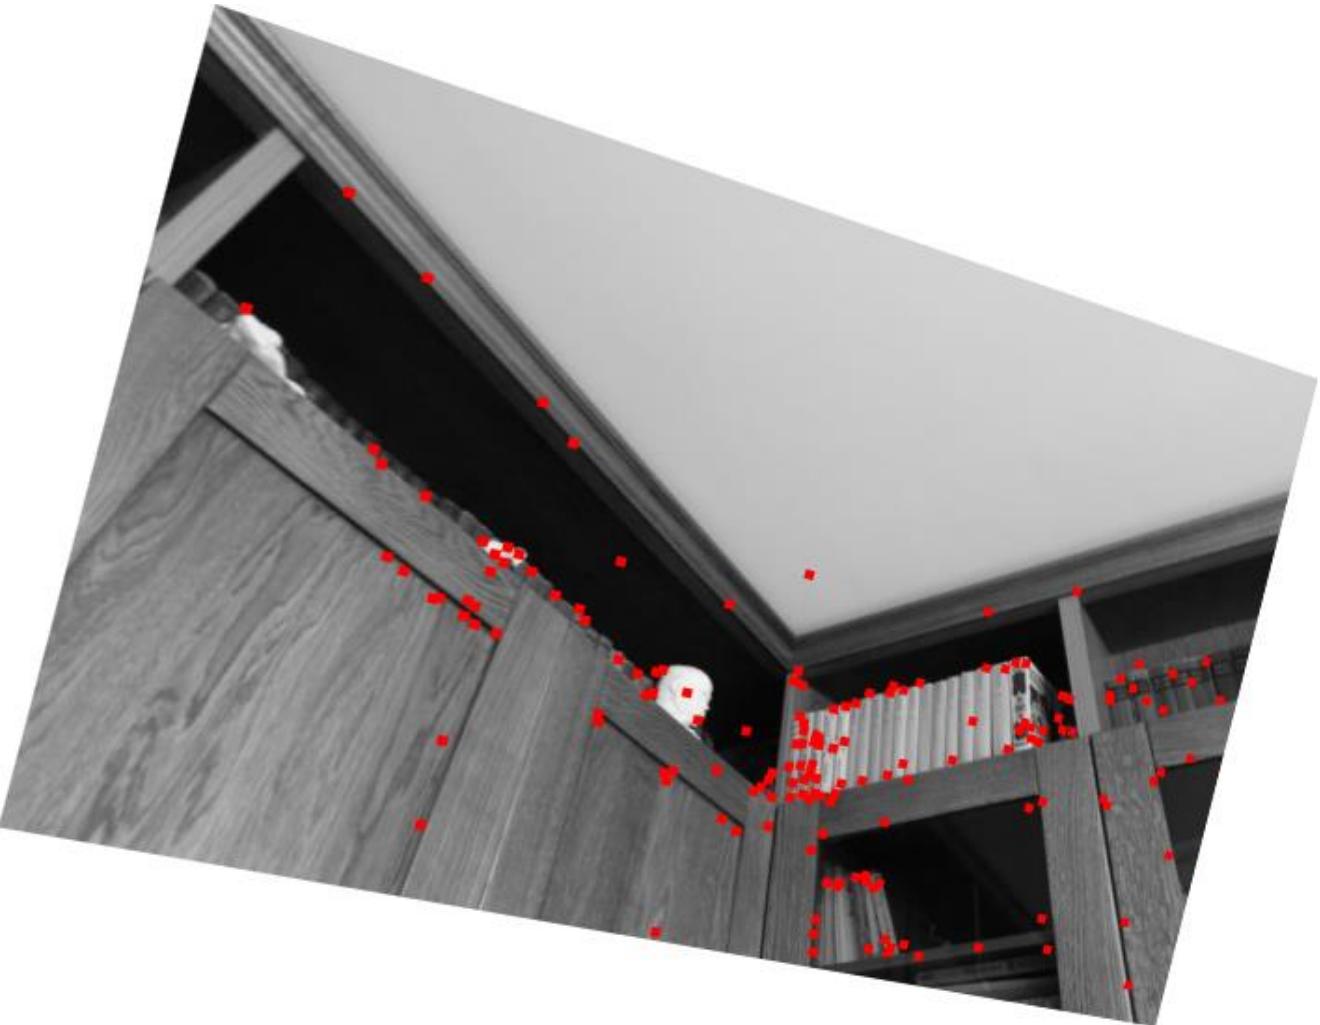
Homography: auto-match



Homography: auto-match



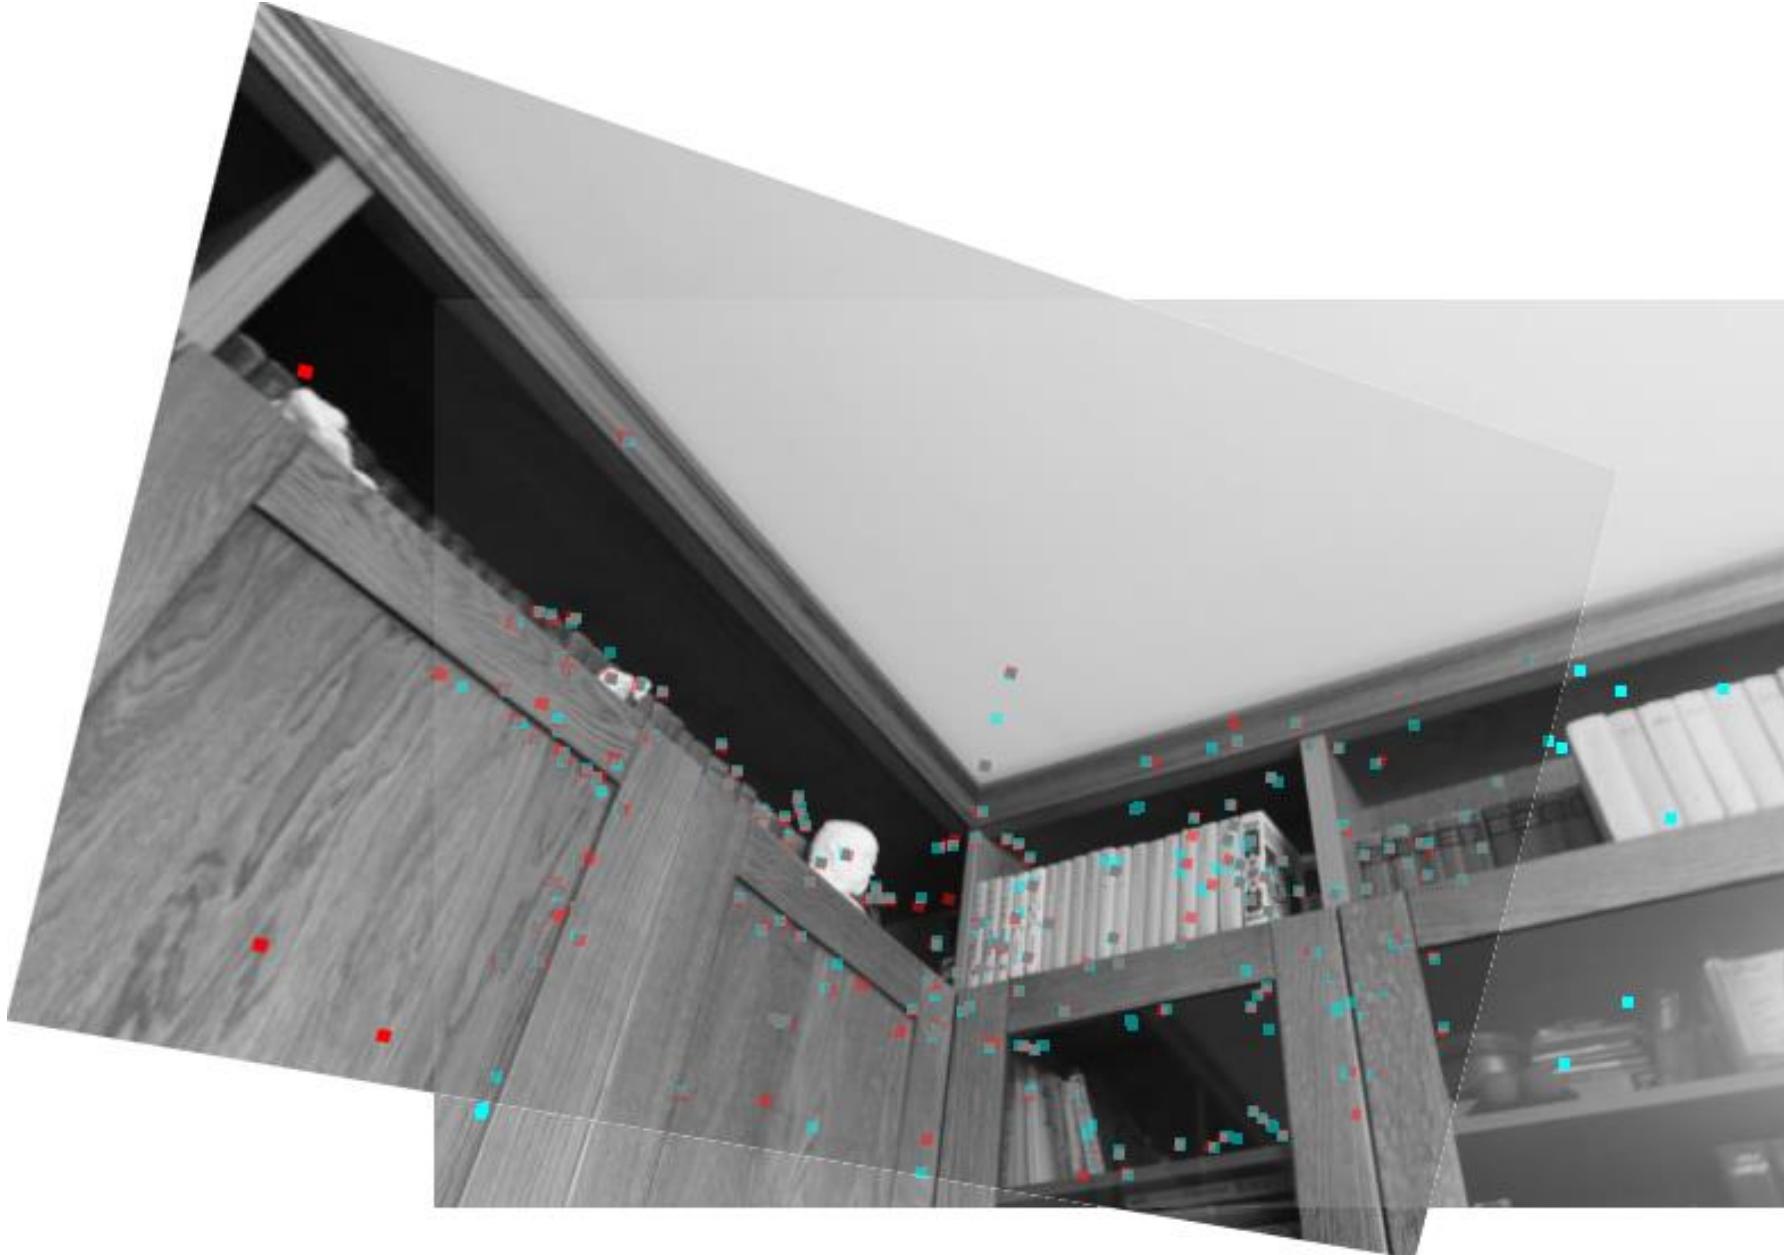
Homography: auto-match



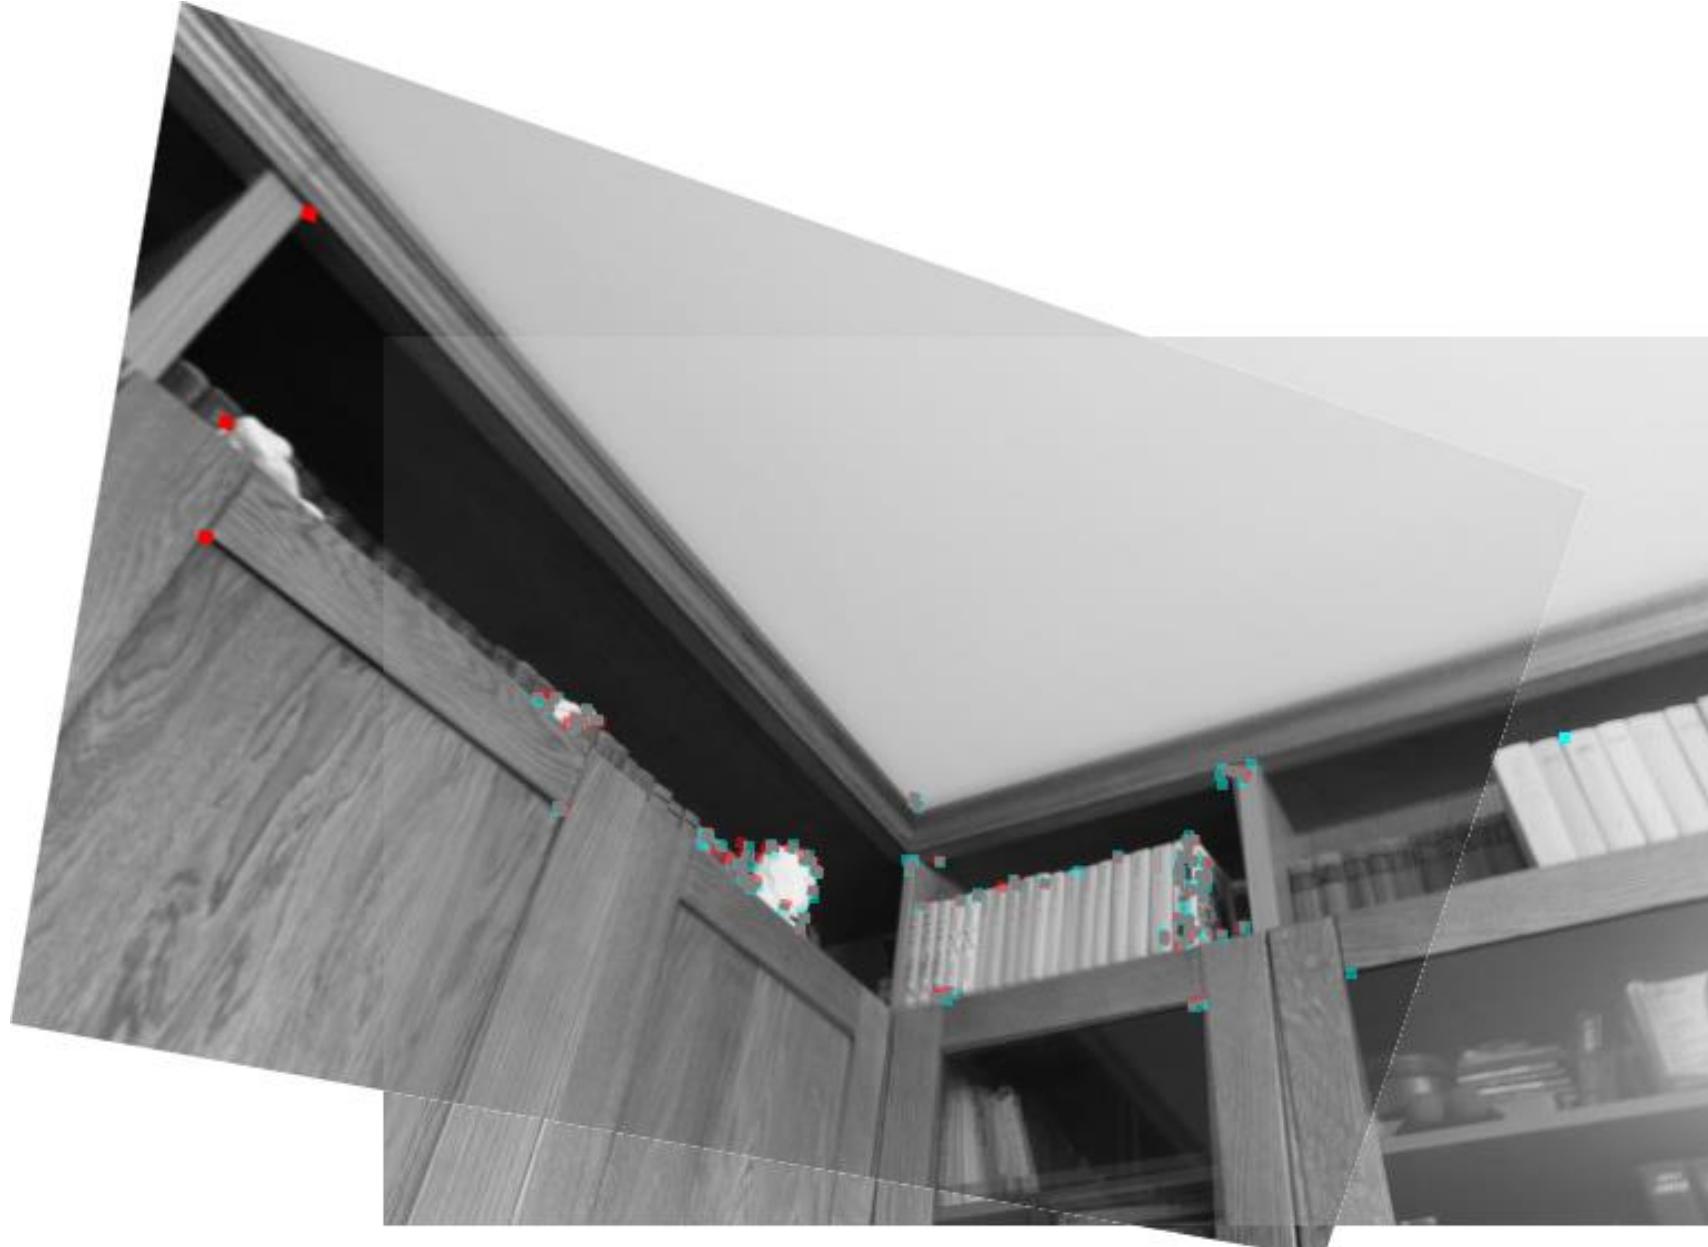
Homography: auto-match SIFT



Homography: auto-match SURF



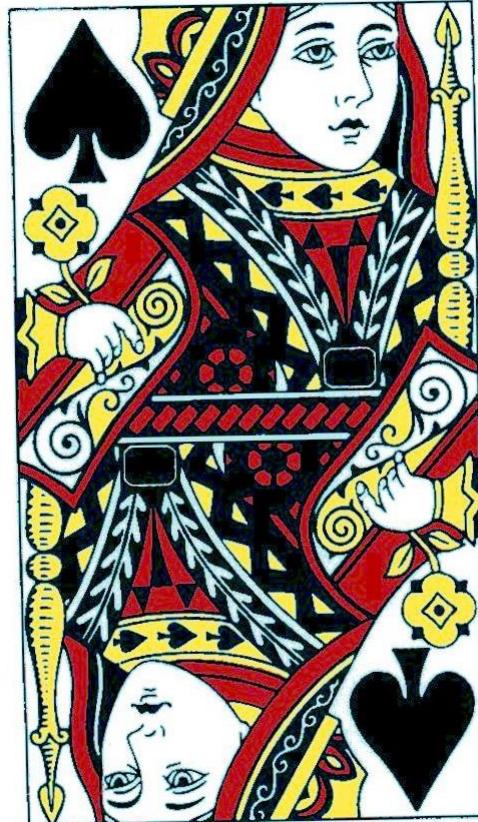
Homography: auto-match ORB



Homography: bitmap transformation

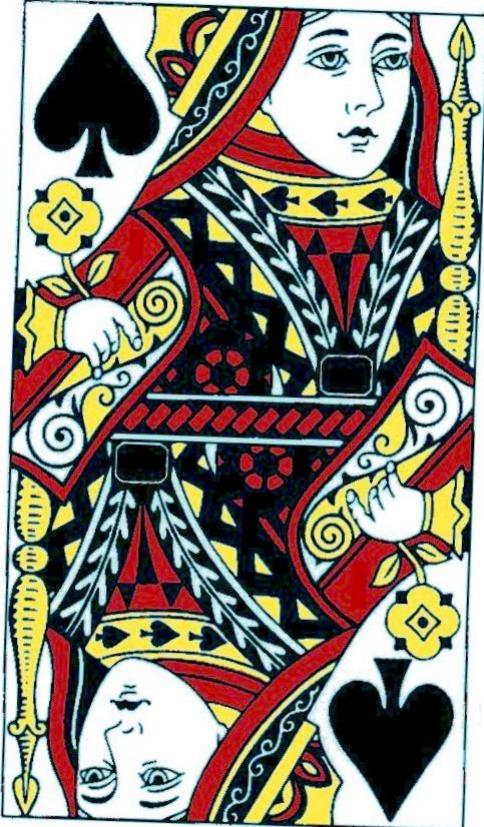


first



second

Homography: bitmap transformation



transformed second



transformed first

4. Blending

BBC  Sign in News Sport Weather Shop Earth Travel More ▾ Search 

capital

NEW SERIES: DISCOVER:
Home Owning Your Time Affording Your Life Level Up Bright Sparks



The video player interface shows a portrait of a smiling man with short, light-colored hair. A black play button icon is visible in the bottom-left corner of the video frame.

The Diversity Box

**What the average
American CEO looks like**

Follow BBC Capital



Blending: average



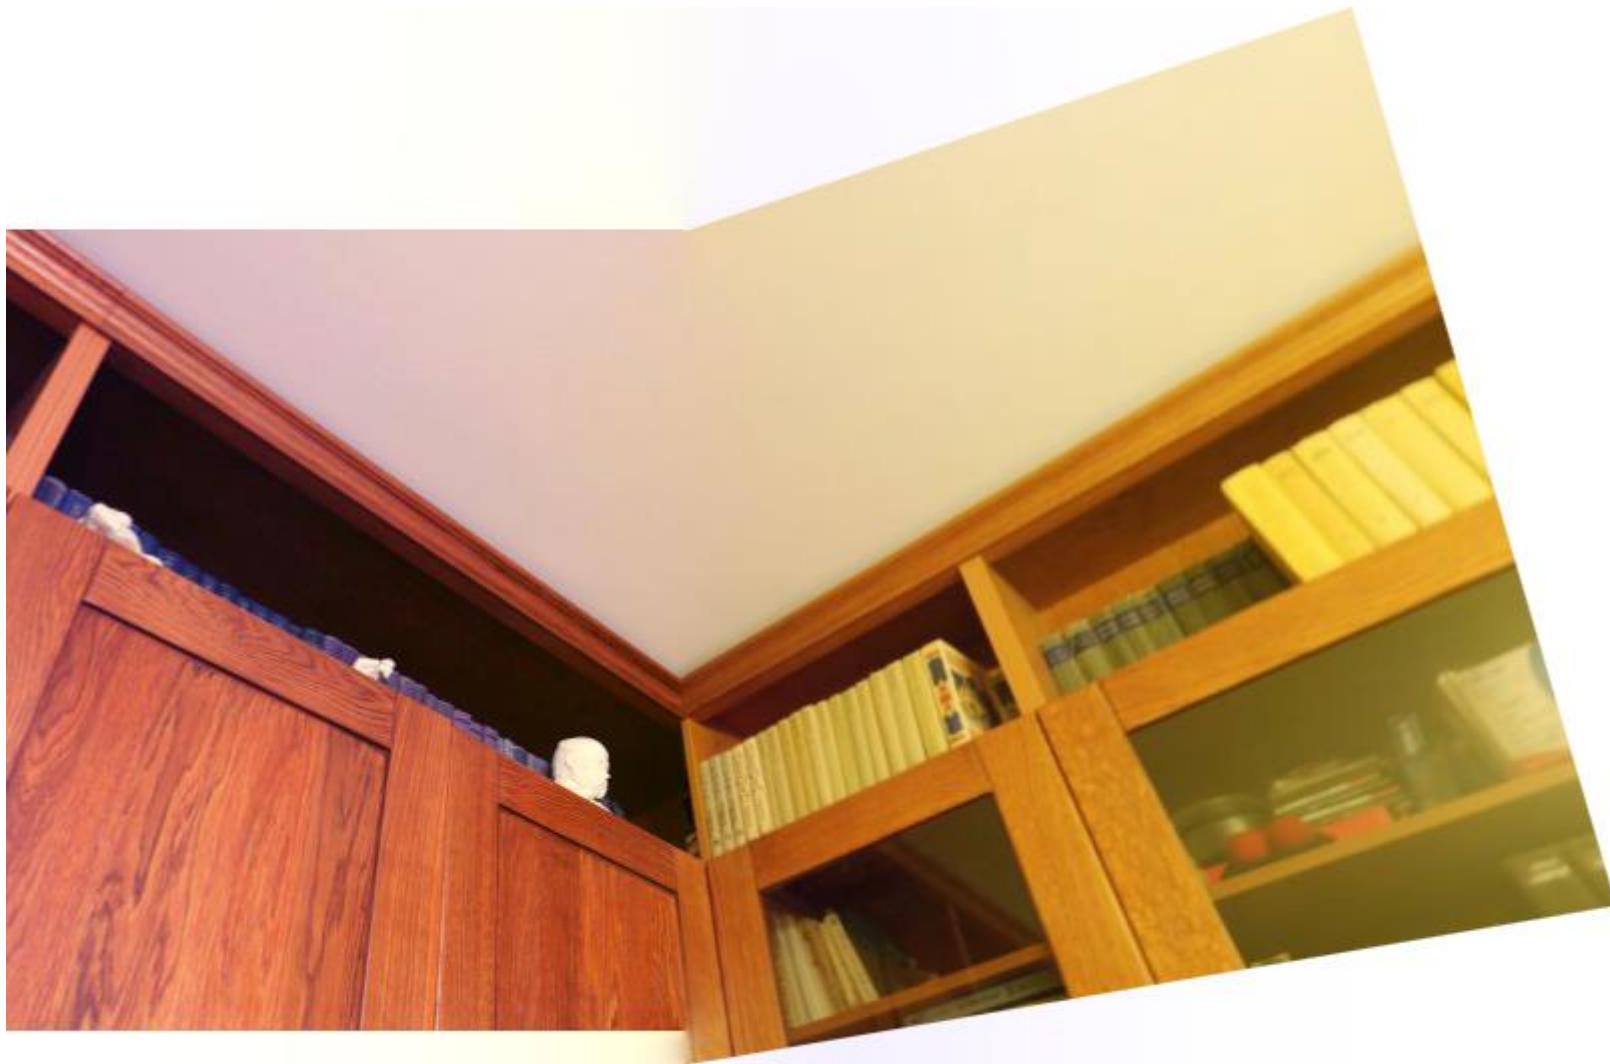
Blending: average



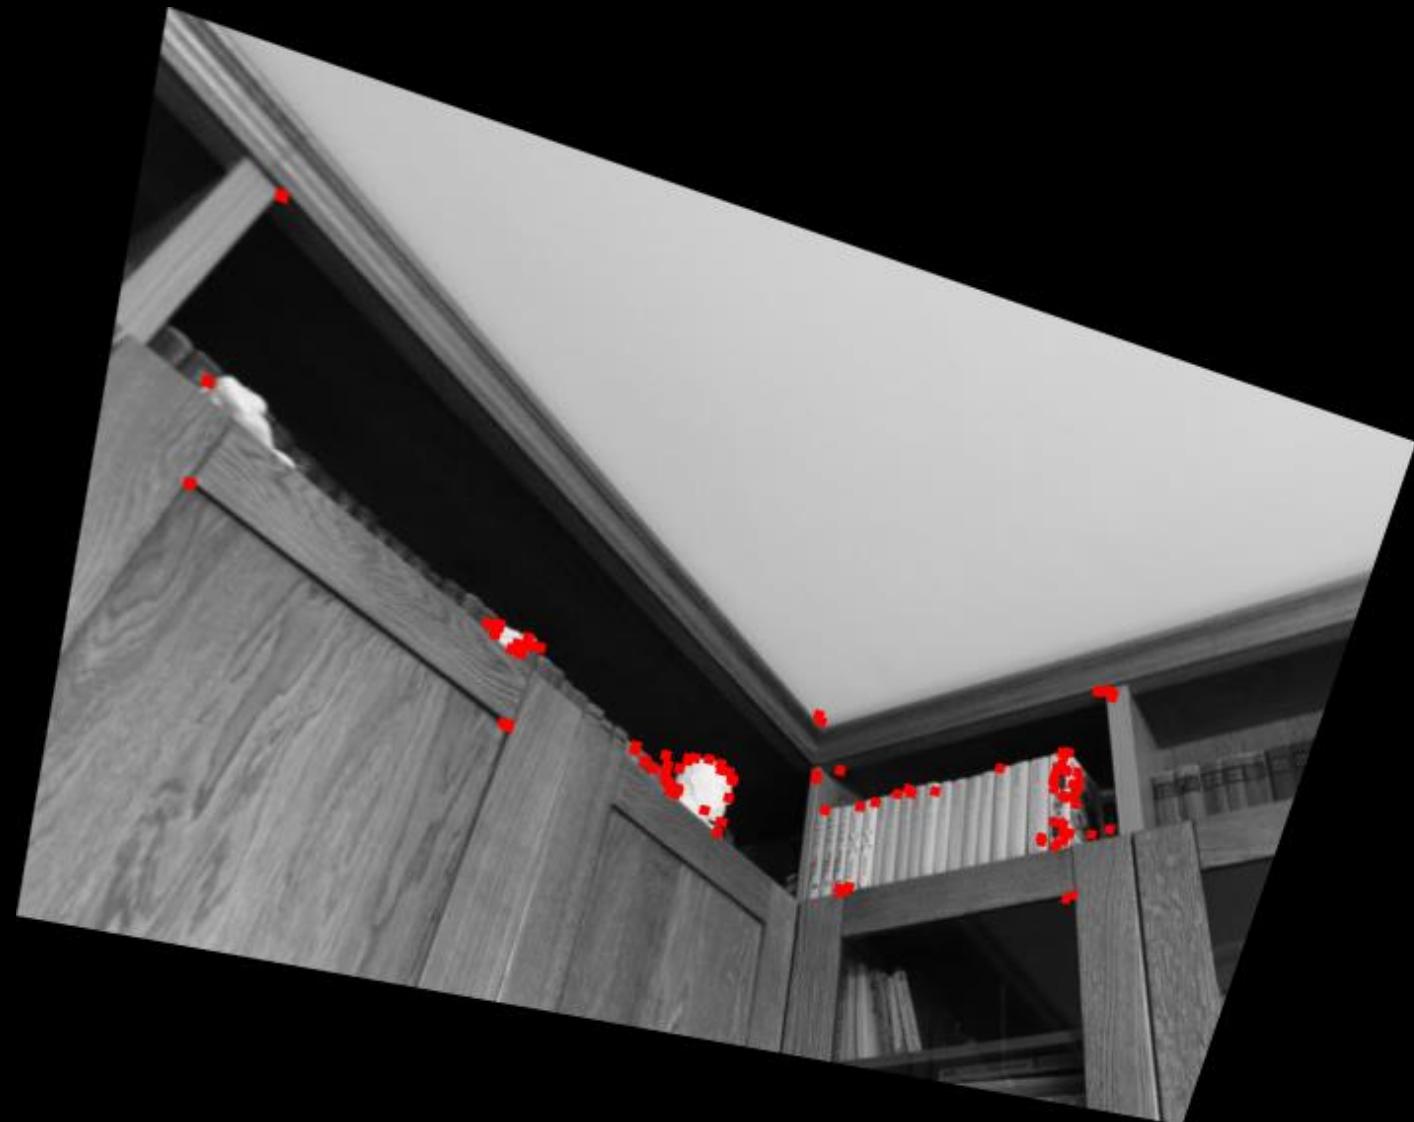
Blending: average



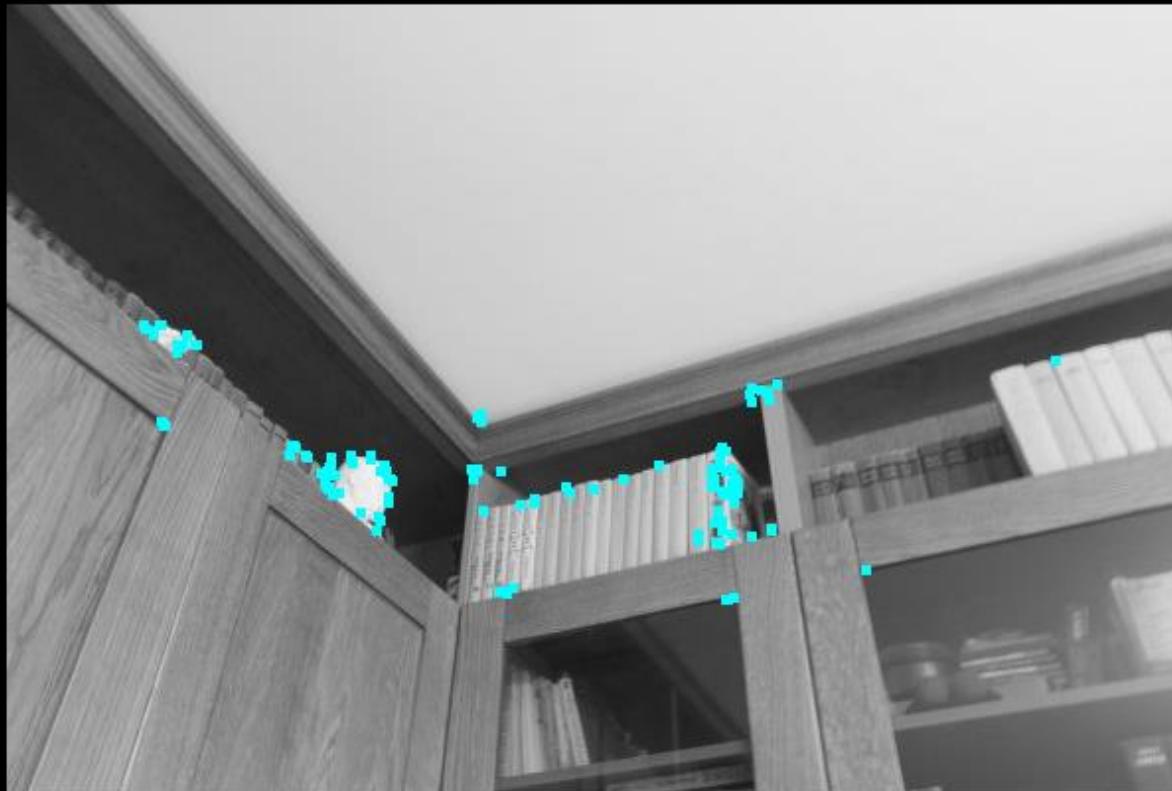
Blending: average



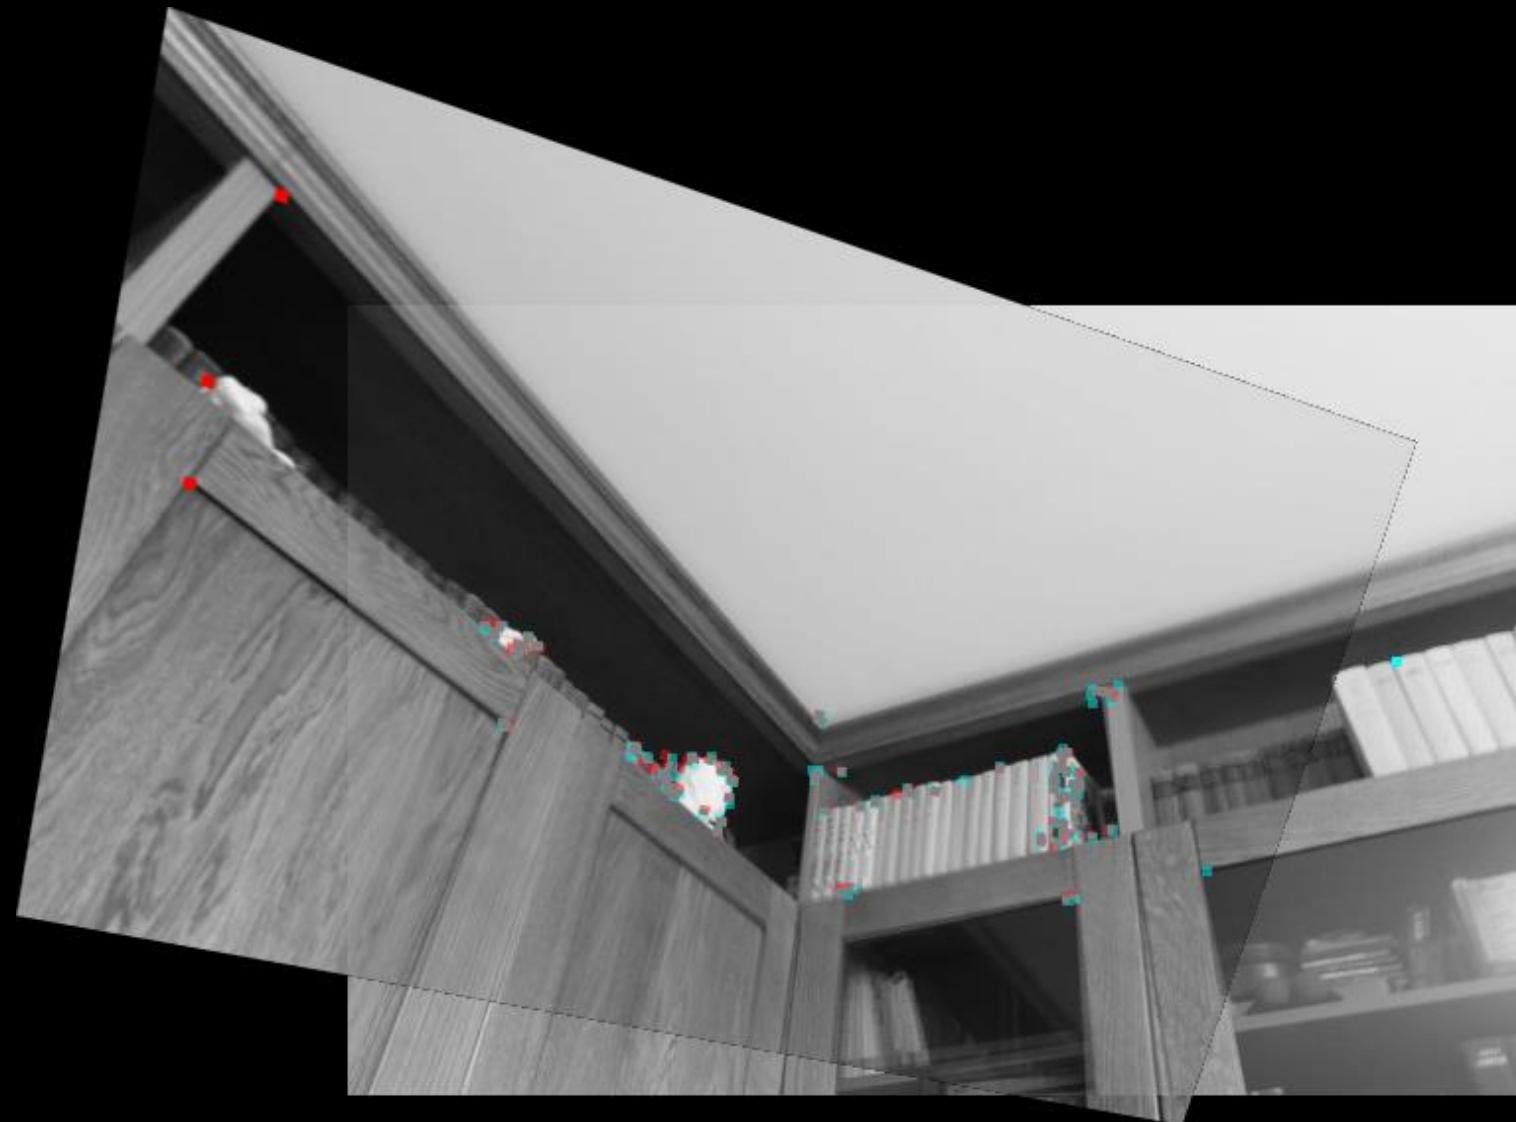
Blending: average



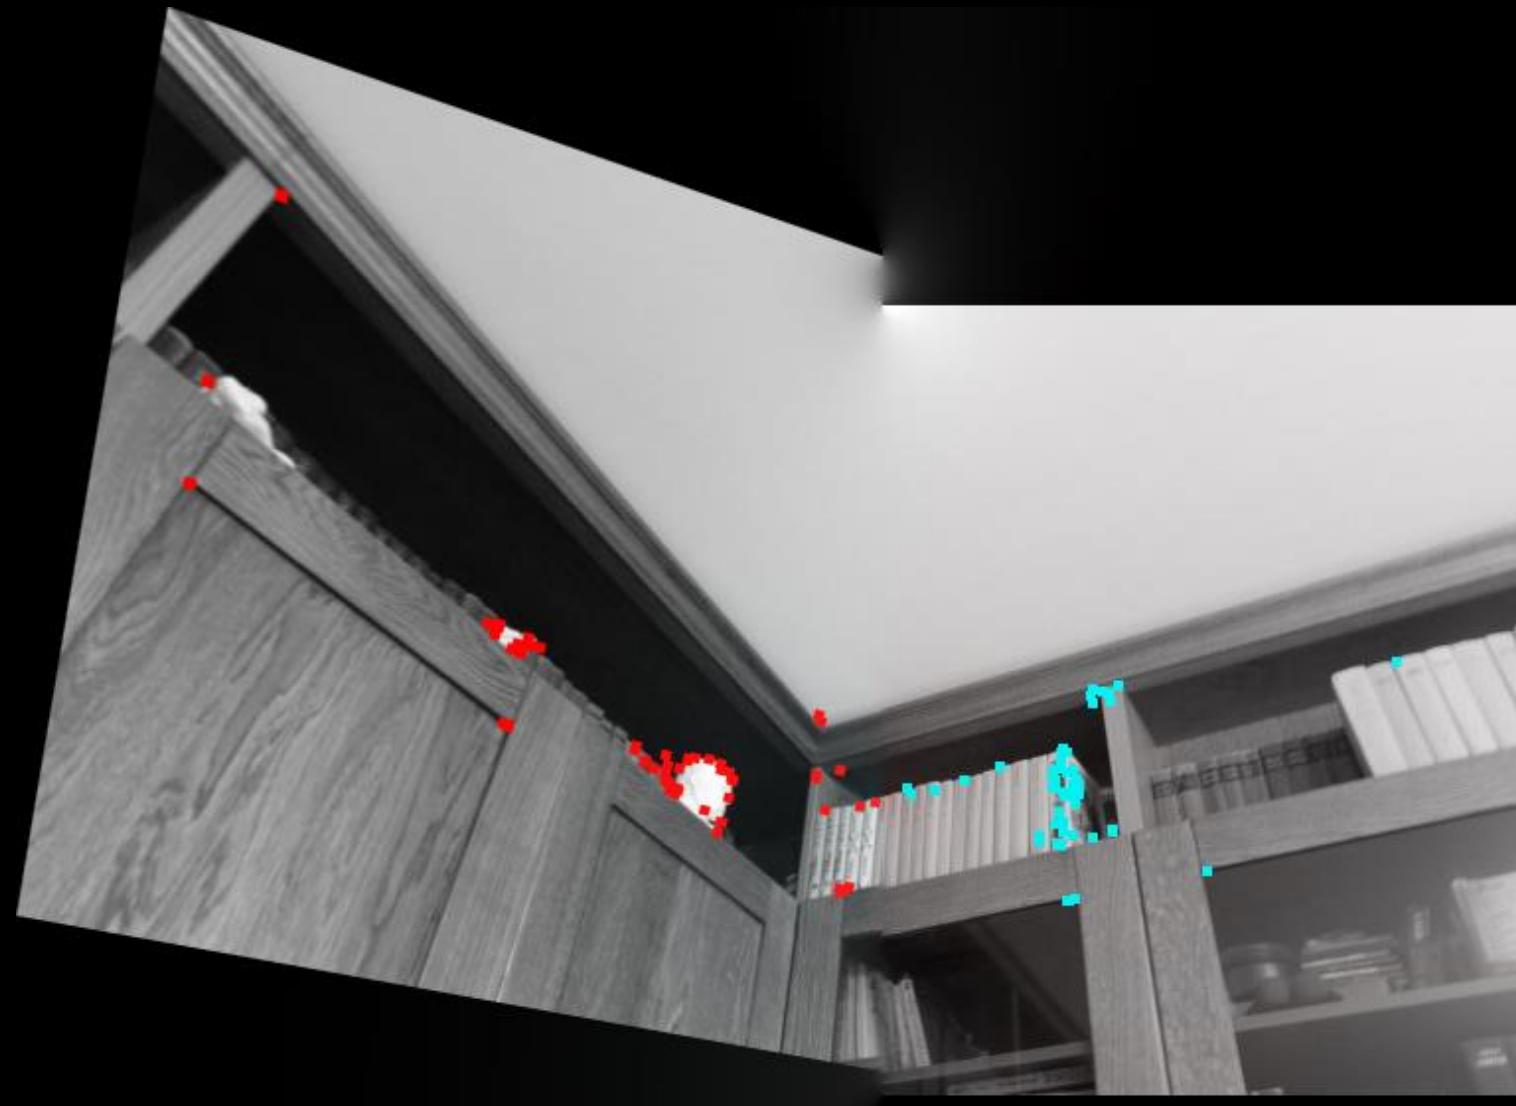
Blending: average



Blending: average



Blending: average

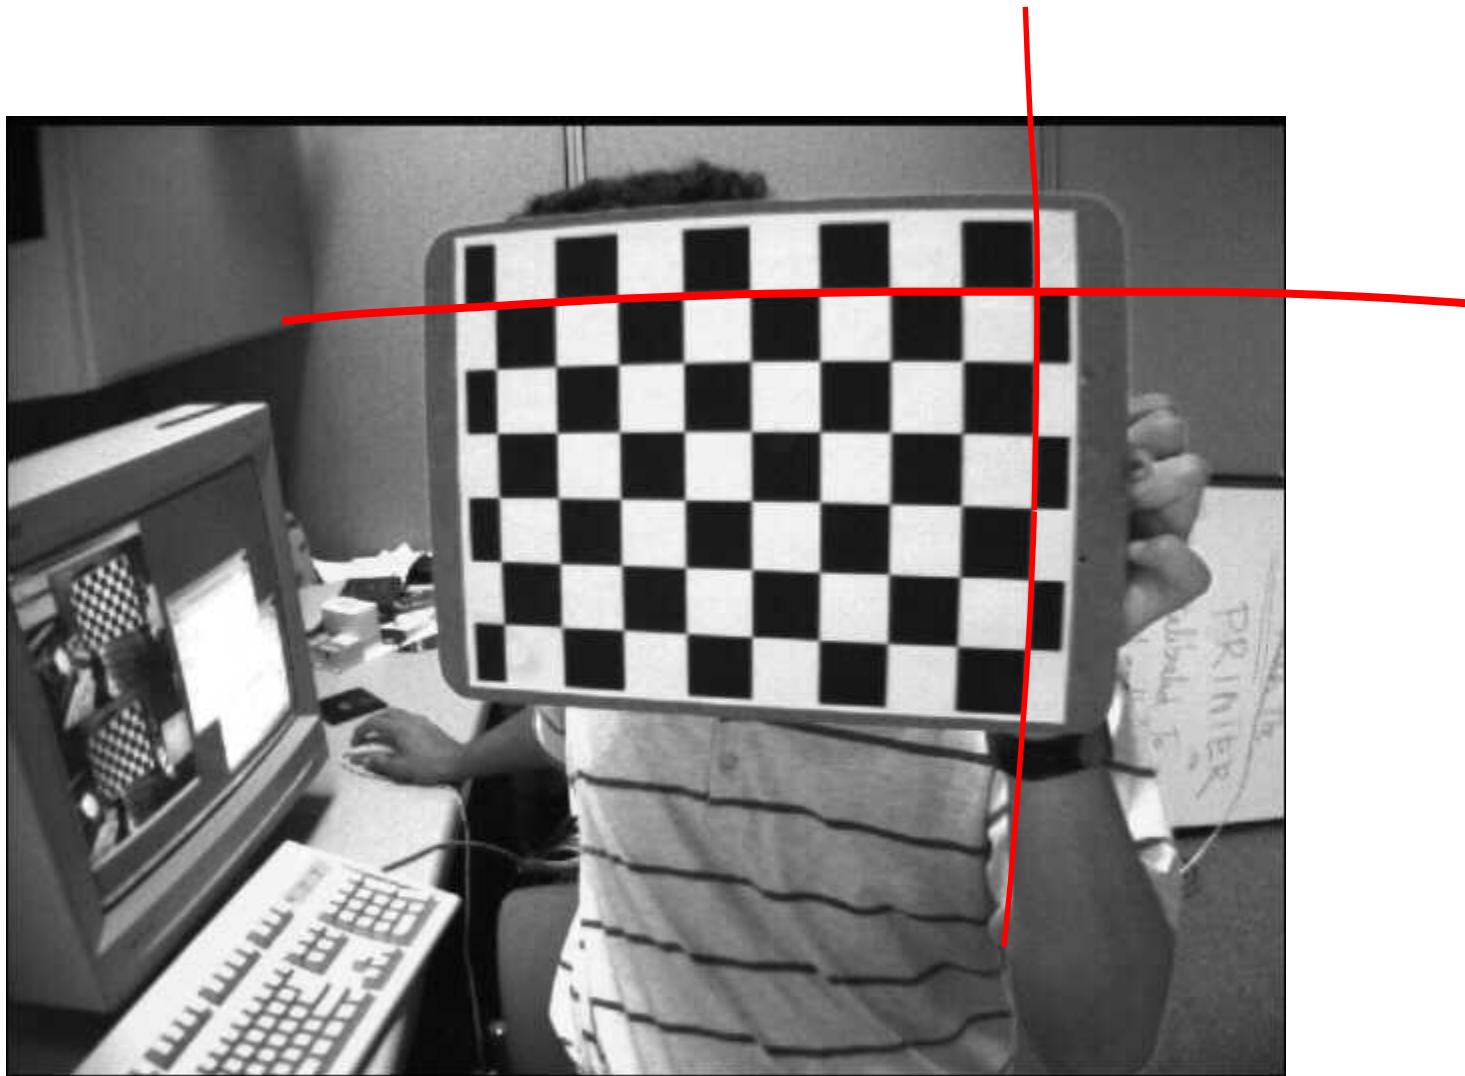


Camera calibration

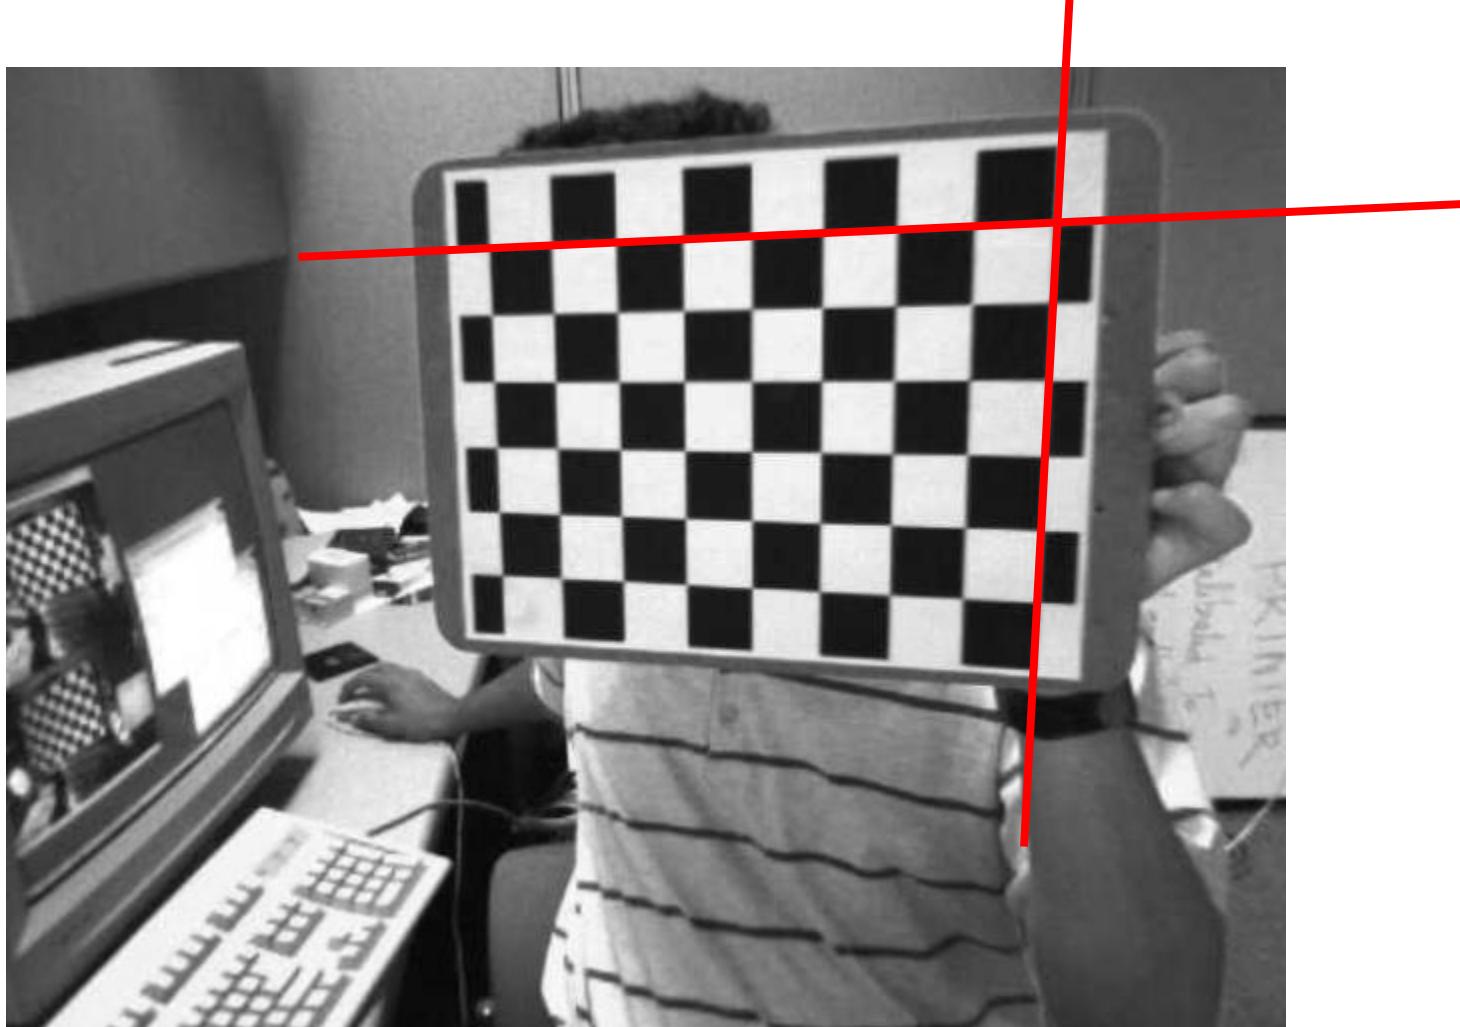
https://github.com/dryabokon/geometry/blob/master/ex06_camera_calibration.py



Camera calibration



Camera calibration

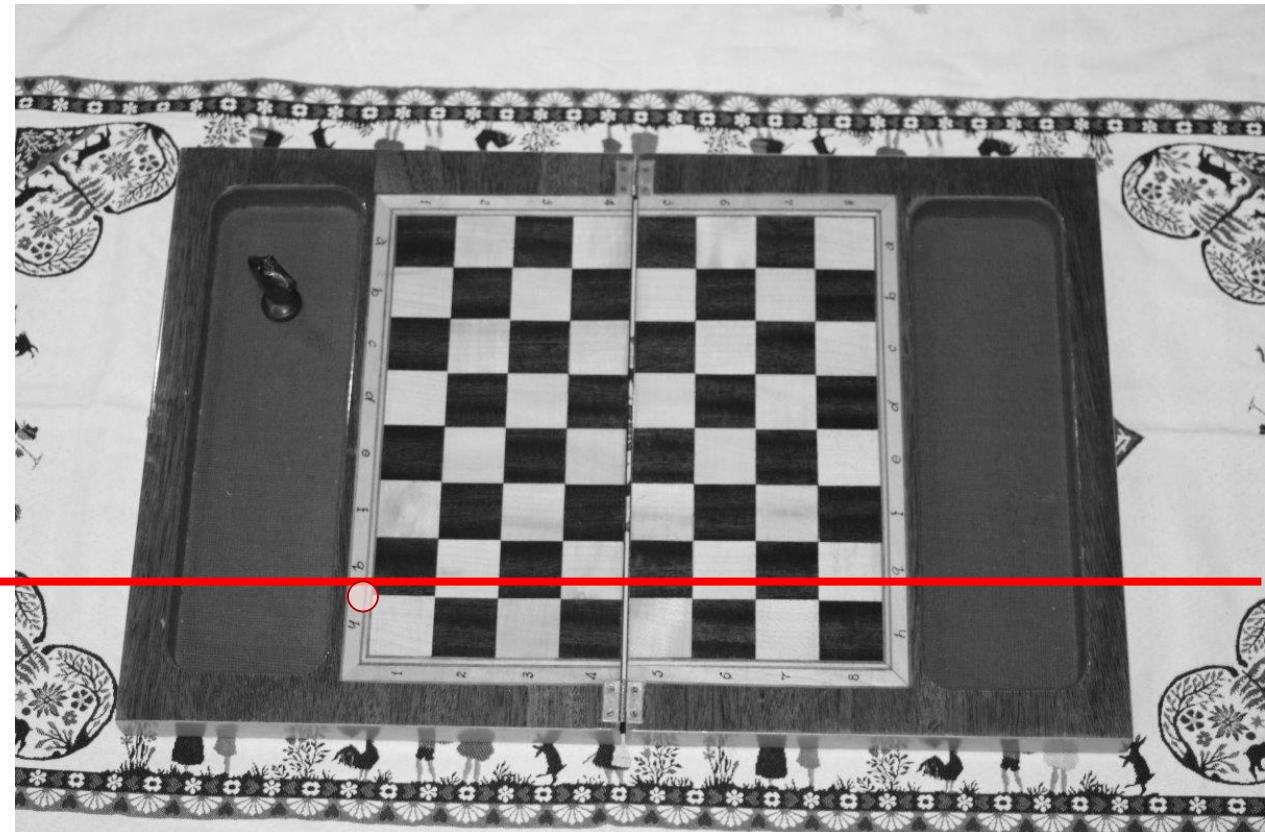
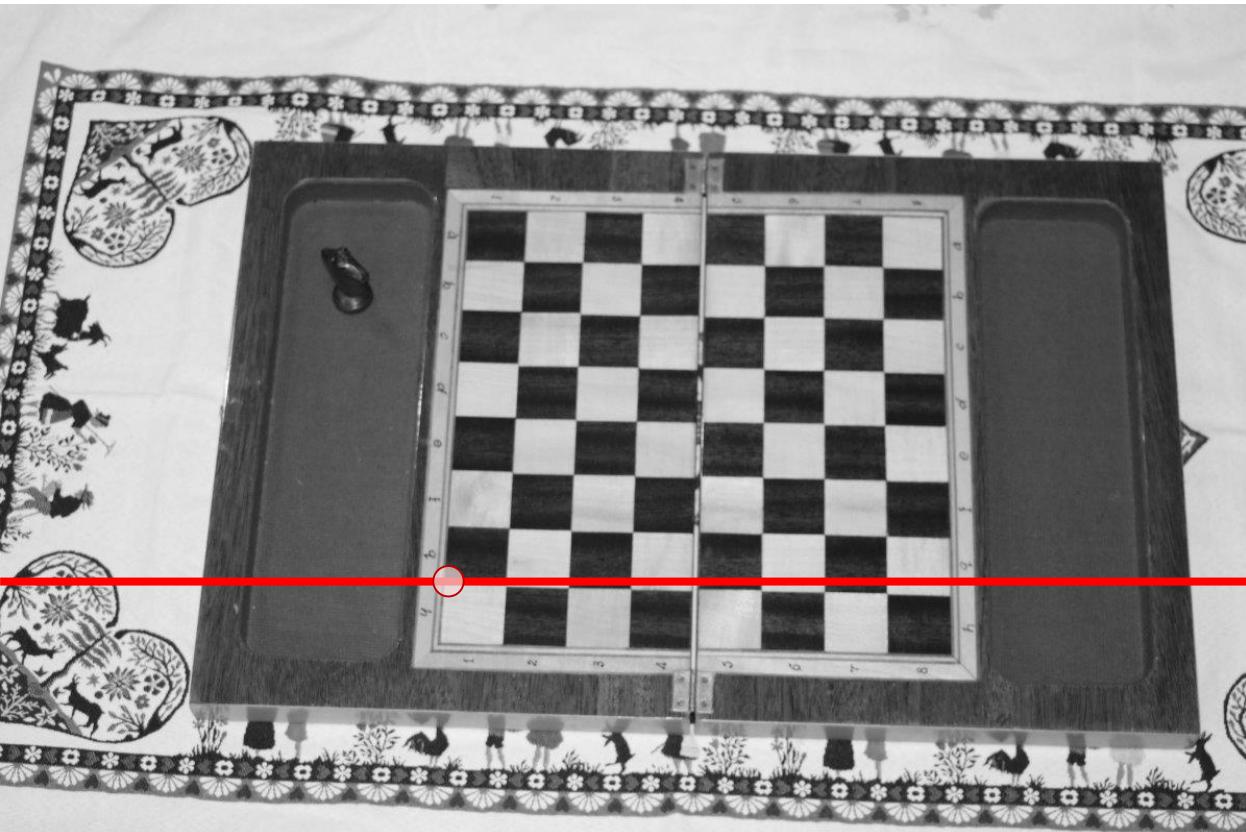


Stereopair rectification

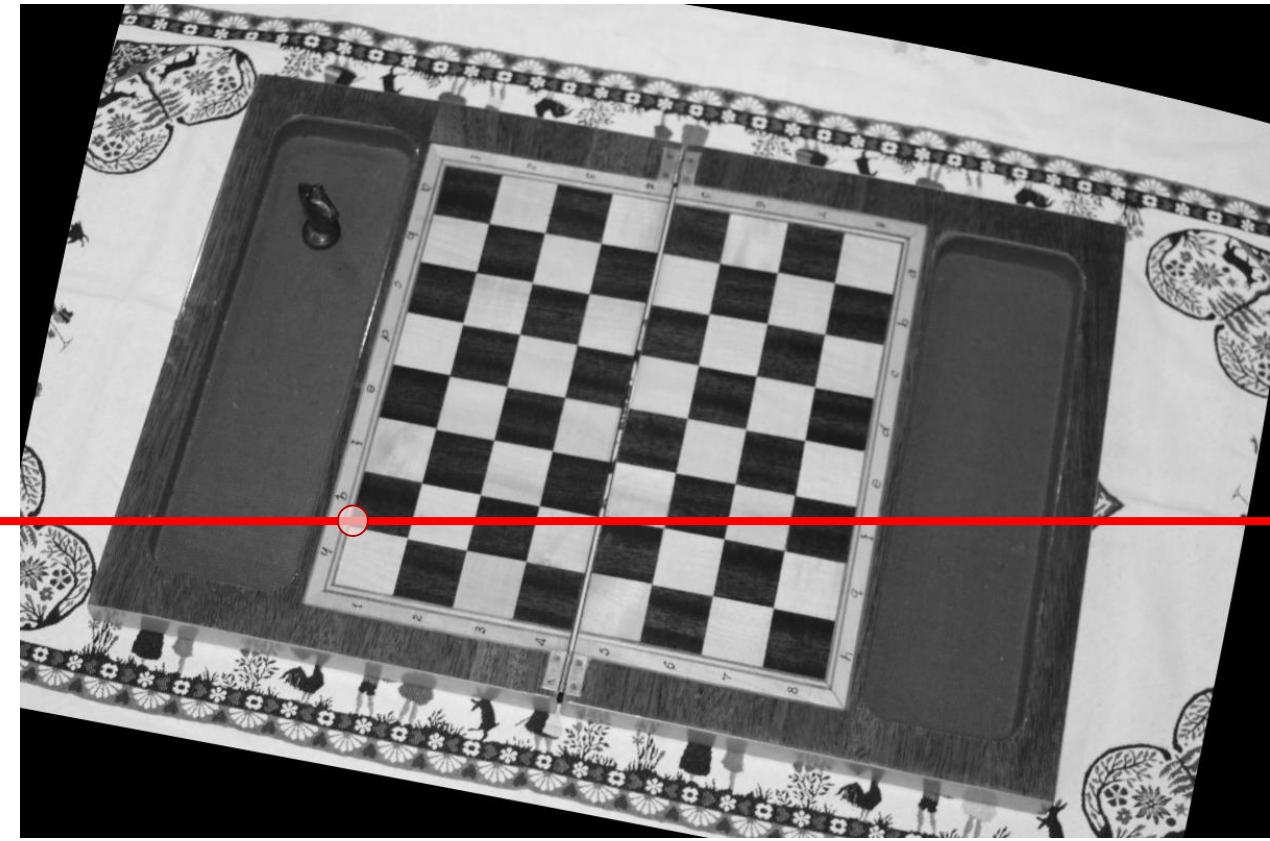
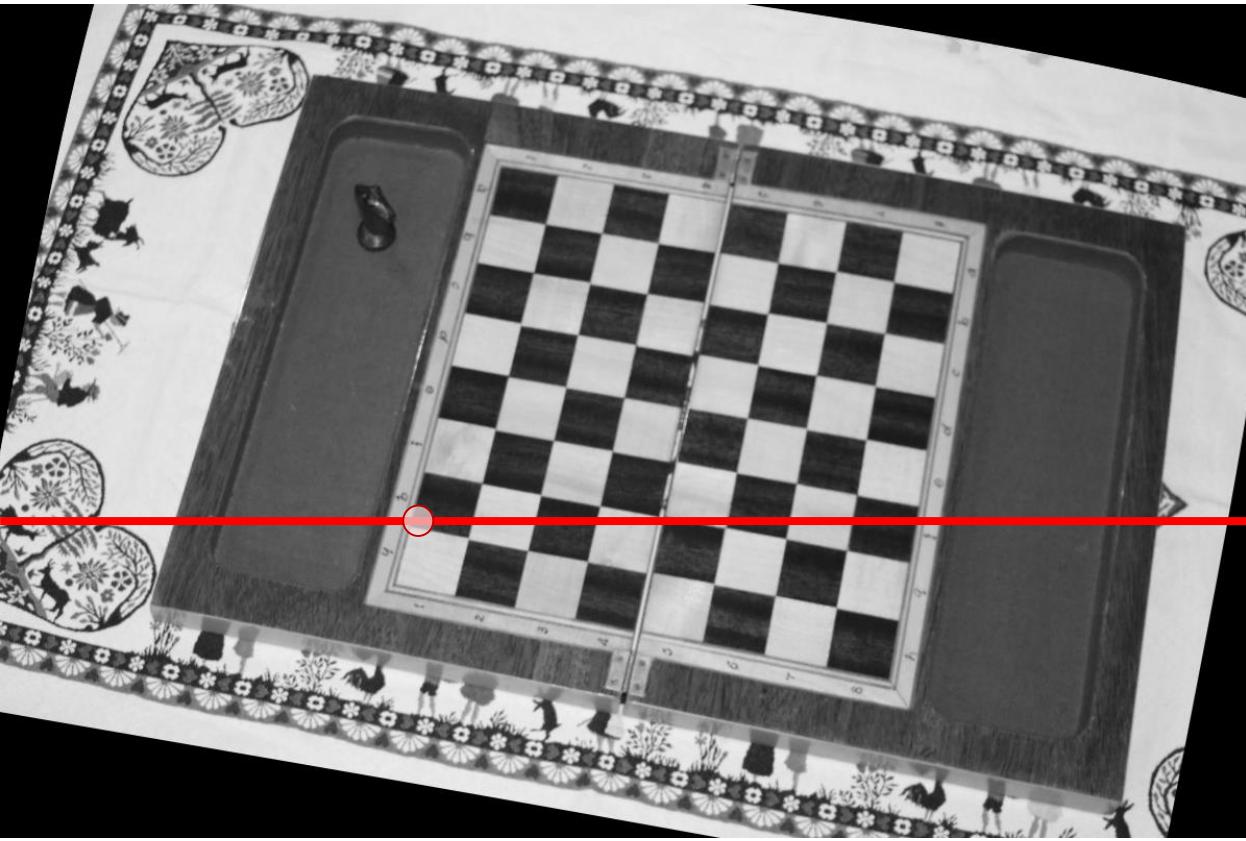


https://github.com/dryabokon/geometry/blob/master/ex07_rectify_stereopair.py

Stereopair rectification



Stereopair rectification



Template matching



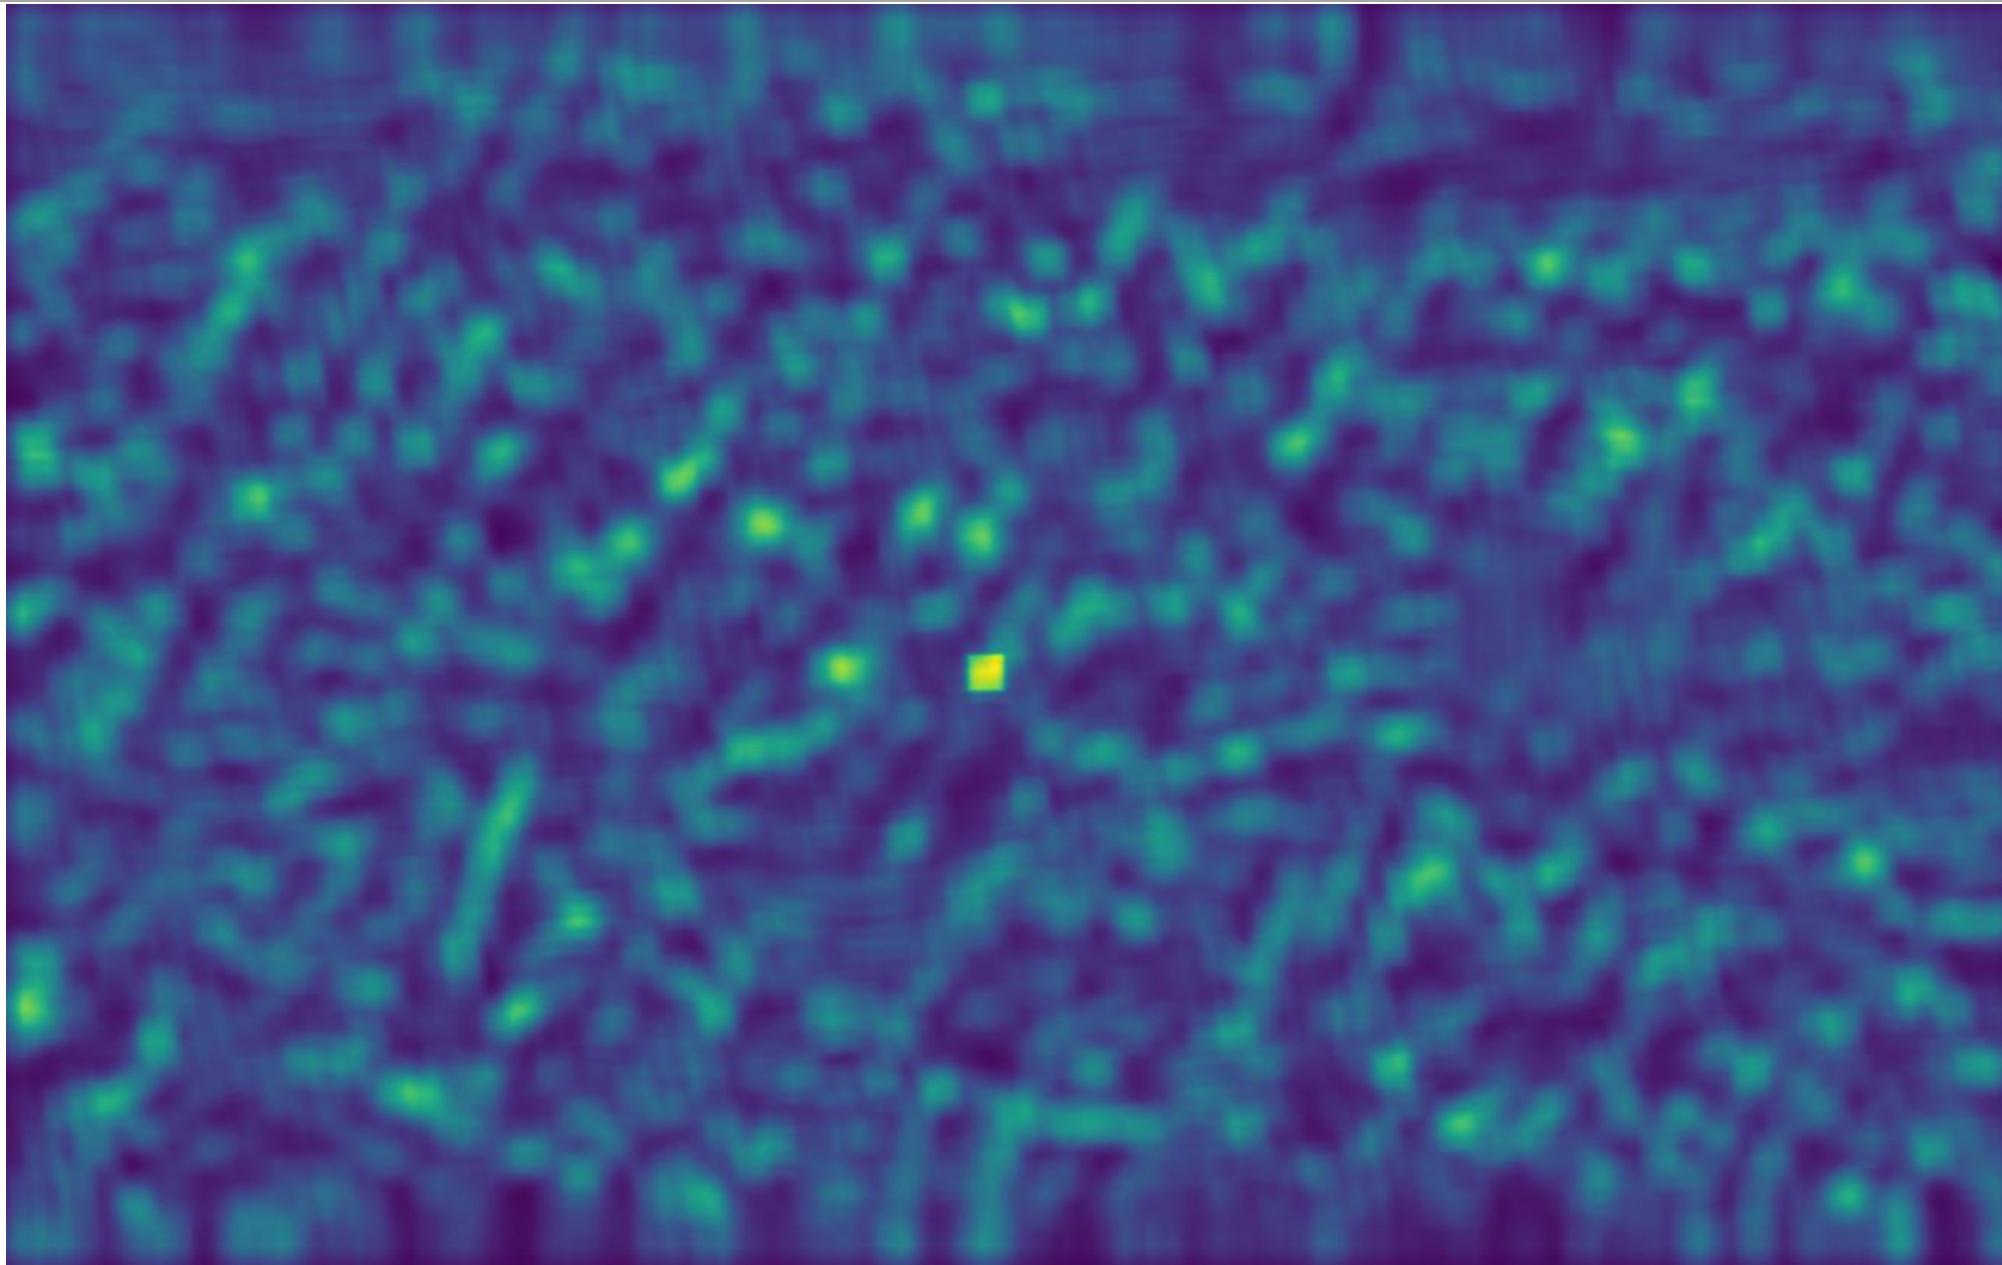
Template Matching



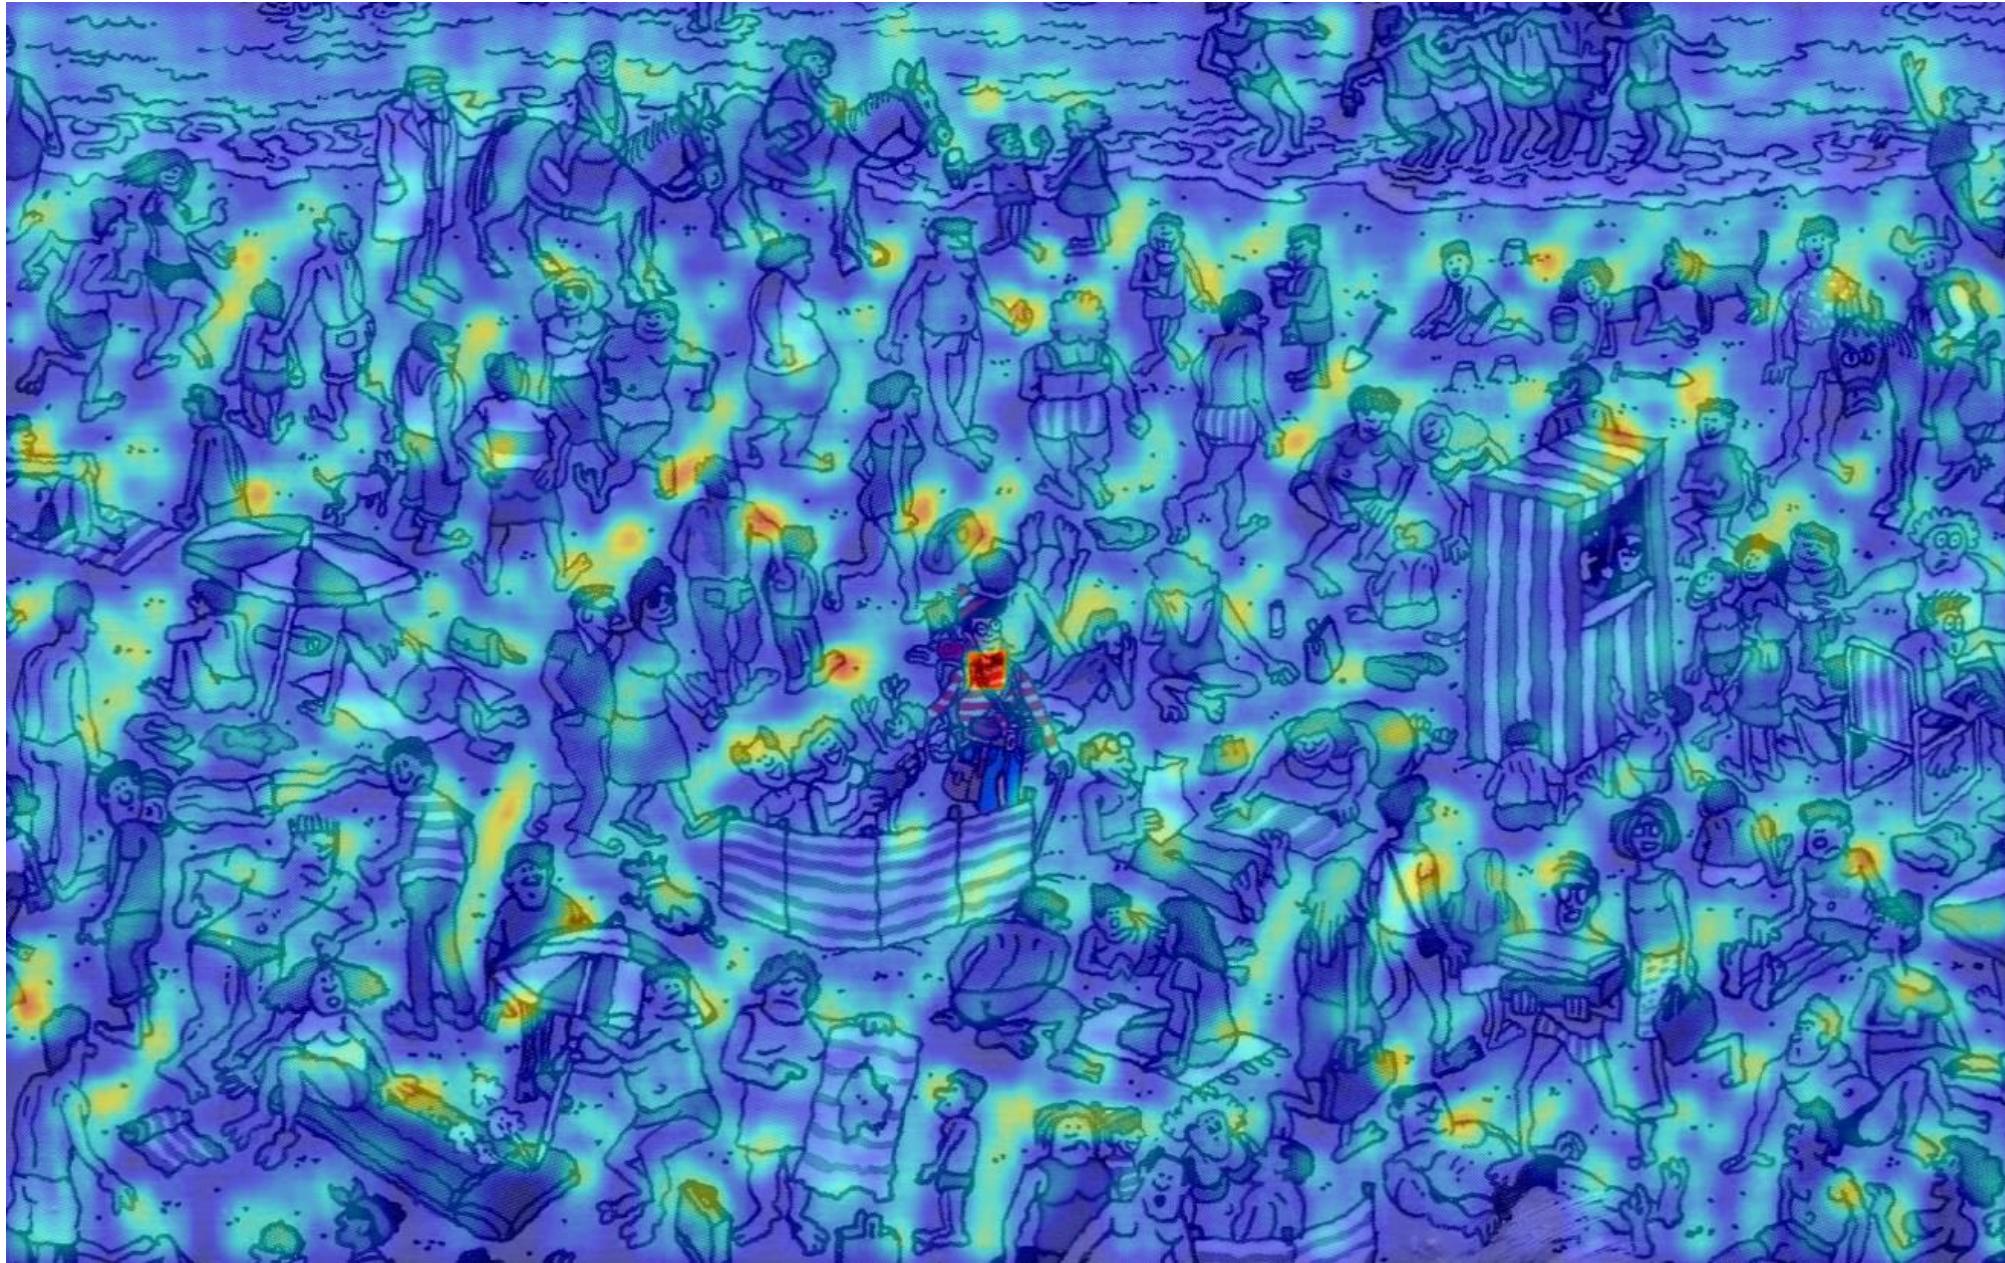
Template Matching



Template Matching



Template Matching



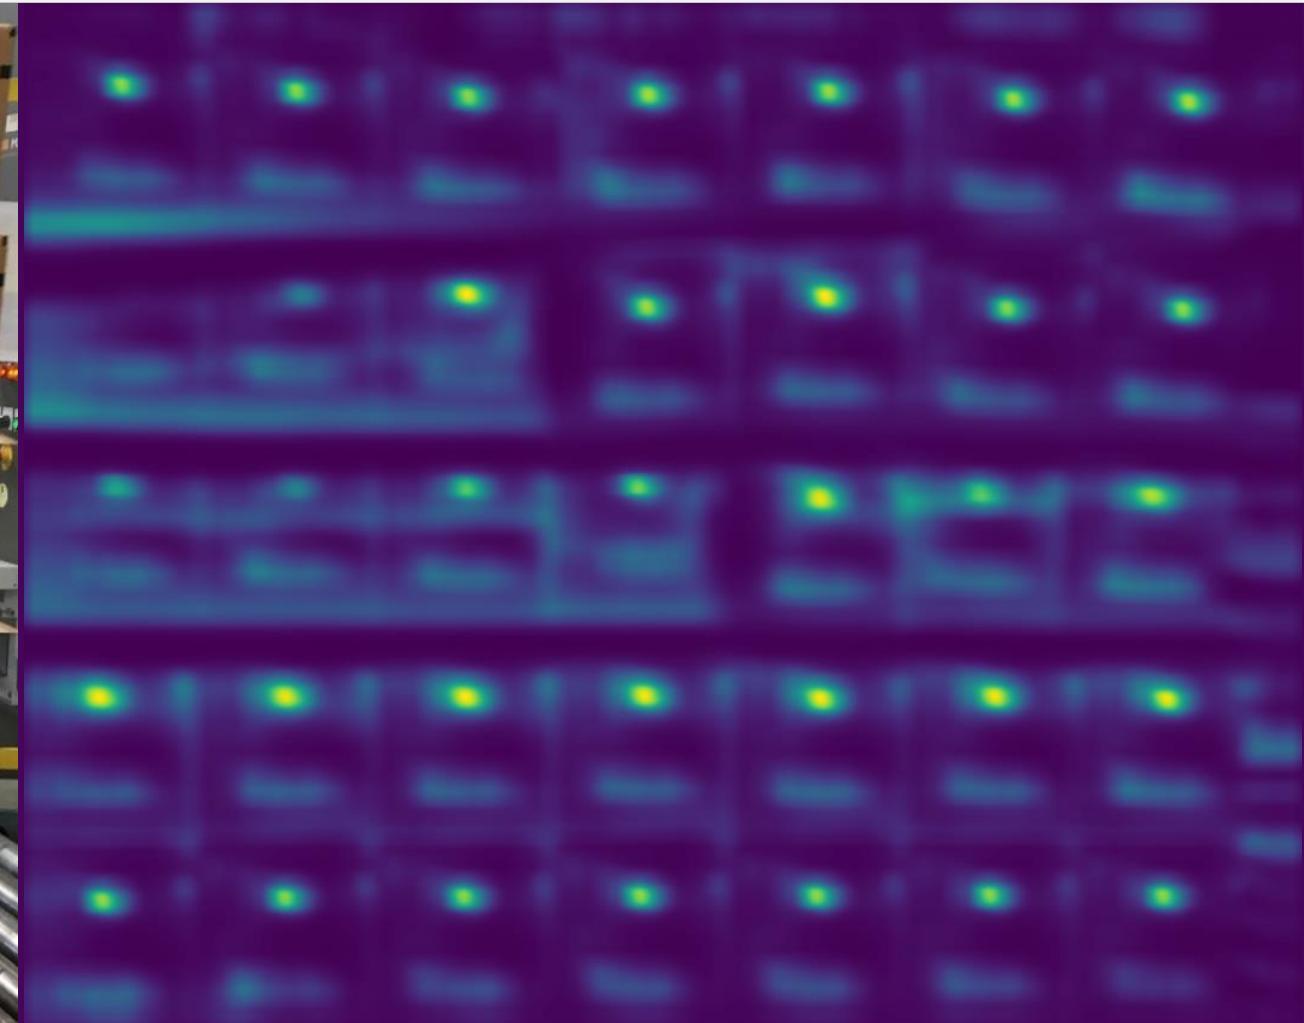
Template Matching



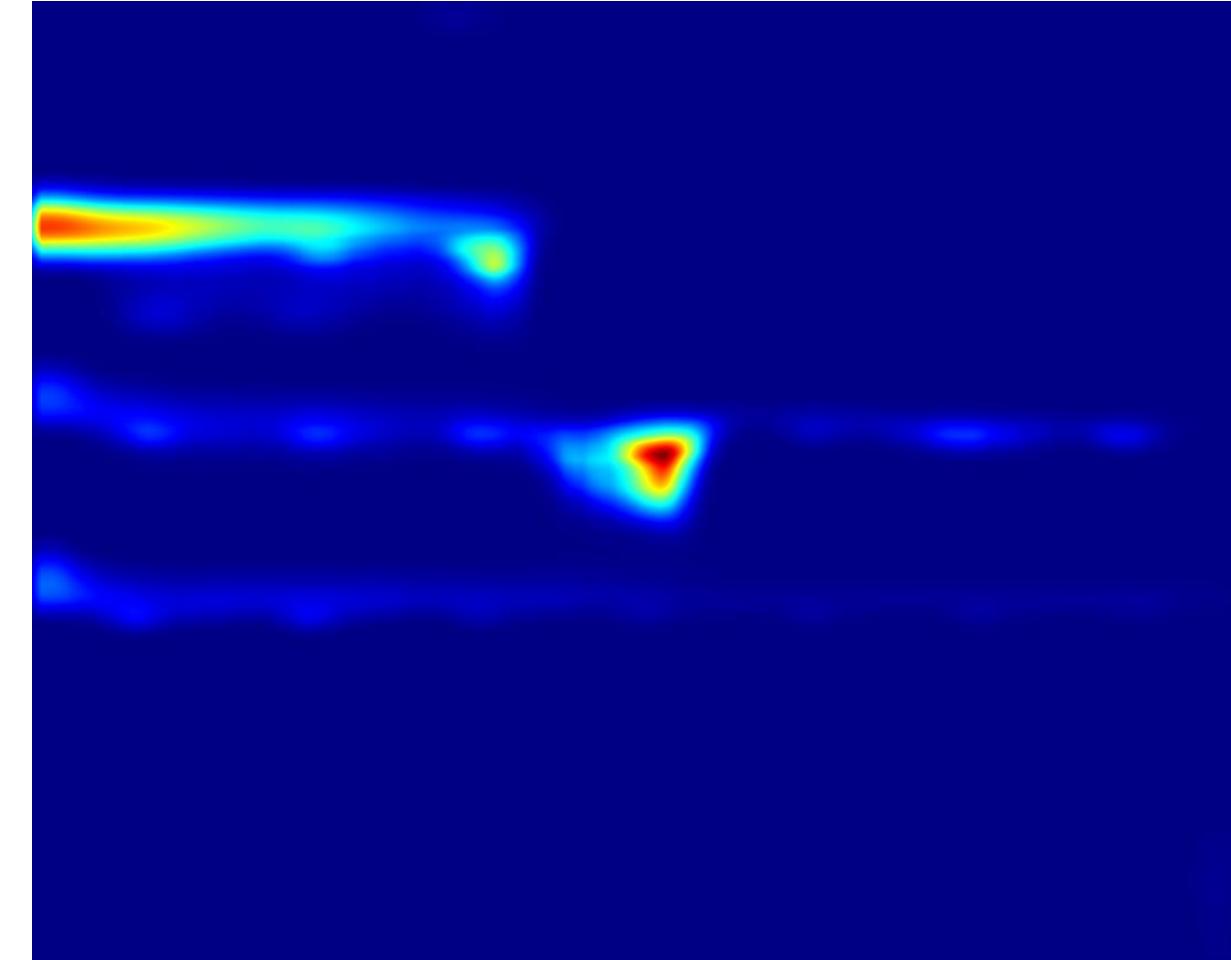
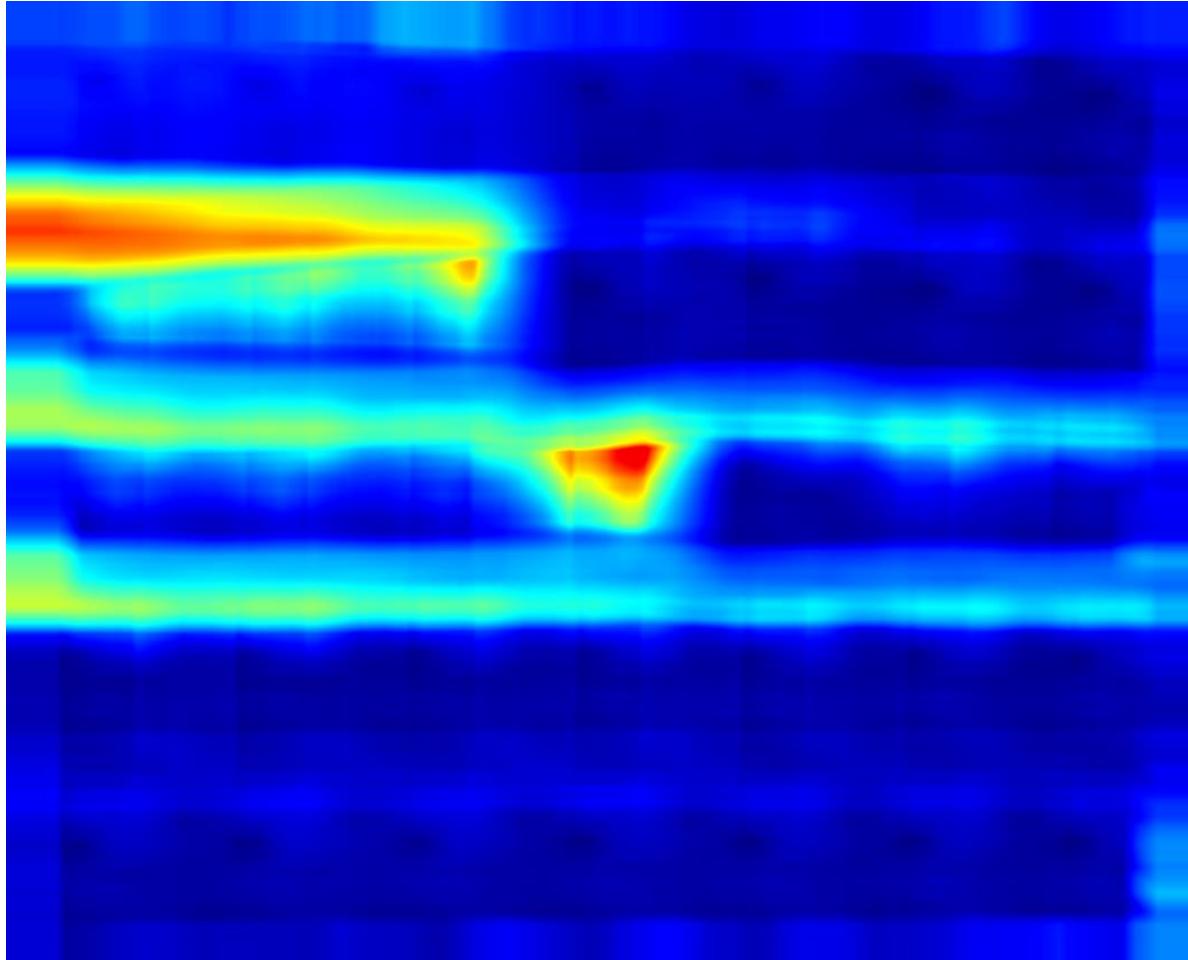
Template Matching



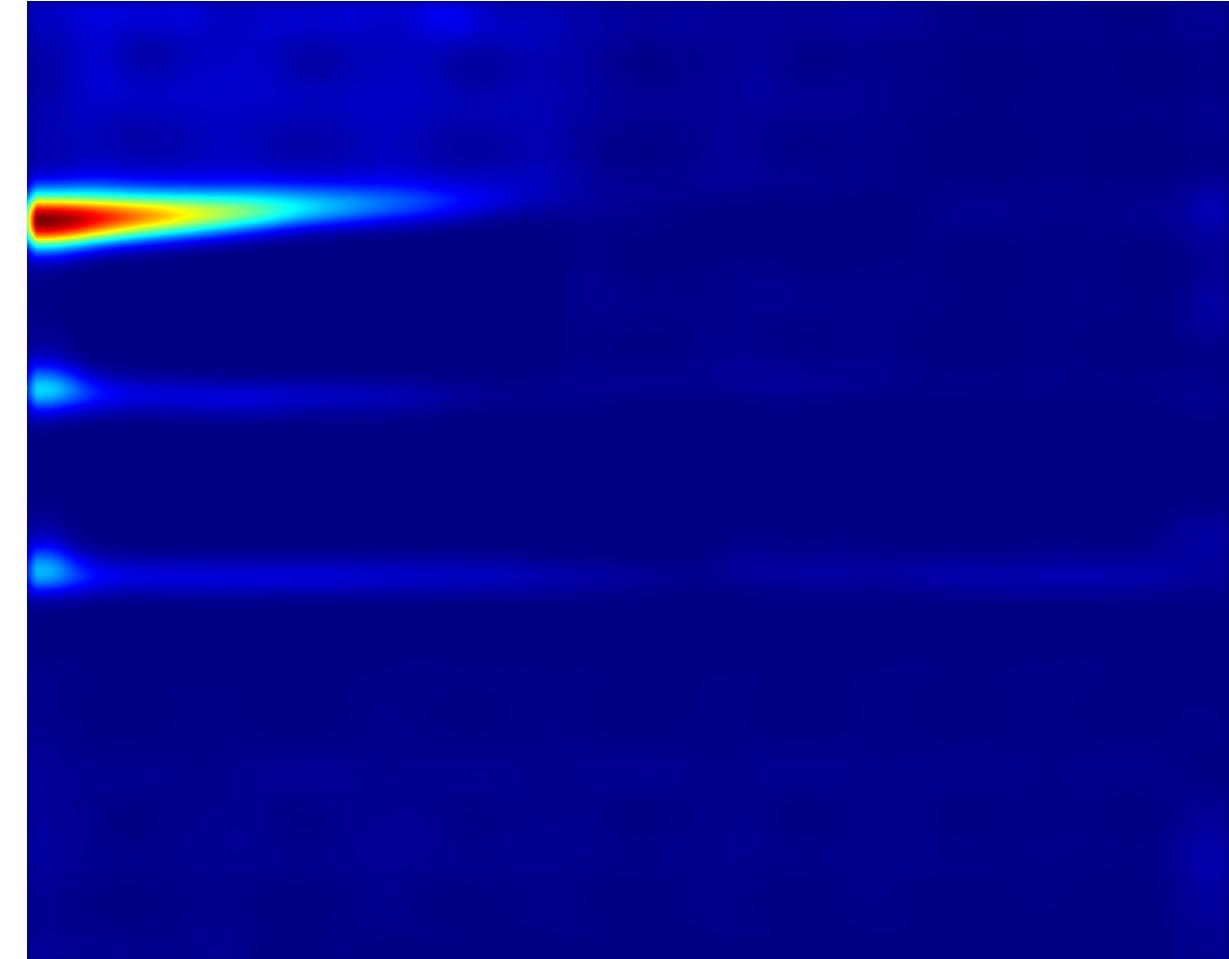
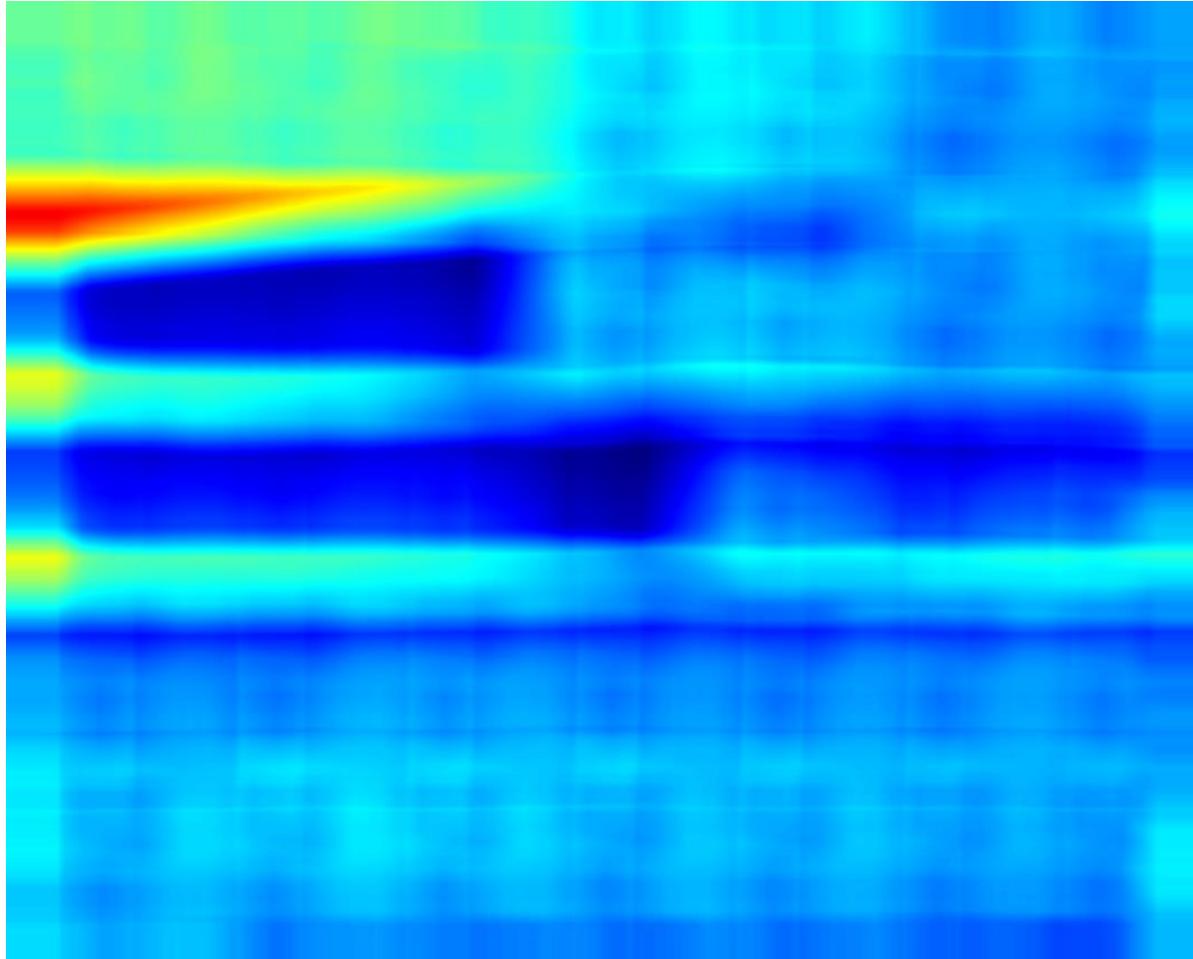
Template Matching



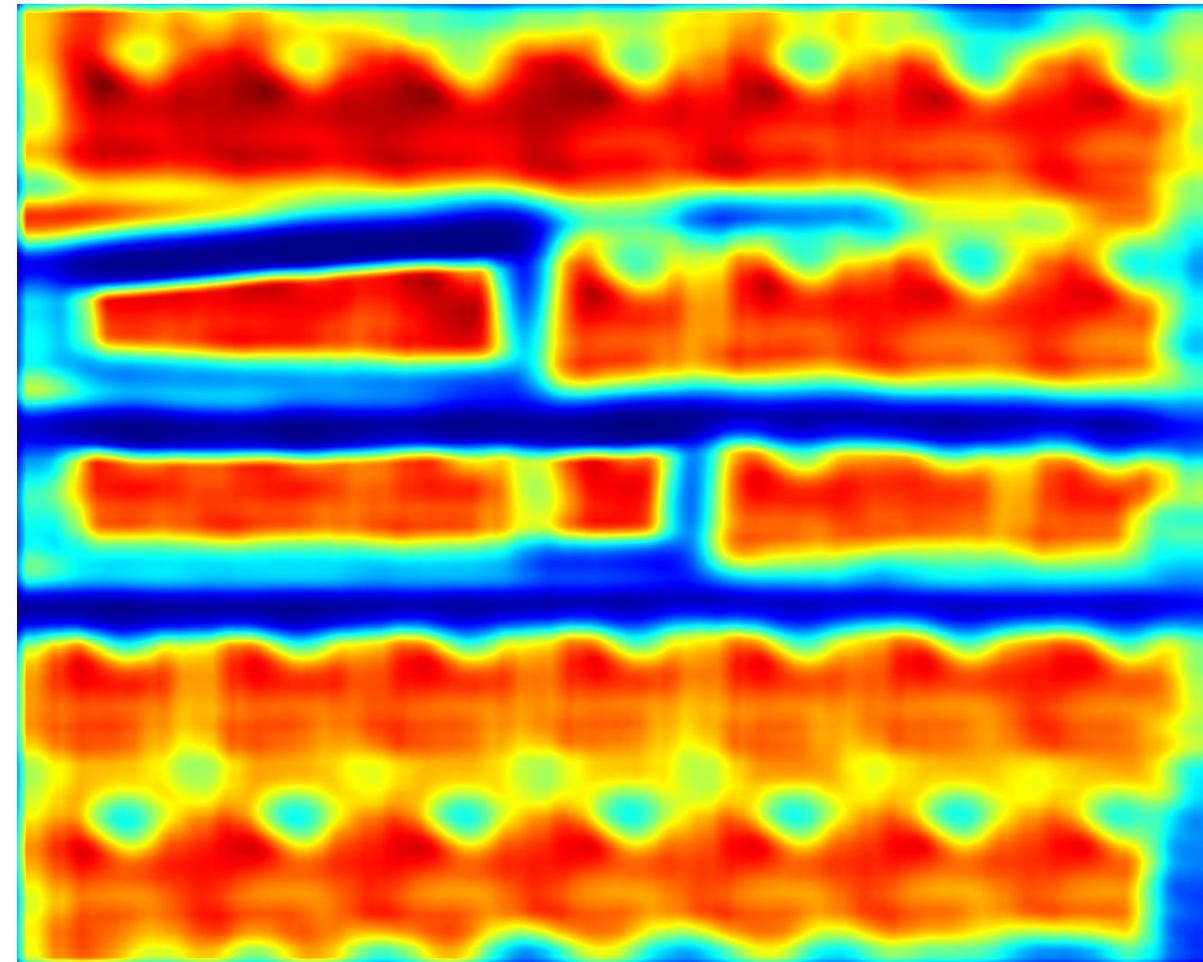
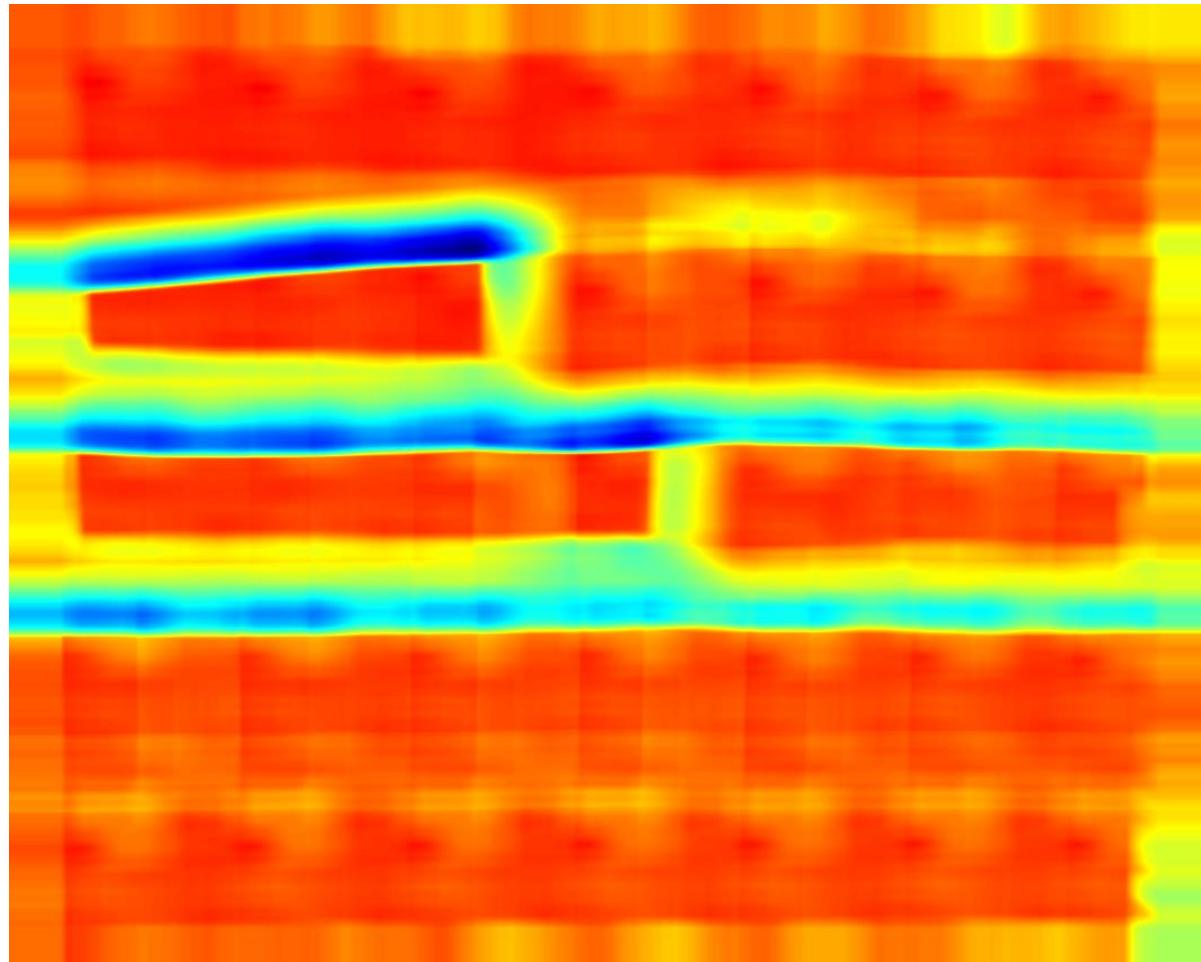
Template Matching: TM_SQDIFF_NORMED



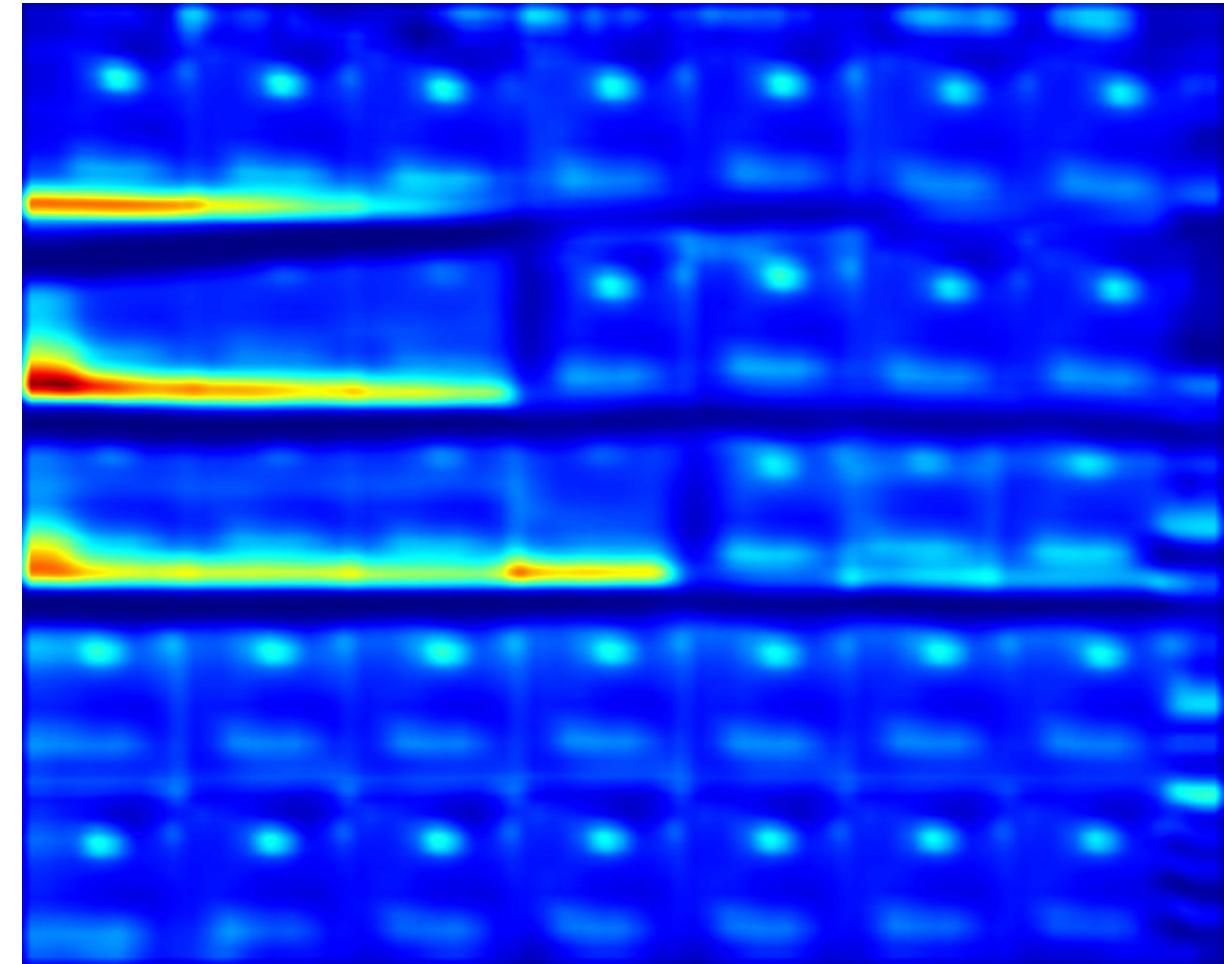
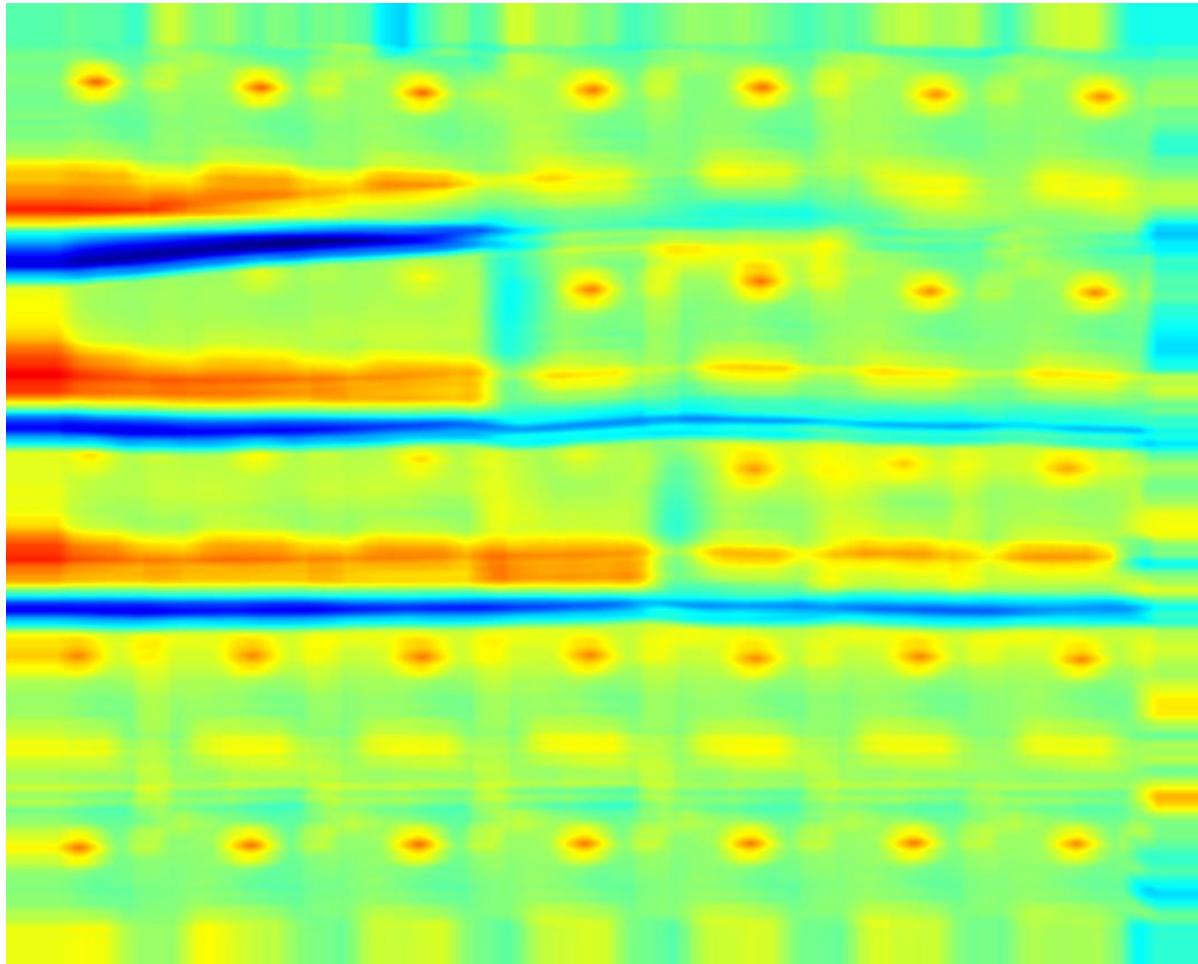
Template Matching: TM_CCORR



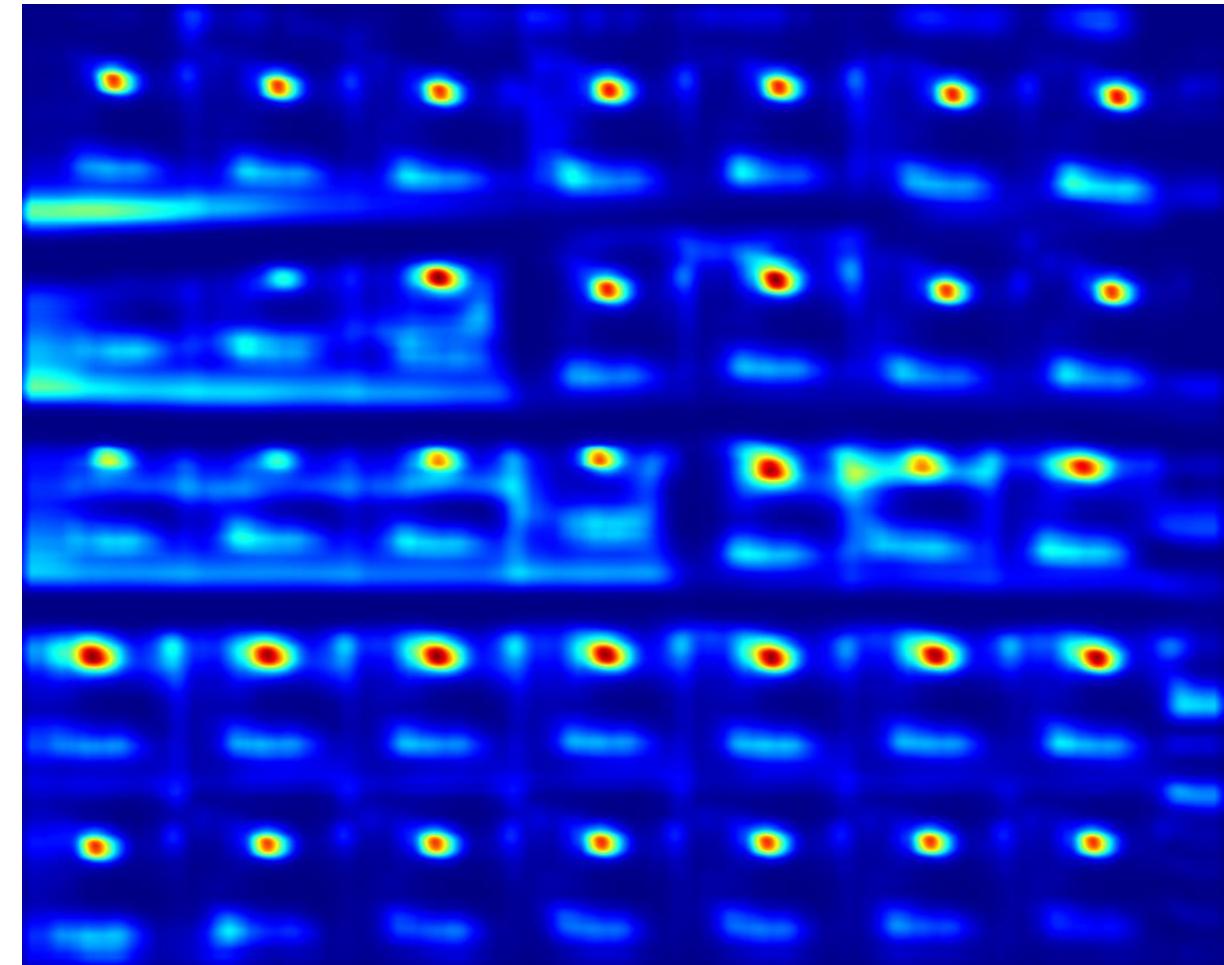
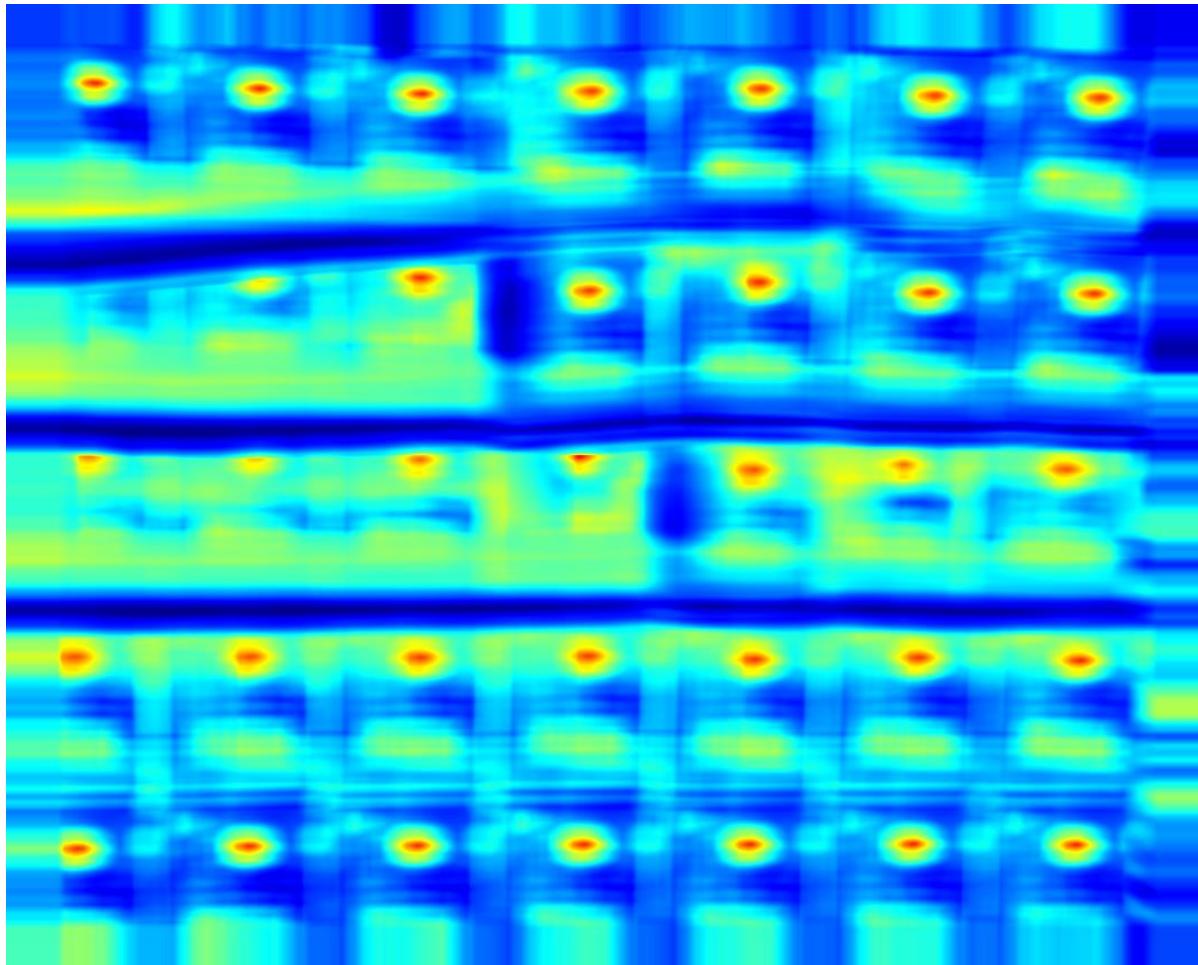
Template Matching: TM_CCORR_NORMED



Template Matching: TM_CCOEFF



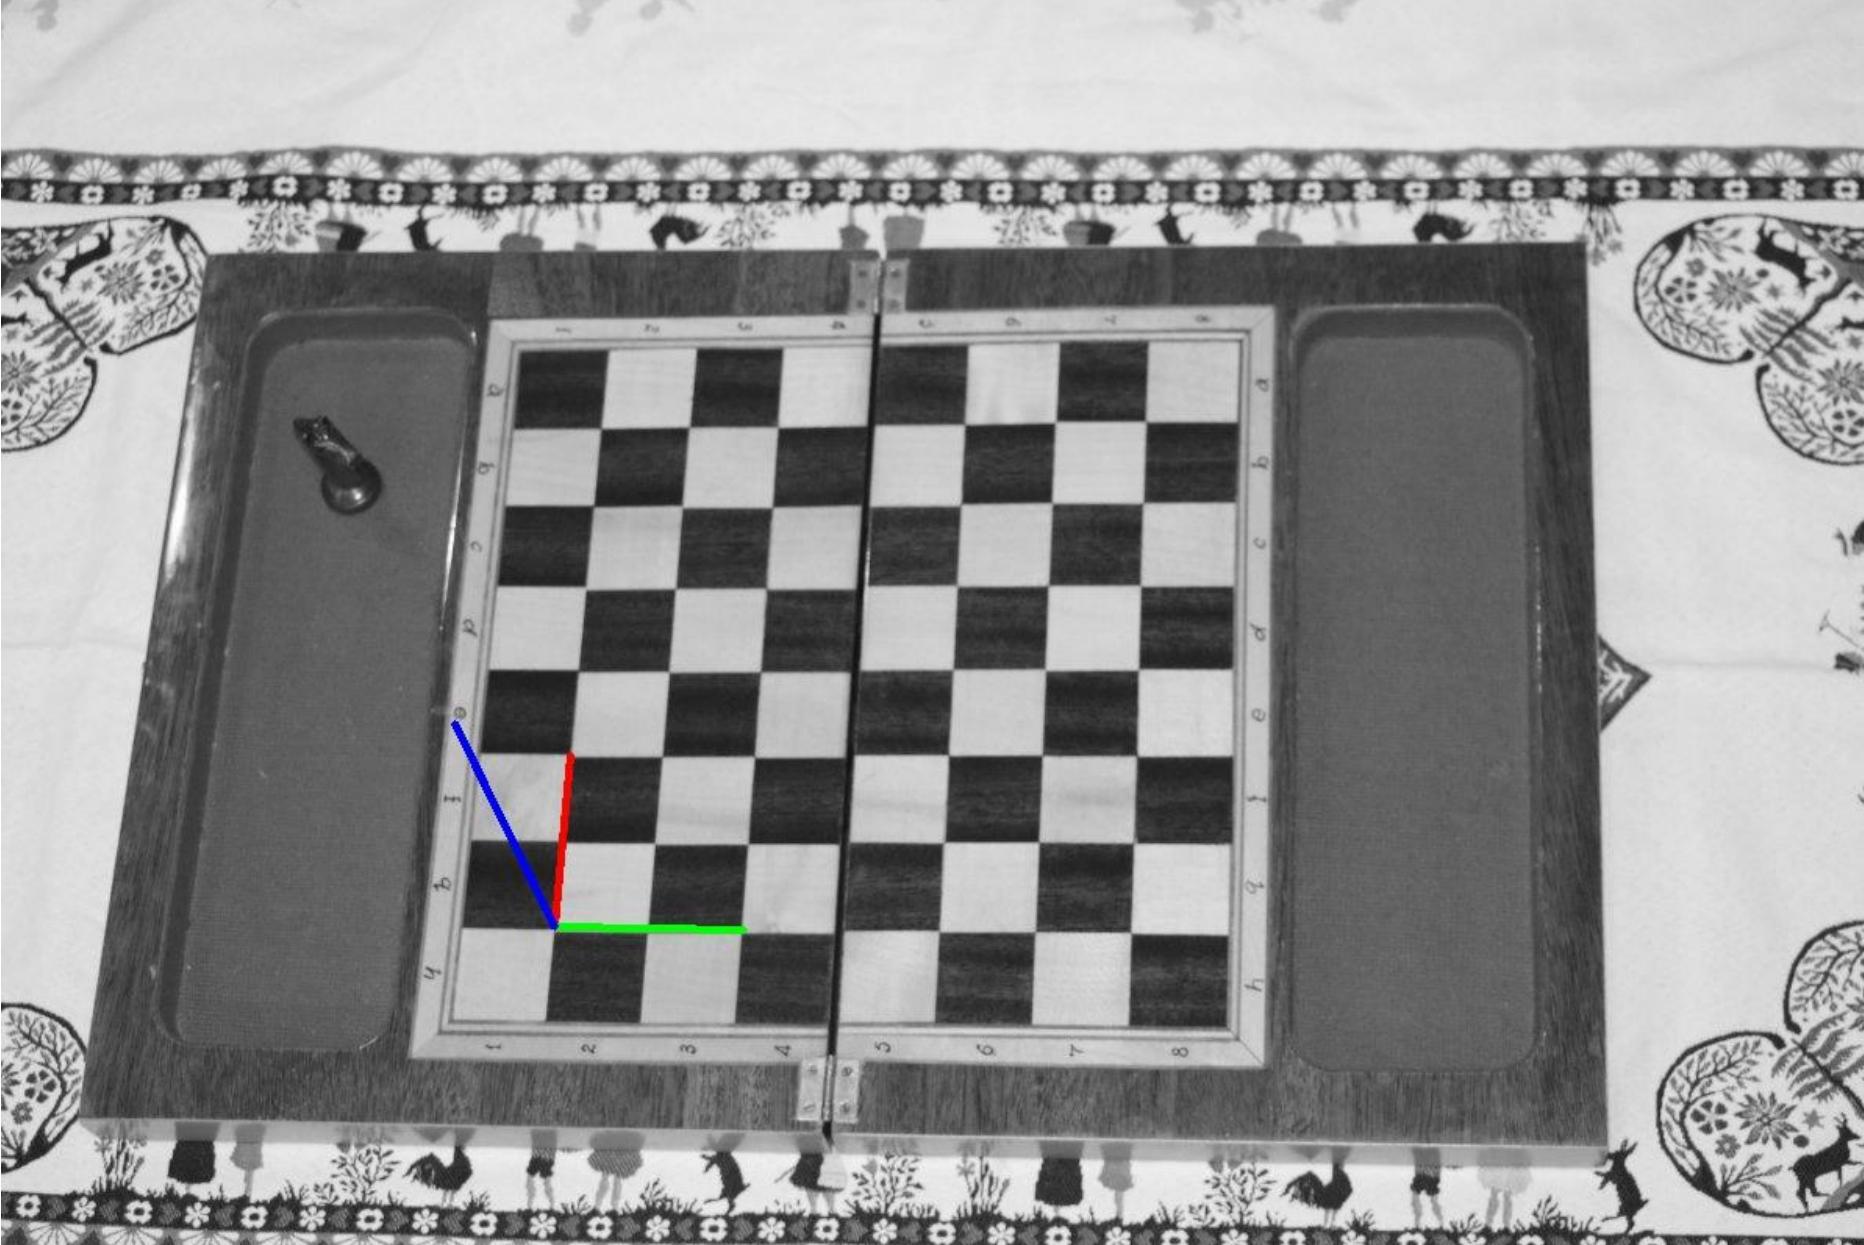
Template Matching: TM_CCOEFF_NORMED



Pose estimation



Pose estimation



Part 5: Computer Vision Applications

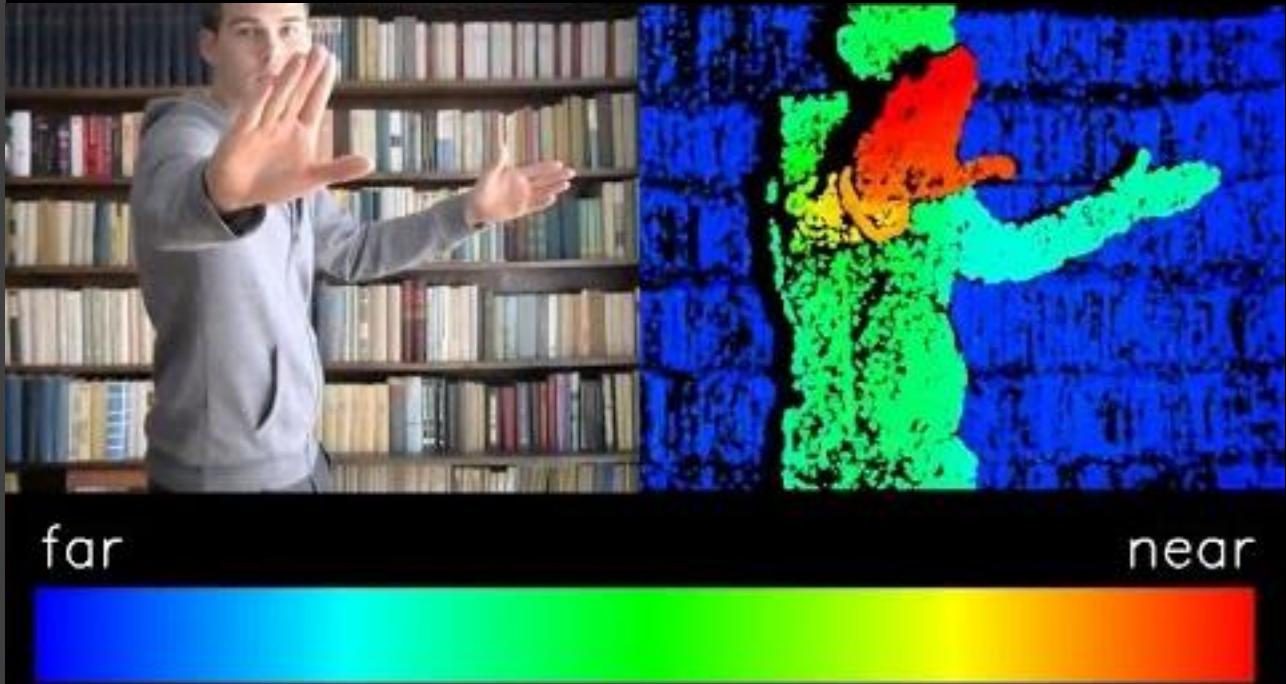
Stereo Vision

Face Swap

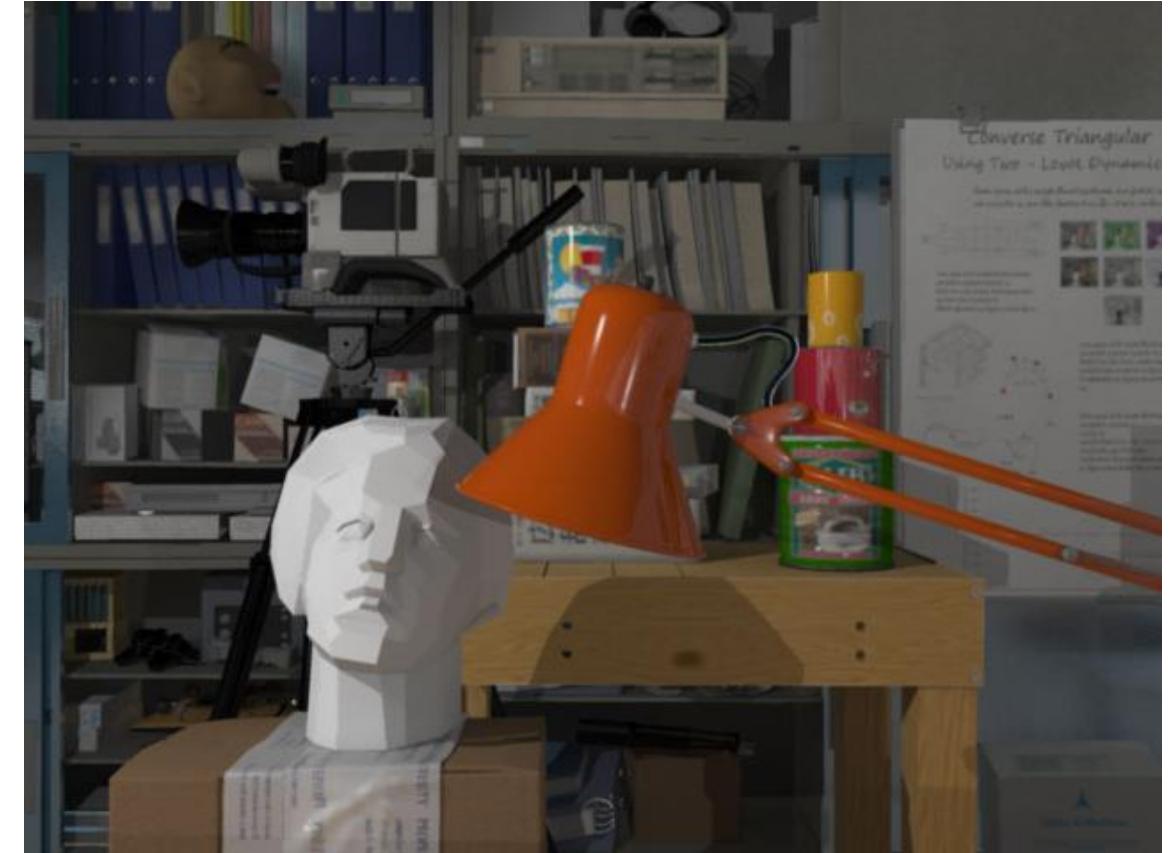
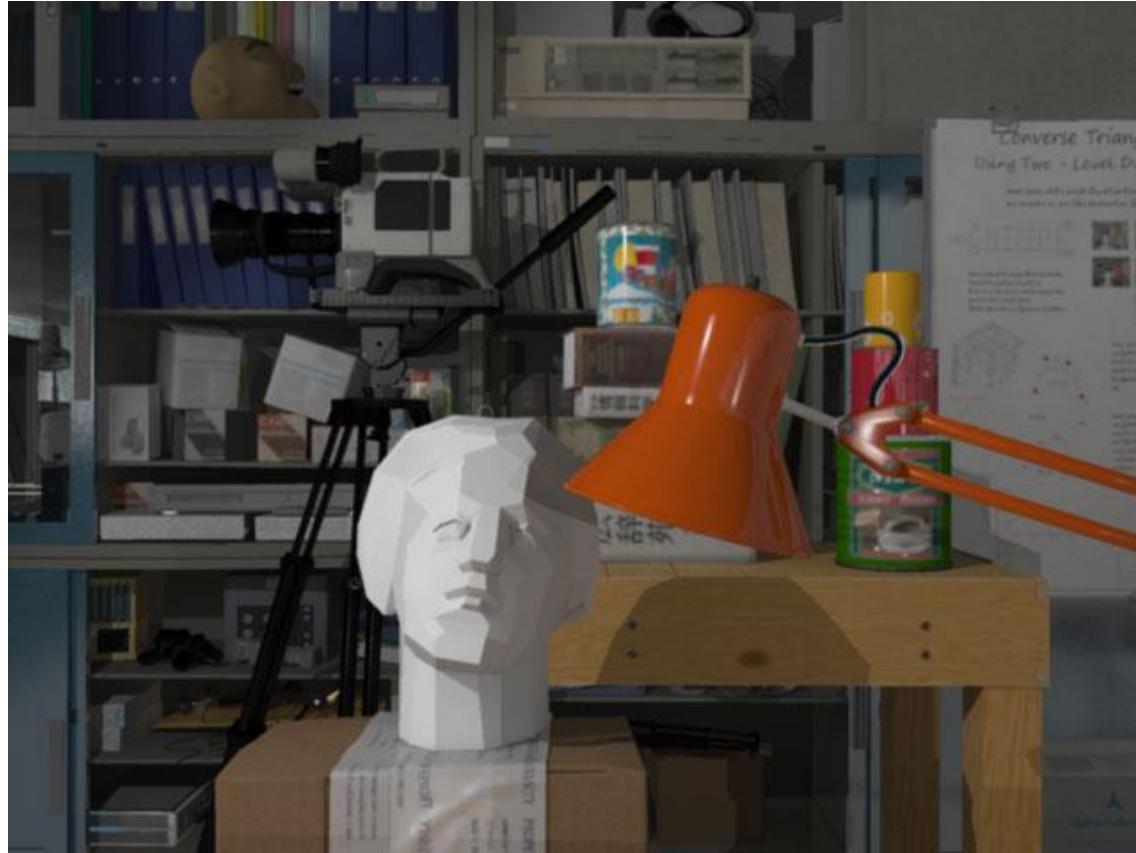
Face Filter

Style Transfer

Stereo Vision



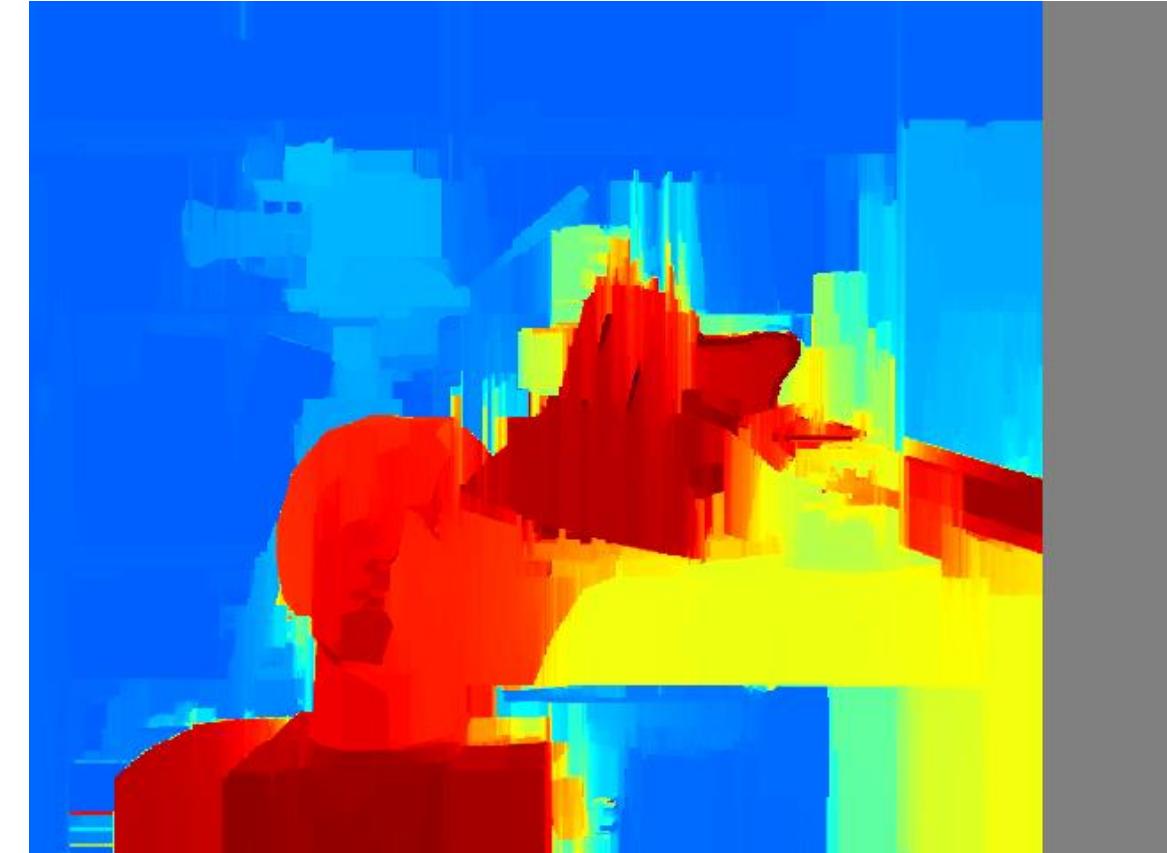
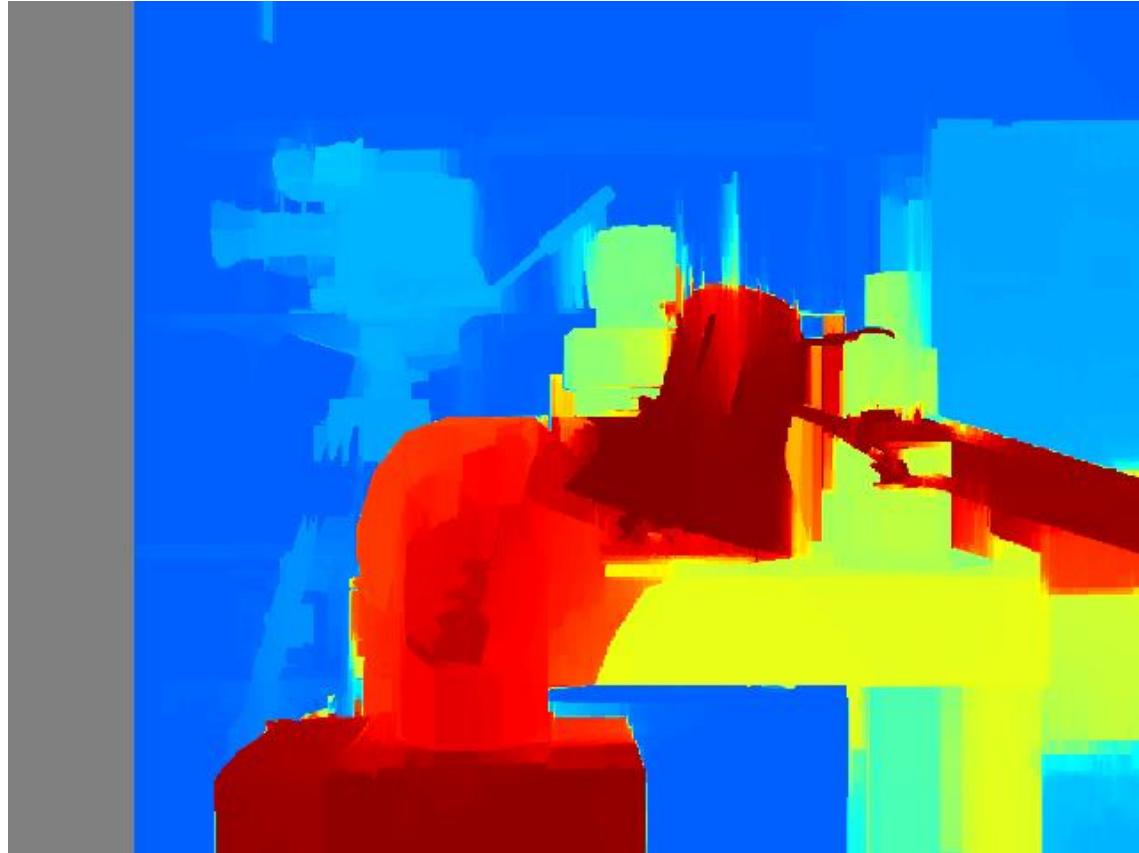
Stereo vision



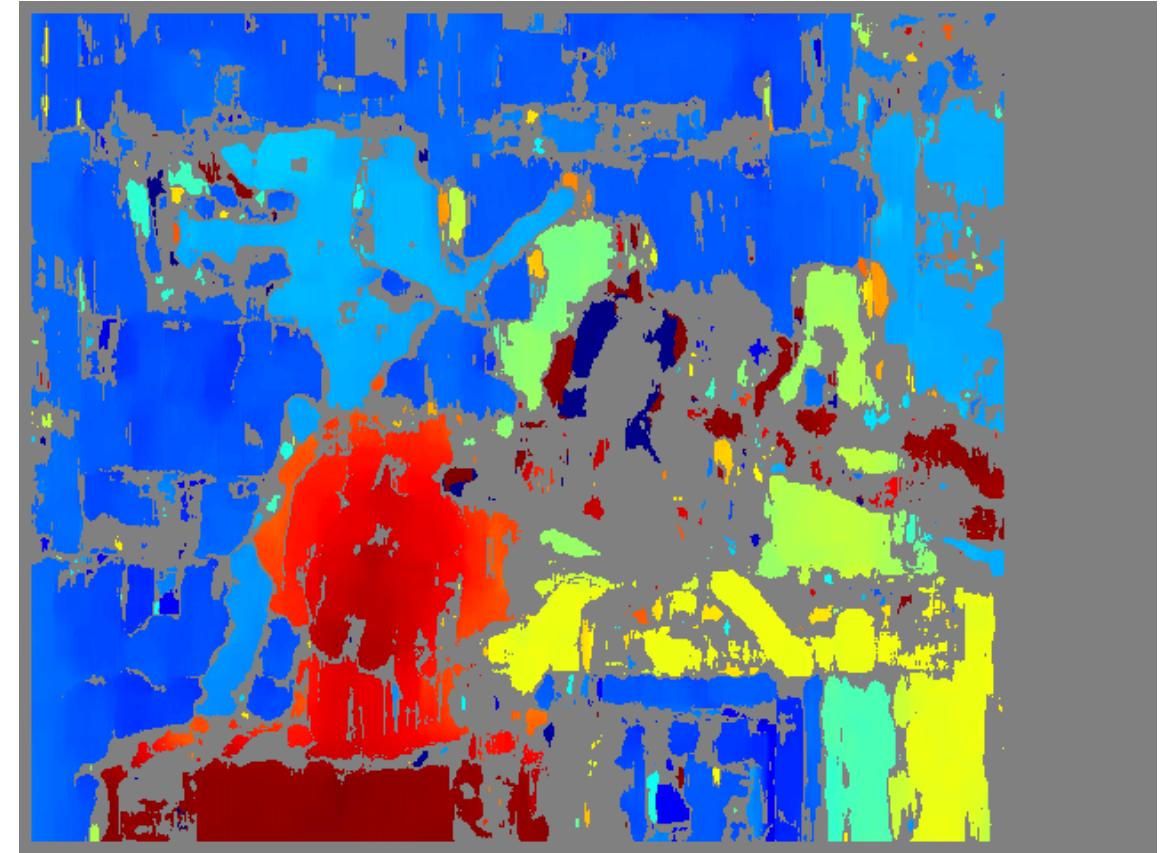
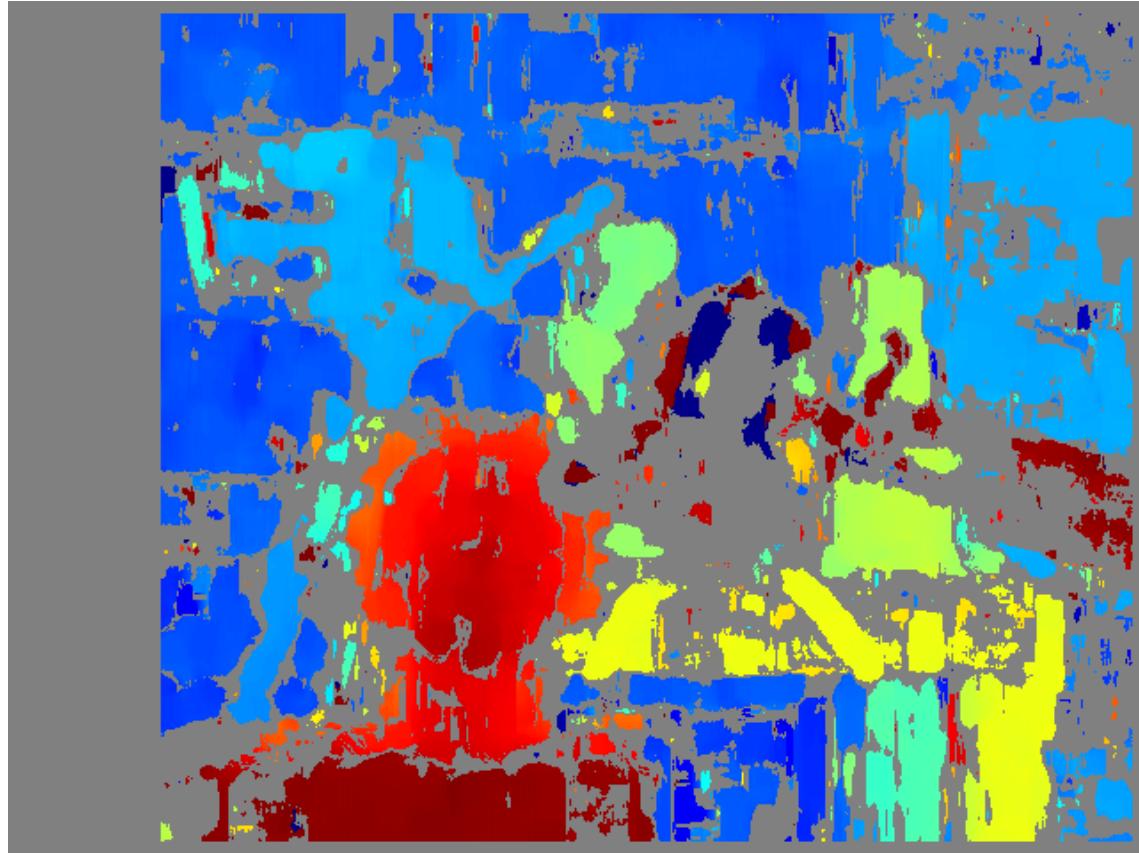
Stereo vision



Stereo vision



Stereo vision



Stereo vision

```
def get_disparity_v_01(imgL, imgR, disp_v1, disp_v2, disp_h1, disp_h2):
    window_size = 7
    left_matcher = cv2.StereoSGBM_create(
        minDisparity=disp_h1,
        numDisparities=int(0 + (disp_h2 - disp_h1) / 16) * 16,
        blockSize=5,
        P1=8 * 3 * window_size ** 2,
        P2=32 * 3 * window_size ** 2,
        disp12MaxDiff=disp_h2,
        uniquenessRatio=15,
        speckleWindowSize=0,
        speckleRange=2,
        preFilterCap=63,
        mode=cv2.STEREO_SGBM_MODE_SGBM_3WAY
    )

    right_matcher = cv2.ximgproc.createRightMatcher(left_matcher)
    dispr = right_matcher.compute(imgR, imgL)
    displ = left_matcher.compute(imgL, imgR)

    wls_filter = cv2.ximgproc.createDisparityWLSFilter(matcher_left=left_matcher)
    wls_filter.setLambda(80000)
    wls_filter.setSigmaColor(1.2)

    filteredImg_L = wls_filter.filter(displ, imgL, None, dispr)
    wls_filter = cv2.ximgproc.createDisparityWLSFilter(matcher_left=right_matcher)
    filteredImg_R = wls_filter.filter(dispr, imgR, None, displ)
```

Stereo vision

```
def get_disparity_v_02(imgL, imgR, disp_v1, disp_v2, disp_h1, disp_h2):
    max = numpy.maximum(math.fabs(disp_h1), math.fabs(disp_h2))
    levels = int(1 + (max) / 16) * 16
    stereo = cv2.StereoBM_create(numDisparities=levels, blockSize=15)
    displ = stereo.compute(imgL, imgR)

    dispR = numpy.flip(stereo.compute(numpy.flip(imgR, axis=1), numpy.flip(imgL, axis=1)), axis=1)
    return -displ / 16, -dispR / 16
```