

Neural Networks for Computer Vision

Dmitry Ryabokon, github.com/dryabokon

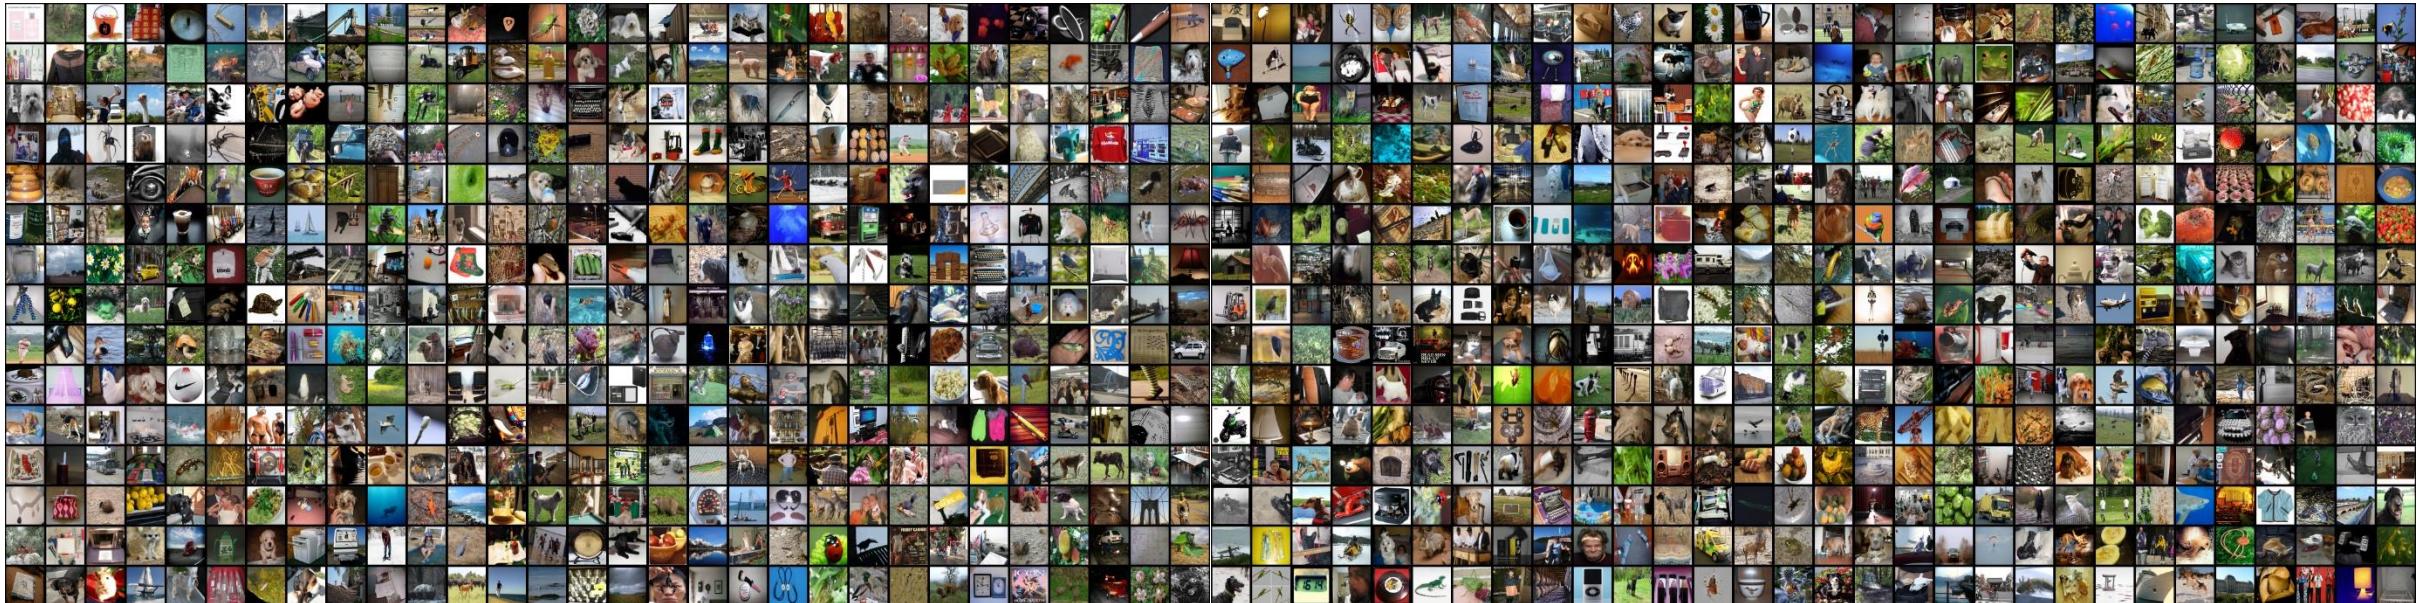
Content

- 1. Intro
- 2. CNNs out of box
- 3. Basic operations
- 4. Feature extraction
- 5. Feature visualization
- 6. Visualization of layers and filters
- 7. Transfer learning
- 8. Fine Tuning
- 9. Architecture of the popular networks
- 10. Build, Train and Test own CNN
- 11. Benchmarking the CNNs
- 12. Autoencoders
- 12. Object detection
- 13. Style transfer
 - Colorization
 - CAM
 - Texture analysis
 - Texture synthesis
 - PixelRNN, GAN

Intro

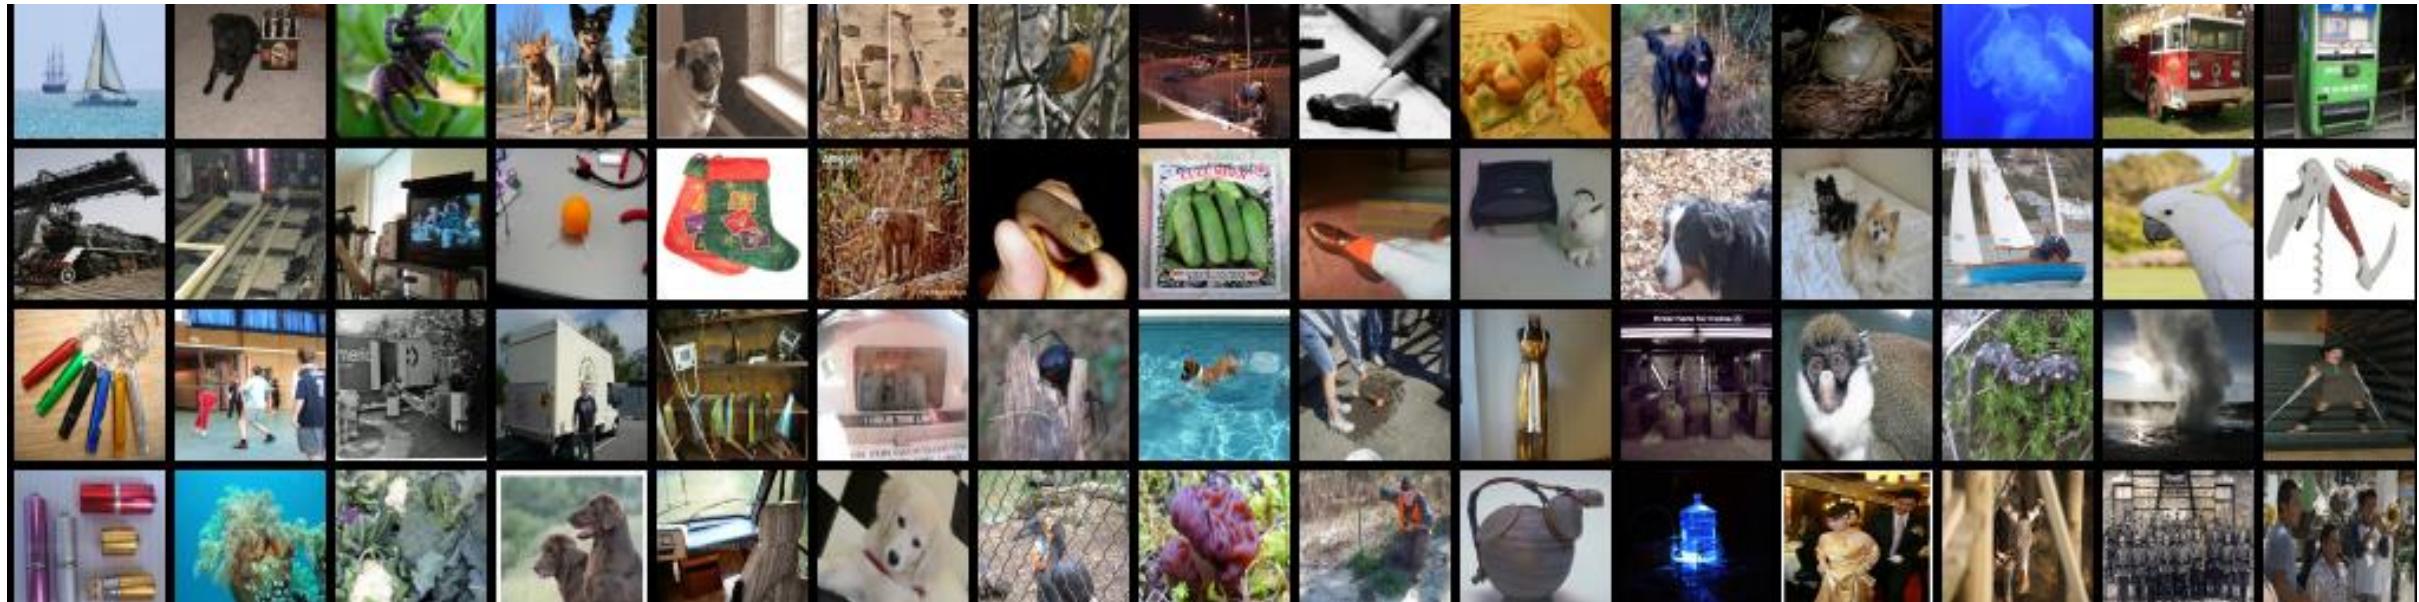
Intro

ImageNet database: 1000 classes, ~100 samples for each class



Intro

ImageNet database



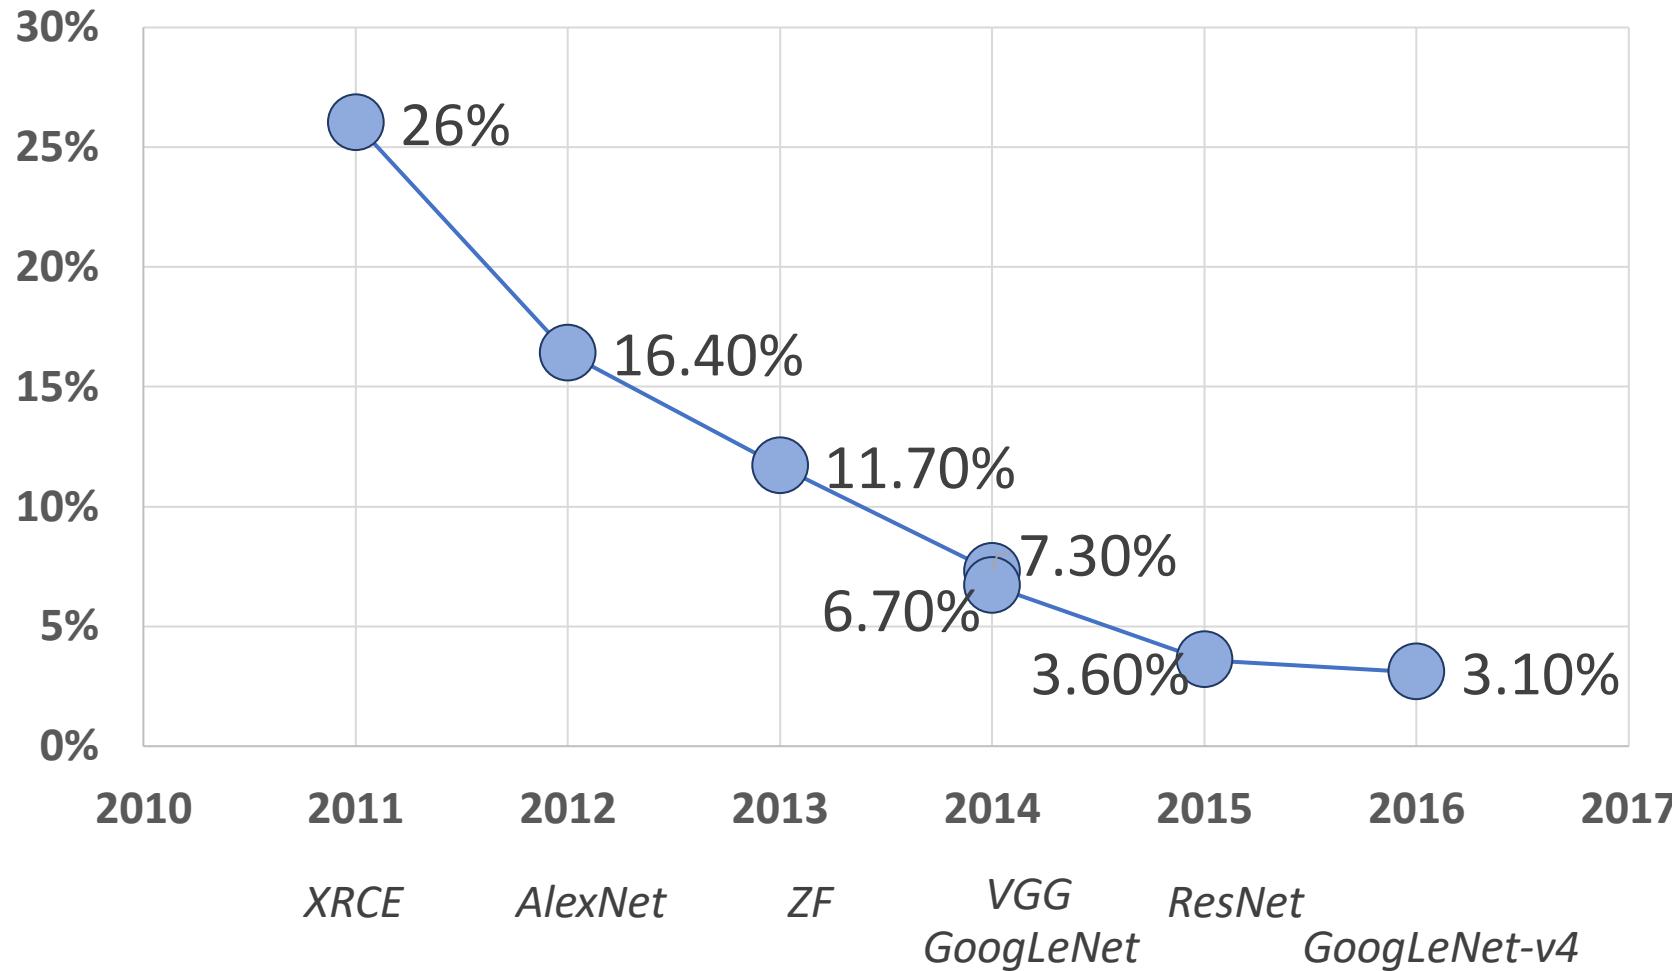
Intro

ImageNet database



Intro

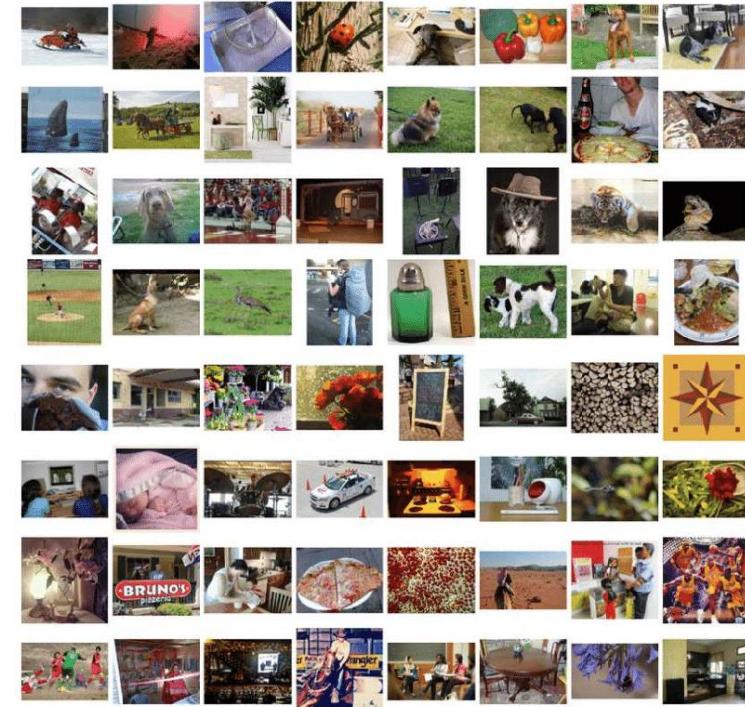
ImageNet Classification error



Intro

Some datasets available for research

MNIST: 10 classes, ~7000 ex. per class



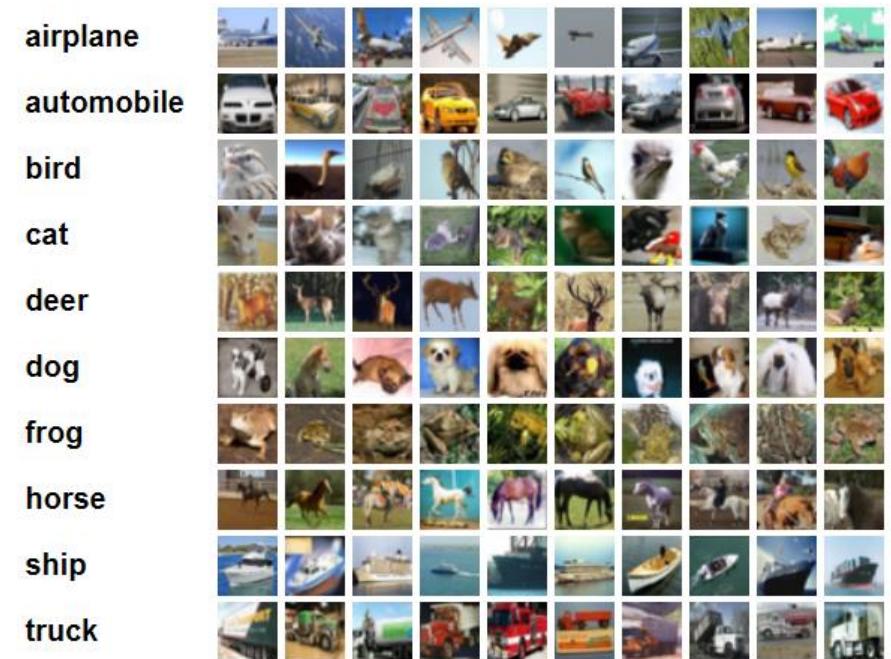
ImageNet: 1000 classes, ~100 ex per class

Intro

Some datasets available for research



The Street View House Numbers
10 classes, ~2000 ex. per class



CIFAR: 10 classes, 6000 ex. per class
100 classes, 600 ex per class

Intro

Some datasets available for research



Olivetti database: 40 classes

The screenshot shows the Kaggle datasets homepage. At the top, there's a navigation bar with back, forward, and refresh buttons, and a URL field showing <https://www.kaggle.com/datasets>. Below the navigation is a blue header bar with the word "Datasets". Underneath the header, there are tabs for "Public", "Your Datasets", and "Favorites". Further down are filters for "Sizes", "File types", "Licenses", and "Tags". The main content area displays four dataset entries:

- Heart Disease UCI** (101 datasets): A dataset for classification, featuring a stethoscope icon. It links to <https://archive.ics.uci.edu/ml/datasets/Heart+Disease> and was updated 7 months ago by ronit.
- Graduate Admissions** (424 datasets): A dataset for regression and classification, featuring a graduation cap icon. It links to "JUMP START ON GE SE" and was updated a month ago by Mohan S Acharya.
- FIFA 19 complete player dataset** (334 datasets): A dataset for sports and data visualization, featuring a soccer player icon. It links to "FIFA 19" and was updated a month ago by Karan Gadiya.
- Suicide Rates Overview 1985 to 2016** (54 datasets): A dataset for world demographics and economics, featuring a bar chart icon. It links to "Suicide Rates Overview 1985 to 2016" and compares socio-economic info with suicide rates by year.

.. and much more @ kaggle

Usage of the neural networks

- identify the name of a street (in France) from an image
- compressing and decompressing images
- semantic image segmentation
- real-world image text extraction.
- image classification
- image-to-text
- unsupervised learning
- 3D object reconstruction
- image matching and retrieval
- automatic speech recognition
- localizing and identifying multiple objects in a single image
- computer vision
- discovery of latent 3D keypoints
- predicting future video frames

The popular networks

Classification

- LeNet [Model](#)
- AlexNet [Model](#)
- VGG [Model](#)
- ResNet [Paper](#)
- YOLO9000 [Paper](#)
- DenseNet [Paper](#)

Segmentation

- FCN8 [Paper](#)
- SegNet [Paper](#)
- U-Net [Paper](#)
- E-Net [Paper](#)
- ResNetFCN [Paper](#)
- PSPNet [Paper](#)
- Mask RCNN [Paper](#)

Detection

- Faster RCNN [Paper](#)
- SSD [Paper](#)
- YOLOv2 [Paper](#)
- R-FCN [Paper](#)

CNNs out of box

CNNs out of box

The screenshot shows a browser window with the URL https://tfhub.dev/google/imagenet/resnet_v2_50/feature_vector/1. The page has a white header with a back arrow, forward arrow, and refresh icon. The main navigation bar is orange with the text "TensorFlow Hub" and a search icon. Below the header, there's a breadcrumb trail with a back arrow and the text "imagenet/resnet_v2_50/feature_vector".

Usage

This module implements the common signature for computing [image feature vectors](#). It can be used like

```
module = hub.Module("https://tfhub.dev/google/imagenet/resnet_v2_50/feature_vector/1")
height, width = hub.get_expected_image_size(module)
images = ... # A batch of images with shape [batch_size, height, width, 3].
features = module(images) # Features with shape [batch_size, num_features].
```

...or using the signature name `image_feature_vector`. The output for each image in the batch is a feature vector of size `num_features = 2048`.

For this module, the size of the input image is fixed to `height x width = 224 x 224` pixels. The input `images` are expected to have color values in the range `[0,1]`, following the [common image input](#) conventions.

CNNs out of box

```
import tensorflow_hub as hub
import urllib.request
import cv2
import numpy
import json
import tensorflow as tf
# -----
URL = 'https://s3.amazonaws.com/deep-learning-models/image-models/imagenet_class_index.json'
data = json.loads(urllib.request.urlopen(URL).read().decode())
class_names = [data['%d'%i][1] for i in range(0,999)]
# -----
def example_predict():

    module = hub.Module("https://tfhub.dev/google/imagenet/resnet_v2_50/classification/1")
    img = cv2.imread('data/ex-natural/dog/dog_0000.jpg')
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = cv2.resize(img, (224, 224)).astype(numpy.float32)
    img = numpy.array([img]).astype(numpy.float32) / 255.0
    sess = tf.Session()
    sess.run(tf.global_variables_initializer())
    sess.run(tf.tables_initializer())

    outputs = module(dict(images=img), signature="image_classification", as_dict=True)
    prob = outputs["default"].eval(session=sess)[0]
    idx = numpy.argsort(-prob)[0]

    print(class_names[idx], prob[idx])
    sess.close()

    return
# -----
if __name__ == '__main__':
    example_predict()
```

CNNs out of box

← → ⌂ https://keras.io/getting-started/faq/#how-can-i-use-pre-trained-models-in-keras

Keras Documentation

Search docs

- Home
- Why use Keras
- GETTING STARTED
 - Guide to the Sequential model
 - Guide to the Functional API

FAQ

- How should I cite Keras?
- How can I run Keras on GPU?
- How can I run a Keras model on multiple GPUs?
- What does "sample", "batch", "epoch" mean?
- How can I save a Keras model?
- Why is the training loss much higher than the testing loss?
- How can I obtain the output of an intermediate layer?
- How can I use Keras with datasets that don't fit in memory?
- How can I interrupt training when the validation loss isn't decreasing

How can I use pre-trained models in Keras?

Code and pre-trained weights are available for the following image classification models:

- Xception
- VGG16
- VGG19
- ResNet50
- Inception v3
- Inception-ResNet v2
- MobileNet v1
- DenseNet
- NASNet
- MobileNet v2

They can be imported from the module `keras.applications`:

```
from keras.applications.xception import Xception
from keras.applications.vgg16 import VGG16
from keras.applications.vgg19 import VGG19
from keras.applications.resnet50 import ResNet50
from keras.applications.inception_v3 import InceptionV3
from keras.applications.inception_resnet_v2 import InceptionResNetV2
from keras.applications.mobilenet import MobileNet
from keras.applications.densenet import DenseNet121
from keras.applications.densenet import DenseNet169
from keras.applications.densenet import DenseNet201
from keras.applications.nasnet import NASNetLarge
from keras.applications.nasnet import NASNetMobile
from keras.applications.mobilenet_v2 import MobileNetV2

model = VGG16(weights='imagenet', include_top=True)
```

CNNs out of box

```
from keras.applications import MobileNet
from keras.applications.xception import Xception
from keras.applications.mobilenet import preprocess_input
import urllib.request
import cv2
import numpy
import json
from keras import backend as K
K.set_image_dim_ordering('tf')
# -----
URL = 'https://s3.amazonaws.com/deep-learning-models/image-models/imagenet_class_index.json'
data = json.loads(urllib.request.urlopen(URL).read().decode())
class_names = [data['%d' % i][1] for i in range(0, 999)]
# -----
def example_predict():

    CNN = MobileNet()
    #CNN = Xception()

    img = cv2.imread('data/ex-natural/dog/dog_0000.jpg')
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = cv2.resize(img, (224, 224)).astype(numpy.float32)

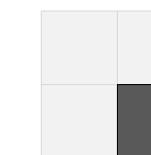
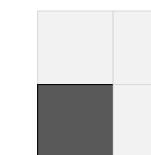
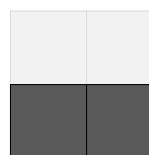
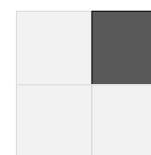
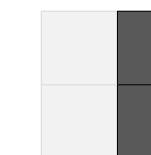
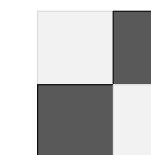
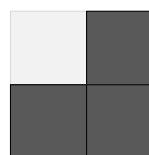
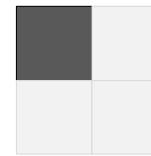
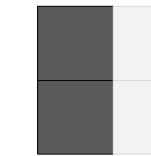
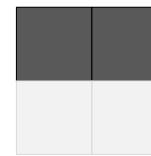
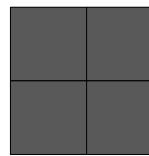
    prob = CNN.predict(preprocess_input(numpy.array([img])))
    idx = numpy.argsort(-prob[0])[0]
    print(class_names[idx], prob[0, idx])

    return
# -----
if __name__ == '__main__':
    example_predict()
```

Basic operations

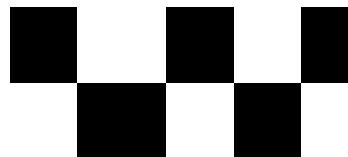
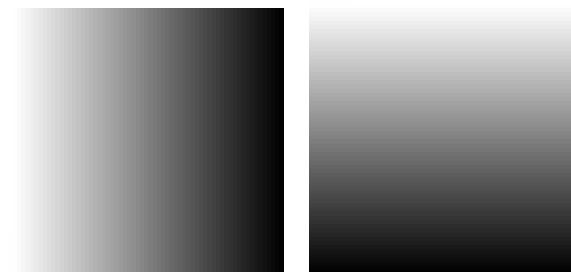
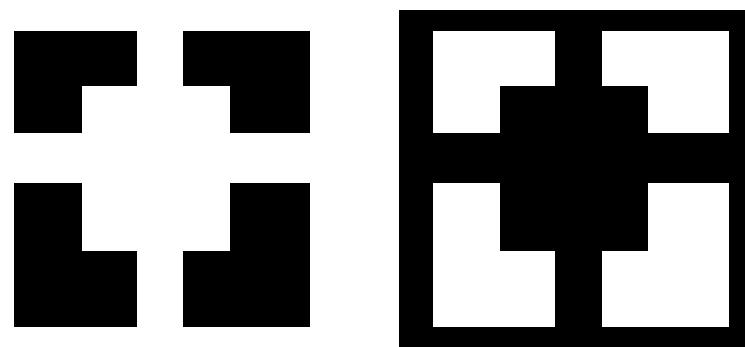
Basic operations

Convolution



Basic operations

Convolution

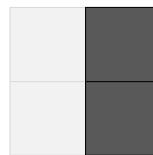
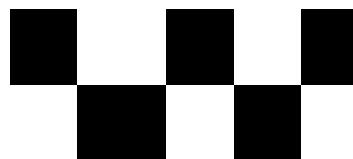
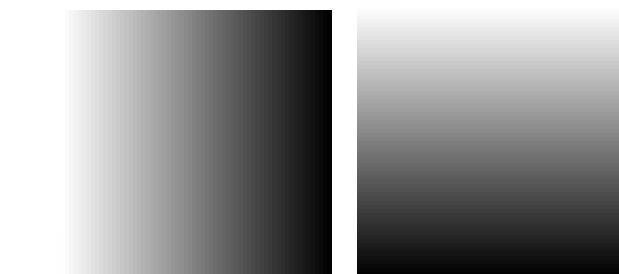
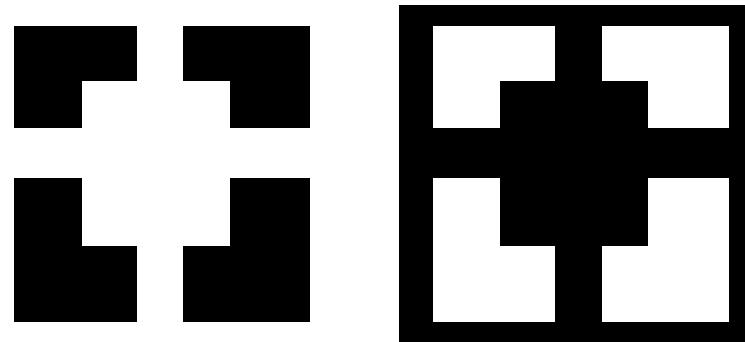


=



Basic operations

Convolution

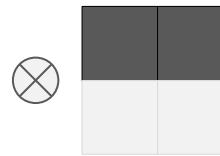
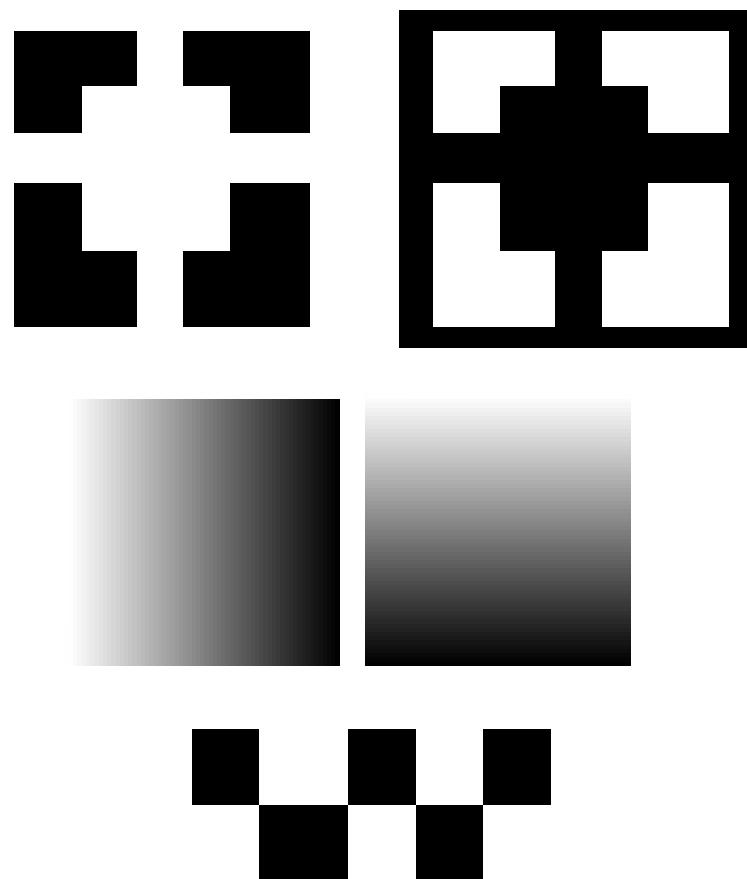


=



Basic operations

Convolution

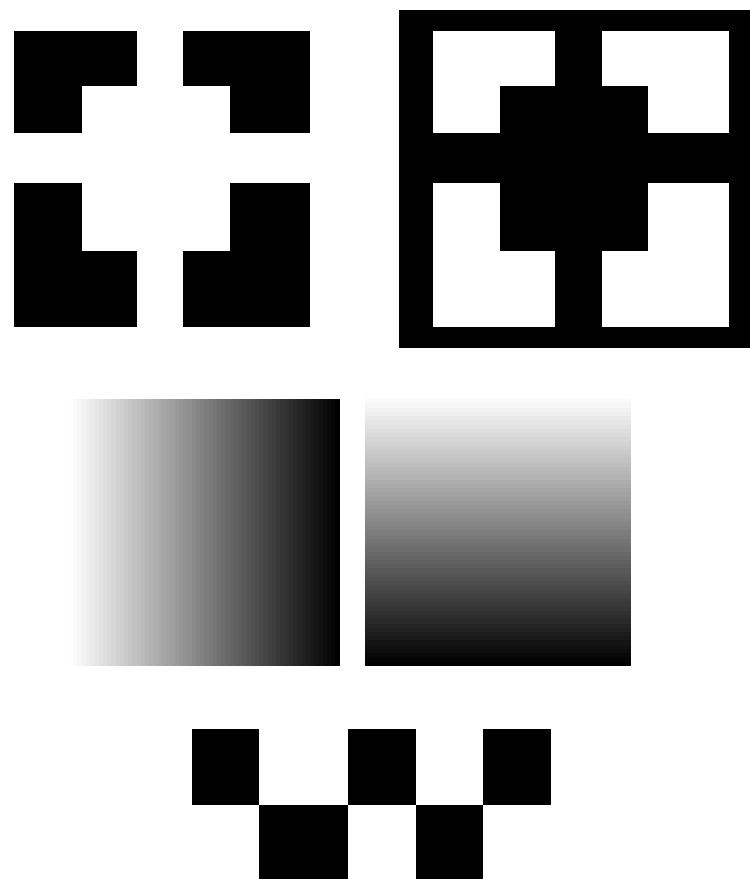


=

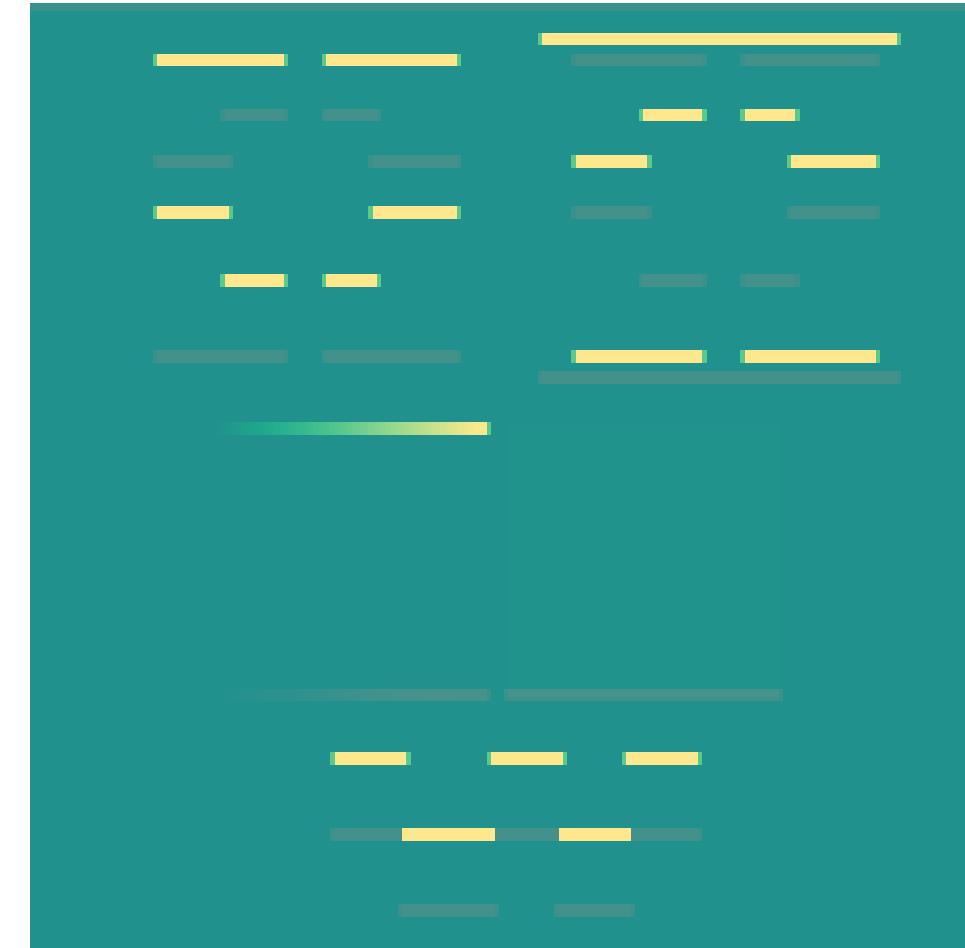


Basic operations

Convolution

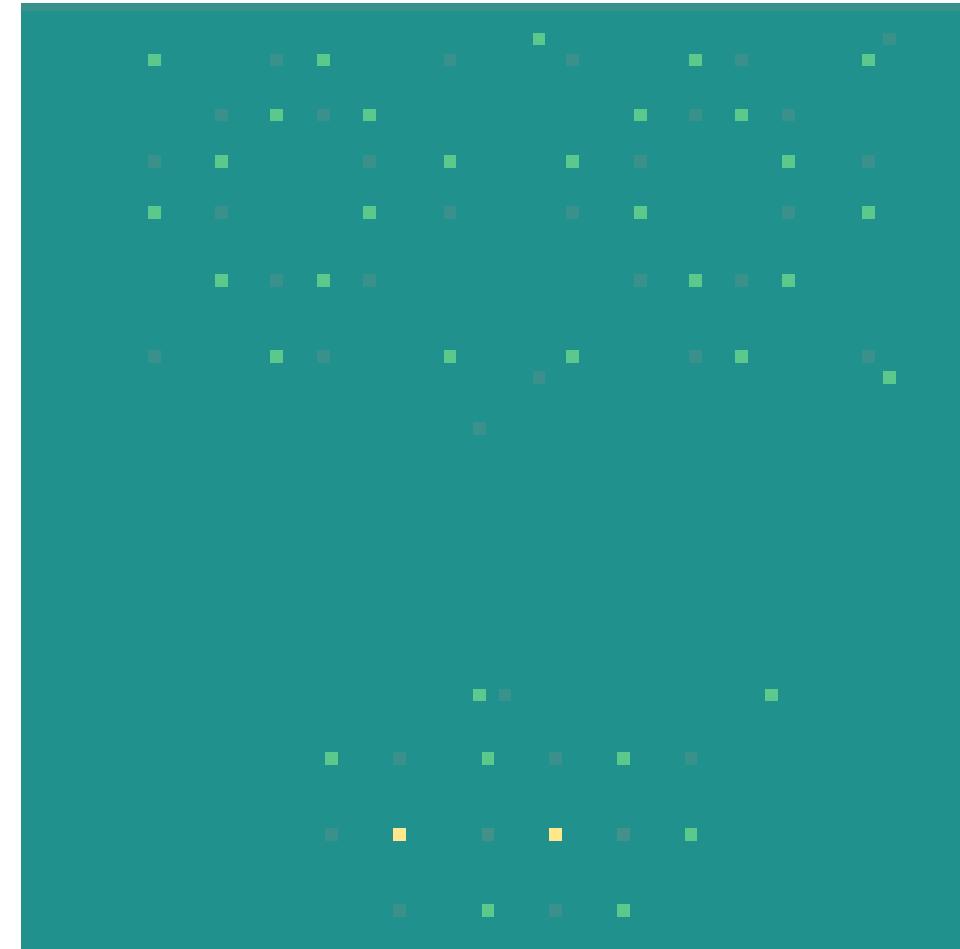
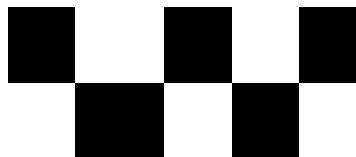
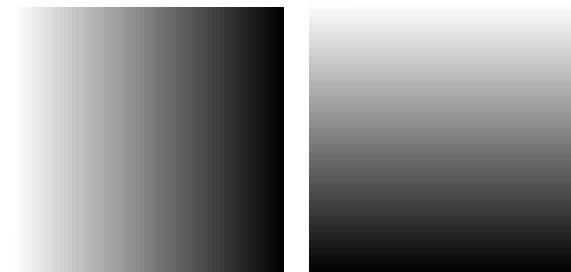
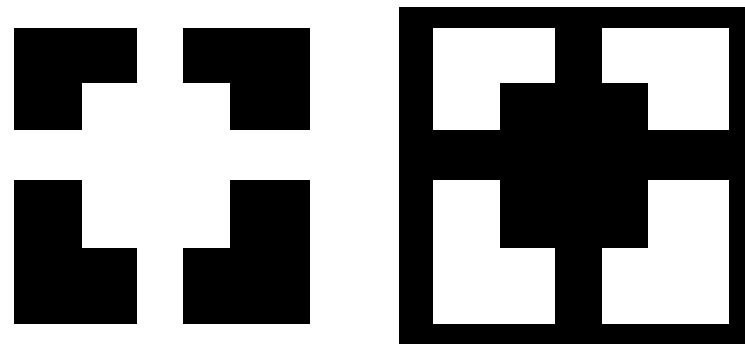


$$\otimes \quad \begin{matrix} & & \\ & & \\ \text{---} & \text{---} & \text{---} \\ & & \end{matrix} =$$



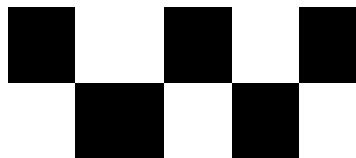
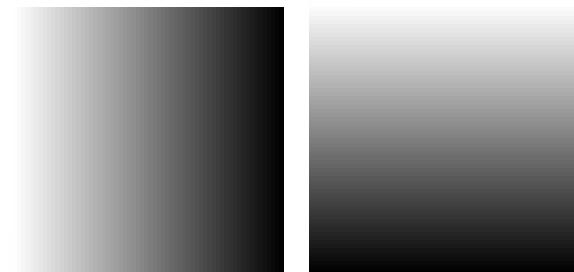
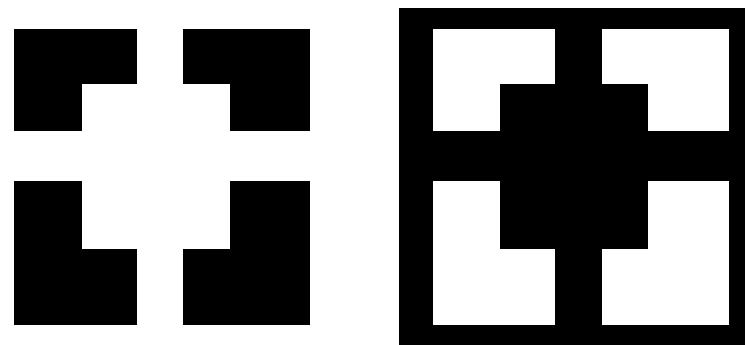
Basic operations

Convolution

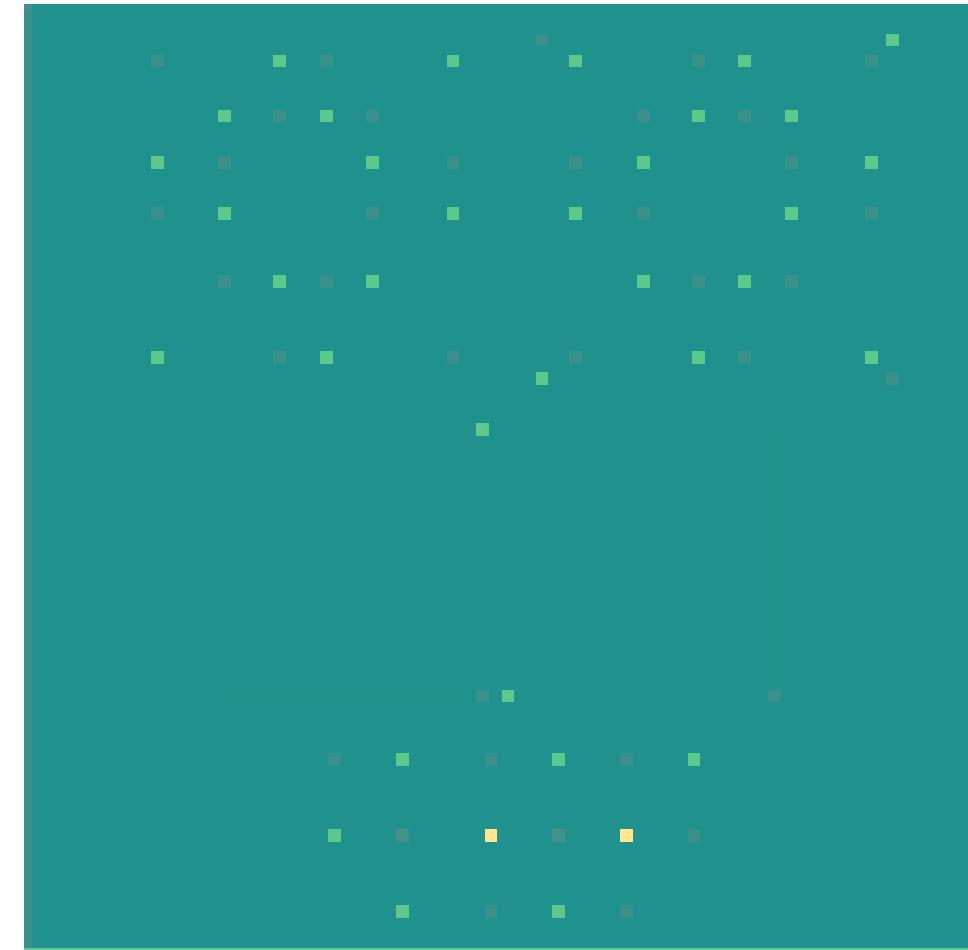


Basic operations

Convolution

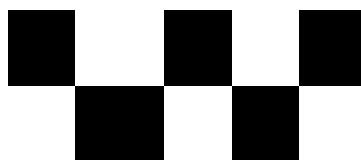
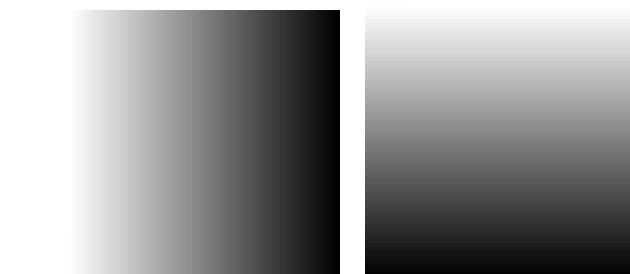
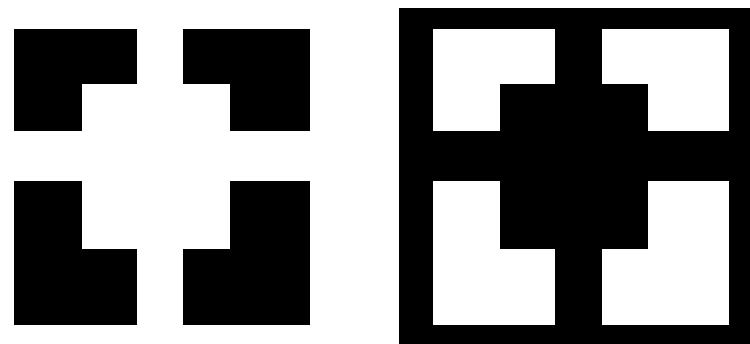
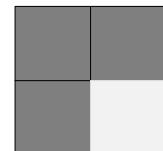
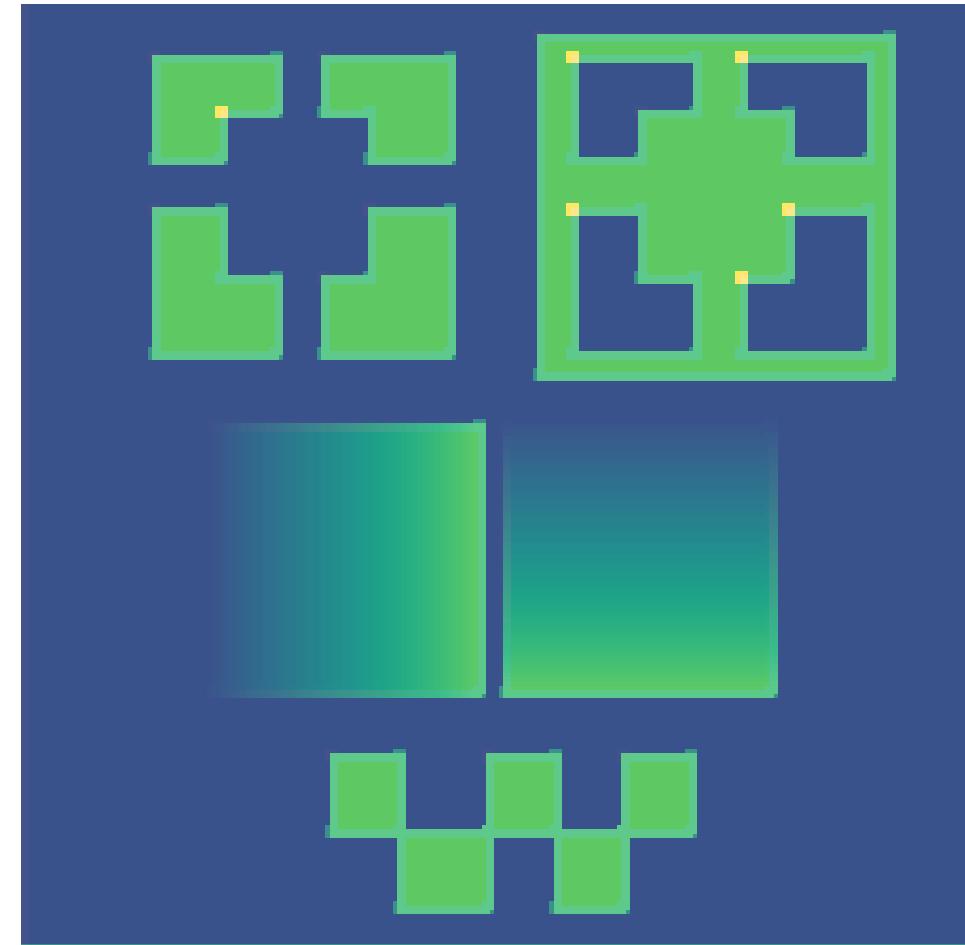


=



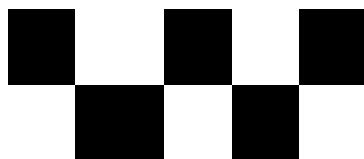
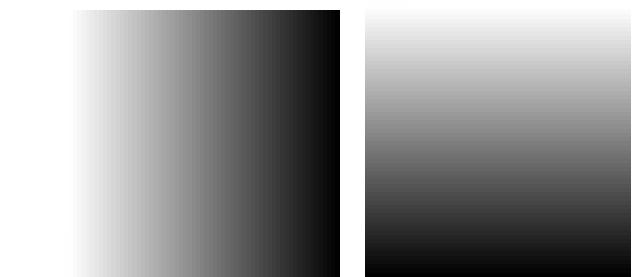
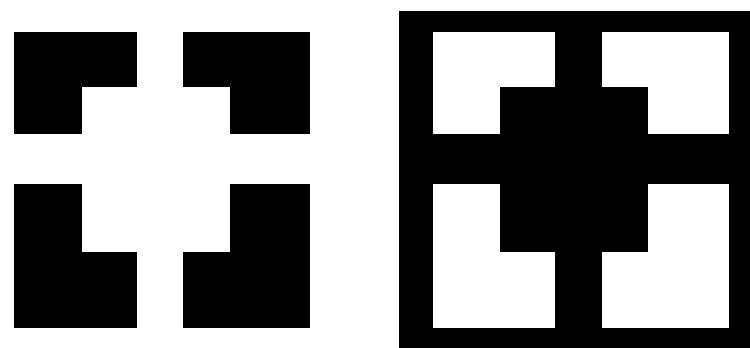
Basic operations

Convolution

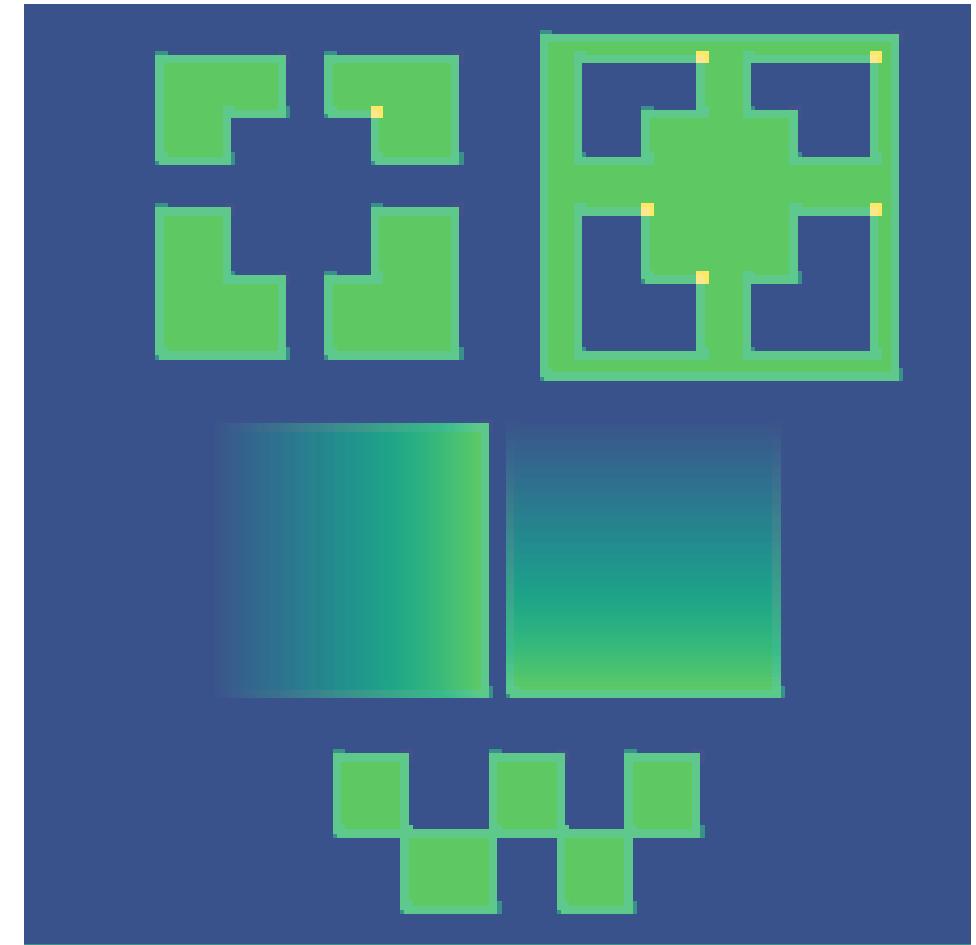
 \otimes  $=$ 

Basic operations

Convolution

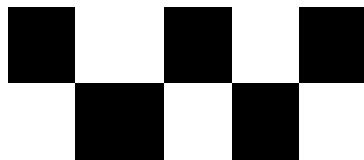
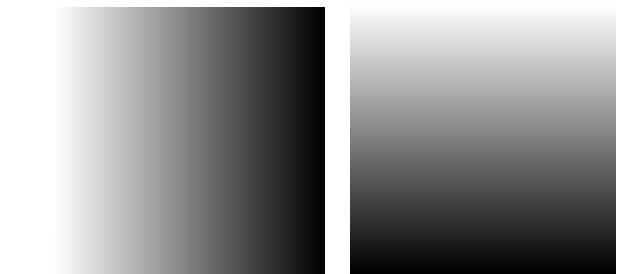
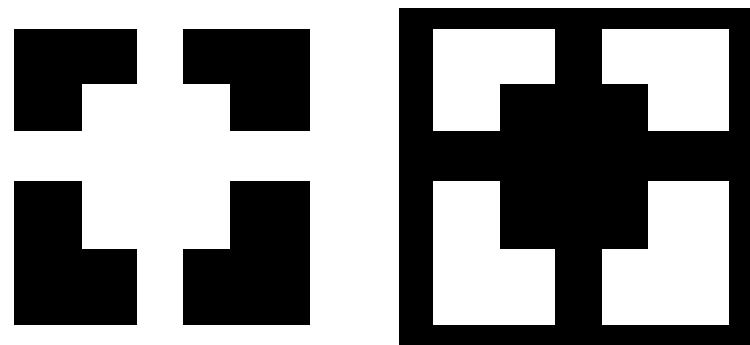


$$\otimes \begin{matrix} & & \\ & \text{dark gray} & \\ & & \end{matrix} =$$

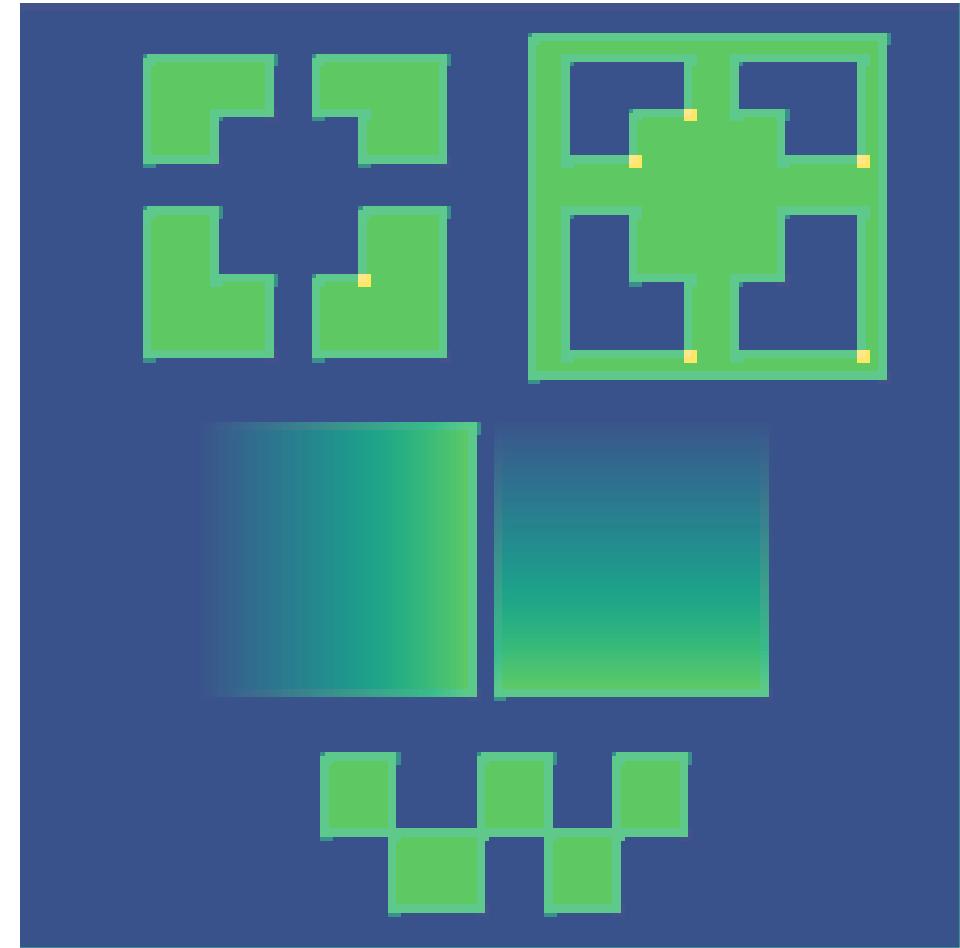


Basic operations

Convolution

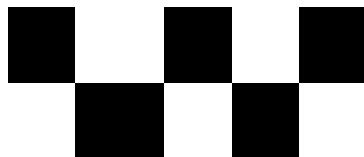
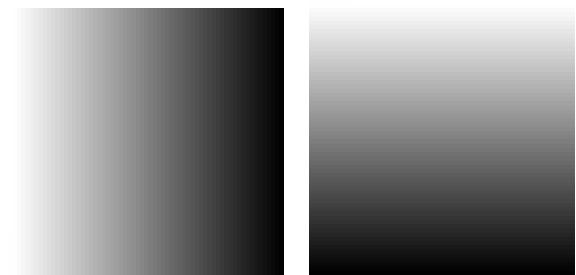
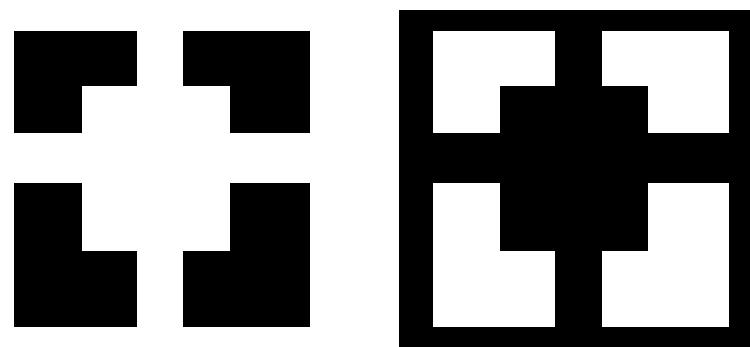


$$\otimes \quad =$$
A 3x3 convolution kernel with a central gray square and three white squares around it.

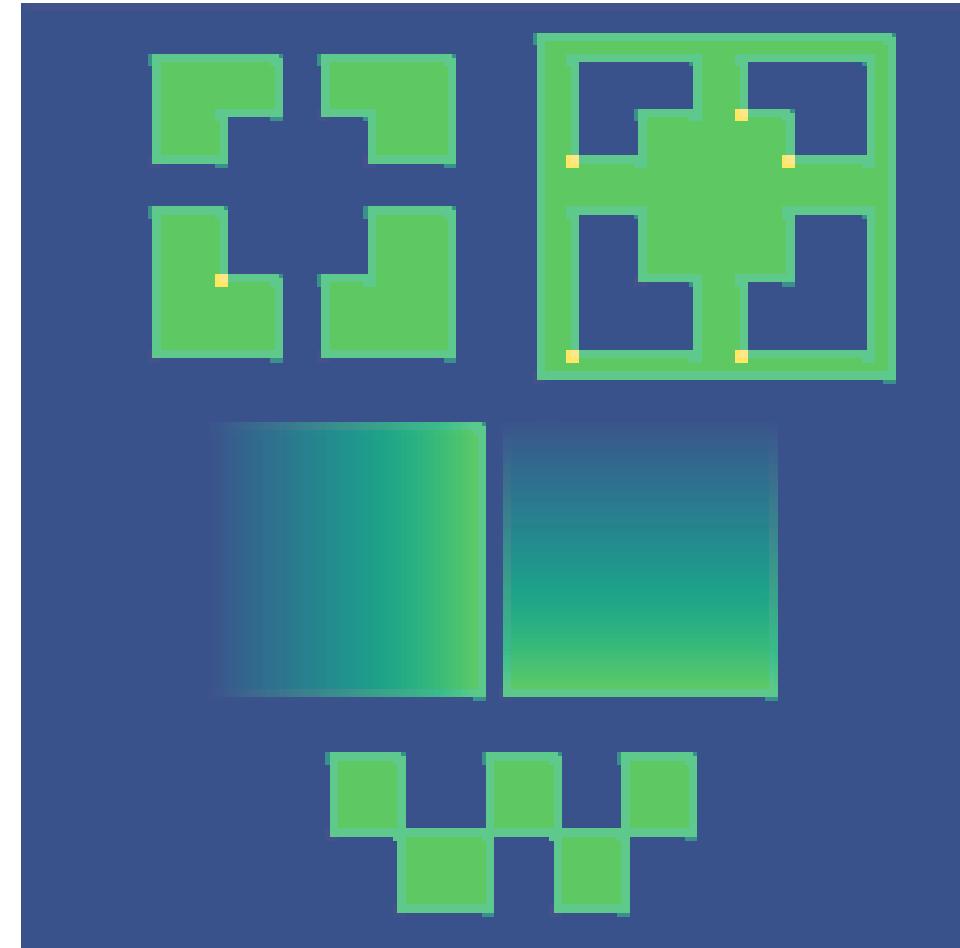


Basic operations

Convolution

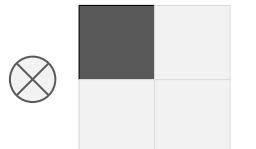
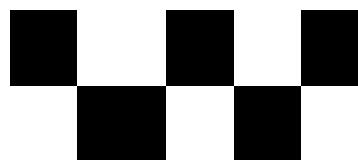
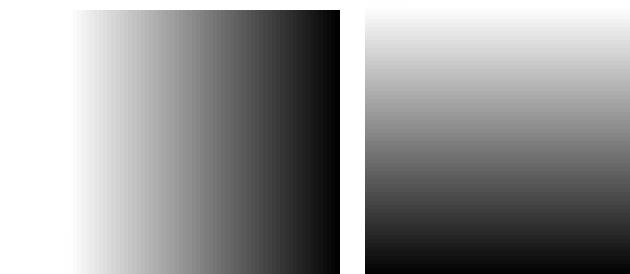
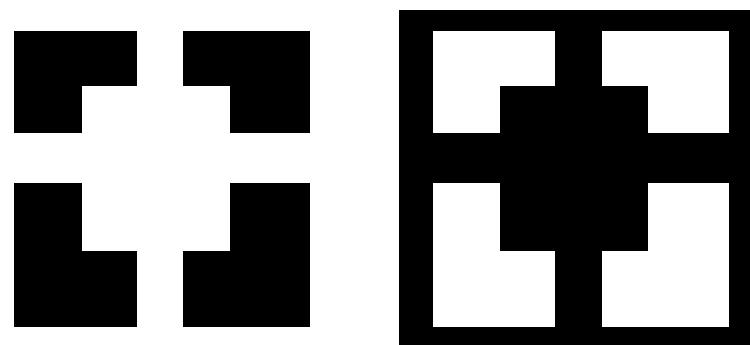


$$\otimes \begin{matrix} & & \\ & \text{gray} & \\ & & \end{matrix} =$$

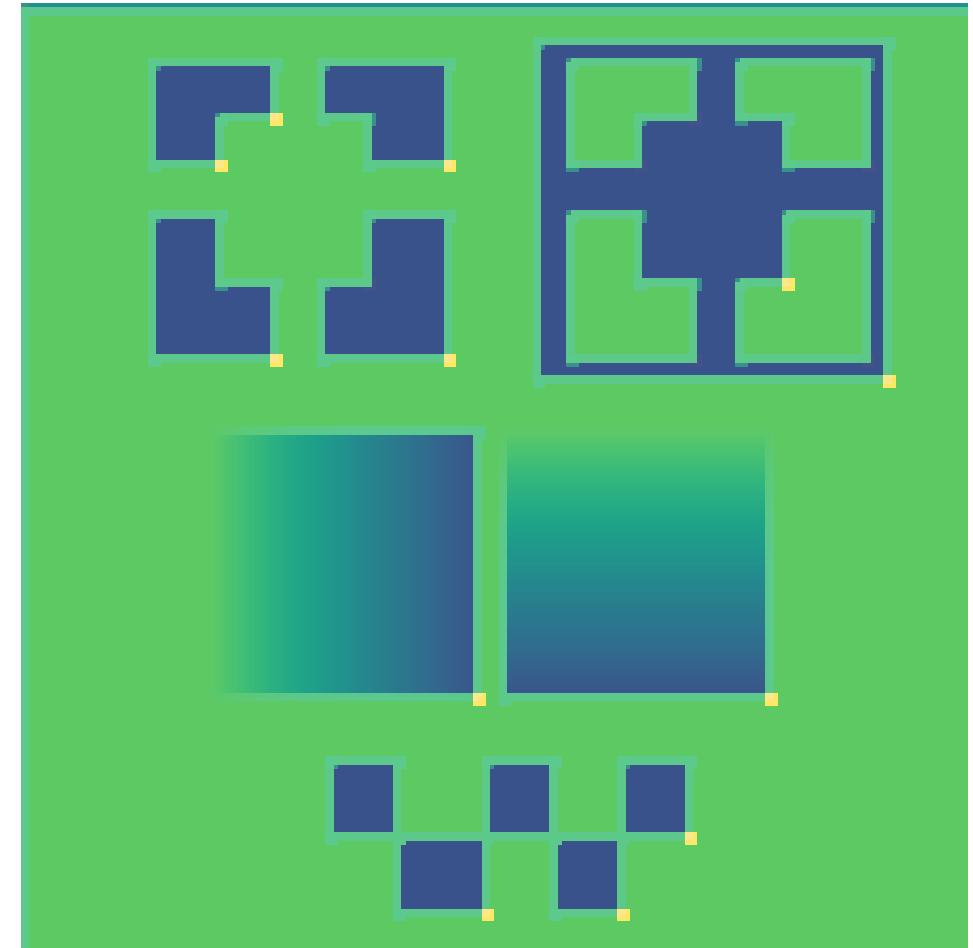


Basic operations

Convolution

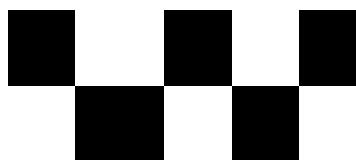
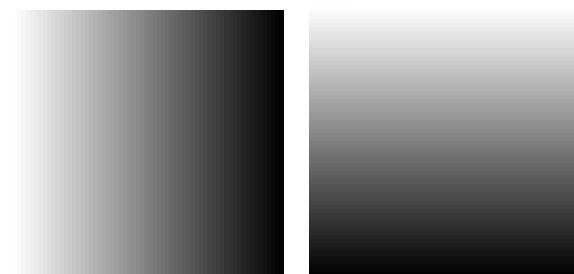
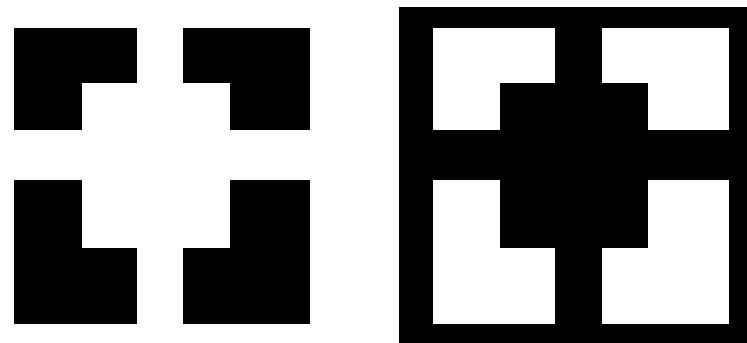


=

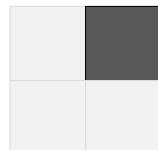


Basic operations

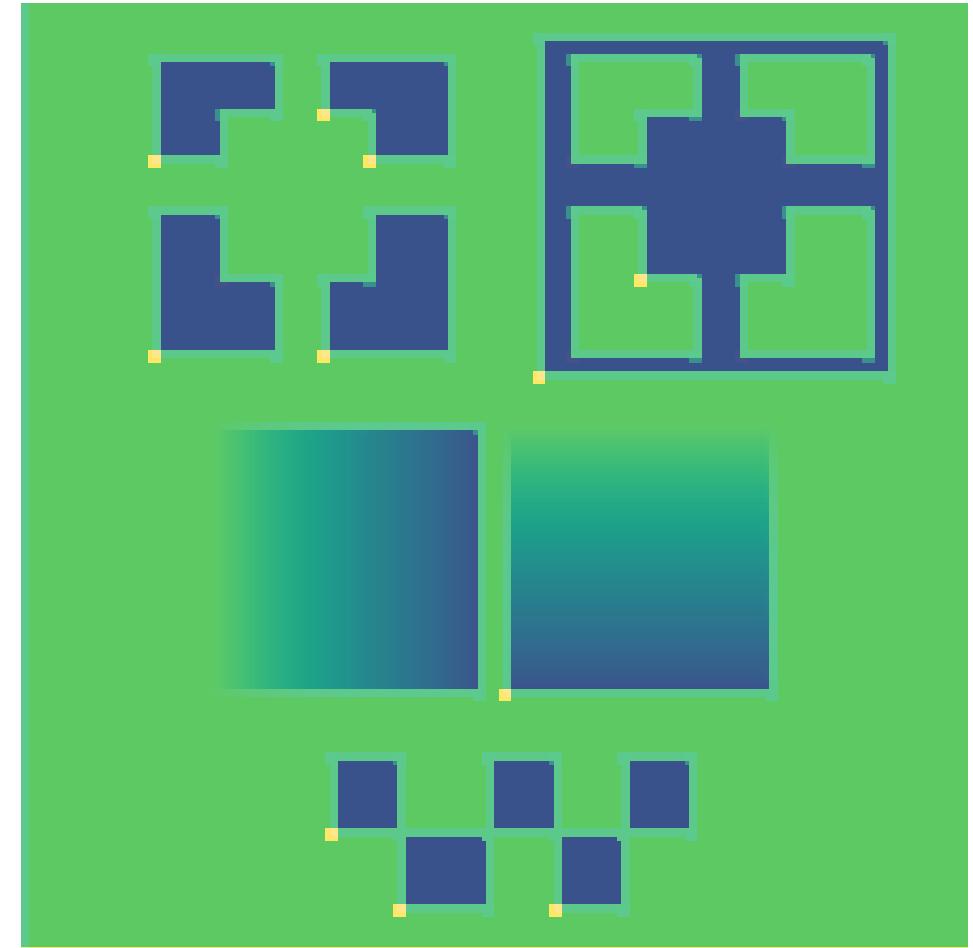
Convolution



\otimes

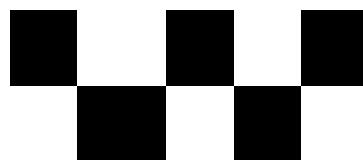
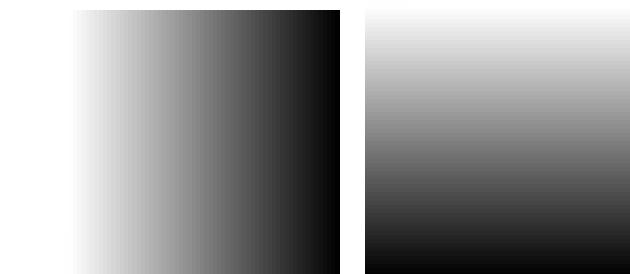
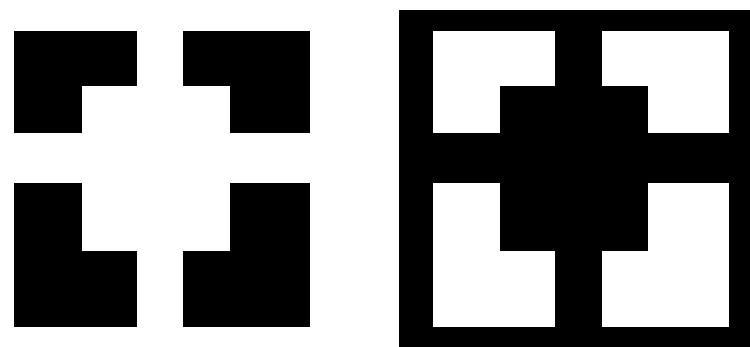


=

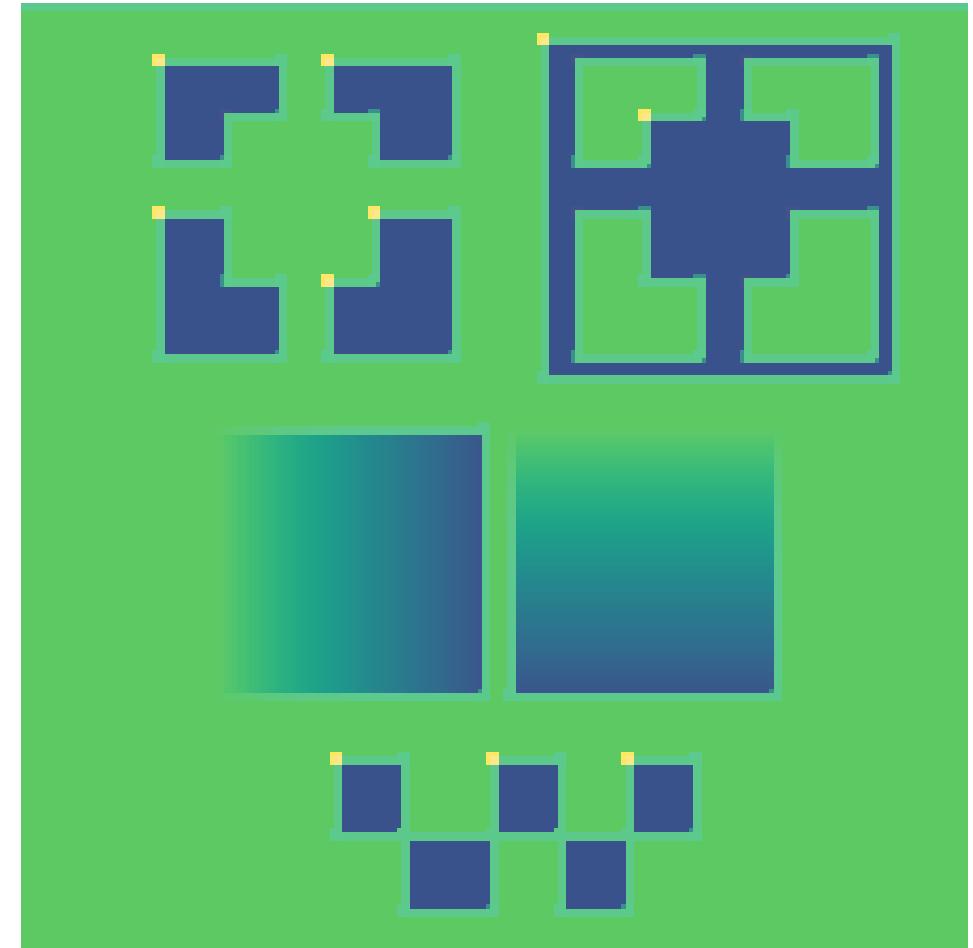


Basic operations

Convolution

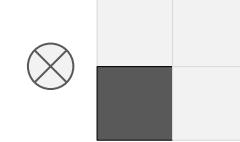
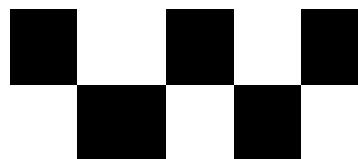
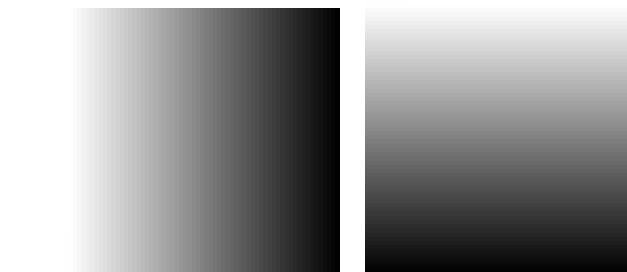
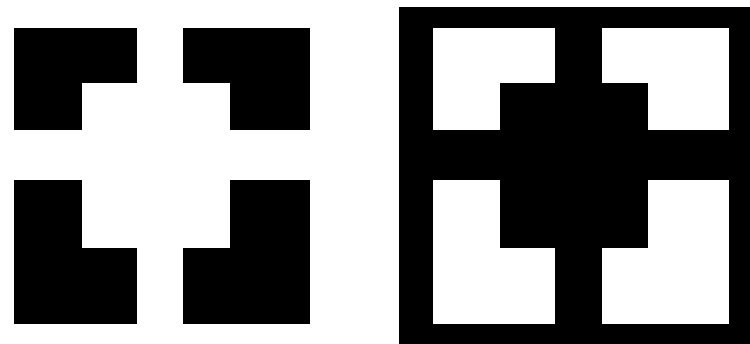


$$\otimes \quad \begin{matrix} & & \\ & & \\ \text{---} & \text{---} & \text{---} \\ & & \end{matrix} =$$

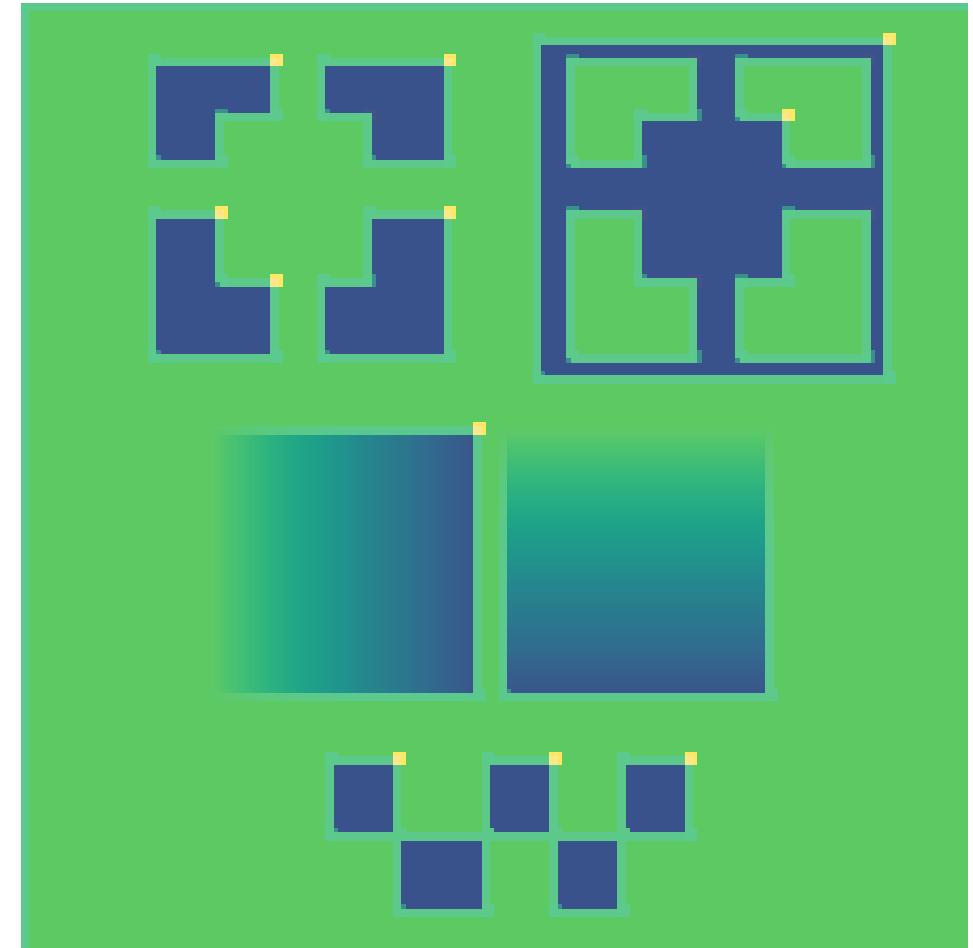


Basic operations

Convolution

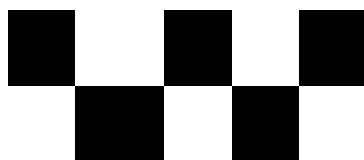
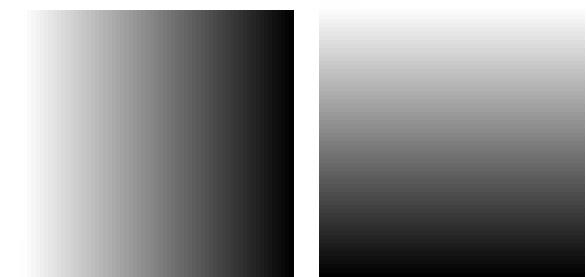
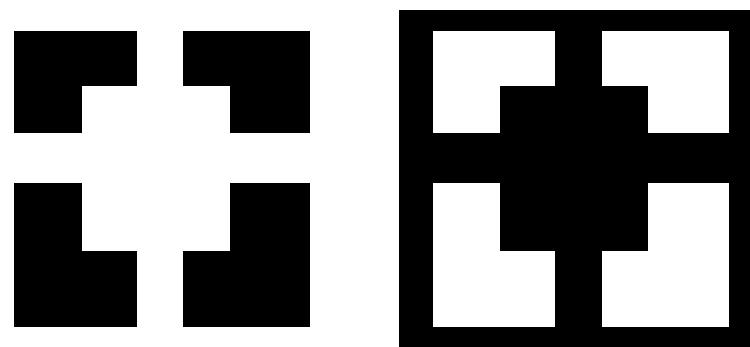


=

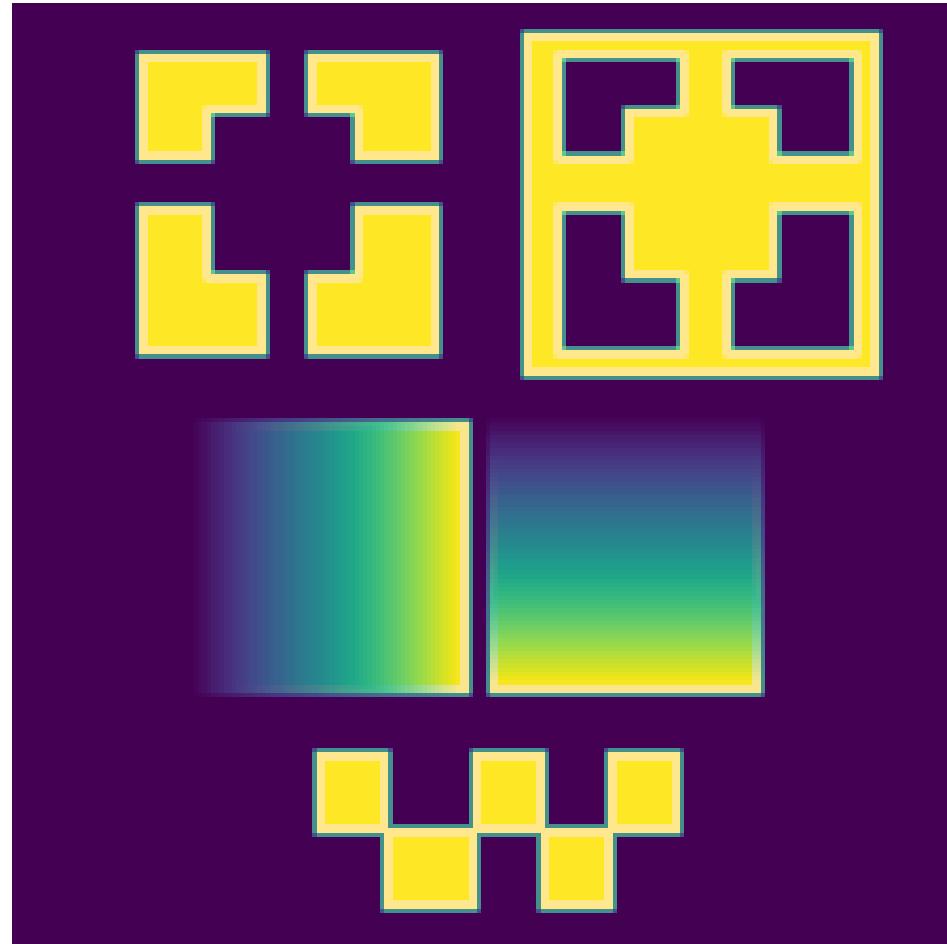


Basic operations

Convolution

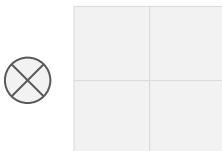
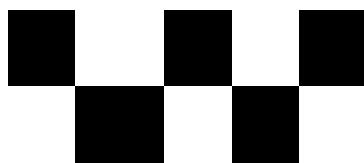
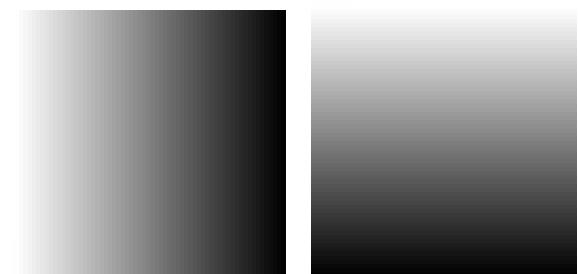
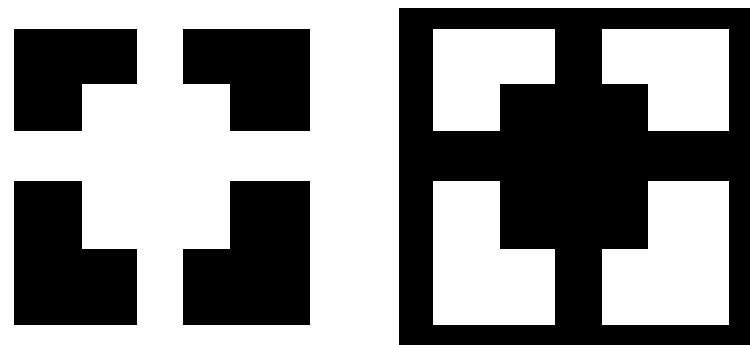


$$\otimes \begin{array}{|c|c|}\hline & & \\ \hline & & \\ \hline & & \\ \hline\end{array} =$$

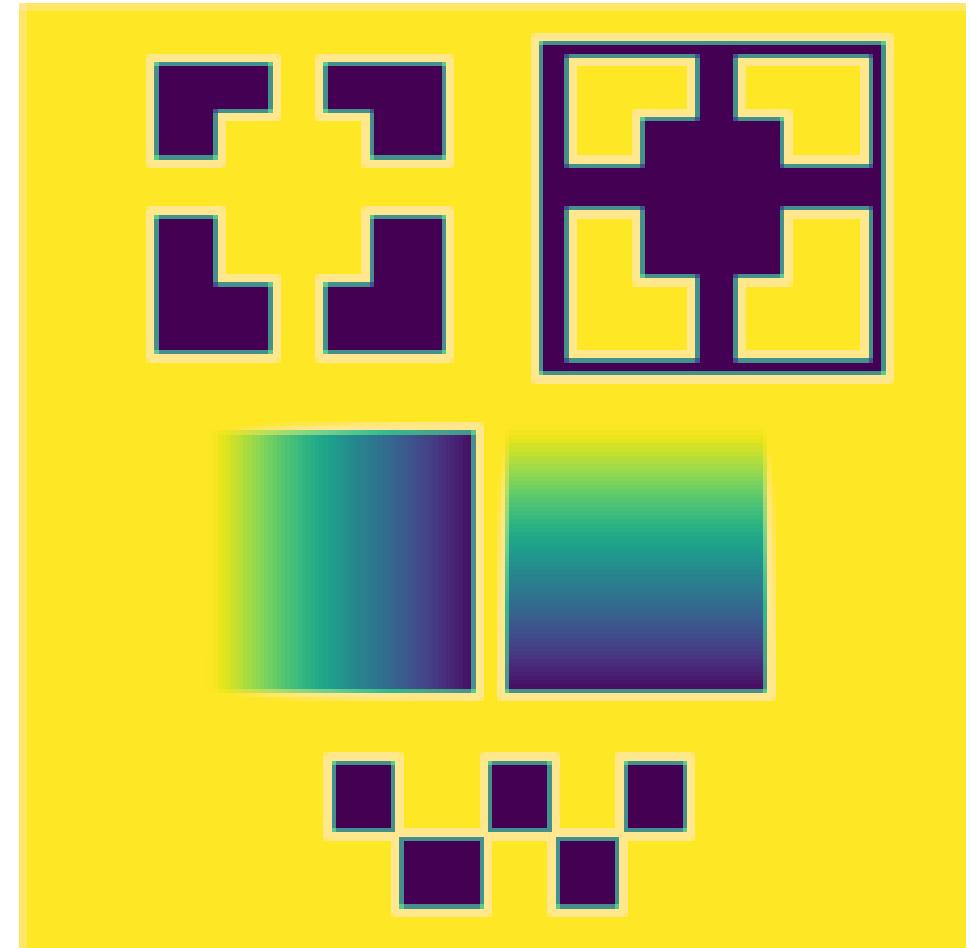


Basic operations

Convolution

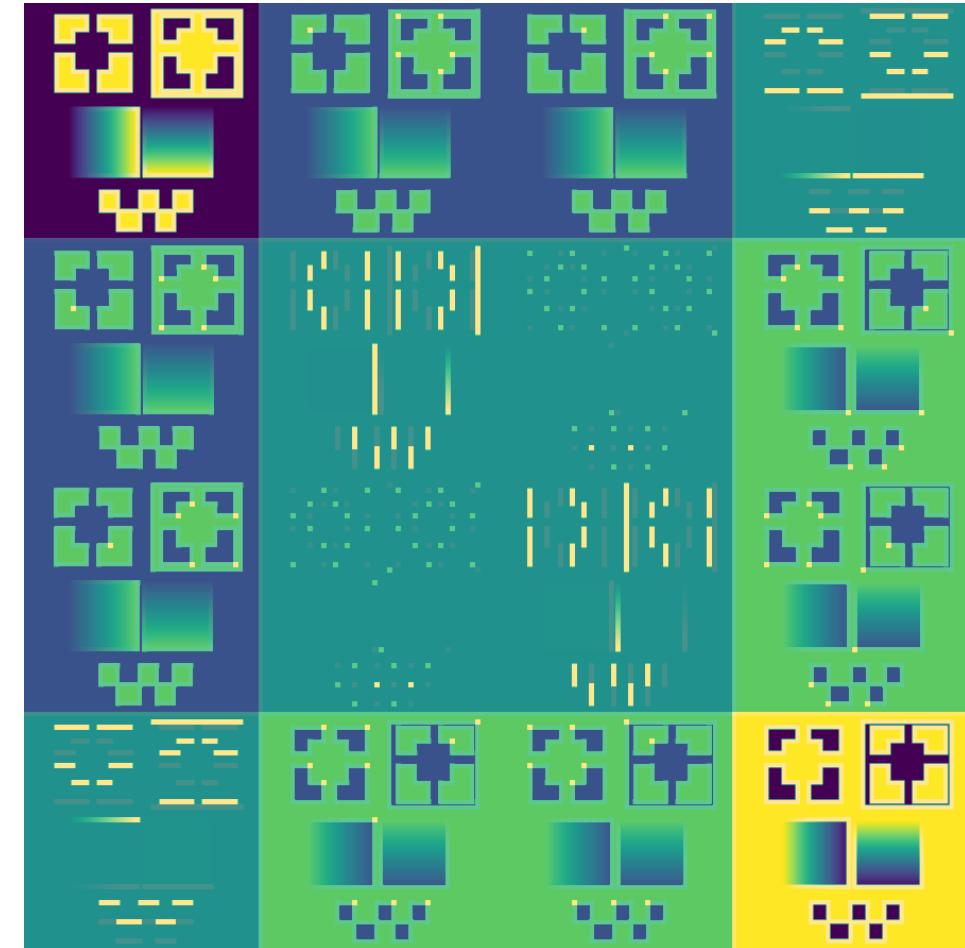
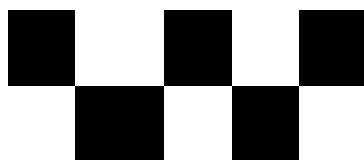
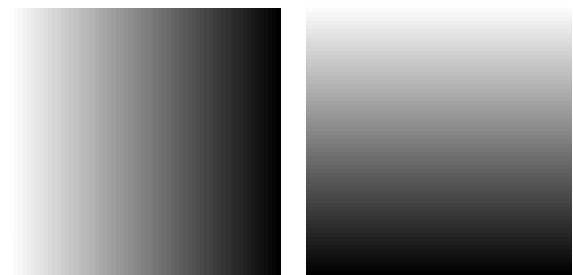
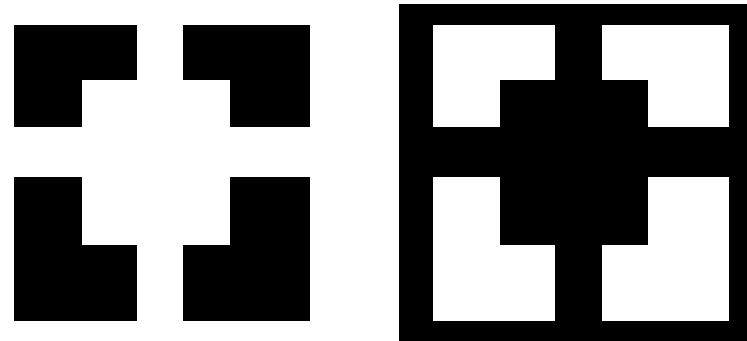


=



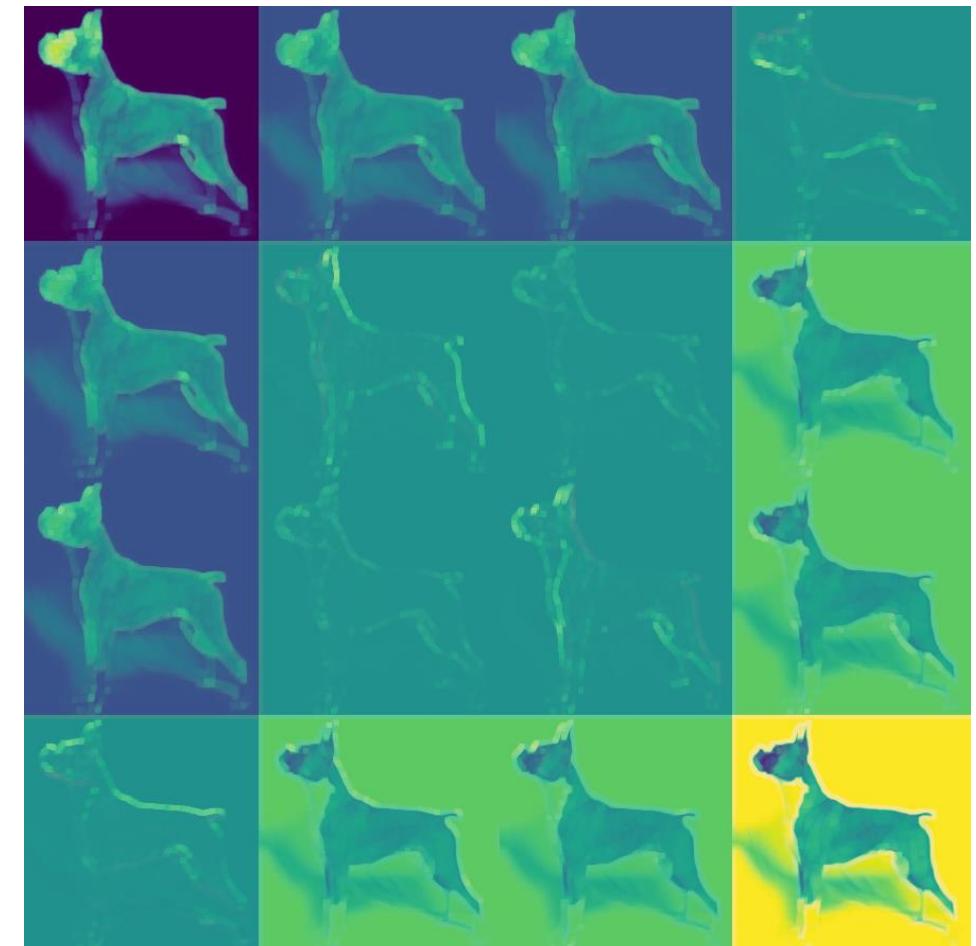
Basic operations

Convolution



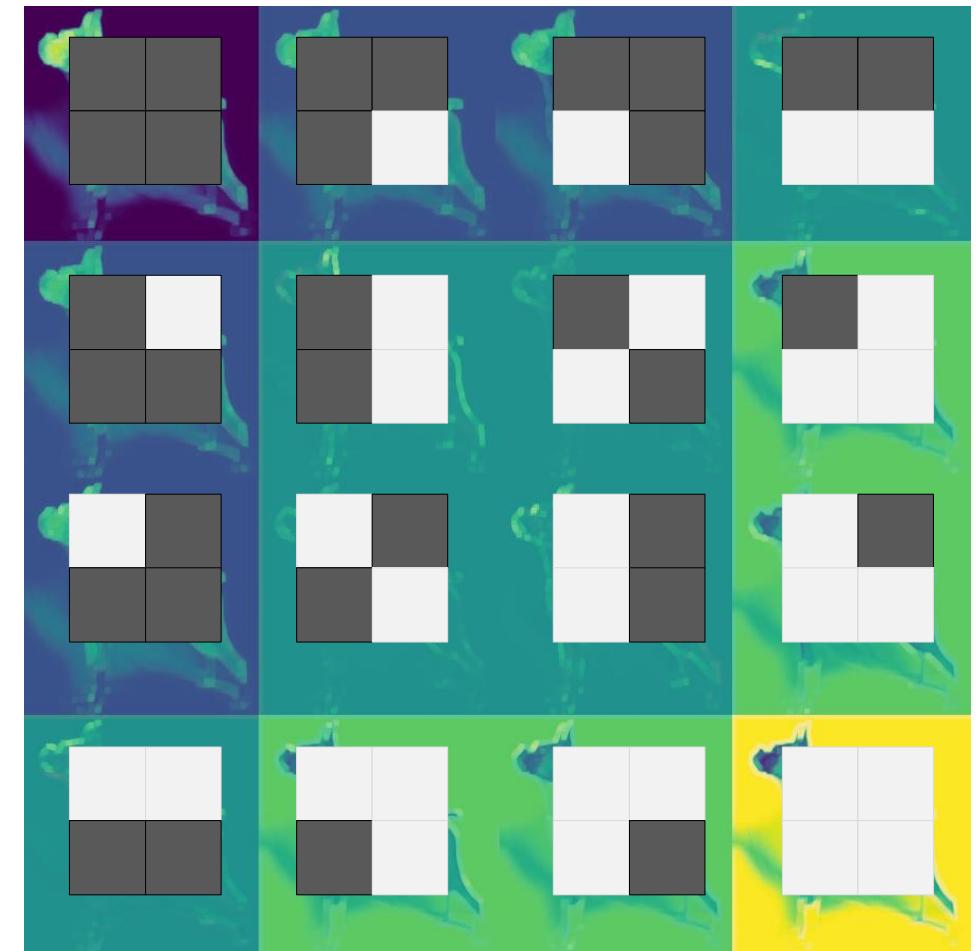
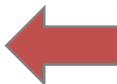
Basic operations

Convolution



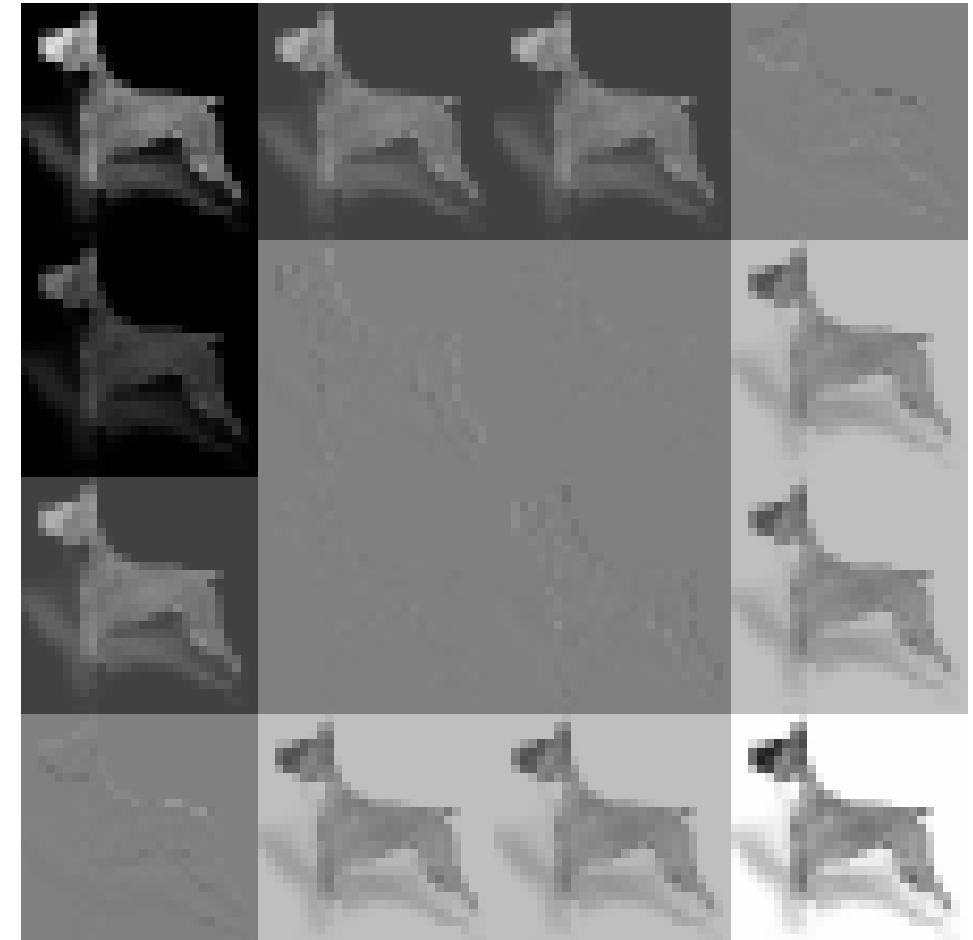
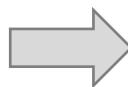
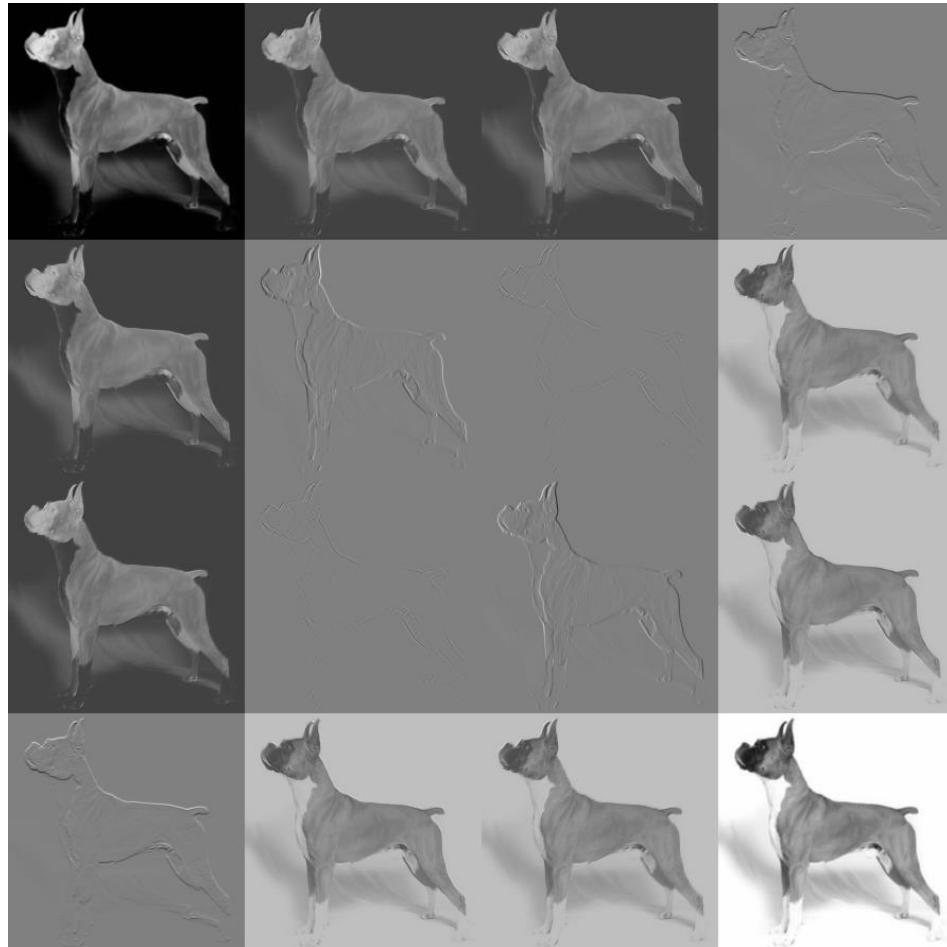
Basic operations

De-Convolution



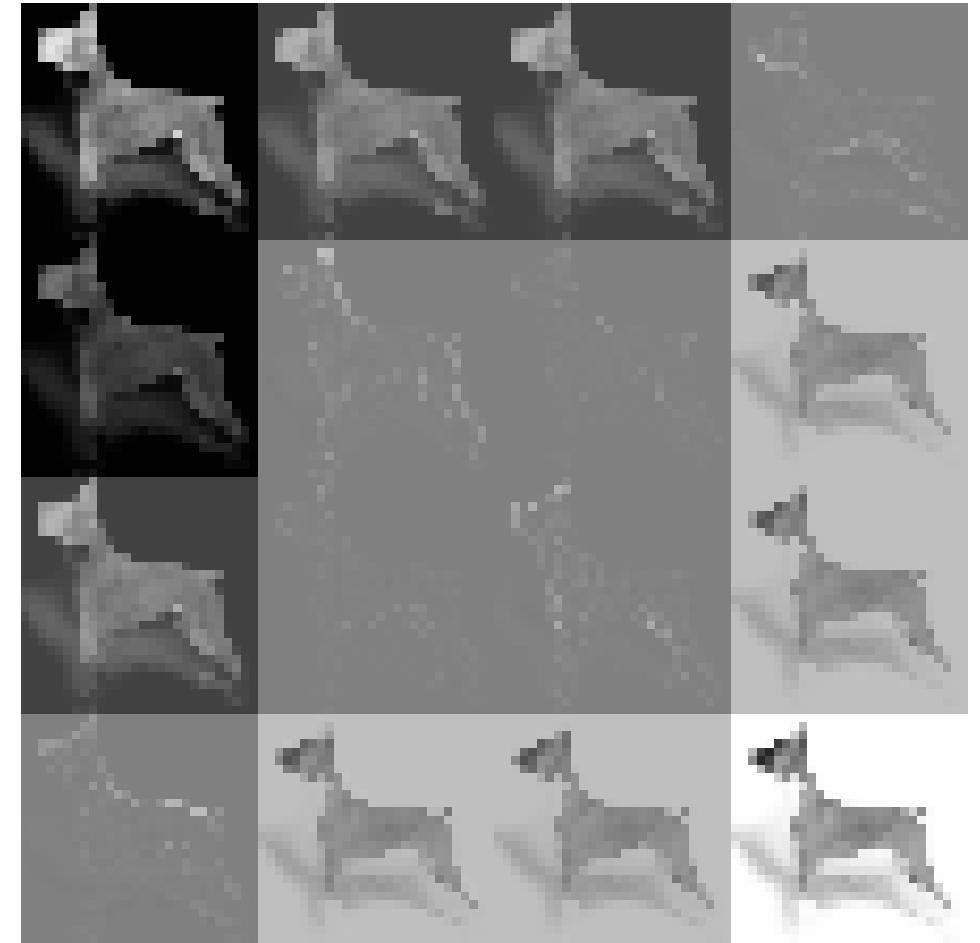
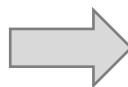
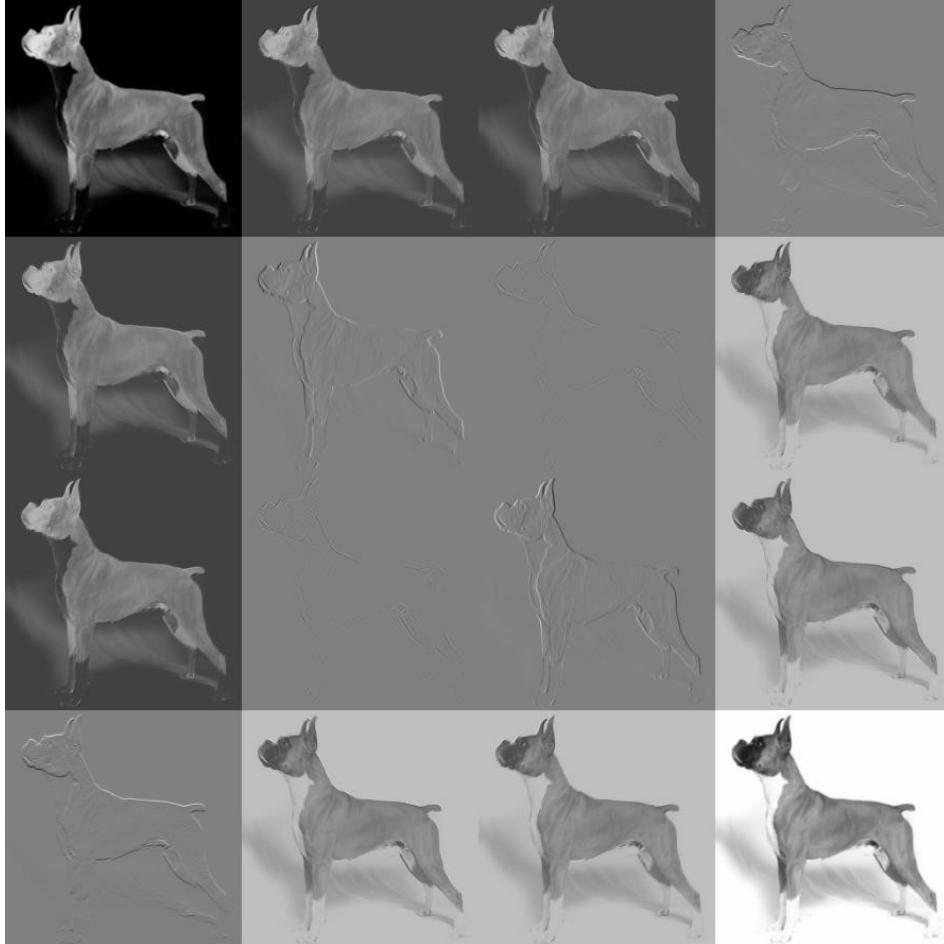
Basic operations

Pool: avg



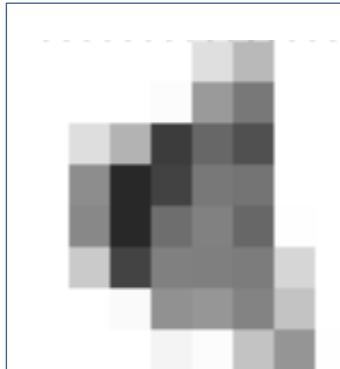
Basic operations

Pool: max



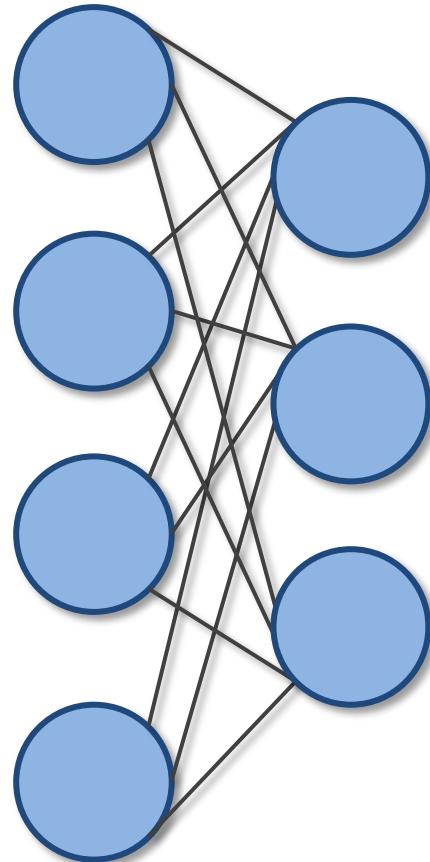
Basic operations

Flatten



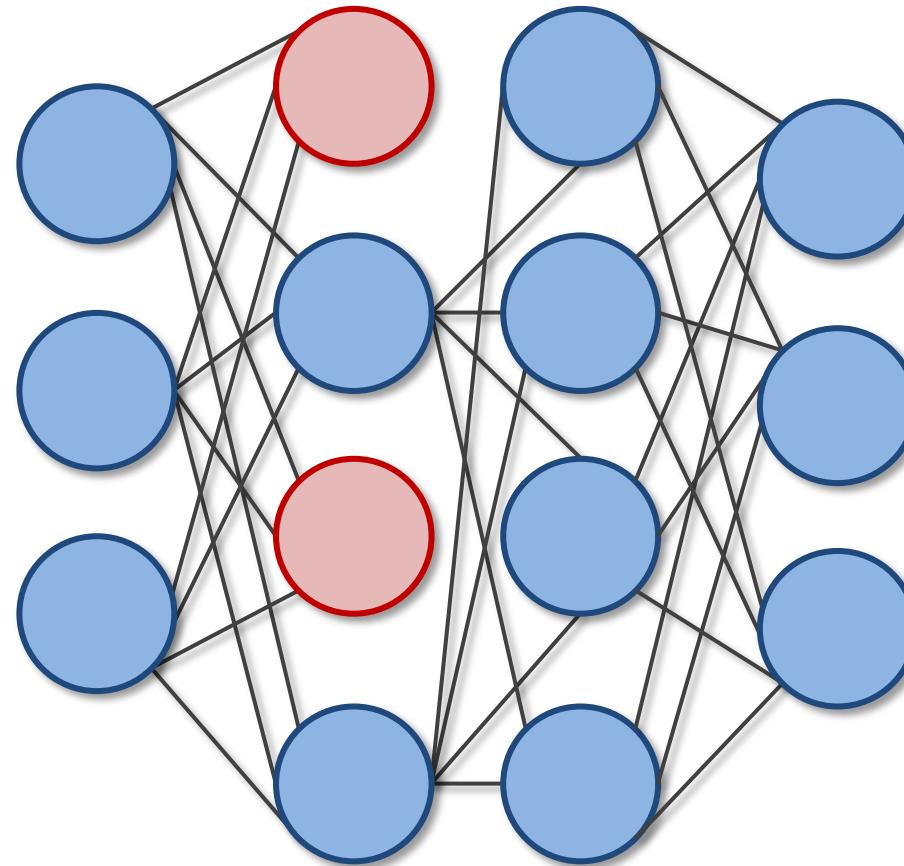
Basic operations

Fully Connected



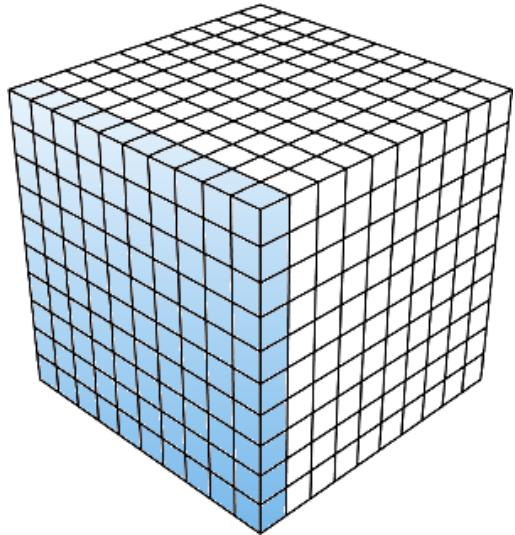
Basic operations

Dropout

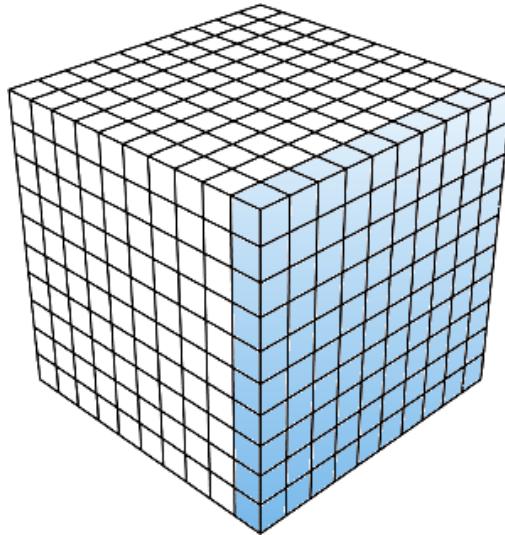


Basic operations

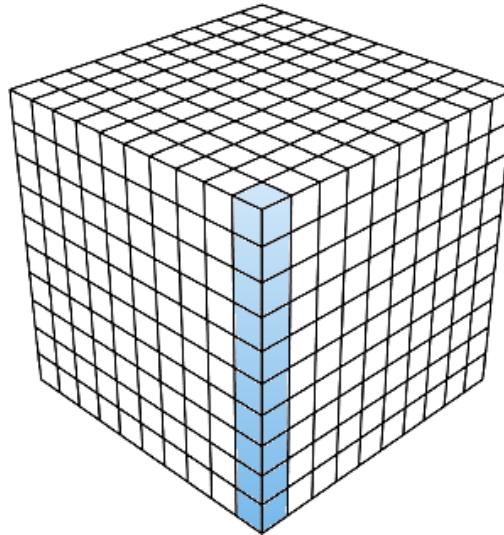
Normalization



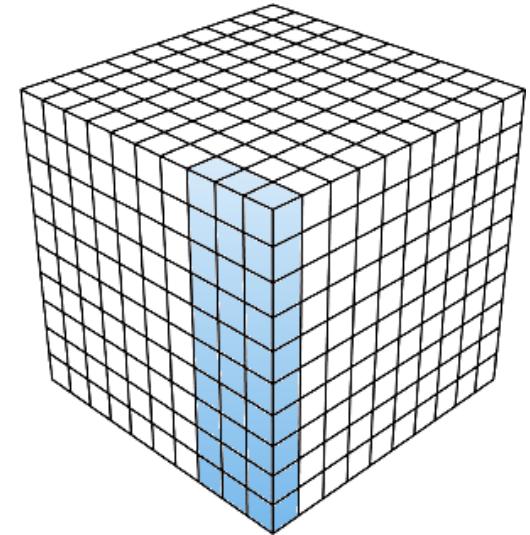
layer



batch



instance



group

Feature extraction

Feature extraction

Illustration

Feature extraction: Keras

```
from keras.applications.xception import Xception
from keras.applications.mobilenet import preprocess_input
from keras.models import Model
import cv2
import numpy
from keras import backend as K
K.set_image_dim_ordering('tf')
# -----
-----
def example_feature_extract_Keras():

    CNN = Xception()

    img = cv2.imread('data/ex-natural/dog/dog_0000.jpg')
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = cv2.resize(img, (224, 224)).astype(numpy.float32)

    model = Model(inputs=CNN.input, outputs=CNN.get_layer('predictions').output)
    prob = model.predict(preprocess_input(numpy.array([img])))

    model = Model(inputs=CNN.input, outputs=CNN.output)
    prob = model.predict(preprocess_input(numpy.array([img])))

    model = Model(inputs=CNN.input, outputs=CNN.get_layer('avg_pool').output)
    feature = model.predict(preprocess_input(numpy.array([img])))

    return
# -----
-----
if __name__ == '__main__':
    example_feature_extract_Keras()
```

Feature extraction: Tensorflow

```
import tensorflow_hub as hub
import cv2
import numpy
import tensorflow as tf
# -----
def example_feature_extract_TF():

    module = hub.Module("https://tfhub.dev/google/imagenet/resnet_v2_50/classification/1")
    img = cv2.imread('data/ex-natural/dog/dog_0000.jpg')
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = cv2.resize(img, (224, 224)).astype(numpy.float32)
    img = numpy.array([img]).astype(numpy.float32) / 255.0
    sess = tf.Session()
    sess.run(tf.global_variables_initializer())
    sess.run(tf.tables_initializer())

    outputs = module(dict(images=img), signature="image_classification", as_dict=True)
    prob = outputs["default"].eval(session=sess)[0]

    feature = module(img).eval(session=sess)

    sess.close()

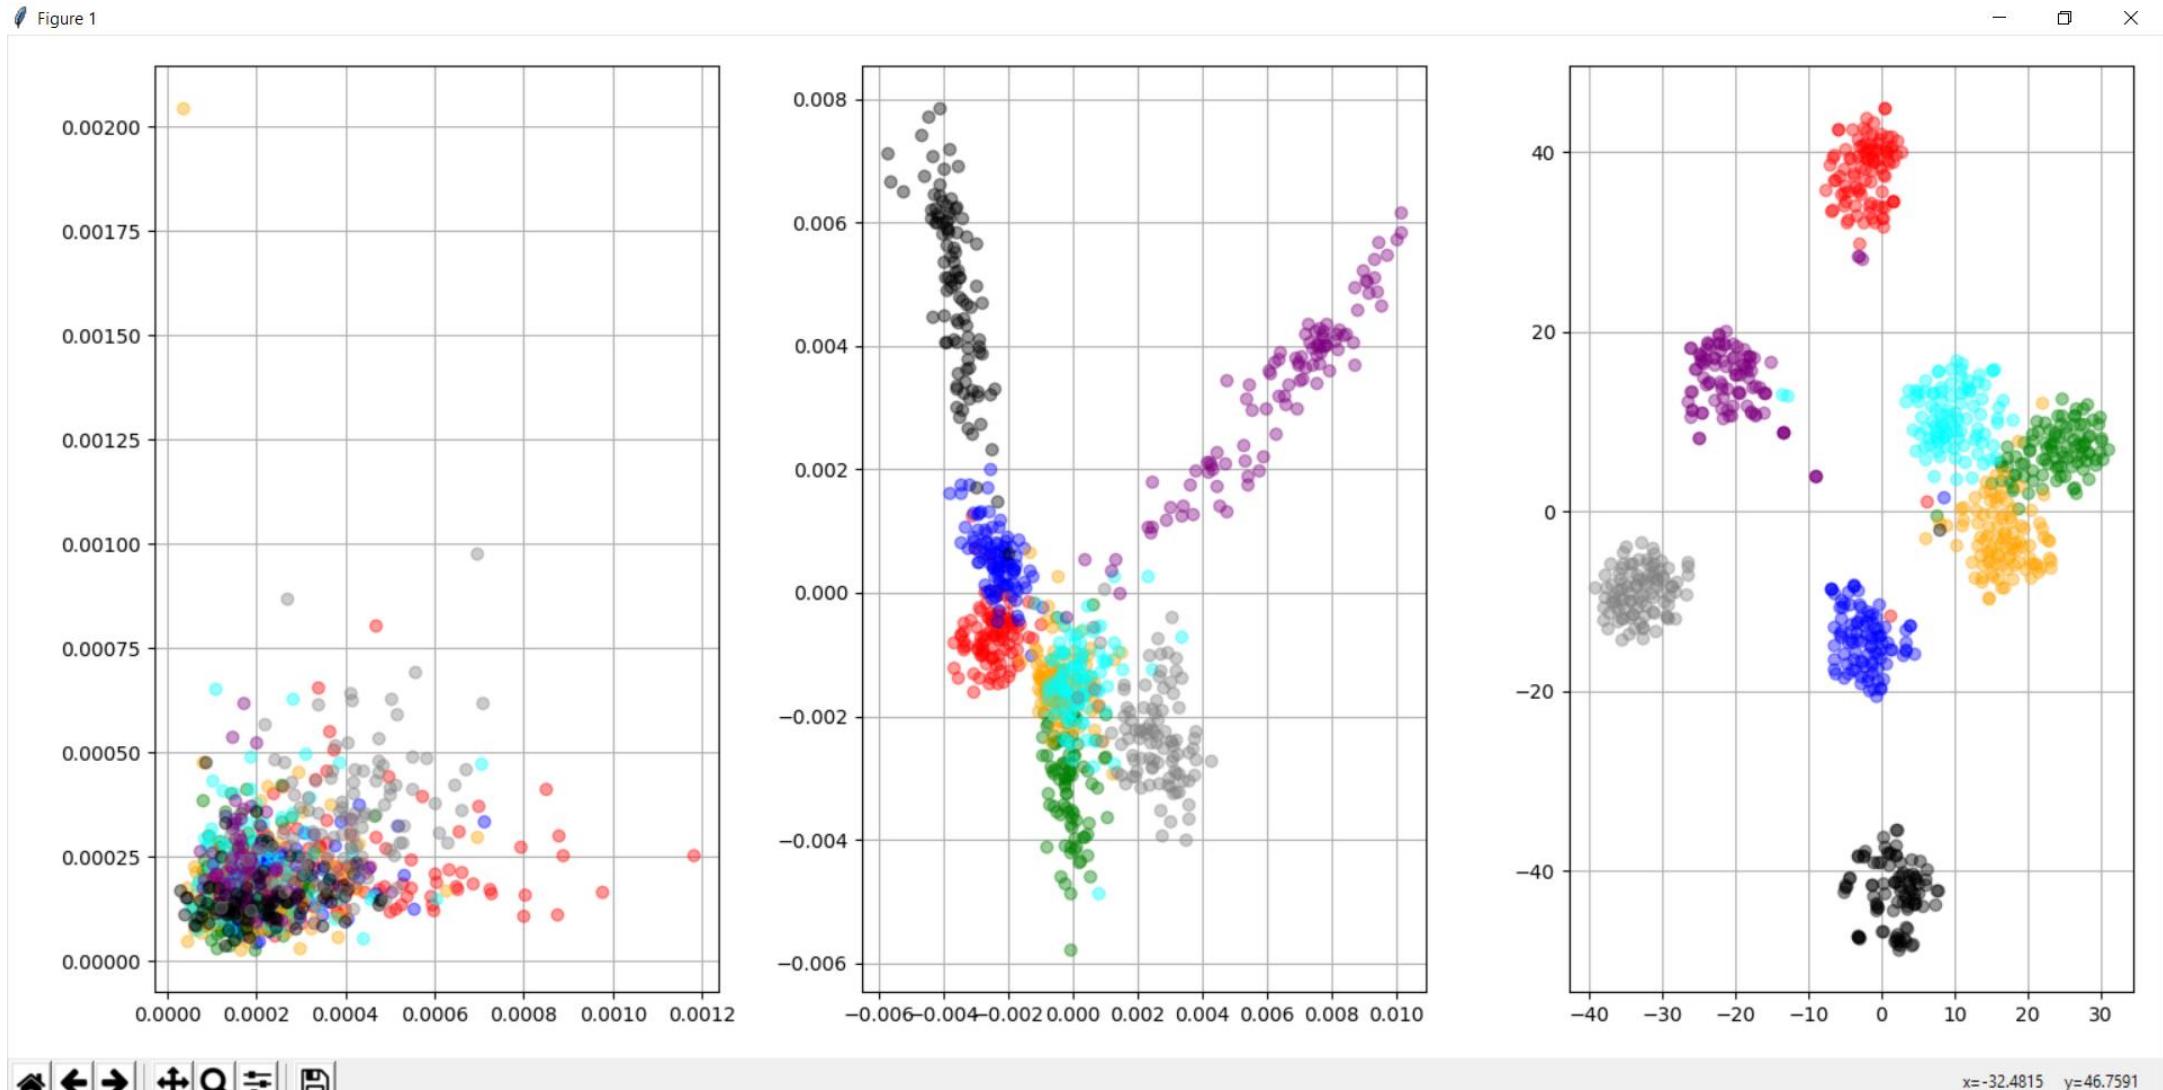
    return
# -----
if __name__ == '__main__':
    example_feature_extract_TF()
```

Feature visualization

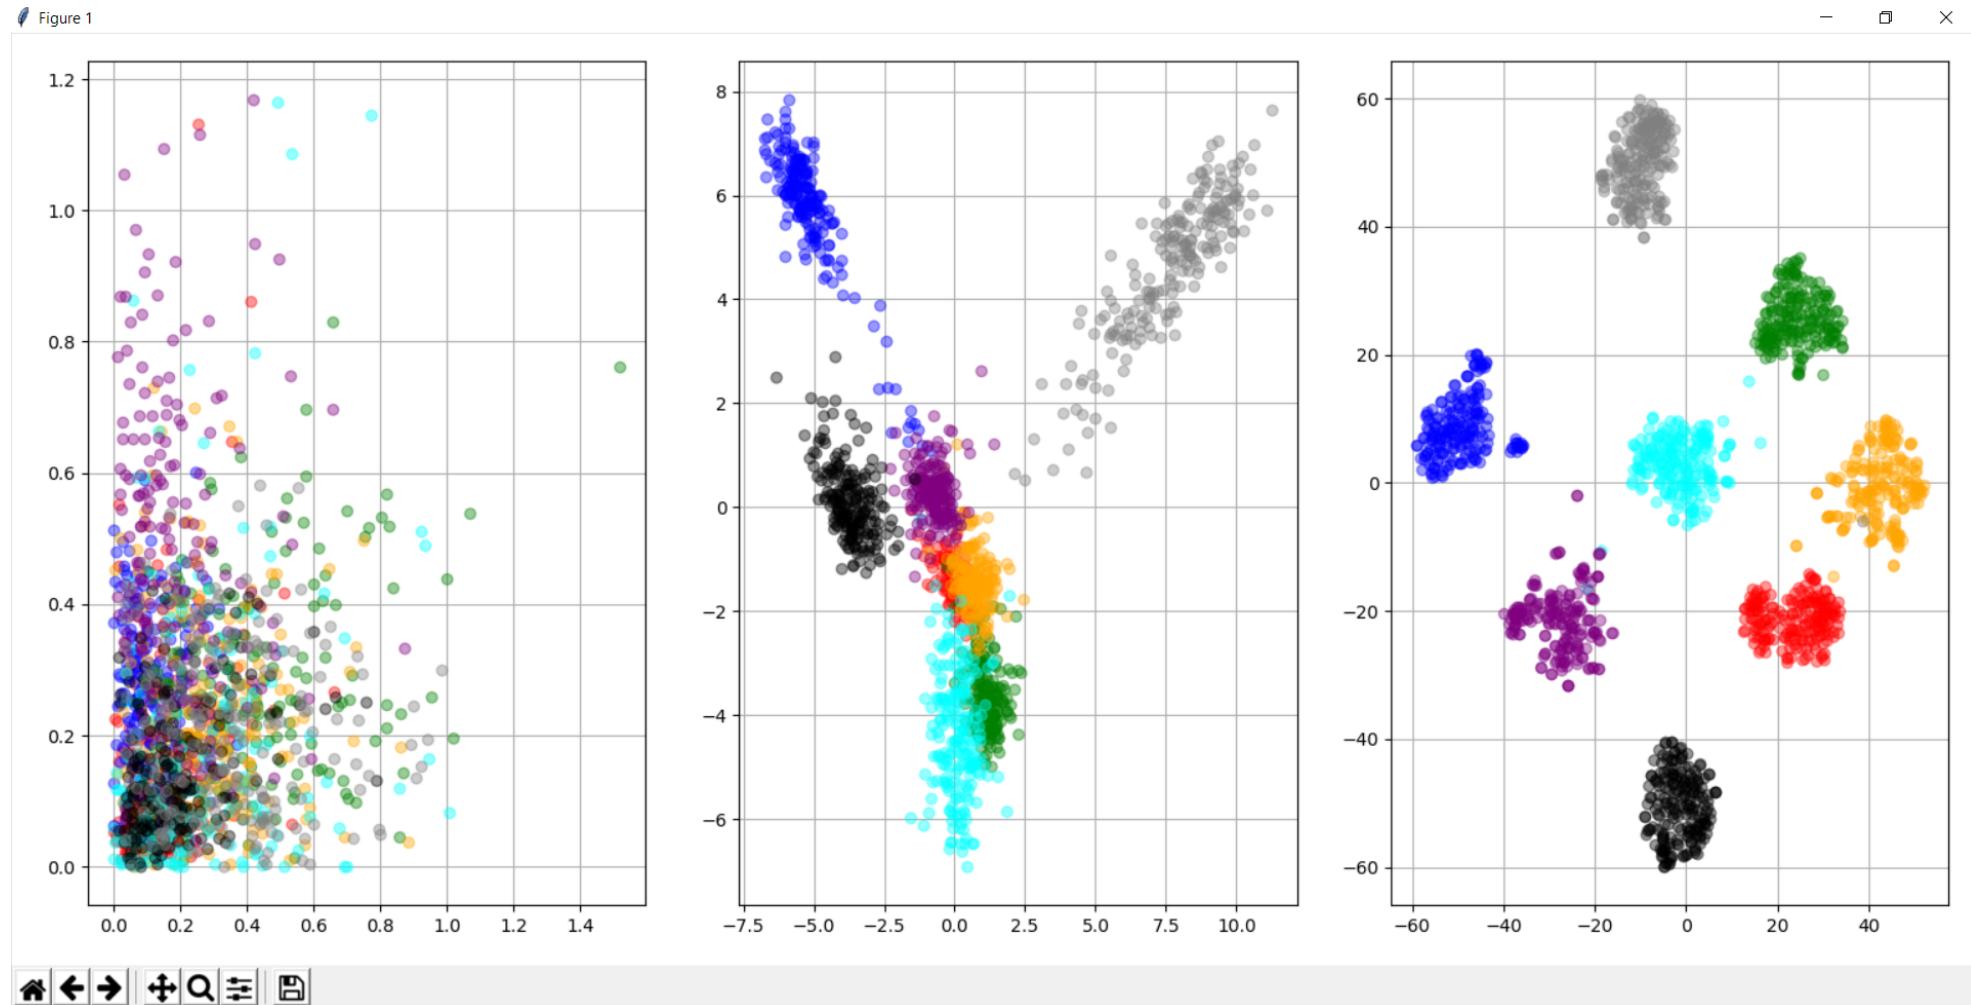
https://github.com/keras-team/keras/blob/master/examples/conv_filter_visualization.py

Feature visualization

Figure 1



Feature visualization



Feature extraction: Tensorflow

```
# -----
# from sklearn.manifold import TSNE
# import matplotlib.pyplot as plt
# import fnmatch
# from os import listdir
# from sklearn import decomposition
# import numpy
#
# -----
# import tools_IO
# import tools_ML
# -----
def example_visualize_features():

    path_input= 'data/features/CNN_Inception_TF/'

    patterns = fnmatch.filter(listdir(path_input),'*txt')
    for i in range (0,len(patterns)):
        patterns[i]=patterns[i].split('.')[0]

    patterns = numpy.array(patterns)
    ML = tools_ML.tools_ML(None)

    X,Y, filenames = ML.prepare_arrays_from_feature_files(path_input, patterns=patterns, feature_mask='.txt')
    X_PC = decomposition.PCA(n_components=2).fit_transform(X)
    X_TSNE = TSNE(n_components=2).fit_transform(X)

    tools_I0.plot_2D_scores_multi_Y(plt.subplot(1,3,1),X, Y)
    tools_I0.plot_2D_scores_multi_Y(plt.subplot(1,3,2),X_PC, Y)
    tools_I0.plot_2D_scores_multi_Y(plt.subplot(1,3,3),X_TSNE, Y)
    plt.tight_layout()
    plt.show()

    return
#
# if __name__ == '__main__':
#     example_visualize_features()
```

Visualization of layers and filters

Visualization of filters in Keras

```
for i in range (0, len(keras_model.layers)):  
    tensor = keras_model.layers[i]  
    if isinstance(tensor, Conv2D):  
        tensor = tensor.get_weights()[0]  
        tensor -= tensor.min()  
        tensor *= 255.0 / tensor.max() #(W,H,3,N)  
  
    if tensor.shape[2]==3:  
        cv2.imwrite(path_out + 'filter_%03d.png'%i, tensor_color_4D_to_image(tensor))  
    else:  
        cv2.imwrite(path_out + 'filter_%03d.png'%i, tensor_gray_4D_to_image(tensor,do_colorize=True))
```

Visualization of layers in Keras

```
layer_outputs = [layer.output for layer in keras_model.layers]
activation_model = Model(inputs=keras_model.input, outputs=layer_outputs)
activations = activation_model.predict(numpy.expand_dims(image, axis=0))

for i in range(0, len(activations)):
    tensor = activations[i][0]
    tensor -= tensor.min()
    tensor *= 255.0 / tensor.max()

    if len(tensor.shape) == 3:
        if need_transpose:
            tensor = tensor.transpose((1, 2, 0))

        if tensor.shape[2] == 3:
            cv2.imwrite(path_out + 'layer_%03d.png' % i, tensor)
        elif tensor.shape[2] != 3:
            tensor_gray_3D_to_image(tensor, do_colorize=True)
    elif len(tensor.shape) == 1:
        cv2.imwrite(path_out + 'layer_%03d.png' % i, tools_image.hitmap2d_to_viridis(tensor_gray_1D_to_image(tensor)))
```

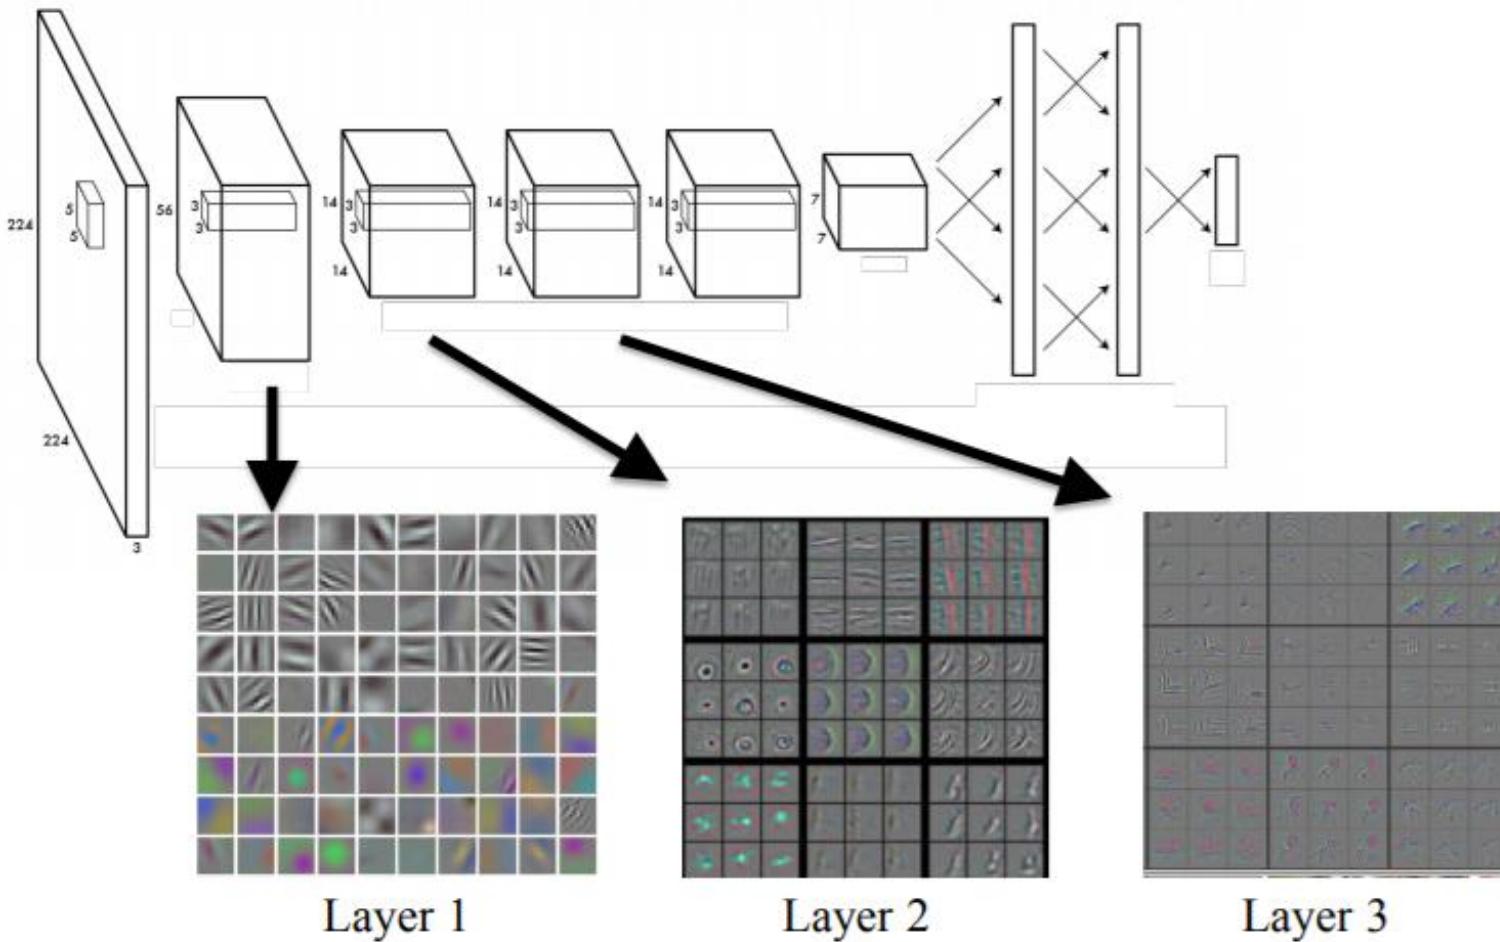
Visualization of layers in Tensorflow

```
init = tf.global_variables_initializer()
sess = tf.Session()
sess.run(init)

image = cv2.resize(cv2.imread(filename_input), (self.input_shape[0], self.input_shape[1]))

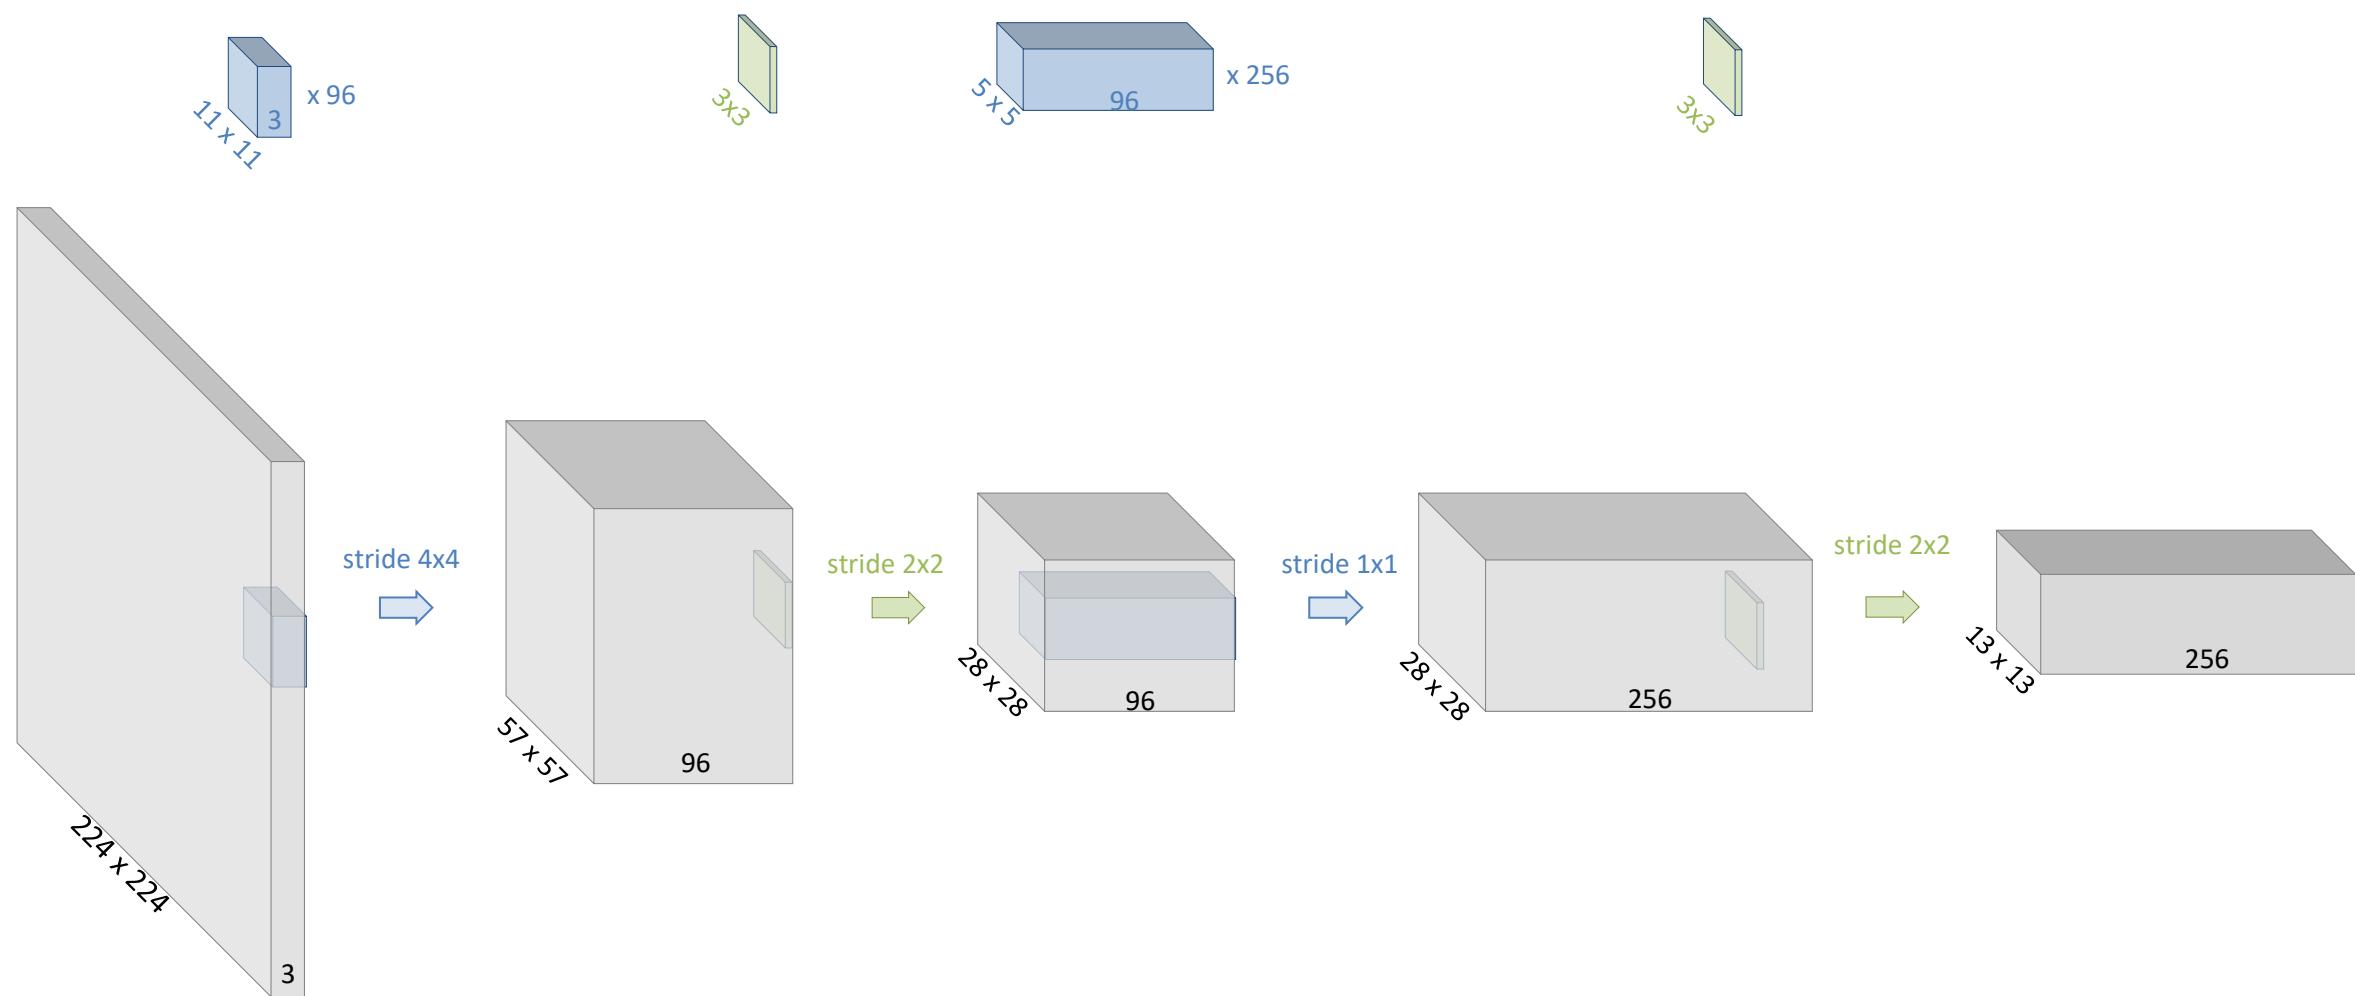
output = sess.run(self.layer_conv1, feed_dict={self.x: [image]})[0] #(57,57,96)
cv2.imwrite(path_output + 'layer01_conv1.png', tools_CNN_view.tensor_gray_3D_to_image(output,do_colorize=True))
```

Visualization of layers and filters

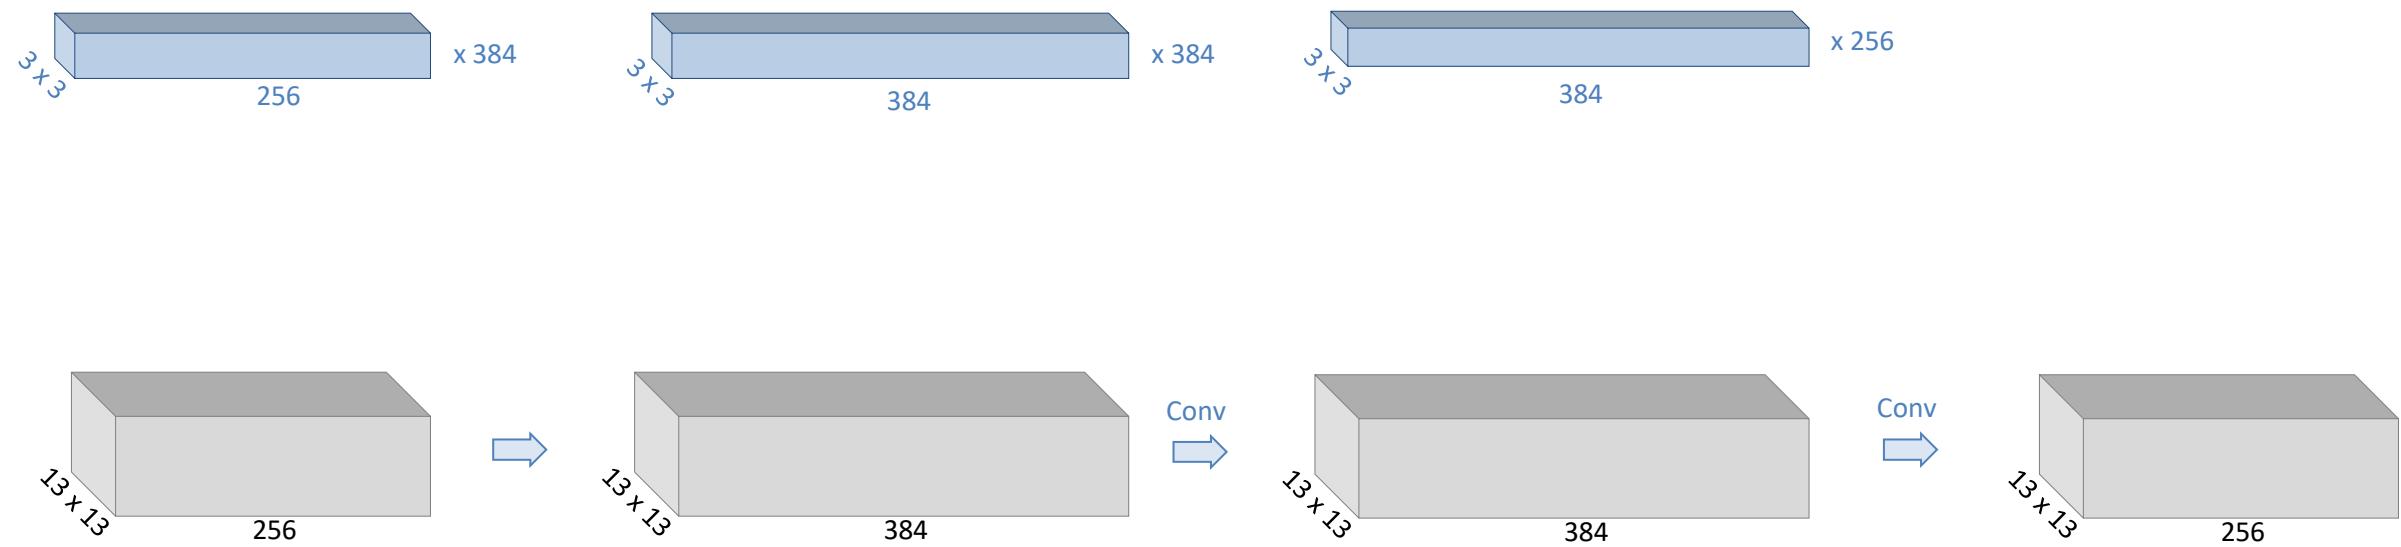


[Zeiler & Fergus ECCV 14]

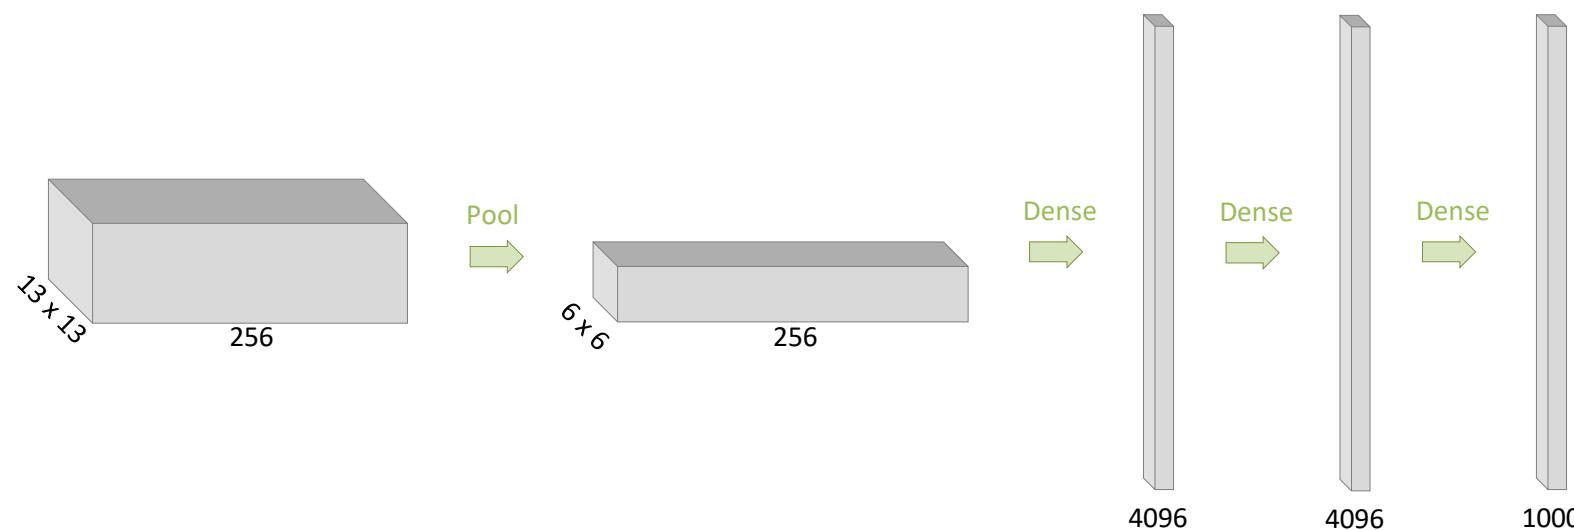
Visualization of layers and filters



Visualization of layers and filters



Visualization of layers and filters



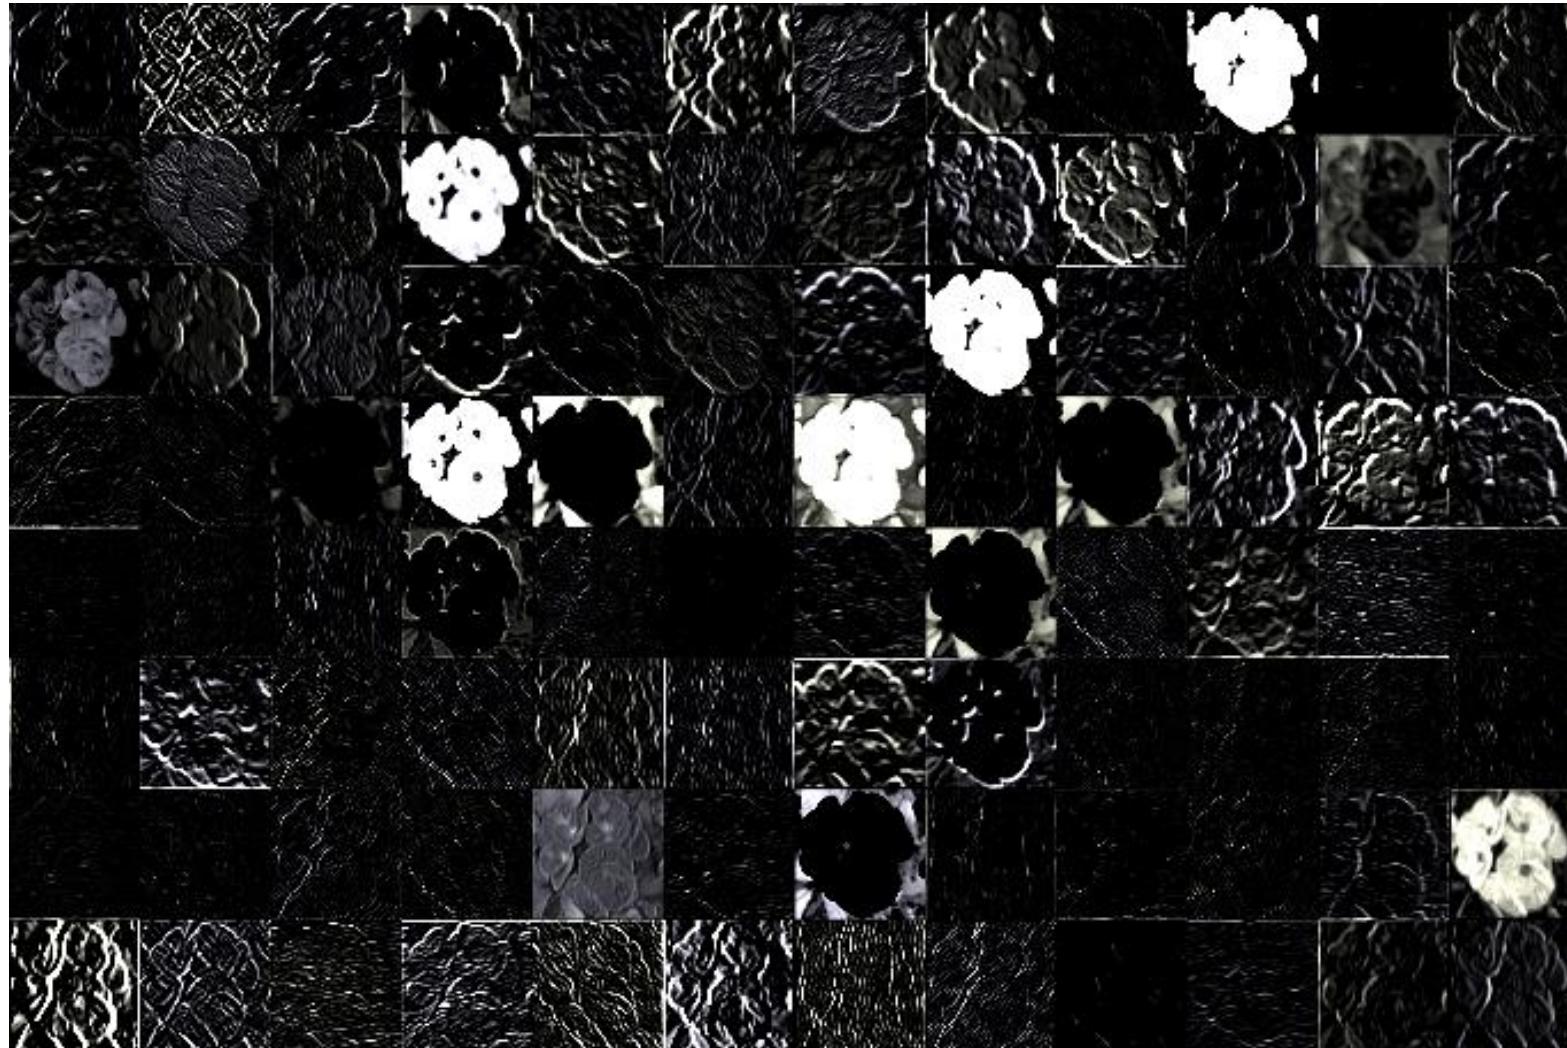
Visualization of layers and filters



Visualization of layers and filters



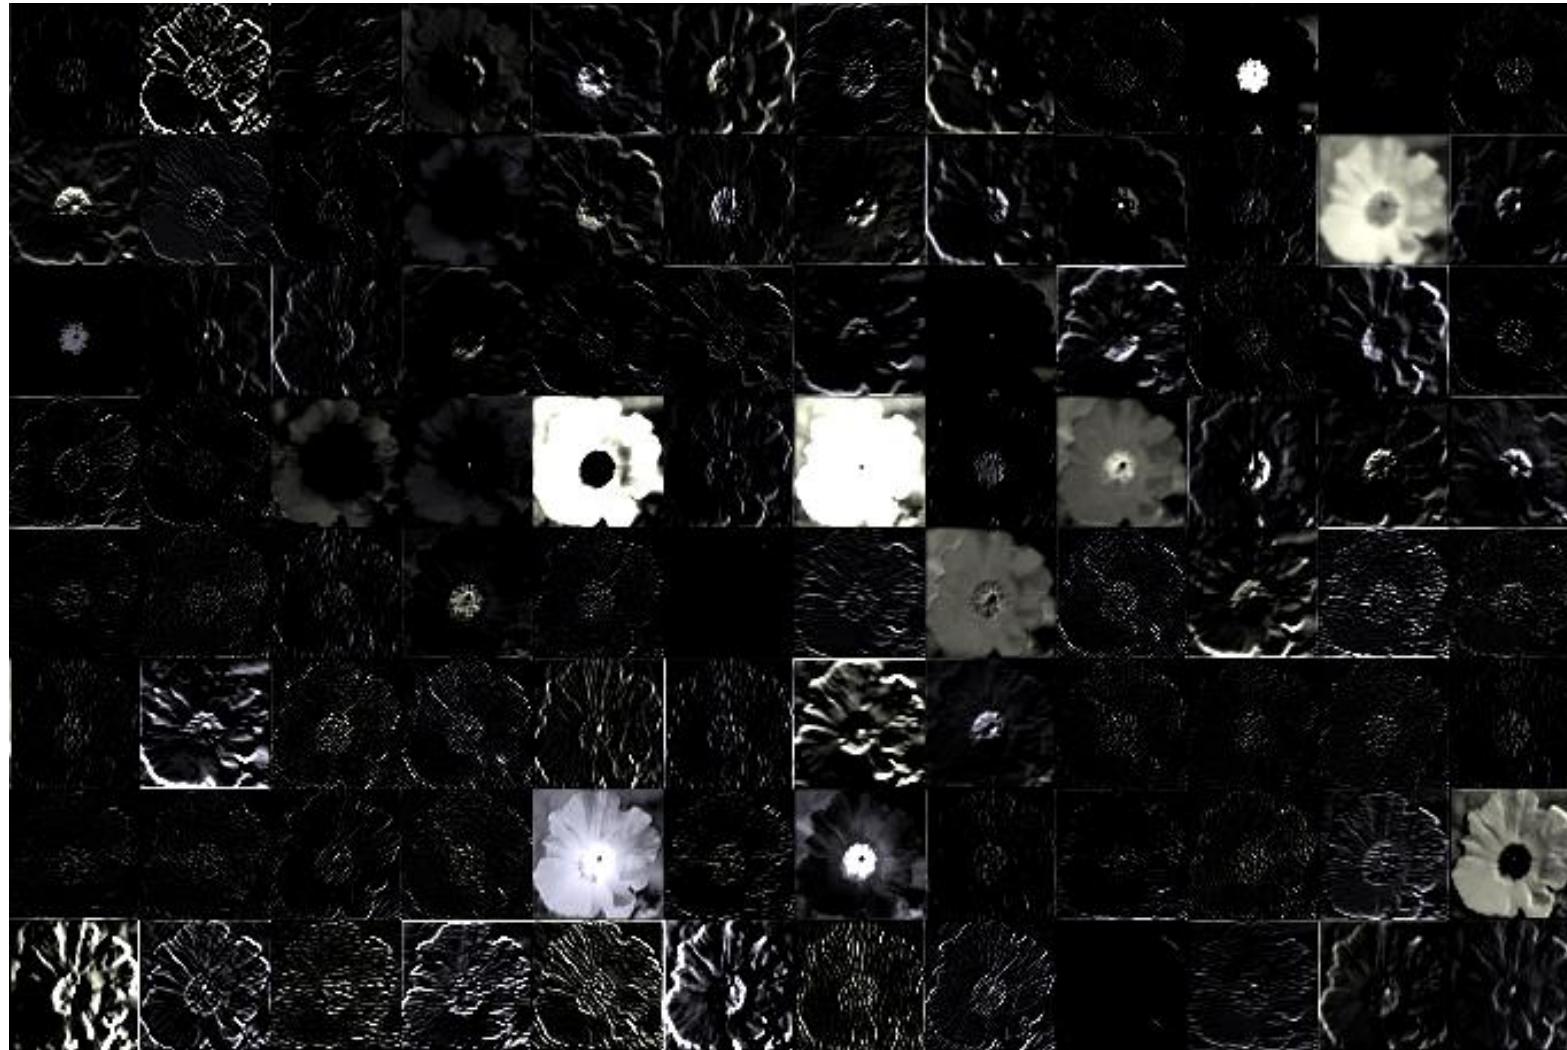
Visualization of layers and filters



Visualization of layers and filters



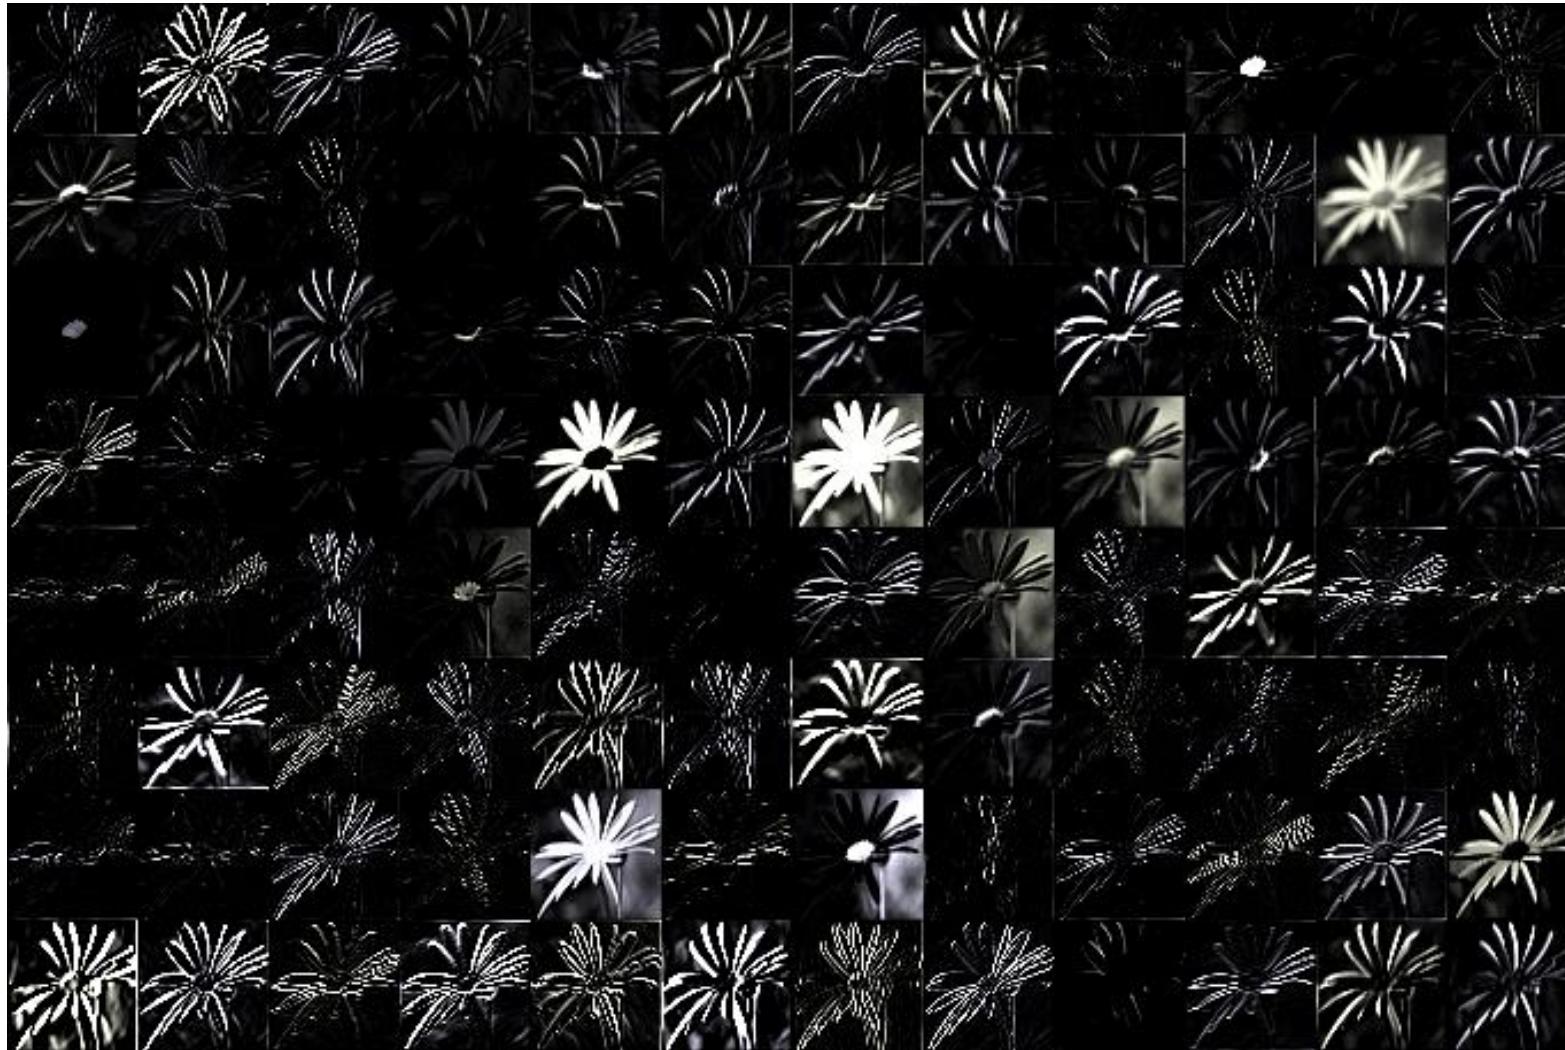
Visualization of layers and filters



Visualization of layers and filters



Visualization of layers and filters



Transfer learning

Training the models

Fine-tuning is the process in which the parameters of a trained model must be adjusted very precisely while we are trying to validate that model taking into account a small data set that does not belong to the train set.

That small validation data set comes from the same distribution as the data set used for the training of the model. The split of the available data to train and validation set is random.

<https://stats.stackexchange.com/questions/343763/fine-tuning-vs-transferlearning-vs-learning-from-scratch>

<https://datascience.stackexchange.com/questions/22302/what-is-the-different-between-fine-tuning-and-transfer-learning>

Transfer Learning or Domain Adaptation is related to the difference in the distribution of the train and test set. So it is something broader than Fine tuning, which means that we know a priori that the train and test come from different distribution and we are trying to tackle this problem with several techniques depending on the kind of difference, instead of just trying to adjust some parameters (usually we are doing this for reasons as preventing overfitting etc.)

Training the models

https://www.tensorflow.org/hub/tutorials/image_retraining

https://github.com/tensorflow/hub/blob/master/examples/image_retraining/retrain.py

<https://medium.com/tensorflow/introducing-tensorflow-hub-a-library-for-reusable-machine-learning-modules-in-tensorflow-cdee41fa18f9>

<https://github.com/tensorflow/hub/tree/master/examples/colab>

Fine tuning

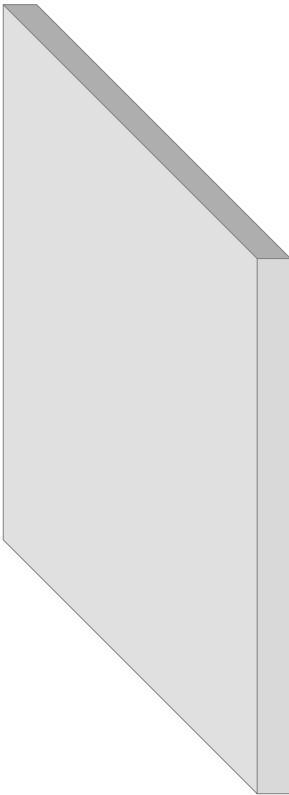
https://www.tensorflow.org/hub/fine_tuning

<https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>

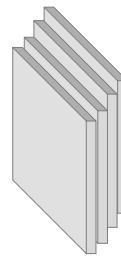
<https://keras.io/applications/#fine-tune-inceptionv3-on-a-new-set-of-classes>

Architecture of the popular networks

Architecture of the popular networks: AlexNet



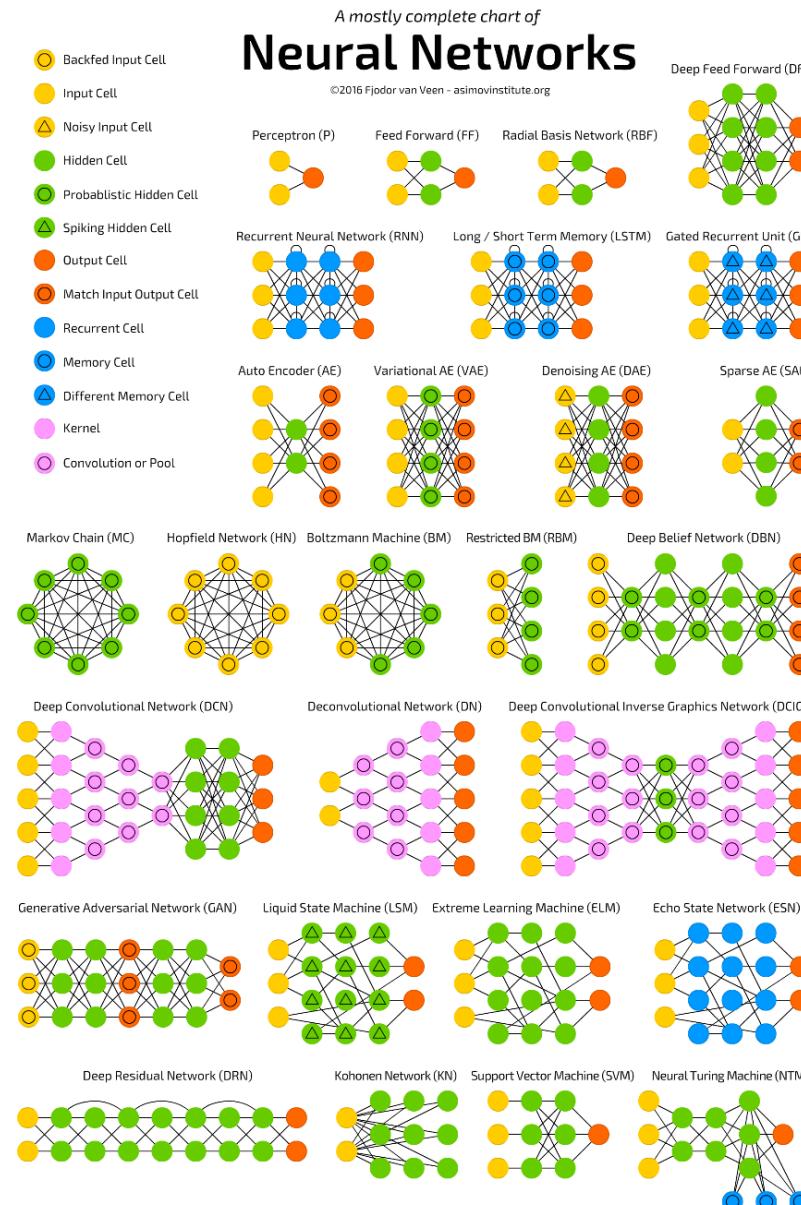
224 x 224 x 3



11 x 11 x 96

<https://lutzroeder.github.io/netron/>

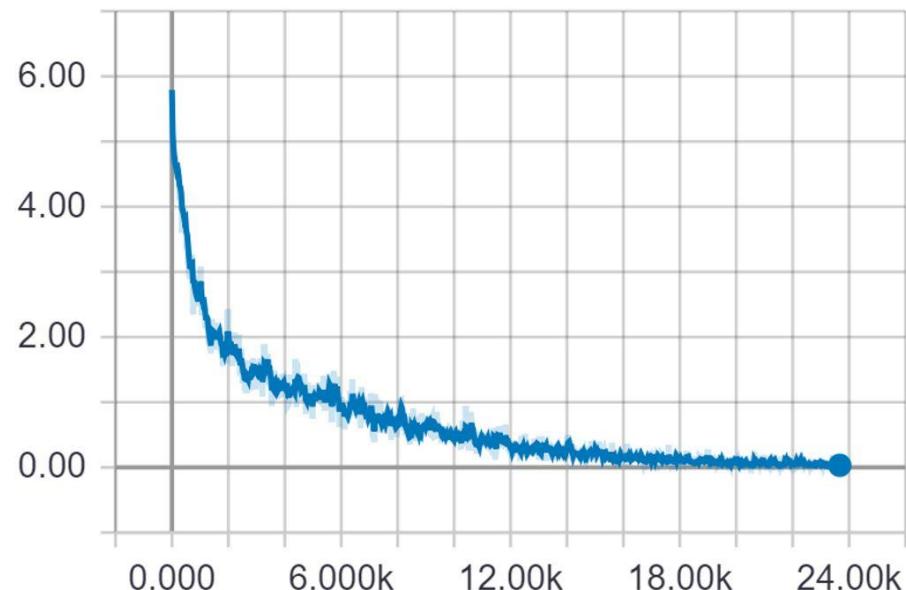
Architecture of the popular networks: AlexNet



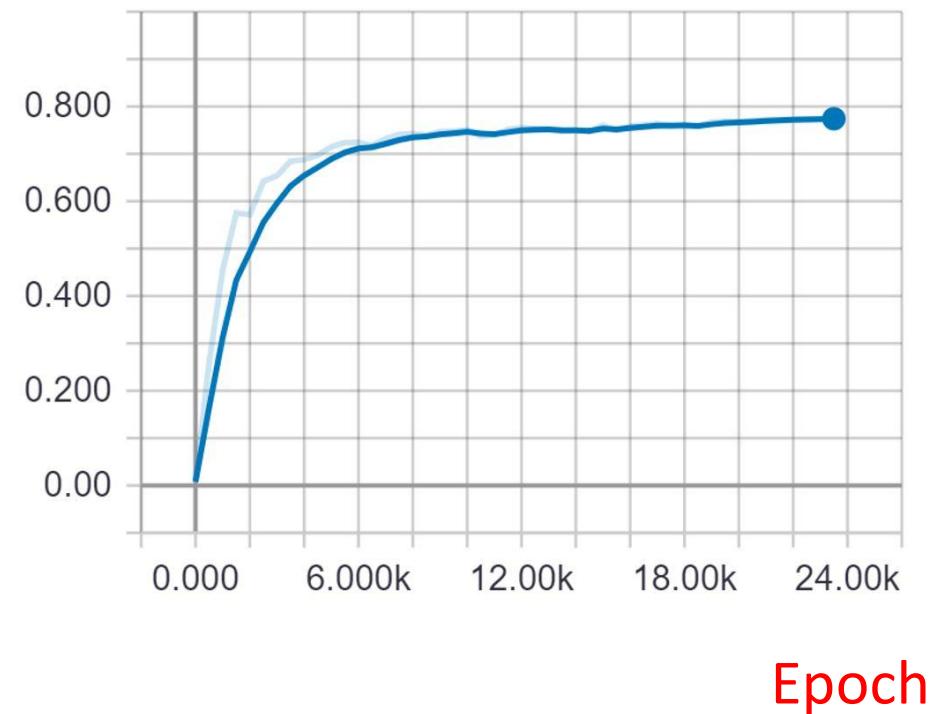
**Build, Train and Test
own CNN**

Build, Train and Test own CNN

Train Loss



Validation Accuracy



Build, Train and Test own CNN

graph

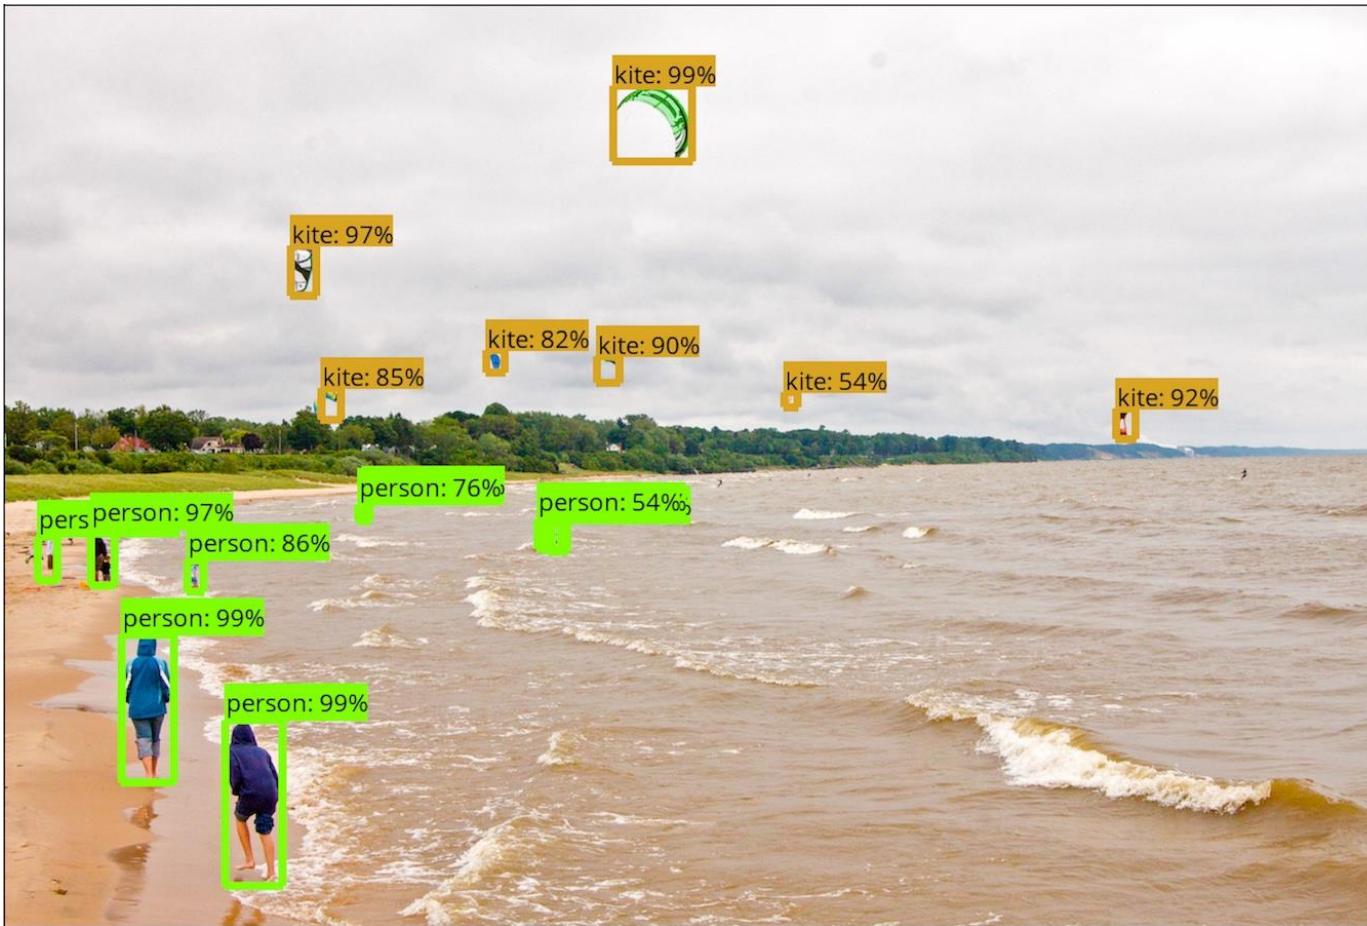
Build, Train and Test own CNN

Tensorflow vs Keras vs PyTorch vs Caffe vs Teano

Benchmarking the CNNs

Object detection

Object detection



https://github.com/tensorflow/models/tree/master/research/object_detection

Auto Encoders

<https://blog.keras.io/building-autoencoders-in-keras.html>

Style transfer

CAM

Class Activation Mapping

CAM

- The class activation map highlights the most informative image regions relevant to the predicted class. This map can be obtained by adding a global average pooling layer at the end of convolutional layers.
- Details of the implementation and more results can be found [here](#). Some results:



References

<https://www.youtube.com/playlist?list=PL3FW7Lu3i5JvHM8ljYj-zLfQRF3EO8sYv>

<http://cs231n.stanford.edu/2017/syllabus.html>

<https://www.jeremyjordan.me/convnet-architectures/>

<https://medium.com/activewizards-machine-learning-company/data-science-for-managers-mindmap-bb95bedc3154>

<https://habr.com/ru/company/cloud4y/blog/346968/>

<https://dou.ua/lenta/digests/ai-ml-digest-9/>

<https://ai.stackexchange.com/questions/5546/what-is-the-difference-between-a-convolutional-neural-network-and-a-regular-neur>

GAN

<https://github.com/junyanz/pytorch-CycleGAN-and-pix2pix>

<https://www.lyrn.ai/2018/12/26/a-style-based-generator-architecture-for-generative-adversarial-networks/>